
CS-E4640 2026

Introduction to Big Data Platforms

Hong-Linh Truong
Department of Computer Science
linh.truong@aalto.fi



07.1.2026

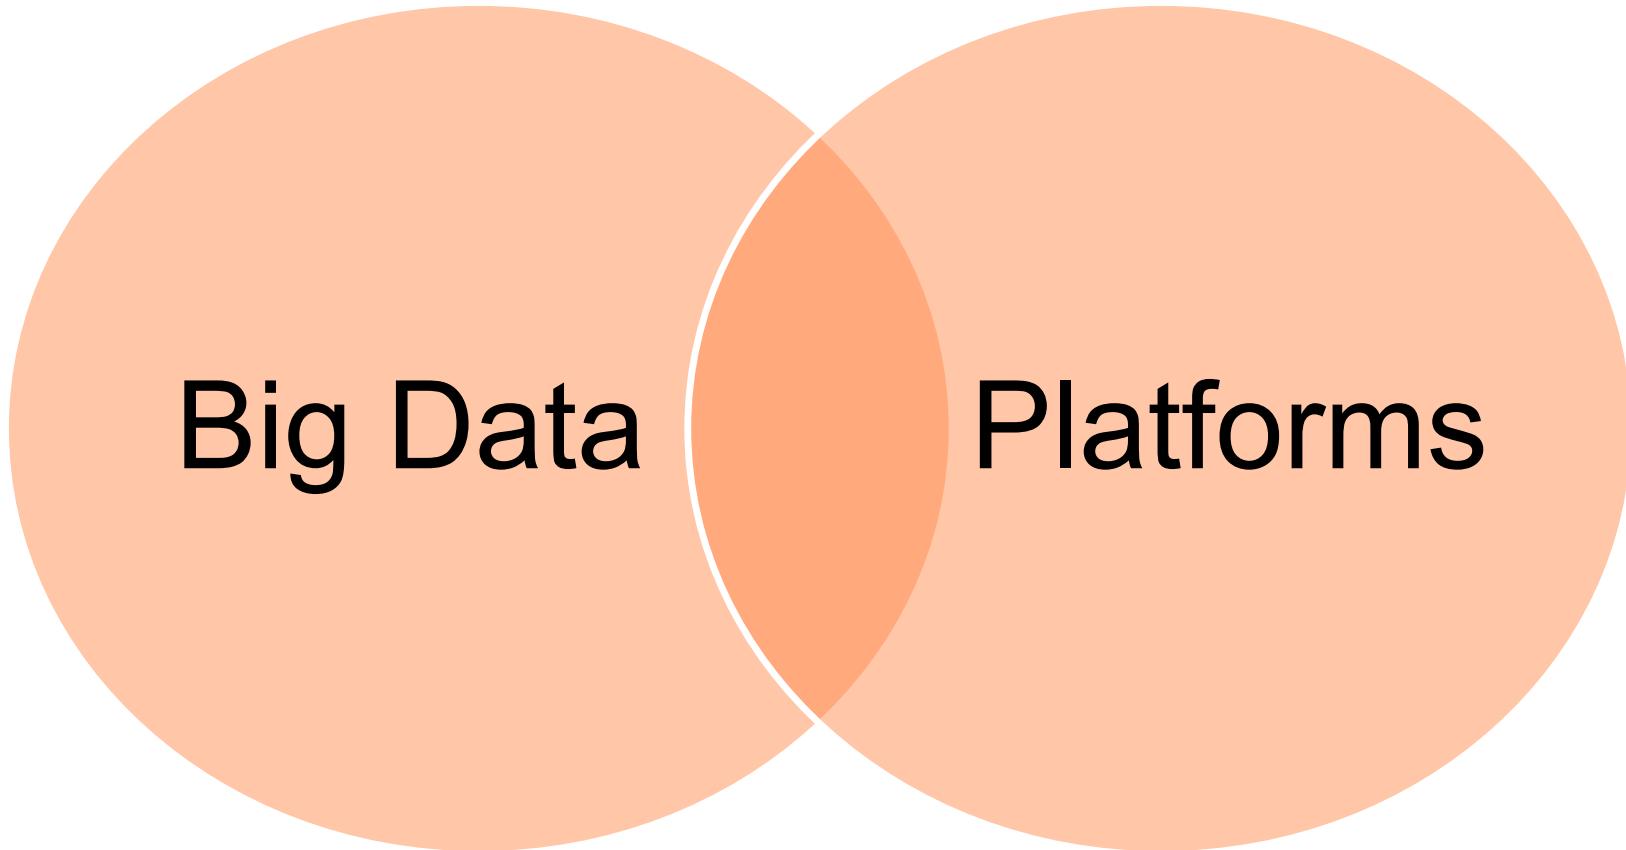


Content is available under
CC BY-SA 4.0 unless otherwise stated

Learning objectives

- Understand “big data” and “platforms” in big data platforms
- Capture high-level views of big data platforms and understand the role of big data platforms
- Understand key aspects in studying big data platforms

What are they?



Emerging factors: IoT/near real-time data, AI/ML requirements and integration, and computing continuum with edge-cloud-HPC

Data: facts, responses, events, measurements, etc.

< amazon_reviews_multilingual_US_v1_00.tsv (3.63 GB)						
Detail	Compact	Column	10 of 15 columns ▾			
▲ marketplace	▲ customer_id	▲ review_id	▲ product_id	# product_parent	▲ product_title	▲ product_c
US	53096384	R63J84G1L0X6R	1563890119	763187671	The Sandman Vol. 1: Preludes and Nocturnes	Books
US	53096399	R1BAL0A11Z06MT	1559947608	381720534	The 22 Immutable Laws of Marketing	Books
US	53096332	R1LLAY5W5PZUS4	0671701800	860650224	Contact	Books
US	53096335	R3R9VTJ82FXECQ	0425132153	624269601	Good Omens	Books
US	51747709	R1P5J3FNWTXY	0517122707	161411385	A Confederacy	Books

Source:
<https://www.kaggle.com/cynthiarempel/amazon-us-customer-reviews-dataset>

Row	unique_key	case_number	date	block	iucr	primary_type
1	7638368	HS442861	2010-08-03 12:10:00 UTC	007XX E 111TH ST	1305	CRIMINAL DAMAGE
2	7658504	HS461920	2010-07-26 05:00:00 UTC	061XX S MELVINA AVE	1305	CRIMINAL DAMAGE
3	7699434	HS447549	2010-08-05 06:12:00 UTC	036XX S ARCHER AVE	2210	LIQUOR LAW VIOLATION
4	7717246	HS524769	2010-09-20 02:30:00 UTC	057XX W AINSLIE ST	2851	PUBLIC PEACE VIOLATION
5	7721499	HS528767	2010-09-22 06:00:00 UTC	031XX S ASHLAND AVE	1121	DECEPTIVE PRACTICE
6	7734397	HS541125	2010-09-29 03:00:00 UTC	132XX S BALTIMORE AVE	1152	DECEPTIVE PRACTICE

Source: Chicago Crime, BigQuery



Sensor	PM2.5 $\mu\text{g}/\text{m}^3$
Median 7 Sens.	23
(+) #20890	25
(+) #26221	23
(+) #31206	13
(+) #31454	28
(+) #34144	20
(+) #43411	1
(+) #59291	36

Source: <https://maps.sensor.community/#8/52.917/8.817>

Big Data

- Extremely large, complex data sets (evolving in time and space)
 - need to be handled with new/different techniques
- Individual data items can be small or big
 - e.g., simple sensor events versus high quality satellite images
- Often characterized by V*
 - e.g., Volume, Variety, Velocity, and Veracity

Is it big data?

Annual Report 2019

annual reports archive 

This report follows directly on from the 2018 report, and analyses the uptake of Copernicus Sentinel data and the performance of the Sentinel Data Access System during the period 1 December 2018 to 30 November 2019 (referred to as Y2019).

By the end of the reporting period, the Sentinel Data Access System was supporting over 280,000 registered users, a daily publication rate of over 30,500 products/day, and an average daily download volume of 214 TiB. A total of 254 million products had been downloaded by users since the start of data access operations, consisting of a total data volume of 158.4 PiB. Over half of these downloads - 128 million - occurred during Y2019 alone. The report provides the detailed statistics behind these numbers, as well as examining the demographics of users, the status of agreements with collaborative and international partners, the challenges and solutions found by the Data Access Operations team in publishing and disseminating such huge volumes of data and evolving the System to cope with them, and the outlook for the future.

The 2019 Copernicus Sentinel Data Access Annual Report is available [here](#).



Source: <https://scihub.copernicus.eu/twiki/do/view/SciHubWebPortal/AnnualReport2019>

Is it big data

“As of today we have 60 PB of query-able event data stored in an S3 based data lake and about 10 PB of raw data is being scanned every day using Presto.”

Source: Aug, 2019

[https://eng.lyft.com/presto-infrastructure
-at-lyft-b10adb9db01](https://eng.lyft.com/presto-infrastructure-at-lyft-b10adb9db01)

e.g., 112M rows (2021 check)

The screenshot shows a dataset page for '2018 Yellow Taxi Trip Data' from the NYC OpenData portal. At the top, there's a navigation bar with 'View Data', 'Visualize', 'Export', 'API', and a three-dot menu. Below the title, a category 'Transportation' is shown. The main content area includes a brief description of the dataset, which contains yellow and green taxi trip records with various fields like pick-up and drop-off dates/times, locations, distances, fares, and passenger counts. It notes that the data was provided by the NYC Taxi and Limousine Commission (TLC). To the right, there's a sidebar with 'Updated' (February 8, 2020) and 'Data Provided by' (TLC). The central part of the page is titled 'About this Dataset'. It lists the last update as February 8, 2020, and provides details about data creation (September 24, 2018), views (41K), and downloads (5,258). It also shows the data provider (TLC) and dataset owner (NYC OpenData). On the right side, there are sections for 'Update' (frequency: historical data, automation: no, date made public: 10/19/2018), 'Dataset Information' (agency: TLC), 'Attachments' (a PDF file), and 'Topics' (category: Transportation, tags: none).

Snapshot from: <https://data.cityofnewyork.us/Transportation/2018-Yellow-Taxi-Trip-Data/t29m-gskq>

Is it big data?

Texts used for training LLMs

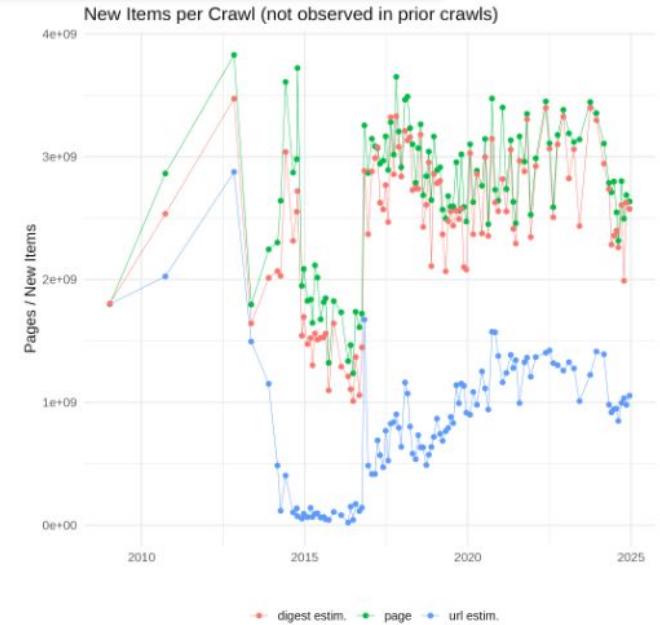
Statistics of Common Crawl Monthly Archives

Number of pages, distribution of top-level domains, crawl overlaps, etc. - basic metrics about Common Crawl Monthly Crawl Archives

Latest crawl: CC-MAIN-2024-51

- [Home](#)
- [Size of crawls](#)
- [Top-level domains](#)
- [Registered domains](#)
- [Crawler metrics](#)
- [Crawl overlaps](#)
- [Media types](#)
- [Character sets](#)
- [Languages](#)

[View the Project on GitHub](#)



Source: <https://commoncrawl.org/>

A!

Characterize big data with V*

- **Volume:**
 - big size, large data set, massive of small data
- **Variety:**
 - complex, different formats/structures, types of data and their links, states/readiness of data
- **Velocity:**
 - generating speed, data movement speed
- **Veracity:**
 - quality is very different (timeliness, accuracy, etc.)

Why do we have big data now?

- Social media data generated by human activities
 - Meta/Facebook, TikTok, Twitter, Instagram, etc.
- Internet of Things (IoT)/Machine-to-Machine (M2M)/Industry 4.0
 - data generated from monitoring of equipment, infrastructures and environments
- Advanced sciences data generated by advanced instruments
 - earth observation from satellites/telescopes (Sentinel, James Webb)
- Personal and disease information (e.g., healthcare)
- Business-related customer data
- Asset management and lodging (e.g., cars, homes)
- Software systems (e.g., traces, logs and test results)

Big “operational” and “analytical” data

- Operational data
 - for business/system operations, reads/writes (update), OLTP (online transaction processing)
- Analytical data
 - “non-operational data”, for understanding and optimization of business/subjects/systems/behaviors, historical/integrated data, write once/read many (no update) , data lakes, lakehouses, data science/ML workflows, OLAP (online analytical processing)
- Both types of data are integrated in data platforms
 - different methods/techniques
 - **bridges** between operational data and analytical data

Why do we need to care?

- Because of the values of data!
 - Current “hot topics”: Large-Language models (LLMs)/Gen-AI
- Top-down: Data economy
 - more data \Rightarrow more insights \Rightarrow better decision making \Rightarrow more business innovation and successes
- Bottom-up
 - understanding \Rightarrow optimizing \Rightarrow saving cost/creating new values
- “The Unreasonable Effectiveness of Data” principle \Rightarrow with **more data**, the same algorithm performs much **better!**

Alon Halevy, Peter Norvig, and Fernando Pereira. 2009. The Unreasonable Effectiveness of Data. IEEE Intelligent Systems 24, 2 (March 2009).
<http://static.googleusercontent.com/media/research.google.com/en//pubs/archive/35179.pdf>

What are platforms?

A!

Hong-Linh Truong

CS-E4640 Big Data Platforms

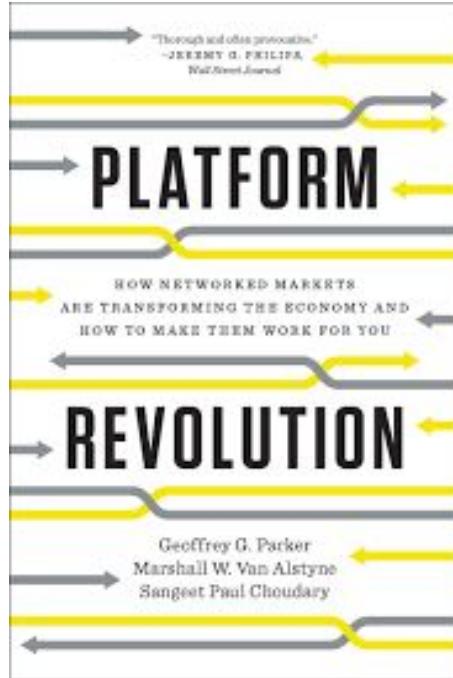
Spring 2026

07.1.2026

13

Example of platforms

Let us see from the business viewpoint from “Platform Revolution”:



Disruptive platforms: Airbnb, Amazon, Uber, Alibaba, Instagram, Meta/Facebook, Youtube, etc.

<https://www.amazon.com/Platform-Revolution-Networked-Markets-Transforming/>

The “Platform Revolution’s definition of a platform (from a business viewpoint)

“A platform is a business based on enabling value-creating interactions between external producers and consumers. The platform provides an open, participative infrastructure for these interactions and sets governance conditions for them. The platform’s overarching purpose: to consummate matches among users and facilitate the exchange of goods, services, or social currency, thereby enabling value creation for all participation”

Source: Geoffrey G. Parker, Van Alstyne, Marshall W. Van Alstyne , Sangeet Paul Choudary, *Platform Revolution: How Networked Markets Are Transforming the Economy - and How to Make Them Work for You*, March 28, 2016

An interpretation of platforms for big data

- Being large-scale service platforms, e.g.
 - On-demand computing platforms for data-centric products
 - On-demand analytics service platforms
 - On-demand data management platforms
- Enabling interactions between big data producers and big data consumers
 - Integration, management, analysis, optimization
- Facilitating the exchange of big data and products centered around data
- Not just a database or data marketplace (even they are big!)

Platforms support data as “asset” and “product” perspectives

- Known views about data
 - “data as an asset”
 - “data as a product”
- “Data as an asset” (see “Design data governance”)
 - something is very valuable that must be managed and exploited for the benefit of the owner
- “Data as a product” (see “Data mesh”)
 - product thinking for how one should process, manage and produce data to be used and satisfied by the data user

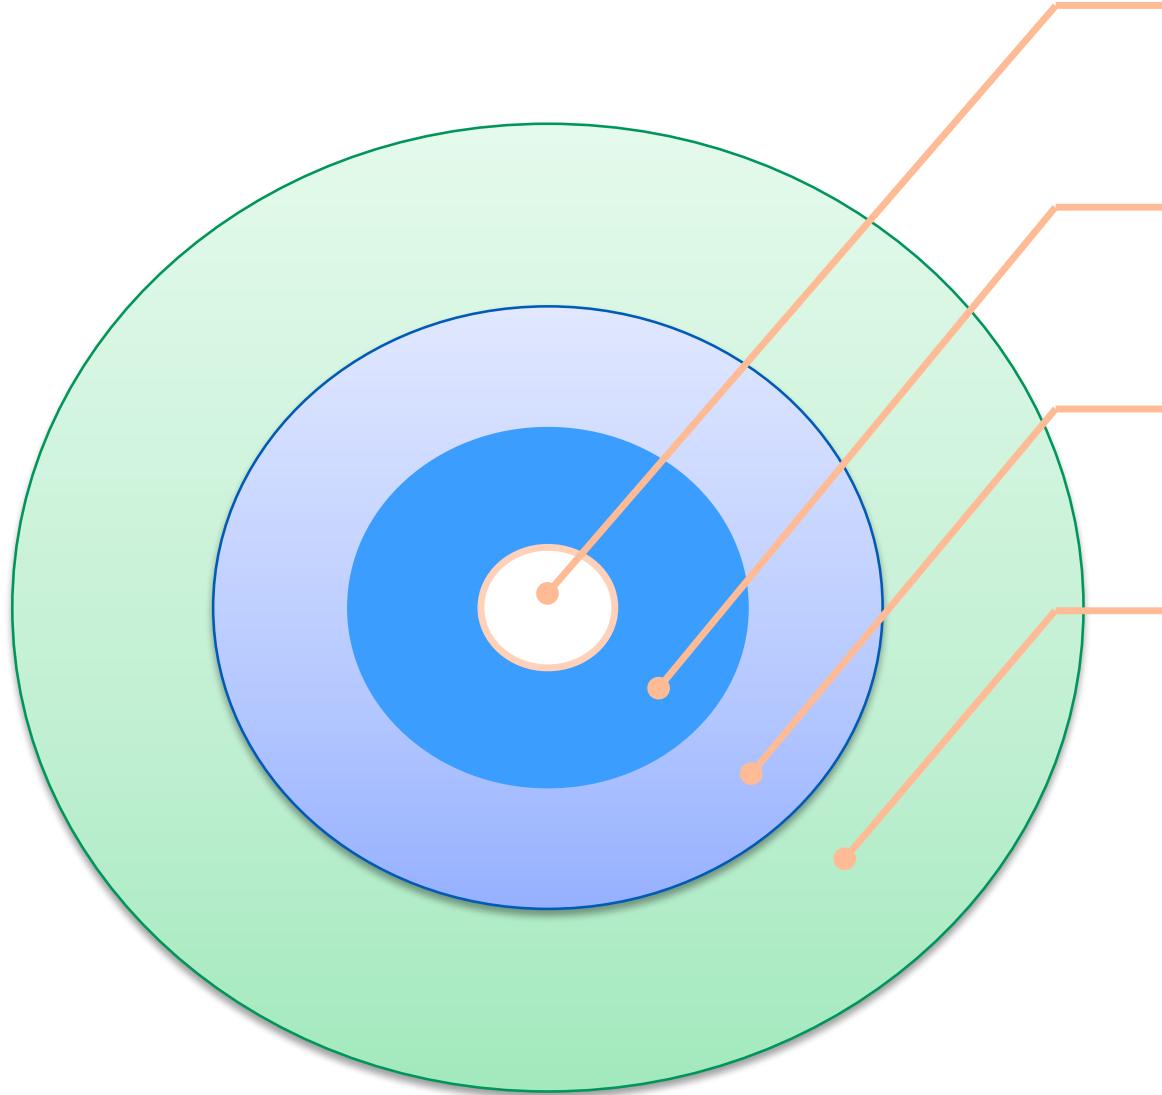
[“Data as an asset”]: Vijay Khatri and Carol V. Brown. 2010. Designing data governance. Commun. ACM 53, 1 (January 2010), 148–152.
<https://doi.org/10.1145/1629175.1629210>

[“Data as a product”]: Zhamak Dehghani, Data Mesh, 2022

Big data platforms

- Data-centric services
 - a lot of data with different types and added continuously
 - complex data governance and technological infrastructures
- Extensibility
 - allowing new data, new services, components to be added and integrated
- With diverse types of stakeholders
 - data consumers, data providers, and data integrators
 - service consumers, service providers and service integrators
 - regulators/auditors, etc.

Big data platforms: Onion architecture



Data-centric Virtualized Infrastructures

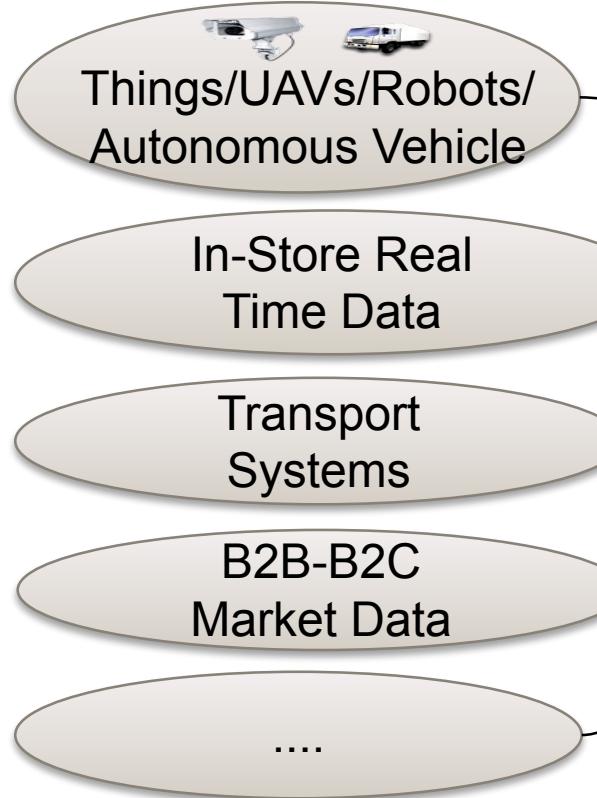
Middleware Platforms: building, deploying, operating and scaling reliable big data services and applications

Big data services & applications

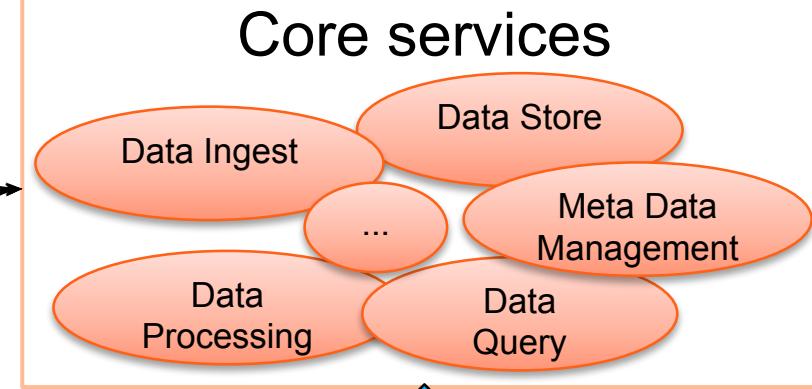
Big data platforms consumers and producers: Sensors, Things, Equipment, Robots, Industrial Processes, People

Analytics and big data platforms

Data sources from different data providers/tenants

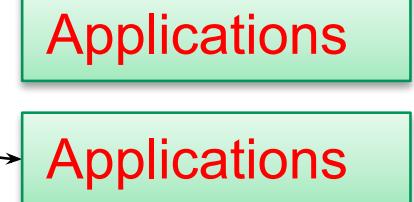


The core part of Big Data Platforms

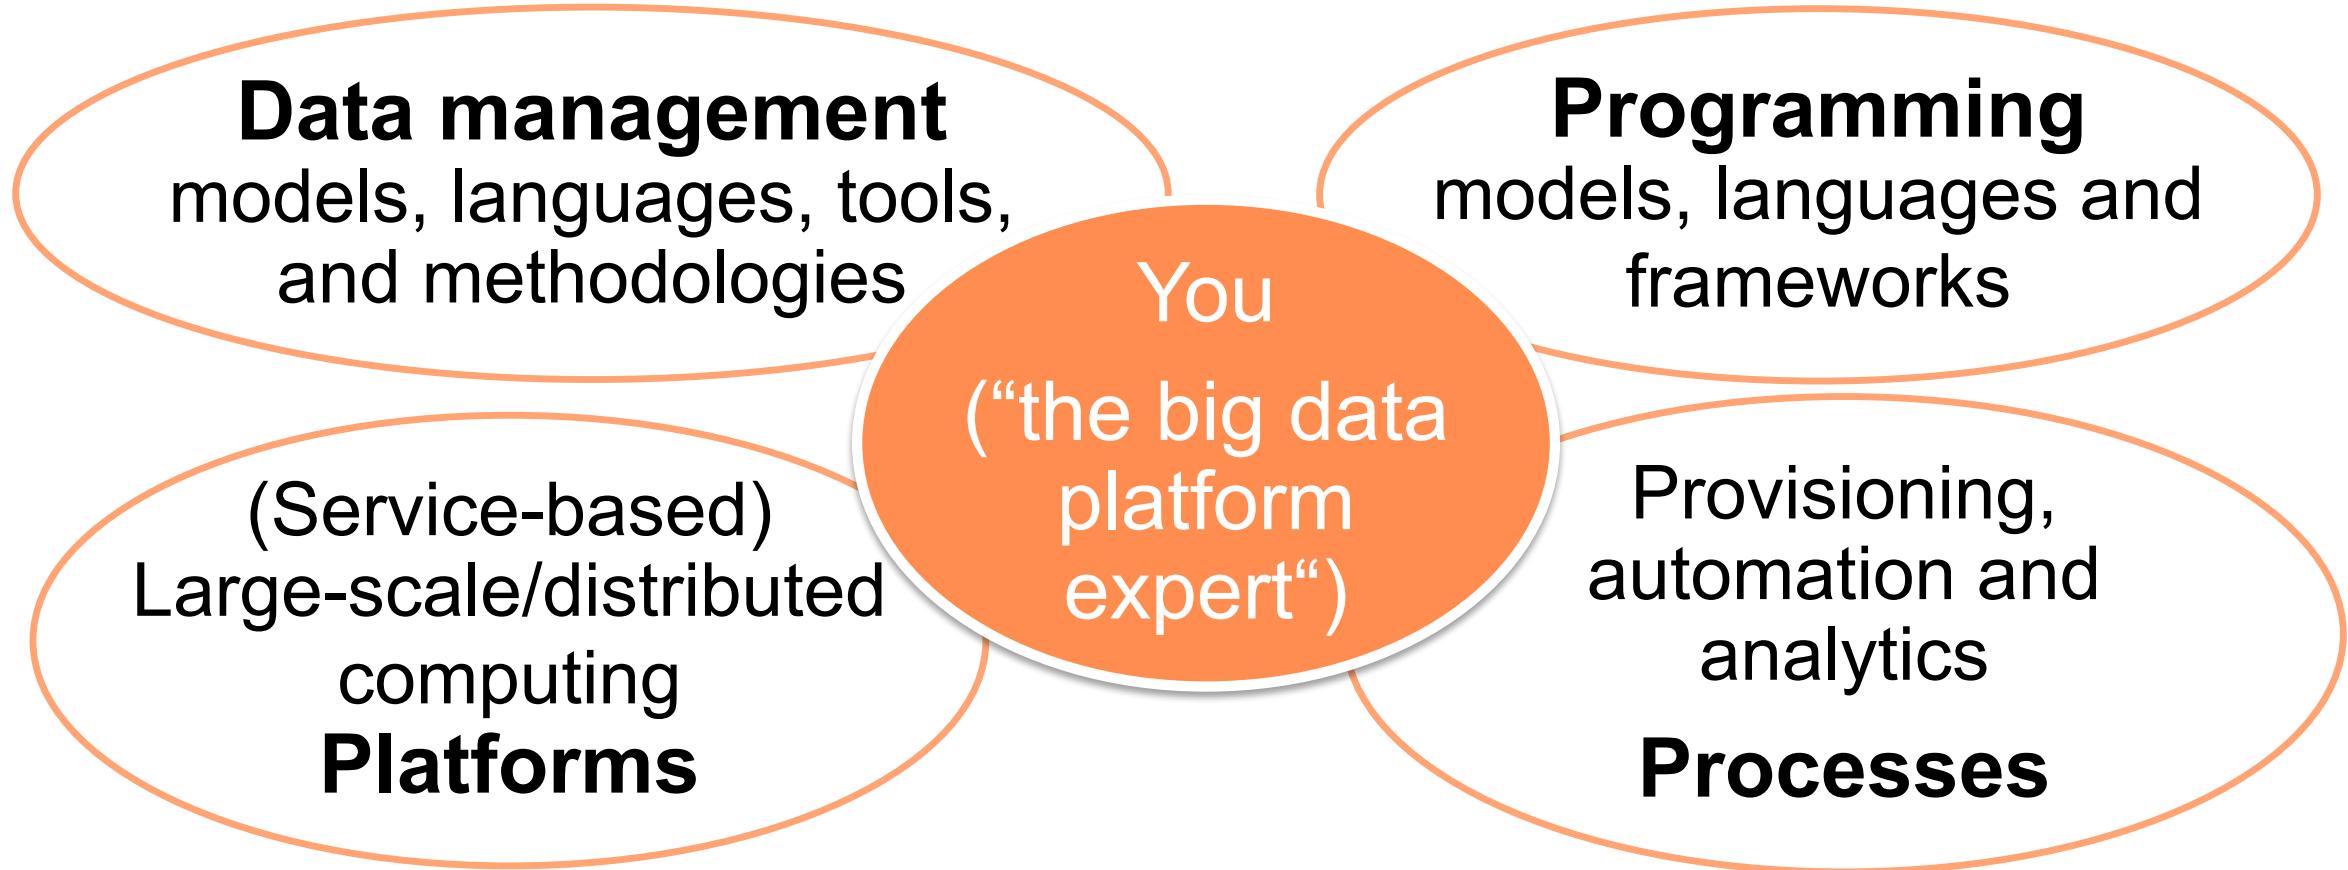


Algorithms & Pipelines

Consuming applications



Core principles/techniques for Big Data Platforms



Focuses in studying big data platforms

- Design/Development vs Operation
- Data-centric vs Service-centric vs Platform-centric activities
- High-level SQL-style analytics vs programmatic data processing models and workflows/processes
- Quality and governance

Stakeholders and roles

- As a user: to use and program atop big data platforms
- As a provider: to deploy and operate big data platforms
- As a designer/architect: to design new (solutions for) big data platforms
- As a developer: to develop services/applications in big data platforms

Business models vs data platform engineering

Business Models



**Engineering
Solutions**

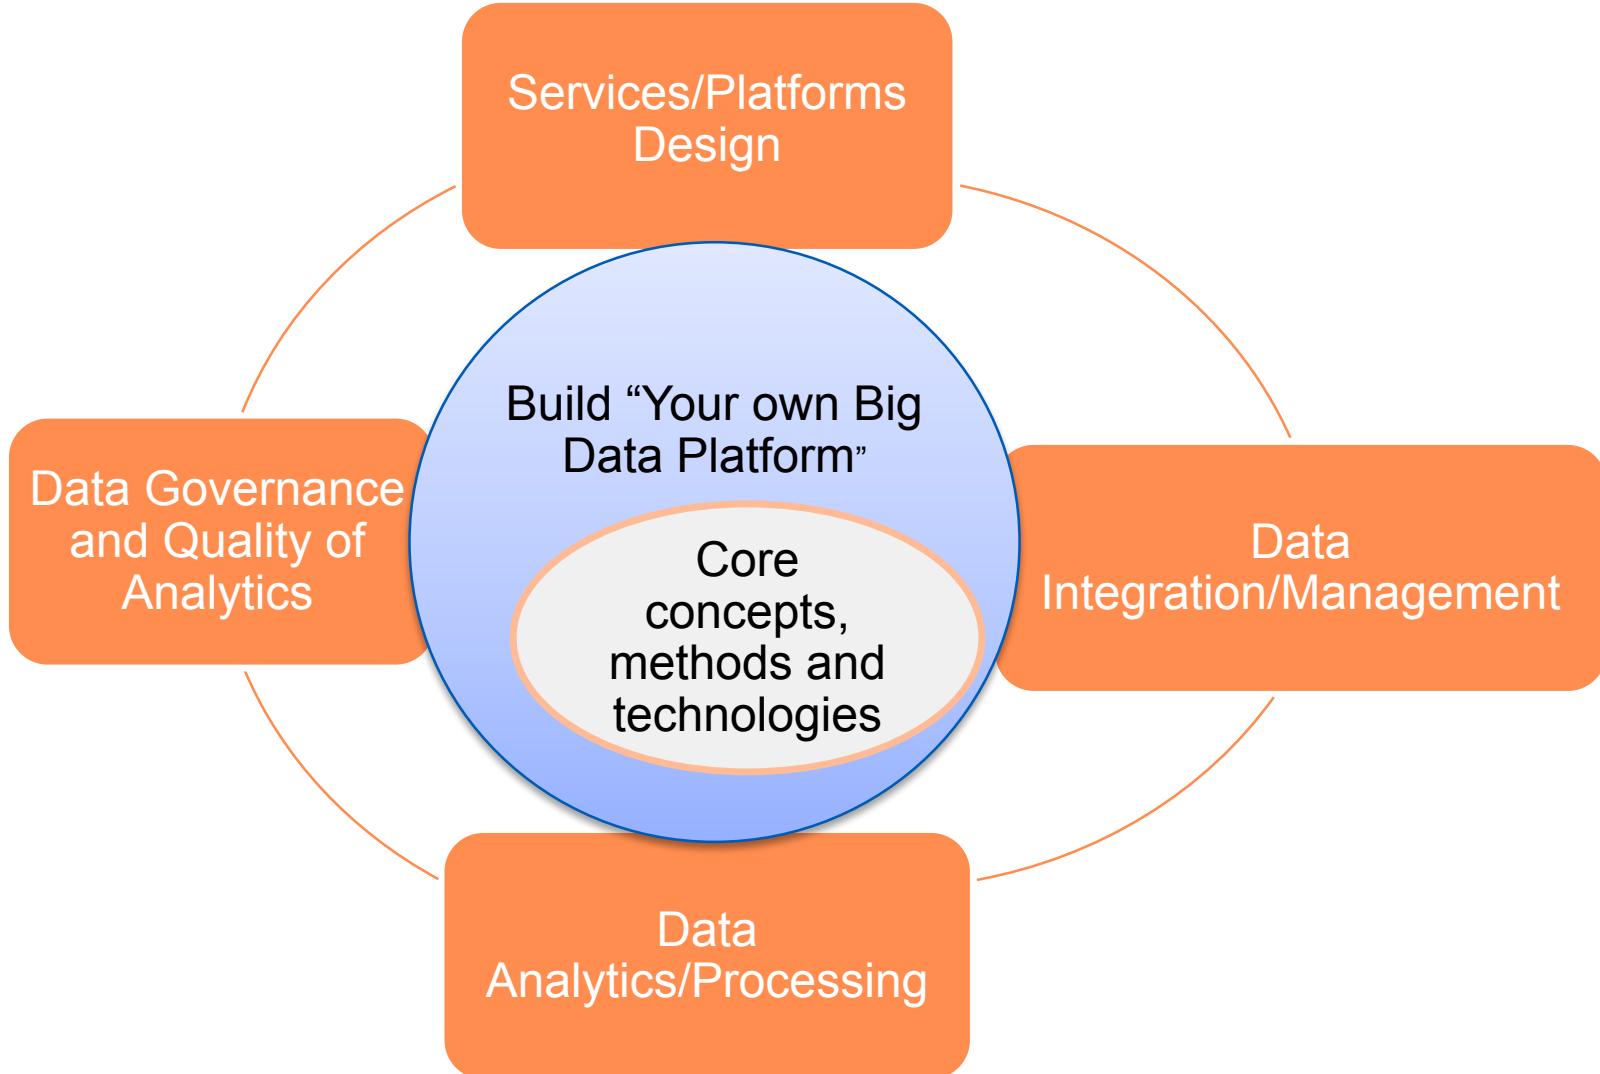
- Distinguish between engineering and business models
 - Engineering: design, operation and governance
 - Business models: stakeholder management, pay-per-use, pricing, tenant models
- **We mainly focus on engineering**
 - Business aspects are reflected in requirements for the engineering

Concepts/techniques vs Technologies

- Concepts/Techniques versus Technologies
 - concepts: e.g., NoSQL data models, sharding techniques, coordination for scalable datastores
 - technologies: e.g., Apache Cassandra, Apache Spark, Airflow, Flink
- **Focus mainly on concepts/techniques**
 - technologies can be very complex or “everything is behind an API”
 - often a well-known technology is packaged and sold under different product names by different vendors
 - but don’t forget key concepts and techniques
 - implement key concepts with state-of-the-art technology in limited but realistic scenarios

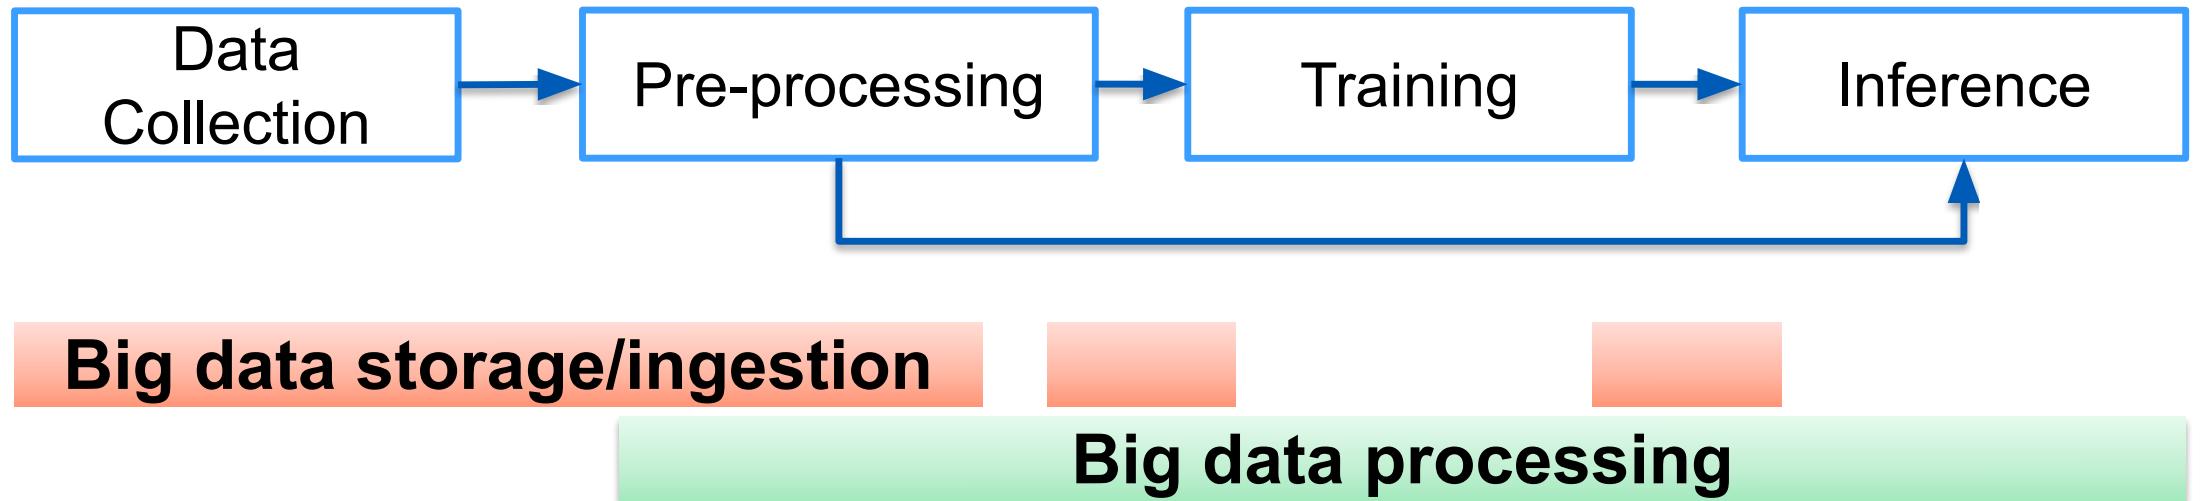
Build your story

Focus on foundations & explore your strengths/interests

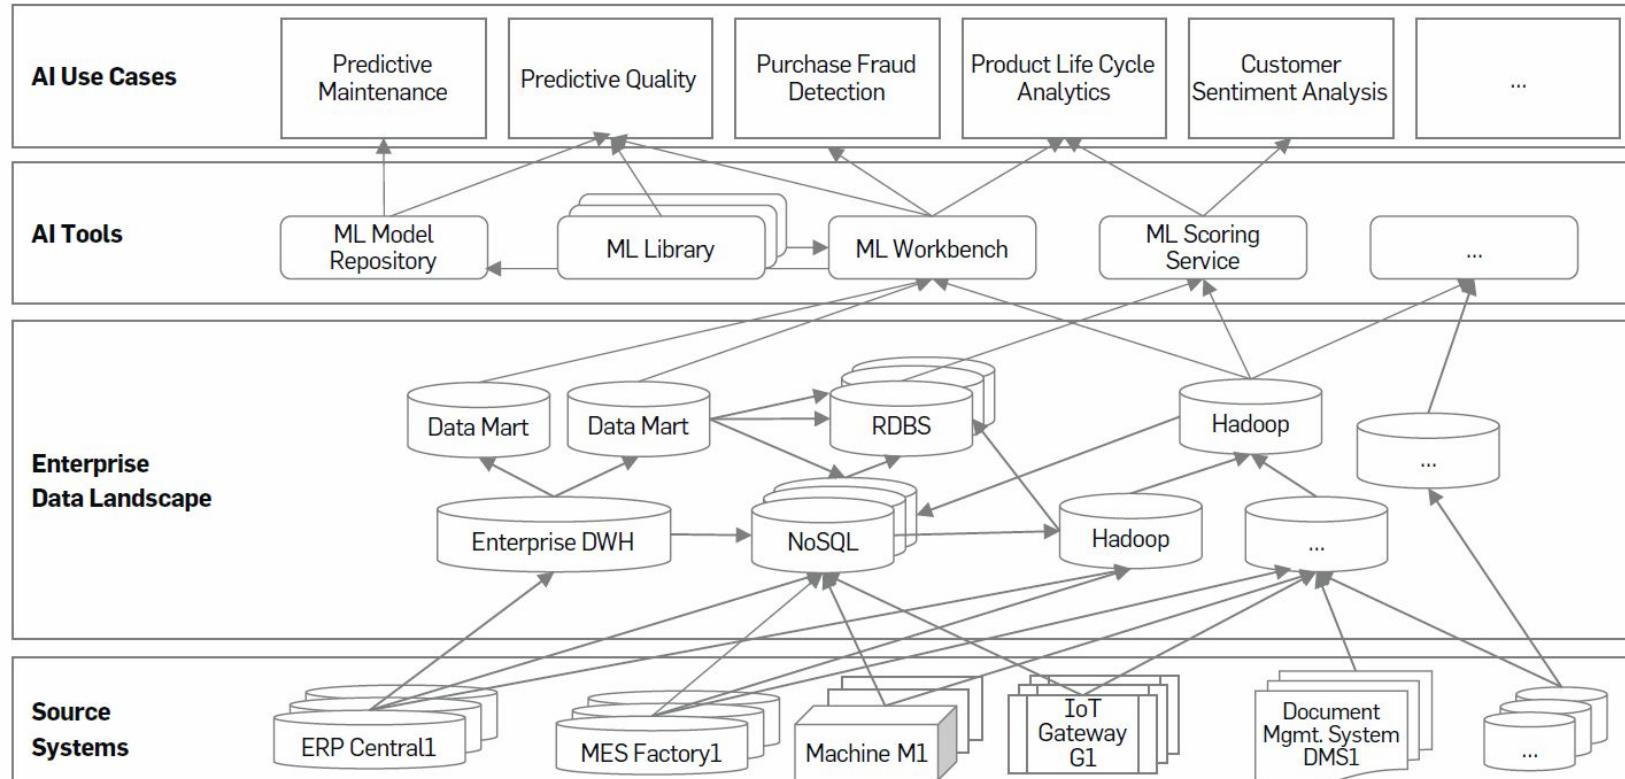


Design and implementation mainly reflect activities of platform developers/providers in the lifecycle of a big data platform based on real-world data sets and scenarios

Example of a typical big data machine learning pipeline



“No AI Without Data”



DWH: Data Warehouse, ERP: Enterprise Ressource Planning, MES: Manufacturing Execution System, ML: Machine Learning, RDBS: Relational Database System

Figure source: There Is No AI Without Data, by Christoph Gröger, Communications of the ACM, November 2021, Vol. 64 No. 11, Pages 98-108, 0.1145/3448247

Summary: the importance of big data platforms

- Foundations and backbones for various research/development areas
- AI/ML/LLMs
 - Data storage/services (feature/training data, vectorized data for RAG), data management, data preparation pipelines, ...
- Data Science
 - Data and data management, computational models, statistics and algorithms
- Enterprise computing
 - 360-degree analytics of customers
- Industrial IoT/Manufacturing/Predictive maintenance
 - Monitor machines and optimizing machines through real-time and predictions
- Smart cities and sustainability

Thanks!

Hong-Linh Truong
Department of Computer Science

rdsea.github.io



—
Kiitos
aalto.fi