

---

# Architecting Big Data Platforms

Hong-Linh Truong  
Department of Computer Science  
linh.truong@aalto.fi



07.1.2026



Content is available under  
CC BY-SA 4.0 unless otherwise stated

# Learning objectives

- Understand key issues in designing a big data platform
- Learn key architecture design issues
- Be familiar with the course' selected big data platform technologies

# Your big data platform story - an example of a scenario

“Your team has to build a big data platform for **X types of data**. Data will be generated/collected from **N data sources**. We expect to have **10+ GBs/day** of data to be ingested into our platform. We will have to serve **K thousands of requests** for different types of analytics – to be determined. Our response time for an analytics request should be **in t milliseconds**. Our services should not be ...”

# You may have several questions?

- Do we have to support multiple types of data?
- How do data pipelines and data load look like?
- How to enable different data processing models?
- Which runtime parameters must be monitored? Which service level and data metrics must be guaranteed?
- Which are the main building blocks/sub systems?
- To where we should distribute/deploy our components?
- Which part of the platform we must manage by ourselves and which part will be fully managed by other providers?
- How to design elastic big data infrastructures?
- Etc.

# **Starts with Big Data Platform architectures!**


**To architect the platform  
centered around big data and  
data intensive activities!**

# **Underlying computing infrastructures is for data intensive tasks**

- Computing resources and services
  - many machines, virtual infrastructures, different types of services
- Distributed infrastructures from different administrative domains
  - in multiple data centers, locations and countries
  - with different security and network policies
- Diverse data contracts, service level objectives (SLO) and service level agreements (SLAs)
  - performance, service failure, cost, privacy/security ...
  - data quality and data governance

# Understanding the underlying infrastructures for big data platforms


large-scale distributed infrastructures ⇒ hybrid and/or multi cloud/edge



Storage (NVMe/SSD, HDD), Interconnects (CXL, RoCE v2, InfiniBand) and Compute resources (CPU/GPU)

**A!**

# Emerging edge-cloud continuum affecting data platform designs: data storage and processing



A!


# Data-centric development & operations

- Data ingestion and ETL (Extract, Transform, Load)
  - from various data sources we move data into the platform
- Data storing and management
  - ingested data stored and managed using different types of storages and databases
- Data analysis and (Machine) learning
  - data processed, analyzed and learned, finding insights and creating ML models
  - data **at rest** vs data **in motion**
- Reporting and visualization
  - patterns/insights in data will be interpreted and presented for making decisions, reporting, and creating stories

**Data Governance at very large-scale**

# Big Data Pipelines

Multiple big data pipelines can be constructed atop a big data platform (and across distributed infrastructures)



# Handling multiple types of data


- Hardly to avoid the support for multiple types of data
  - tenant requirements and application needs, given data as “products”/“assets”
- Multiple types of data
  - different characteristics and values
- Any elastic solution that ensures minimum changes to support **generalization and extensibility**
  - multi-model databases, microservices or multiple of databases or data lake
  - new workloads, new types of customers

# Data governance: concerns & SLAs

- Key common problems:
  - data quality and data lineage
- Ingesting data: ingestion of data under V\*
  - mapping and transforming data, data validation/quality control
- Storing data
  - data sharding and consistency, data backup, retention, etc.
  - the impact of **the rights to remove data**
- SLA multitenancy versus single tenancy
  - security, privacy, performance, reliability and maintenance?

The volume of data is **increasing** but its **usefulness may not** because of the bad data quality


# Data concerns: example




- Risk and ethical consequences?
- Regulation-compliant platforms: e.g., GDPR, AI Act

# Fast/slow reliable processing

big data but not near real-time, e.g., *take customer transaction files from companies and move to data centers for analytics*



fast, small IoT data in near real-time flows, e.g. *position of cars*



# Software services design goals


For dealing with V\*

- Responsive: guarantee quality of services
- Resilient: deal within failures
- Elastic: deal with different workload and quality of analytics
- Loosely coupling: support reusability, composition, and extensibility

# Designs must address various aspects for big data

- **Responsive:**
  - distributed computing, multi layer optimization
- **Resilient:**
  - replication, containment, isolation
- **Elastic:**
  - sharding, replication, load balancing, scale up/out
- **Message-driven:**
  - loosely coupling with messages, non-blocking protocols, location-independent

## Reactive systems



Source: <https://www.reactivemanifesto.org/>

# Efficiency and sustainability

- Highly efficient might not be good for sustainability or resilience
- Sustainability within big data platforms
  - energy consumption, reusability, extensibility


- Reliability
- Resilience
- Elasticity

How to have them for



■ *Sustainability*

## Design and implementation



**Distributed systems of components are used to manage,  
ingest data and process data**

# Partitioning and composability: splitting functionality & data

- Breakdown the complexity and responsibility
  - easy to implement, replace and compose
  - deal with performance, scalability, security, data quality, etc.
  - support teams in DevOps and data products
  - data responsibility and ownership
  - cope with technology changes
- Many things are related to the current trends
  - microservices, domain-driven designs, data mesh, data platform modernization (storage, query, and analysis/AI couplings)
- But tradeoffs with management w.r.t data resources

# Example of functional and data partitioning and composability (1)

**Service-oriented components**

**Microservices and domain-oriented microservices**

**Serverless functions/function as-a service**




Figure source: <http://queue.acm.org/detail.cfm?id=1971597>

# Example of functional and data partitioning and composableity (2)

**Multi (hardware) storage capabilities**

**Data sharding**

**Multi data spaces**

**Multi data services**

**Multiple data infrastructures**

**Data products**







Figure source: <http://queue.acm.org/detail.cfm?id=1971597>

# Interaction: multiple models

- Protocols
  - REST, gRPC, Message Passing, Stream-oriented Communication
- Models
  - one-to-many, many-to-one, many-to-many
  - synchronous/asynchronous calls
  - internal data exchange versus open/external exchange




# Coordination: Orchestration and Reactiveness/Choreography



# Distribution: Edge or Cloud Data Centers?

- Data & components can be distributed in different places!
  - performance, security, regulation, energy efficiency
- Global deployment or not?
- Move analytics/work or move data?



An outage can lead to a huge problem. Example:

<https://www.thousandeyes.com/blog/aws-outage-analysis-dec-7-2021>

# Scalability & elasticity

Presto: <https://prestodb.io/>




Figure source: *Presto: SQL on Everything*

<https://ieeexplore.ieee.org/stamp/stamp.jsp?tp=&arnumber=8731547&tag=1>

# Lyft Presto Gateway

*"As of today we have 60 PB of query-able event data stored in an S3 based data lake and about 10 PB of raw data is being scanned every day using Presto"*





Figure and text source:

<https://eng.lyft.com/presto-infrastructure-at-lyft-b10adb9db01>

# API for Data Platform as a Service

- APIs are very important! Why?
  - enable consumers access to data and analytics functions from your big data platforms without worrying about changes within your platforms
  - enable virtualization and management (hide internal, control access, throttling)
  - establish protocols for data exchanges and governances



A!

# Delivery data as “asset”/”product”

- FAIR principles
  - Findable, Accessible, Interoperable and Reusable
- Product/asset characteristics


⇒ The platform must support

- Metadata and governance
- Service-level objectives and quality controls for data
- Search
- Access control policies

[FAIR principles]: <https://www.nature.com/articles/sdata201618>

# **Common, high-level architecture view with popular state-of-the art technologies for our study**

# Our big data at large-scale: the big picture in this course



Thanks!

Hong-Linh Truong  
Department of Computer Science

[rdsea.github.io](https://rdsea.github.io)



—  
Kiitos  
aalto.fi