



Aalto University
School of Science

Big Data Storage and Database Services – common systems & integration problems

Hong-Linh Truong

Department of Computer Science

linh.truong@aalto.fi, *<https://rdsea.github.io>*

Common data models and data storage/database systems

Common data models

- **File**
- **Relational data model**
- **Key-Value data model**
- **Document-oriented model**
- **Column family model**
- **Graph model**
- **Vectorization model**

Special file/table formats in big data

- **For analytical big data, remember the patterns about data**
 - write once, read many
 - analytics queries often access data based on “columns” (e.g., sum of all “trip payments”)
- **File/table formats**
 - compression, columnar representation for column-based queries/accesses, encryption
 - suitable with big data analytics (e.g., Spark, Hadoop)
- **Examples**
 - Apache ORC (<https://orc.apache.org/>)
 - Apache Parquet (<https://parquet.apache.org/>)
 - Apache Iceberg (<https://iceberg.apache.org/spec/>)
- **They are the file/table formats under many big data systems**

Blob data

Big files:


- Pictures, documents, big log files, images, video, backup data

Storage

- File systems or blob storage

Implementations

- File systems: NFS, GPFS, Lustre (<http://lustre.org/>), Hadoop File systems
- Storage: Amazon S3, Azure Blob storage, OpenStack Swift, Minio
- Simple API for direct access



Example - Amazon S3

Store blob files and their metadata

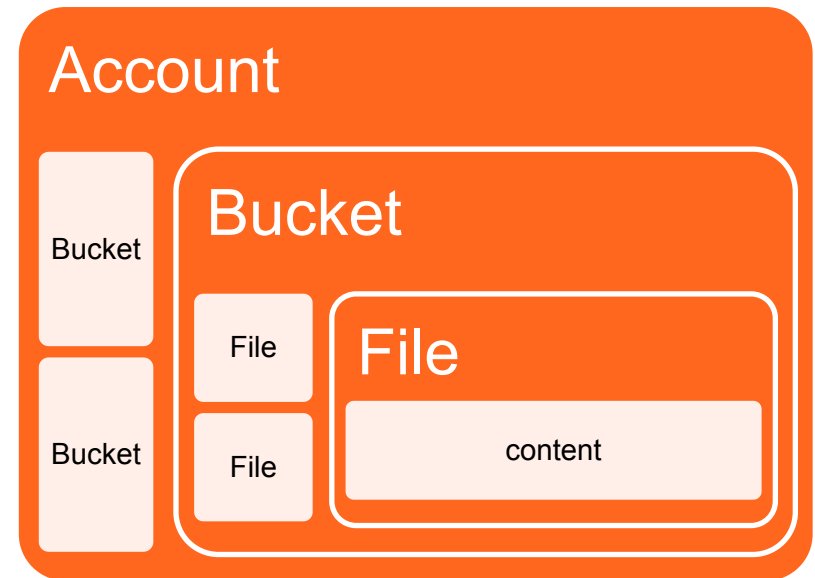
- Max 5TB per file
- A file is identified by a key

Structure

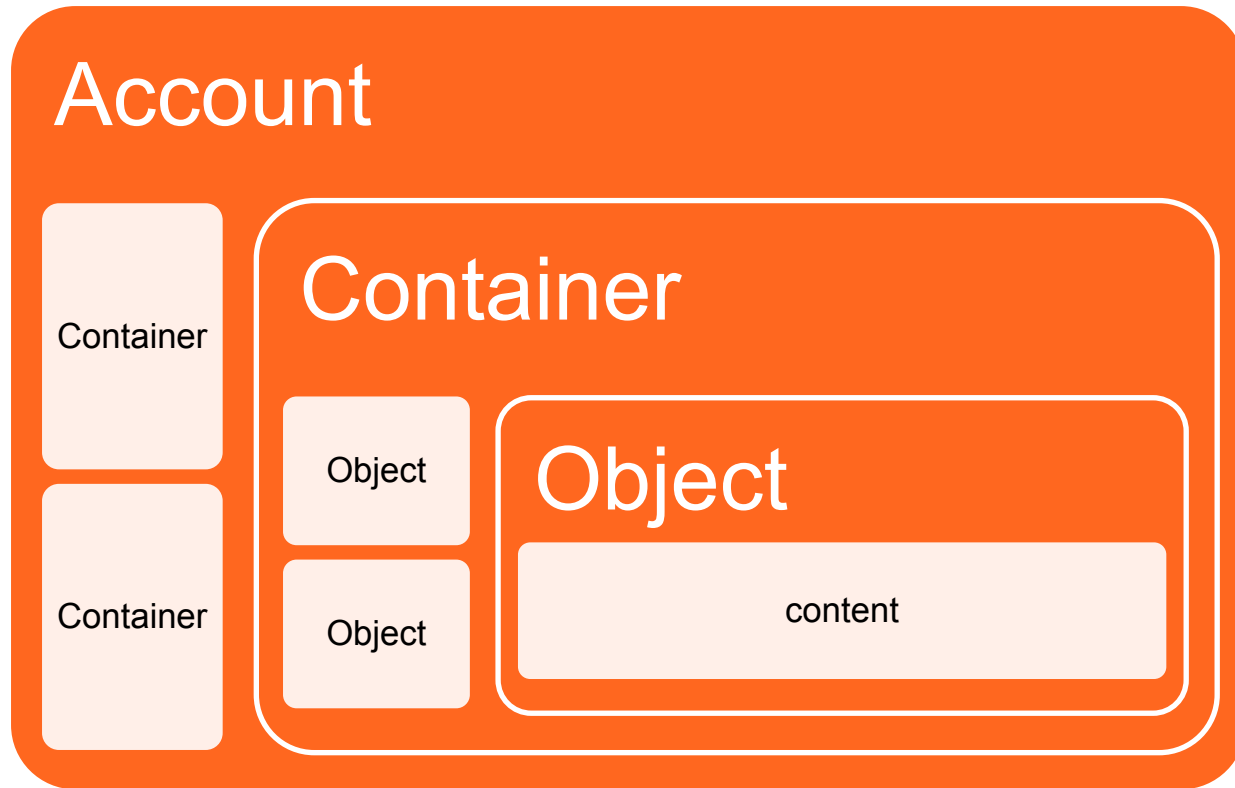
- File = Object
- Object: name and metadata
- Objects are organized into Buckets

Simple APIs

- REST



OpenStack Swift



<http://docs.openstack.org/developer/swift/>

Minio

- Check <https://min.io/>
- For different deployment models: Kubernetes, VMs, edge-cloud
- S3 compatibility

Relational Model

- **Well-known, long history**
- **Tables with rows and columns**
 - Strict schema requirements
- **Powerful querying & strong consistency support**
 - E.g.: Oracle Database, MySQL Server, PostgreSQL, MariaDB

Example: Alarm in BigQuery

stationdescription

Schema Details Preview

Field name	Type	Mode
station_id	INTEGER	NULLABLE
code	STRING	NULLABLE
name	STRING	NULLABLE
address	STRING	NULLABLE
description	STRING	NULLABLE
latitude	STRING	NULLABLE
longitude	STRING	NULLABLE

stationparameters

Schema Details Preview

Field name	Type	Mode
reading_time	TIMESTAMP	NULLABLE
value	FLOAT	NULLABLE

station_id

parameter_id

stationalarms

Schema Details Preview

Field name	Type	Mode
station_id	INTEGER	NULLABLE
alarm_id	INTEGER	NULLABLE
parameter_id	INTEGER	NULLABLE
start_time	TIMESTAMP	NULLABLE
end_time	TIMESTAMP	NULLABLE
value	FLOAT	NULLABLE
threshold	INTEGER	NULLABLE



Aalto University
School of Science

CS-E4640 Big Data Platforms, Spring 2025, Hong-Linh Truong
22/01/2025

Big data: relational databases

- **Tables with rows and columns**

- Strict schema requirements, powerful querying & strong consistency support
- E.g.: Oracle Database, MySQL Server, PostgreSQL, CockroachDB

- **Relational database at very large-scale**

- Amazon Aurora, Microsoft Azure SQL Data Warehouse

- **ACID (atomicity, consistency, isolation, durability) is hard with big data**

- relational big database must address replication, distribution, and scalability issues

Relational Databases for big data scenarios

- **Relational database at very large-scale**
 - Amazon Aurora, CockroachDB, Microsoft Azure SQL Data Warehouse
- **Examples of Amazon Aurora (reading list)**
 - based on MySQL/InnoDB but change the architecture, separate storage from engine, support cloud scale and replication, etc.

Key-Value Model

- **Tuple = (key, value)**
 - Values can be based on different structures
- **Scalable and performance**
- **Primary use case: caching (pages, sessions, frequently access data, distributed lock)**
 - Simple, very efficient but limited querying capabilities
- **Implementation:**
 - Memcached, Riak, Redis, Valkey, Apache Accumulo

Example: Redis

- **<http://redis.io/>**
 - see also <https://github.com/valkey-io/valkey>
- **In-memory cache service**
 - *Store (key,value) tuples in memory but persistent back to database*
 - Simple APIs
 - Well supported with many programming languages
 - Widely used in big data ecosystems
- **Learning**
 - <https://app.redislabs.com/#/login> provides a free account

Example: Redis

<http://redis.io/topics/benchmarks>

Document-oriented Model


- **Documents**

- flexible schema (schemaless) with flexible content
- data fields can be complex for sub documents
- use collections, each collection is a set of documents

- **Primary use cases**

- large amounts of semi-structured data
- collection of data with different structures

Document-oriented model – simple analogy



Examples: MongoDB.Atlas

<https://www.mongodb.com/atlas/database>

Graph-oriented model

- **Data is represented as a graph**
 - nodes or vertices represent objects
 - an edge describes a relationship between nodes
 - properties associated with nodes and edge provide other information
- **Use cases**
 - when searching data is mainly based on relations (social networks, asset relationship, knowledge graph)

Graph-oriented model

- **Data is represented as a graph**

- nodes or vertices represent objects, an edge describing a relationship between nodes
- properties associated with nodes and edge provide other information

- **Use cases**

- when searching data is mainly based on relations (social networks, asset relationship, knowledge graph)

Working with graph databases

- **Graph databases**

- Azure CosmosDB, ArangoDB, Titan, TypeDB, Neo4J, OrientDB

- **Query languages:**

- Gremlin, SPARQL, Cypher

- **Graph computing frameworks (analysis)**

- Apache TinkerPop, Apache Spark GraphX

Example

<https://github.com/vaticle/typedb>

Column-family data model

Many situations we aggregate and scan few columns of million rows of data \Rightarrow store big data in columns enable fast scan/retrieval/aggregation

Column Family = (Column, Column, ...): for similar type of data

Column Key = Family: qualifier

Data = (Key, Value) where Key =(Row Key, Column Key, Timestamp)




Figure source: Fay Chang, Jeffrey Dean, Sanjay Ghemawat, Wilson C. Hsieh, Deborah A. Wallach, Mike Burrows, Tushar Chandra, Andrew Fikes, and Robert E. Gruber. 2006. Bigtable: a distributed storage system for structured data. In Proceedings of the 7th symposium on Operating systems design and implementation (OSDI '06). USENIX Association, Berkeley, CA, USA, 205-218.

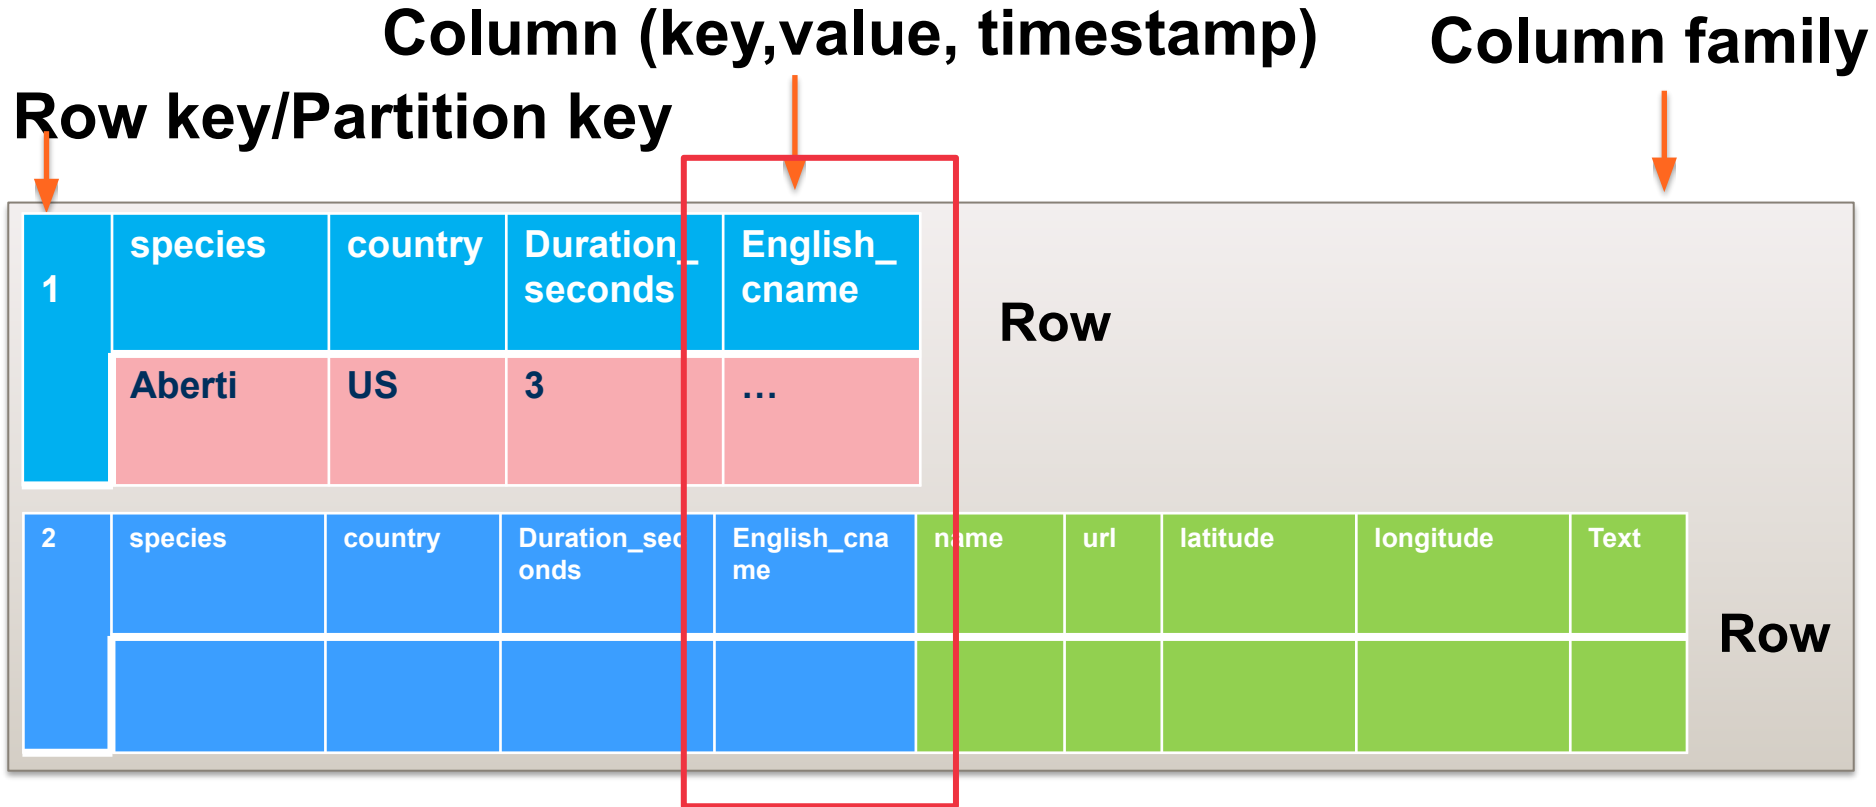
Examples: Cassandra, HBase

Column-family data model

- **Data Model**

- Table consists of rows
- Row consists of a key and one or more columns
- Columns (column name, value, timestamp)
- Columns are grouped into column families
- Columns can be different in rows
 - flexible, wide columns → save spaces

Example of a data model in Cassandra



Examples

Examples of rows

Column (name, value, timestamp)

english_cname	writetime(english_cname)
Black-tailed Gnatcatcher	1569966171073228

(1 rows)

```
cassandra@cqlsh> select * from tutorial12345.bird2;
```

@ Row 1

species	melanura
country	Mexico
duration_seconds	29
english_cname	Black-tailed Gnatcatcher
file_id	71907
latitude	32.156
longitude	-115.793

@ Row 2

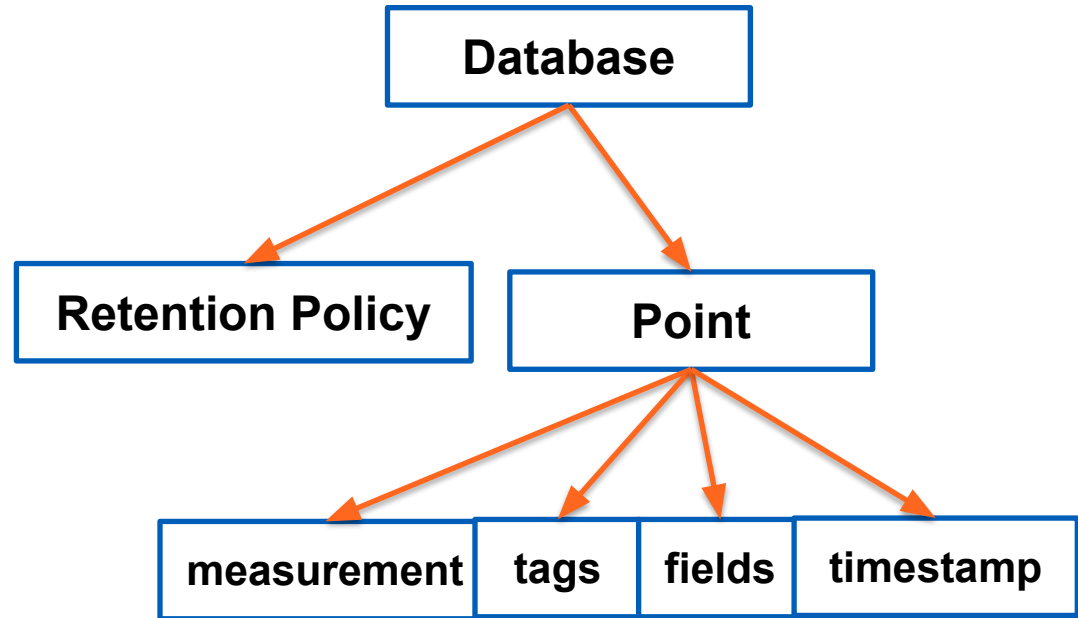
species	melanura
country	United States
duration_seconds	29
english_cname	Black-tailed Gnatcatcher
file_id	358907
latitude	33.7329
longitude	-115.8023

Time Series Database

- **So many types of data in big data are time series**
 - IoT measurements, session data, log, etc.
- **Of course you can also use other databases**
 - e.g., Cassandra, ElasticSearch, BigTable
- **Time Series Databases specially designed for time series data**
 - *examples: Riak TS (Time Series), InfluxDB, Apache Druid*

Example: InfluxDB

- <https://www.influxdata.com/>
- High-level query, SQL-alike Language
- Retention policy for data storage, sharding and replication



An example of InfluxDB

```
> show measurements
```

```
name: measurements
```

```
name
```

```
----
```

```
stationalarm
```

```
stationparameter
```

```
> select * from stationalarm;
```

```
name: stationalarm
```

time	alarm_id	datapoint_id	station_id	value	valueThreshold
----	-----	-----	-----	----	-----
1487444343000000	308	121	1161115016	240	240

In-memory databases

- **Databases use machine memory for storage**
 - Persist data on disks
 - Require very powerful machines
- **In principle it is not just about data models but also data management, data processing, software and hardware optimization, e.g.,**
 - SAP HANA, VoltDB: in memory relational databases
- **Why are in-memory databases important?**



Aalto University
School of Science

Interfaces between a data storage/databases system and its external analysis systems



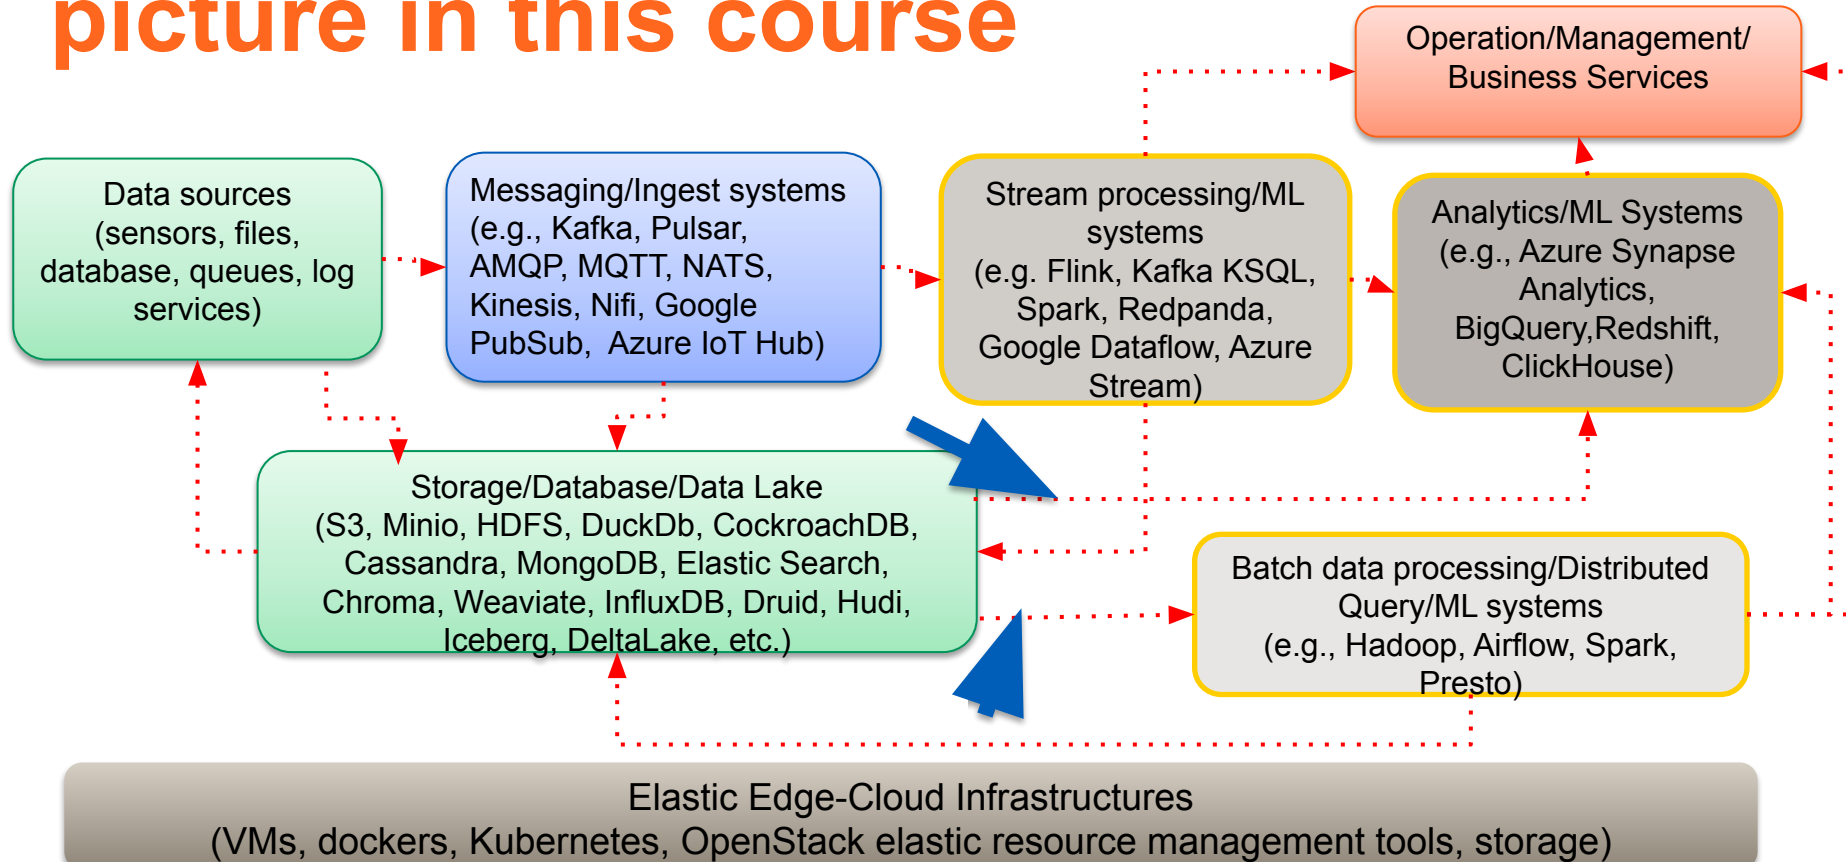
Aalto University
School of Science

In many cases: the data in data storage/database service must be made available for large-scale analysis:

large-scale analytics and data are managed by different systems

an important consideration in big data platforms design!

Big data at large-scale: the big picture in this course



Making data available to the analytics

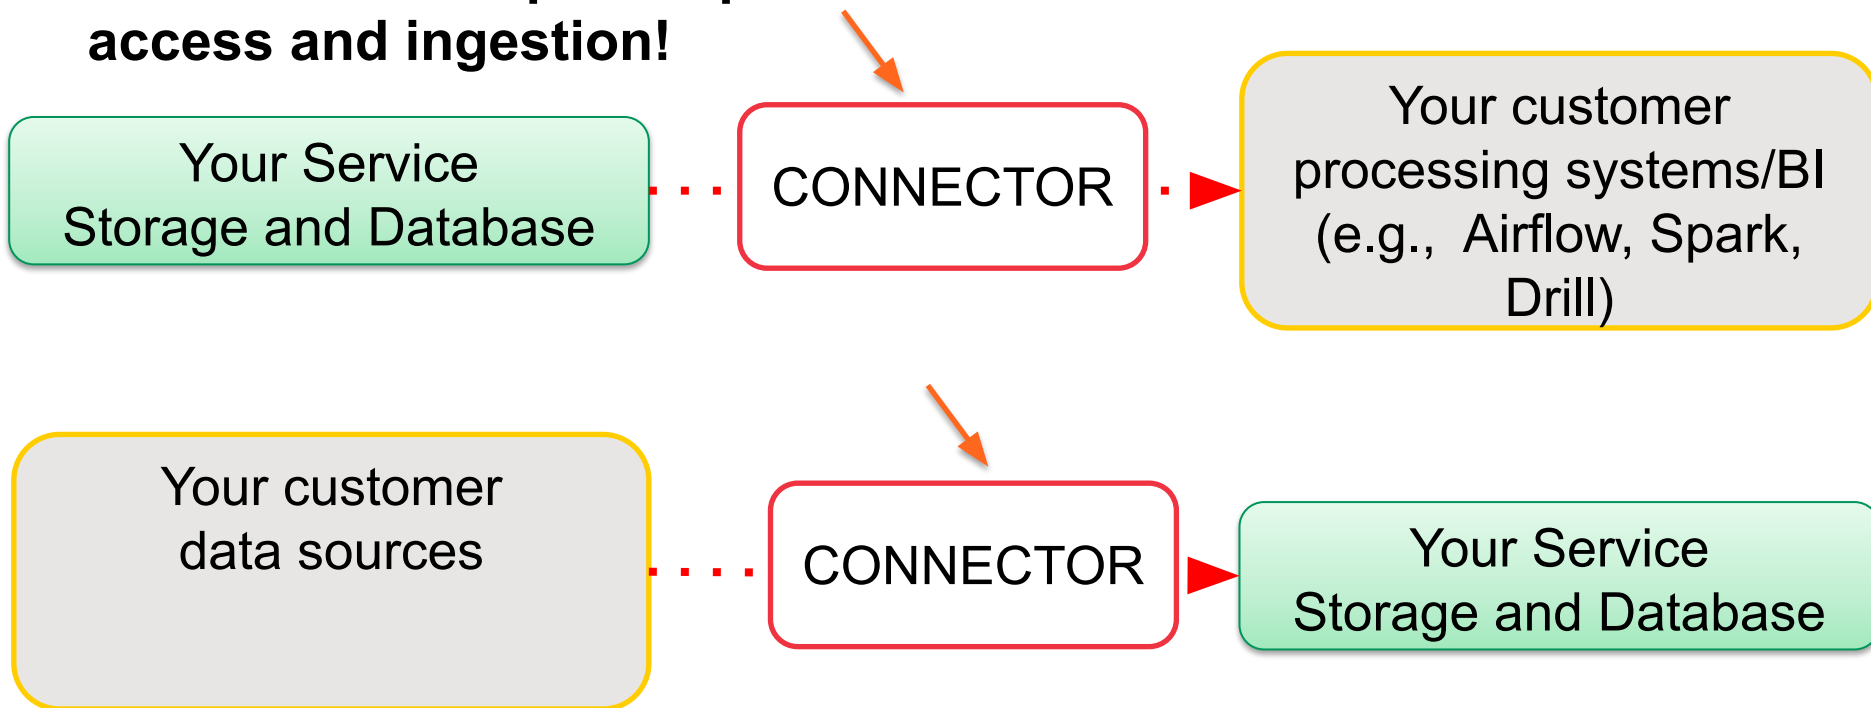
- **Data layer must map/provide data to processing layer**
 - maximize the analytics possibilities
- **Key issues**
 - avoid data movement as much as possible
 - avoid contention between the data management and the data analytics system
- **Techniques**
 - “mount”, specific connectors/drivers, copy-process-remove activities

Mount/"Fuse"

- **Mapping a remote storage as a local file system**
 - Blobfuse (Microsoft Azure), gcsfuse (Google Storage)
 - the network performance is important

Connectors


ODBC or other specific protocol connectors to enable data access and ingestion!



Example: Superset connectors

<https://superset.apache.org/>

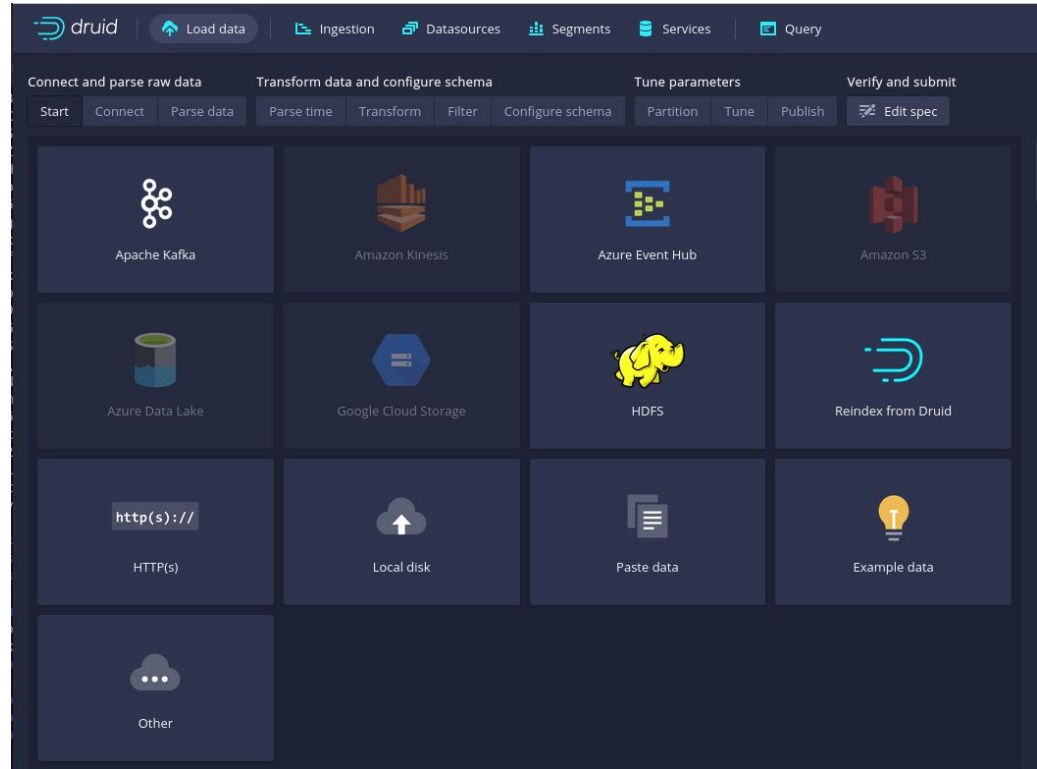
Connectors to different types of databases/datasets to retrieve and analyze data



Example: Druid

<https://druid.apache.org/>

Different types of connectors (e.g., Kafka, Files, S3, etc.) to allow data ingestion into the database



Analytics and Cloud Storage

- **Various connectors for making data in cloud storages available for analytics**
- **Apache Hadoop/Spark (data analysis) can work with Amazon S3, OpenStack Swift, Google Cloud Storage**
- **Examples:**
 - <https://github.com/GoogleCloudDataproc/hadoop-connectors>
 - <https://spark.apache.org/docs/latest/cloud-integration.html>

“Copy and Process”

Client libraries are used to move data from storages and databases to processing places

Examples:

```
import pandas as pd
from sklearn.tree import DecisionTreeClassifier
from cassandra.cluster import Cluster

cluster = Cluster(contact_points=hosts, port=9042, auth_provider=auth_provider)
session = cluster.connect("tutorial12345")
sql_query = "SELECT * FROM tutorial12345.bird1234;"
df = pd.DataFrame()
rows = session.execute(sql_query)
df = rows._current_rows
print(df)
```


Distributed SQL queries: Presto + other as an example

- **Presto (used by Facebook and many others)**
 - distributed query engine
 - decoupled from storage
 - integration with different databases
 - very large-scale with many nodes
- **Analytics: interactive analytics, seconds – minutes**
 - SQL style

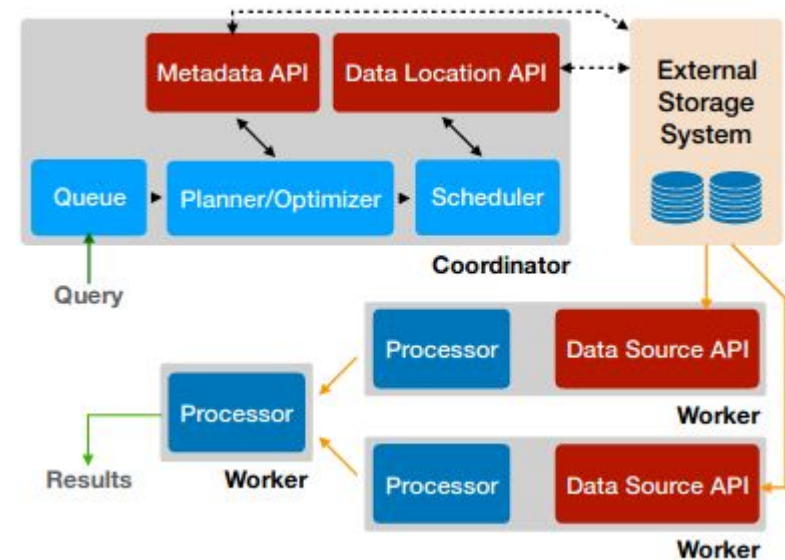


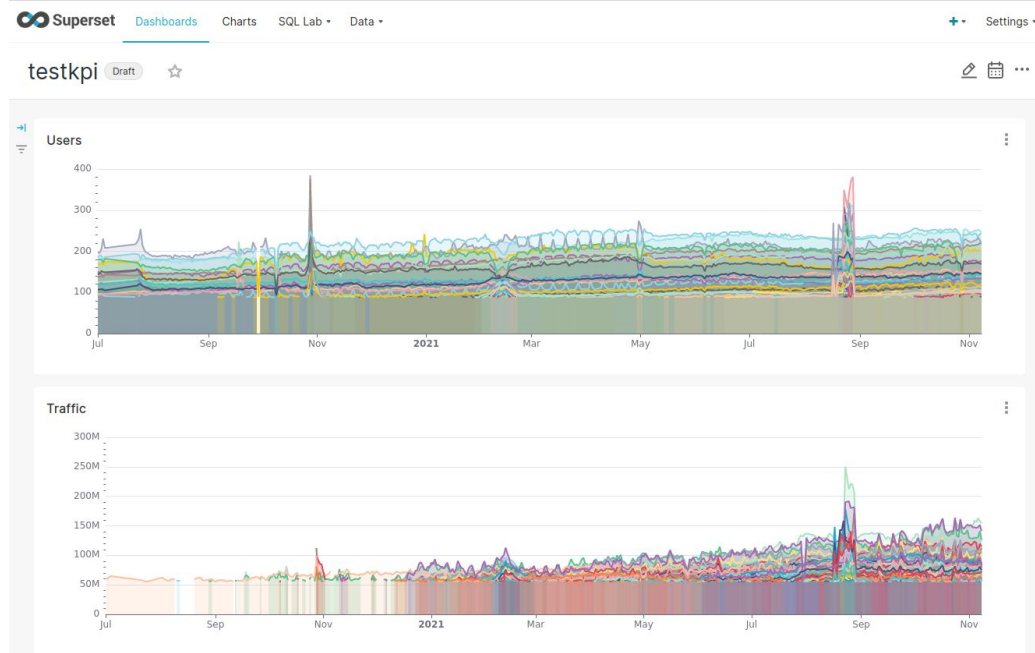
Figure source: *Presto: SQL on Everything*
<https://ieeexplore.ieee.org/stamp/stamp.jsp?tp=&arnumber=8731547&tag=1>

Examples

Analytics: write SQL

```
presto> explain (format text, type logical) SELECT sum(trip_total) FROM chicago_taxi_trips.taxi_trips;

- Output[_col0] => [sum:double]
  _col0 := sum (1:44)
    - Aggregate(FINAL) => [sum:double]
      sum := "presto.default.sum"((sum_4)) (1:44)
        - LocalExchange[SINGLE] () => [sum_4:double]
          - RemoteStreamingExchange[GATHER] => [sum_4:double]
            - Aggregate(PARTIAL) => [sum_4:double]
              sum_4 := "presto.default.sum"((trip_total)) (1:44)
                - TableScan[TableHandle {connectorId='bigquery', connectorHandle='BigQueryTableHa
                  Estimates: {rows: ? (?), cpu: ?, memory: 0.00, network: 0.00}
                  LAYOUT: BigQueryTableHandle{projectId=bigquery-public-data, schemaName=ch
                    trip_total := BigQueryColumnHandle{name=trip_total, type=FLOAT, mode=NULL
```



Thanks!

Hong-Linh Truong
Department of Computer Science

rdsea.github.io