

---

## Apache Kafka for Streaming Data Ingestion - the Core

Hong-Linh Truong  
Department of Computer Science  
linh.truong@aalto.fi




23.1.2026



Content is available under  
CC BY-SA 4.0 unless otherwise stated

# Our big data at large-scale: the big picture in this course




# Abstraction of Data Streams

- **Data stream:** a sequence/flow of *data units*
  - continuously, unbounded or bounded data
- **Data units** are defined by applications:
  - a data unit can be data described by a primitive data type or by a complex data type
  - “small” or “big” w.r.t. size
  - events, messages (requests/responses, information), small documents, etc.
- Usually we encapsulate a data unit in a **record/message** of data
  - record in the application view != record in the system view

# Some use cases - the diversity!

- Data producers generate a lot of near real-time events
- Data producers and data consumers have different processing speeds
  - E.g. logging activities (fast producers but slow consumers)



- Diverse types of data to be produced and consumed
- Dealing with cases when consumers might be on and off (fault tolerance support)
- Asynchronous producing and consuming data

A!


# Examples of log-based/event-based messaging systems

- Apache Kafka
  - <https://kafka.apache.org/>
- Apache Pulsar
  - <https://pulsar.apache.org/>
- RedPanda
  - <https://redpanda.com/>

# Apache Kafka

- <https://kafka.apache.org/>
  - originally from LinkedIn, not a protocol!
- Some components are commercialized by Confluent (bought by IBM)
  - <https://www.confluent.io/>
- Widely used for big data use cases, including message processing in large-scale enterprise service platforms
  - **data messages** (e.g., logs, records, historical events)
    - in the focus of big data platforms
  - request/command messages (e.g., payment/database update)
  - event messages (e.g., notification of a payment due)

# Kafka messaging design



- Use a cluster of brokers to deliver messages
  - e.g., within single data center or on-premise
- Durable messages, ordered delivery via partitions
- Online/offline consumers
- Using underlying file systems **heavily** for message storage and caching


A!

# More than a message broker

- In **Apache Kafka**: the basic data element is **<Key,Value>** tuple
  - also with timestamp and metadata
  - called “Kafka record”
- Data streaming features
  - for near real-time transferring data
- Stream processing
  - streaming applications handle data from streams
  - read and write data back to Kafka message brokers
  - other important frameworks in the ecosystem:
    - Apache Flink and Apache Spark
- High-level SQL-style: KSQL/ksqlDB
  - other possibilities: SQL-like + Java in Apache Flink

In the context of big data: we examine  
Apache Kafka for transferring, ingesting  
and processing **messages of (event) data**

# Kafka design



- A topic consists of different partitions
- Partitions
  - enable parallel processing  
⇒ performance
  - fault-tolerance via replication
- Durable messages, ordered delivery via partitions
- Messages with the same partition key will go to the same partition

# Messages, topics and partitions

- Ordered, immutable sequence of messages
- Messages are kept in a period (regardless of the consumed state)
- Support **total order** for messages within a partition
- Partitions are distributed among server

Anatomy of a Topic




Figure source:

<https://kafka.apache.org/23/getting-started/introduction/>

# Consumers

- Consumer **pulls the data** by sending requests to the broker
- The consumer **keeps a single pointer** indicating the position in a partition to keep track the offset of the next message being consumed
- Why?
  - ⇒ allow customers to design their speed
  - ⇒ support/optimize batching data
  - ⇒ easy to implement total order over message
  - ⇒ easy to implement reliable message/fault tolerance

# Example of a producer

[https://github.com/rdsea/bigdataplatforms/blob/master/tutorials/basic\\_kafka/code/simple\\_kafka\\_producer.py](https://github.com/rdsea/bigdataplatforms/blob/master/tutorials/basic_kafka/code/simple_kafka_producer.py)


# Example of a consumer

[https://github.com/rdsea/bigdataplatforms/blob/master/tutorials/basic kafka/code/simple\\_kafka\\_consumer.py](https://github.com/rdsea/bigdataplatforms/blob/master/tutorials/basic kafka/code/simple_kafka_consumer.py)

# Message delivery

- Message delivery guarantees are important for different use cases/requirements
- Some delivery models
  - At most once
  - At least once
  - Exactly once



# What does it mean exactly one?



- Producer: idempotent delivery  $\Rightarrow$  no duplicate entry in the log
- Transaction-like semantics:
  - within a transaction, EITHER messages to ALL partition topics OR not at all
- Consumer behavior management

# Scalability and fault tolerance

- Topic can be replicated
- Partitions are distributed and replicated among broker servers
- Consumers are organized into groups
- Each message is delivered to **only one consumer instance** in a group
- One partition is assigned to one consumer




Figures source: <https://kafka.apache.org/23/getting-started/introduction/>

# Partitions and partition replication

- Why partitions?
  - support scalability
    - enable arbitrary data types and sizes for a topic
    - enable parallelism in producing and consuming data
  -
- But partitions are replicated, why?
  - for fault tolerance

# Partition replication

Replication model: the leader-follower (primary-secondary) model!  
The leader handles all read and write requests



**Figure source:** <http://de.slideshare.net/junrao/kafka-replication-apachecon2013>

# Consumer group


- Consumer group: a set of consumers
  - used to support scalability and fault tolerance
  - allows multiple consumers to read a topic
- In one group: each partition is consumed by only consumer instance
  - combine „queuing“ model and „publish/subscribe“ model
- Enable different applications receive data from the same topic.
  - different consumers in different groups can retrieve the same data

# Key questions/thoughts


- Why do we need partitions per topic?  
⇒ arbitrary data handling, ordering guarantees, load balancing
- How to deal with high volume of near real-time messages for online and offline consumers?  
⇒ partition, cluster, message storage, batch retrieval, etc.
- Queuing or publish-subscribe model?  
⇒ check how Kafka delivers messages to consumer instances/groups

# Kafka vs RabbitMQ

**Figures source:** Philippe Dobbelaere and Kyumars Sheykh Esmaili. 2017. Kafka versus RabbitMQ: A comparative study of two industry reference publish/subscribe implementations: Industry Paper. In Proceedings of the 11th ACM International Conference on Distributed and Event-based Systems (DEBS '17). ACM, New York, NY, USA, 227-238. DOI: <https://doi.org/10.1145/3093742.3093908>



**Figure 1: Kafka Architecture**



**Figure 2: RabbitMQ (AMQP) Architecture**

# Hands-on

- Understanding the message broker systems and message delivery are key for stream processing
- Check our tutorial:
  - <https://github.com/rdsea/bigdataplatforms/tree/master/tutorials/basickafka>

Thanks!

Hong-Linh Truong  
Department of Computer Science

[rdsea.github.io](https://rdsea.github.io)



—  
Kiitos  
aalto.fi