



Aalto University  
School of Science

# Big Data Ingestion, Transformation and Orchestration

*Hong-Linh Truong*

*Department of Computer Science*

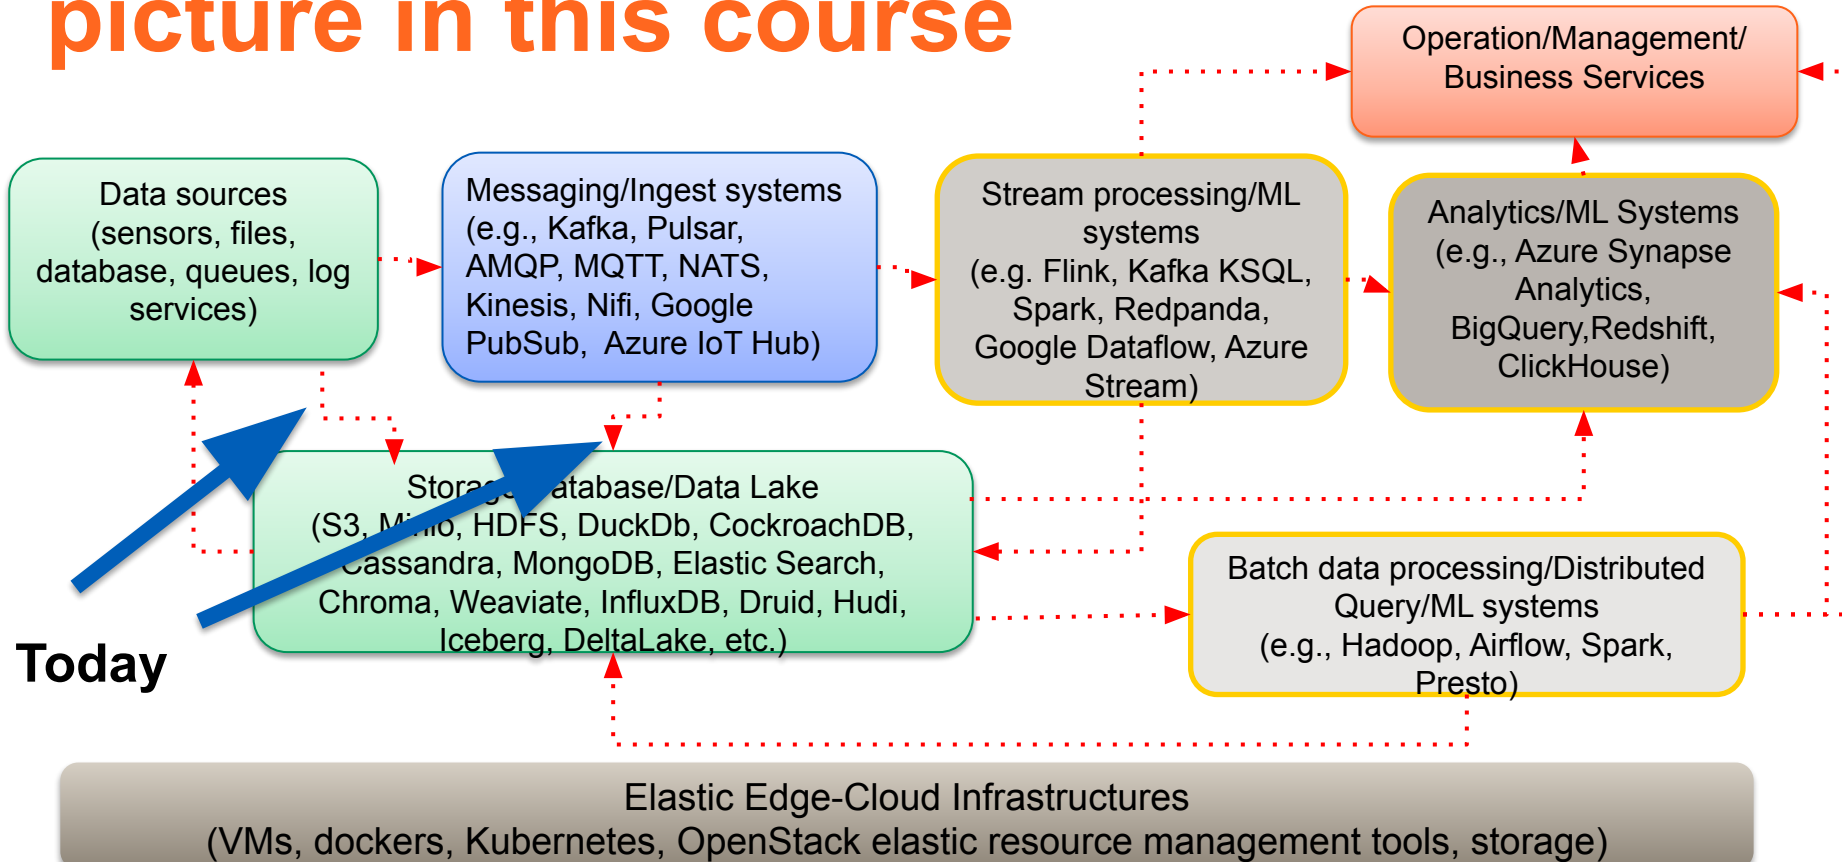
*[linh.truong@aalto.fi](mailto:linh.truong@aalto.fi)*, *<https://rdsea.github.io>*

CS-E4640 Big Data Platforms, Spring 2025, Hong-Linh Truong  
29/01/2025

# Learning objectives

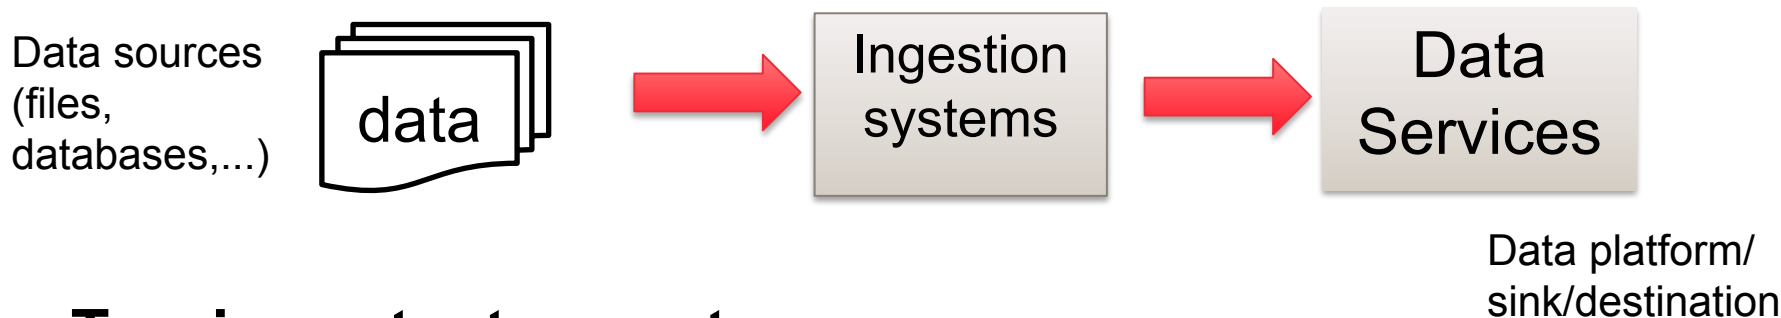
- **Understand the overall design of data ingestion**
- **Study common tasks in data ingestion**
- **Understand and design efficient, robust data ingestion pipelines/processes**
- **Learn existing technologies/frameworks for your own design**

# Big data at large-scale: the big picture in this course



# Ingestion systems

Data ingestion: move data from different sources into data platforms or selected data sinks/destinations



## Two important aspects:

- tasks and non-functional requirements
- architectures, pipelines and service models

**Reusability and extensibility**

# Data sources and sinks


## Data sources



Ingestion Pipelines

Rich, diverse types of **Connectors (libraries)** for source/sink **Connections (runtime)**

## Data sinks



## Big data platform examples:

Hadoop File systems  
Google Storage  
Amazon Storage  
Druid  
Google BigQuery  
Hive  
MongoDB  
ElasticSearch  
Cassandra  
InfluxDB  
Hudi  
Kafka, Pulsar

# Diverse requirements from $V^*$ of big data

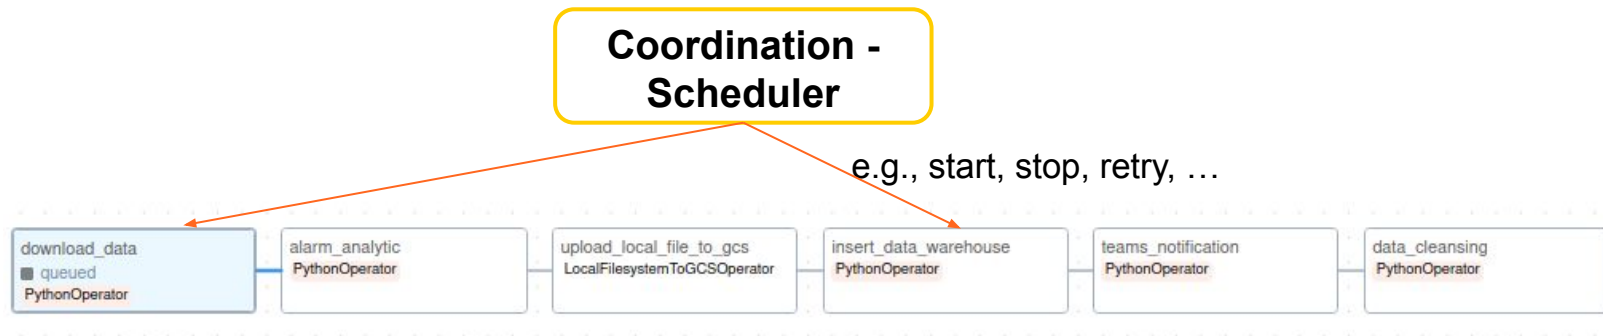
- **Requirements based on data characteristics**
  - structured, unstructured and semi-structured
  - speed, volume, accuracy, confidentiality, data regulation
- **Interact with data sources:**
  - Access APIs and protocols
    - *REST API, ODBC, SFTP, specific client libs*
    - *MQTT, AMQP, CoAP, NATS, Kafka,...*
  - Connection management:
    - *Performance, reliability and security*
- **How deep can a platform support complex requirements?**
  - e.g., able to go into inside of data elements (understanding the syntax and semantics of data)?

# Data transformation

- **Relation with ETL (Extract, Transform, Load)**
  - during ingestion, data transformation tasks might be needed
  - ETL has many operations to deal with the semantics/syntax of data and the business of data
- **Data transformation within ingestion**
- **Data transformation done after, within the (target) platform**
  - called ELT (load and then transformation)

**Performance, correctness and quality assurance**

# Ingestion needs task coordination



- **Big data ingestion involves**
  - many tasks
  - multiple tenants/users
  - ad-hoc, on-demand vs scheduled task pipelines
  - data movement in single vs cross data centers
- **Complex coordination techniques are used**
  - tasks are not in the same machine (executors) → data exchange among tasks (data dependencies)



# Ingestion tasks: common tasks and requirements

# Main tasks in ingestion

- **Key categories of tasks**

- data access and extraction
- data routing
- data wrangling
- data storing
- lineage and observability for quality assurance/governance (quality check)


- **Consumer/user defined tasks vs platform tasks**

- **Other supports: compression, end-to-end security**

**They are different for batch vs near real time ingestions**

# Ingestion tasks implemented as extensible, composable connectors

- Basic tasks for big data ingestion can be (re)used in different cases
- Support end-user goals
  - enables the user to do many tasks through configurations and extensions
- Enable pluggable approaches is important



# Data access and extraction tasks

- **Access**

- obtaining/copy data from sources
  - including change data capture (CDC)
- often built based on common protocols and APIs
  - **connector** library: strongly related to data storage/database/datalake sink/source)
  - runtime **connection** management: maintain list of connections created from connectors
- reusability is important!

- **Encryption, masking/anonymization**

- might need to be done when accessing and extracting data
- also during transfers of data
- data security requirements, personally identifiable information

# Example with NETACT Log

29869;10/01/2017

00:57:56;;Major;PLMN-PLMN/BSC-xxxxxx/BCF-xxx/BTS-xxx;XYZ01N;ABC08;D  
EF081;BTS OPERATION DEGRADED;00 00 00 83 11 11;Processing

## Simple Grok

```
1 input {
2   file {
3     path => "/tmp/alarmtest2.txt"
4     start_position => "beginning"
5   }
6 }
7 filter {
8   grok {
9     match => {"message" => "%{NUMBER:AlarmID};%{DATESTAMP:Start};%{DATESTAMP:End};%{WORD:Severity};%{NOTSPACE:NetworkType};%{NOTSPACE:BSCName};%{NOTSPACE:Sta}
10  }
11 }
12 output {
13   stdout {}
14   csv {
15     fields => ['AlarmID', 'Start', 'Stop', 'Severity', 'NetworkType', 'BSCName', 'StationName', 'CellName', 'AlarmInfo', 'Extra', 'AlarmStatus']
16     path => "/tmp/test-%{+YYYY-MM-dd}.txt"
17   }
18 }
```

# Change data capture (CDC)

- **The principles:**

- capture and ingest only **new data** by listening data changes
  - “**new**”: application-specific, e.g., based on time, value, and version.
- leverage many features of databases (update, query, insert operations), data stream offsets and status notification (e.g., the availability of new files)

- **Implementation in different tools like Redhat Debezium, Hudi DeltaStreamer, Kafka Connect**

# Data wrangling

- **Convert/transform data from one form to another**
  - cleansing, filtering, merging, enriching, inferring, and reshaping data
- **Require access to the data content!**
- **Key design choices**
  - do you support it during the ingestion or after the ingestion?
  - as a platform provider: are you able to do this?

# Data wrangling

- **In the context of big data platforms**
  - define or discover data schemas
  - automatic data wrangling: write pipelines/programs which do the wrangling
- **Wrangling programs provided by customers**
  - needs the platform to support debugging, monitoring and exception handling
  - runtime management for wrangling
- **Wrangling programs provided by platforms**
  - constraints in dealing with customer data



# Examples

Write your own  
code with  
Pandas/Dask  
and Dataframe?



Automatically  
generate code  
for wrangling?

```
Alarms={}
with open(sys.argv[1], 'rb') as csvfile:
    reader = csv.DictReader(csvfile)
    for row in reader:
        try:
            #print row['Started']
            alarm_time = datetime.strptime(row['Started'], '%d.%m.%Y %H:%M:%S')
            #diff = start_time - alarm_time
            #print "different time is ",diff
            if alarm_time >= start_time:
                #print(row['RNW Object Name'], row['Severity'])
                typeOfAlarm = 0
                cleanSeverity = re.sub('\W+', '', row['Severity'])
                if (cleanSeverity in mobifone.AlarmSeverity.keys()):
                    typeOfAlarm = mobifone.AlarmSeverity[cleanSeverity]
                #print ("Type of Alarm: ", typeOfAlarm)

                if row['RNW Object Name'] in Alarms:
                    #print "Again"
                    severies = Alarms[row['RNW Object Name']];
                    severies[typeOfAlarm] = severies[typeOfAlarm] + 1
                else:
                    severies = [row['RNW Object Name'], 0, 0, 0, 0, 0, 0]
                    severies[typeOfAlarm] = severies[typeOfAlarm] + 1
                    Alarms[row['RNW Object Name']] = severies;

        except:
            print "Entry has some problem"
            print row
            #timestamp = long(row['TIME'])
            #times.append(datetime.datetime.fromtimestamp(timestamp/1000))
            #times.append(long(row['TIME']))
            #signals.append(float(row['GSM_SIGNAL_STRENGTH']))
dataframe = pd.DataFrame(Alarms, index=mobifone.AlarmSeverityIndex).transpose()
alarmdata = dataframe.as_matrix();
#TODO print Alarms to file
#only for debugging
print dataframe
dataframe.to_csv(outputFile, index=False)
```

# Data Wrangling: complex data transformation and processing

- **More complex data processing**
  - extract only important data
    - *feature engineering*
  - enrich data on the fly with external sources

## Example: extract vectors from images

```
class TowheeExtractor(BaseExtractor):
    def __init__(self):
        self.towhee_feature_extractor = (
            pipe.input('file_path')
                .map('file_path', 'img', ops.image_decode.cv2_rgb())
                .map('img', 'embedding', ops.image_embedding.timm(model_name='resnet50'))
                .map('embedding', 'embedding', ops.towhee.np_normalize())
                .output('embedding')
        )

    def get_model_name(self):
        return "resnet50"


    def feature_extractor(self, image_file):
        ##currently only a single figure so we have to get the first element
        embedding=self.towhee_feature_extractor(image_file).get()[0]
        return embedding
```

# Behind the scene: complex code & libraries hide low level distributed/parallel tasks

- **Complex distributed and parallel tasks for ingestion**
- **Complex coordination**
- **Underlying, internal task models:**
  - MapReduce model
  - embarrassingly parallel model
  - full direct acyclic graph (DAG) task model


**Embarrassingly Parallel**

Hadoop/Spark/Dask/Airflow/Prefect




**MapReduce**

Hadoop/Spark/Dask



**Full Task Scheduling**

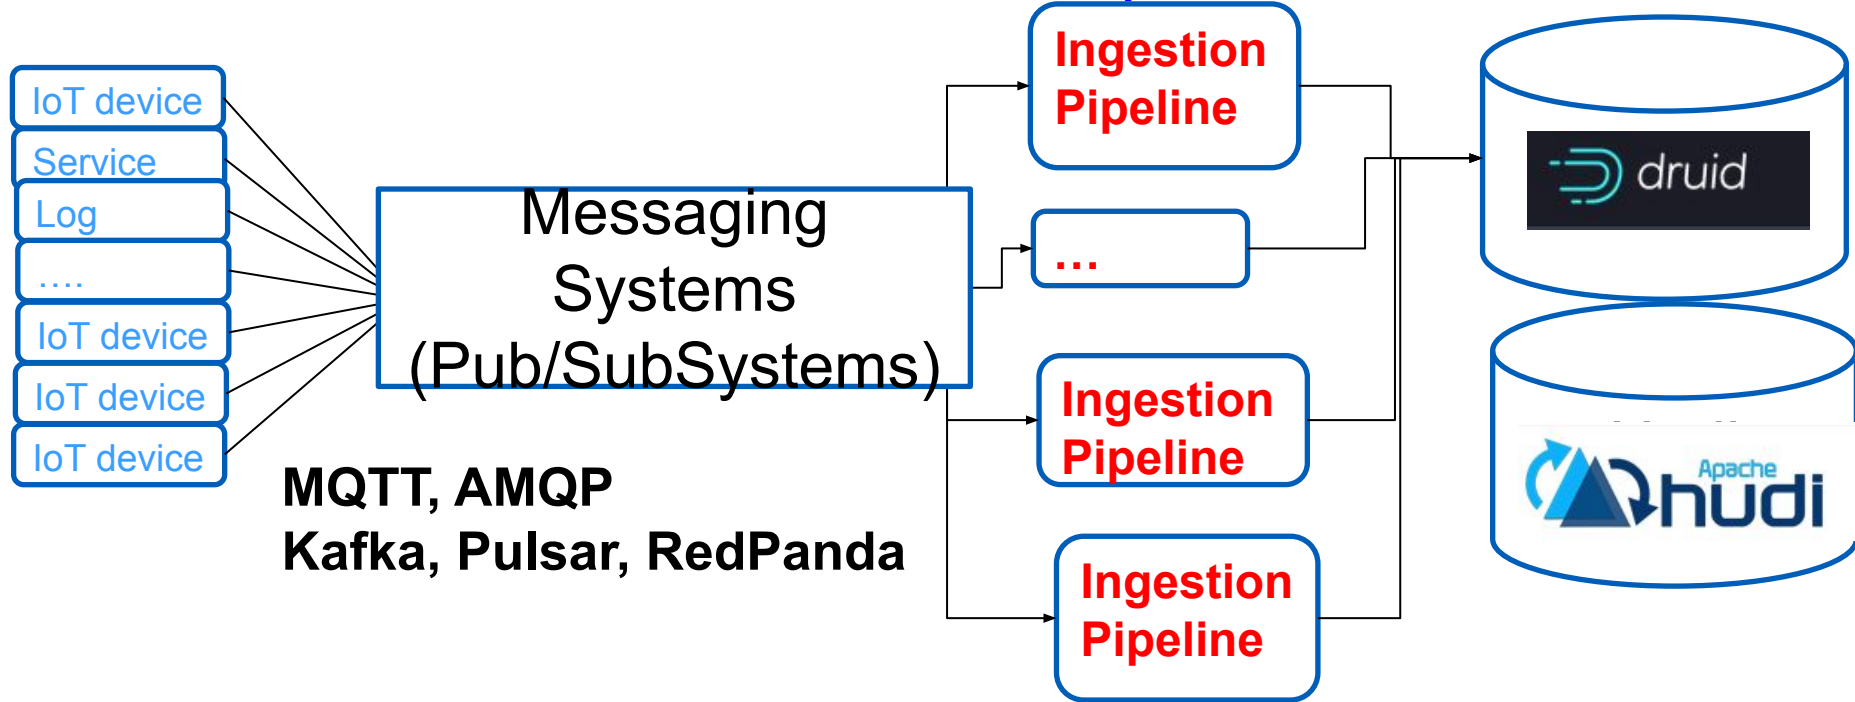
Dask/Airflow/Prefect



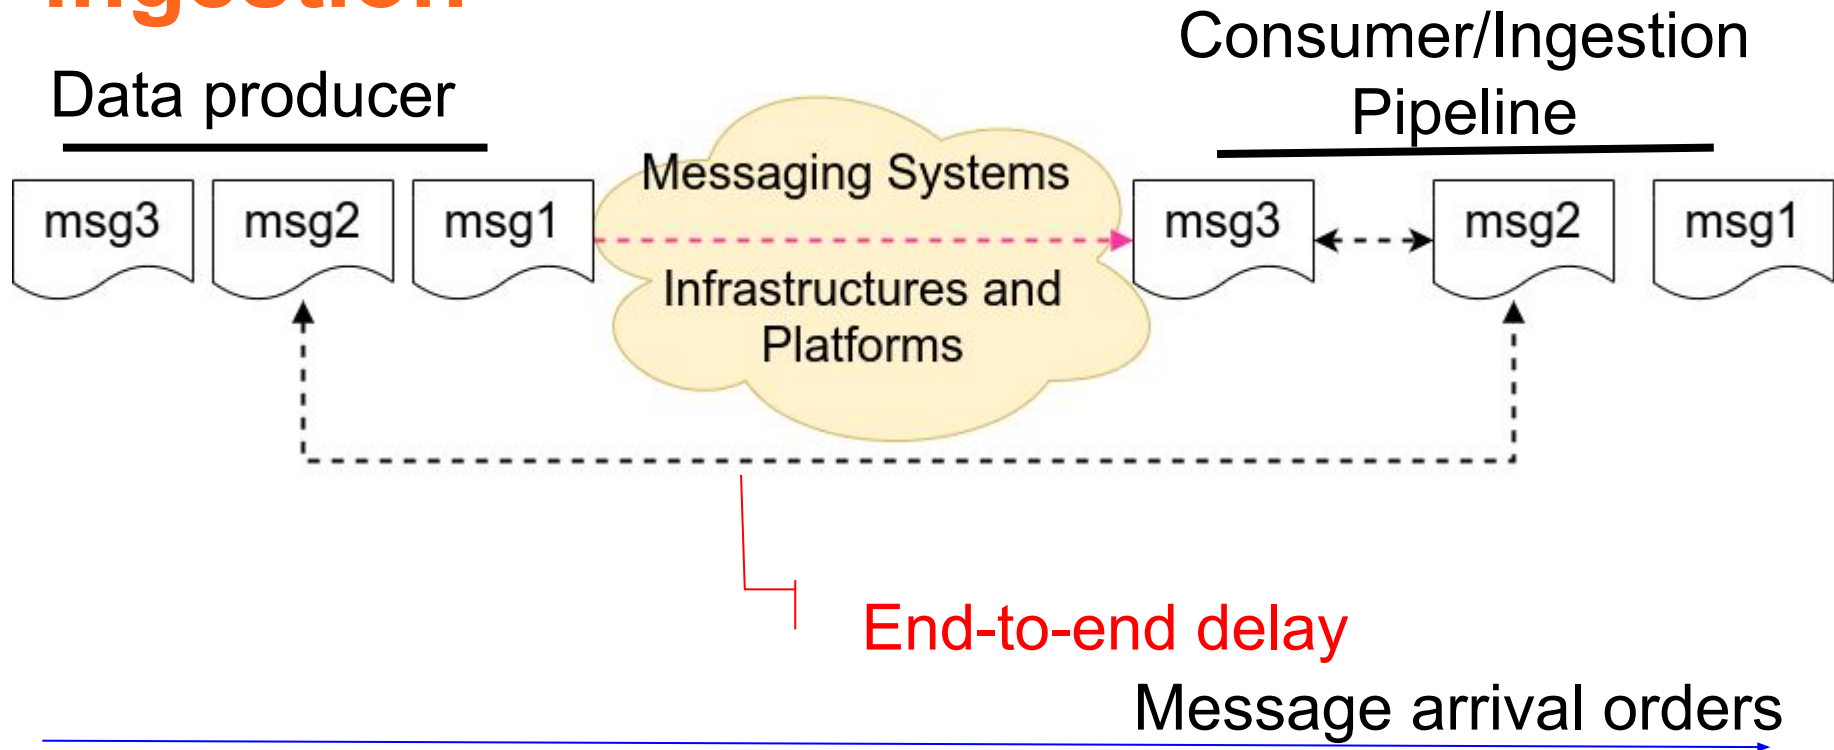
**Figure source:** <https://docs.dask.org/en/stable/graphs.html>

# Near real time ingestion

Run as a service or arbitrary and  
can be complex




# Key issues in streaming data ingestion



# Some key issues for ingestion of streaming data

- **Late data, data out of order?**
- **Exactly once?**
- **Back pressure and retention**
  - for individual components or the whole pipelines
- **Scalability and elasticity**
  - changes in data streams can be unpredictable
- **Data quality**
  - fast processing, overhead

# Split (pub/sub) and partition with ingestion




# Some key issues

- **Multiple topics/streams of data**
  - speed and volume of data per topic vary
  - should not have duplicate data in data store
- **How to distribute topic/data to ingestion pipelines?**
  - application logics and performance
- **Where should we run the messaging system?**
- **Where and when should the elasticity be applied?**
  - computing resources + data resources (streams)



# Processing data before ingestion requires some streaming techniques



E.g., basic processing for data rollup/summarization



Aalto University  
School of Science

# Pipeline designs and execution models

# Dealing with diverse data structures

- the data sender/producer and the receiver/consumer are **diverse**
  - implemented with different languages and software technologies
    - *need to guarantee the message syntax and semantics*
  - performance overhead due to data format conversion
- **Solutions:** don't assume! agreed in advance
  - agreed in advance  $\Rightarrow$  within the implementation or with a standard
  - know and use tools to deal with **syntax differences**
- **Understanding the syntax allows some automatic transformations/quality checks**
- **But semantics are domain/application-specific**

# Example of interoperability in data transfer: Arvo

Syntax specification  
<https://avro.apache.org/>

```
{  
  "namespace": "bdp.courses.aalto.fi",  
  "type": "record",  
  "name": "event",  
  "fields": [  
    {"name": "station_id", "type": "string"},  
    {"name": "datapoint_id", "type": "int"},  
    {"name": "alarm_id", "type": "int"},  
    {"name": "event_time", "type": "int"},  
    {"name": "value", "type": "float"},  
    {"name": "valueThreshold", "type": "float"},  
    {"name": "isActive", "type": "boolean"}  
  ]  
}
```



Also: Protobuf, JSON Schema

# Interoperability in data processing in ETL

## Data sources and formats

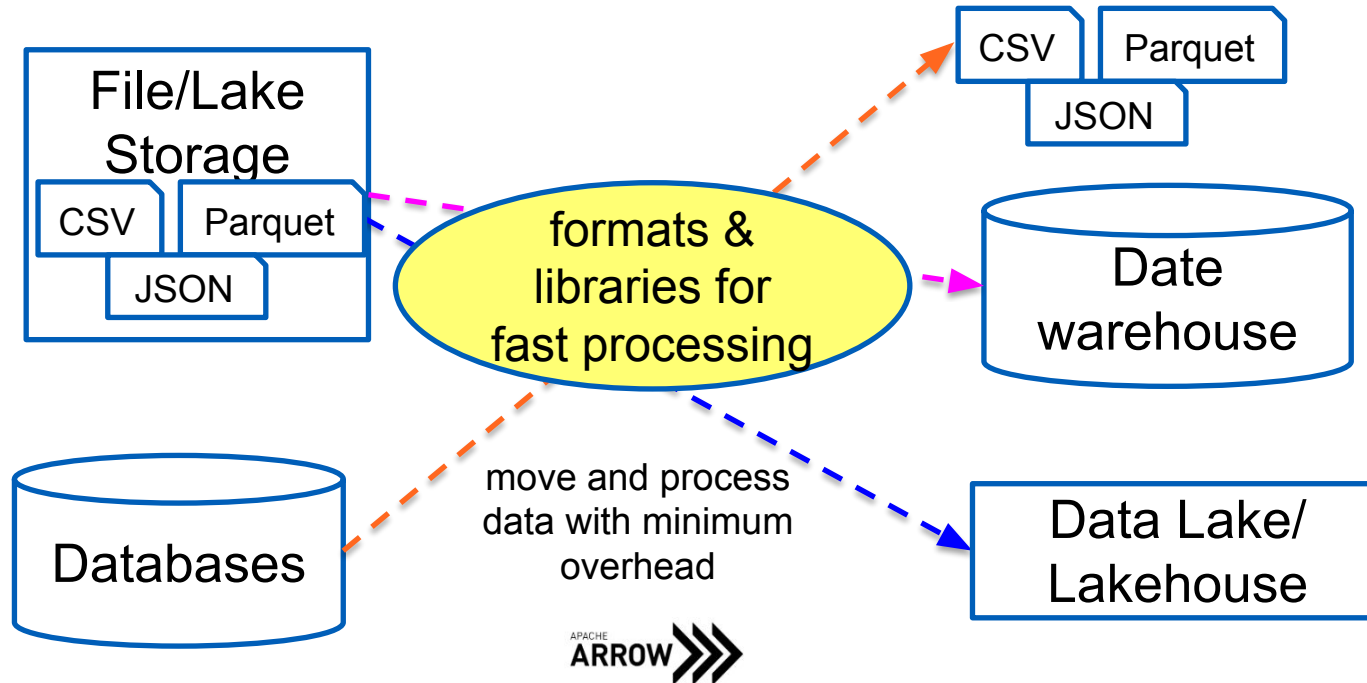
## Data sinks and formats

## Example: Apache Arrow

In-memory  
columnar format

rich ecosystems

<https://arrow.apache.org/>



# Ingestion is not a single task!

## Ingestion pipelines/processes: architectures, composition, coordination, and tools

# Complex deployment and composition models

- Understanding strong dependencies between protocols/APIs, **security, performance and management**

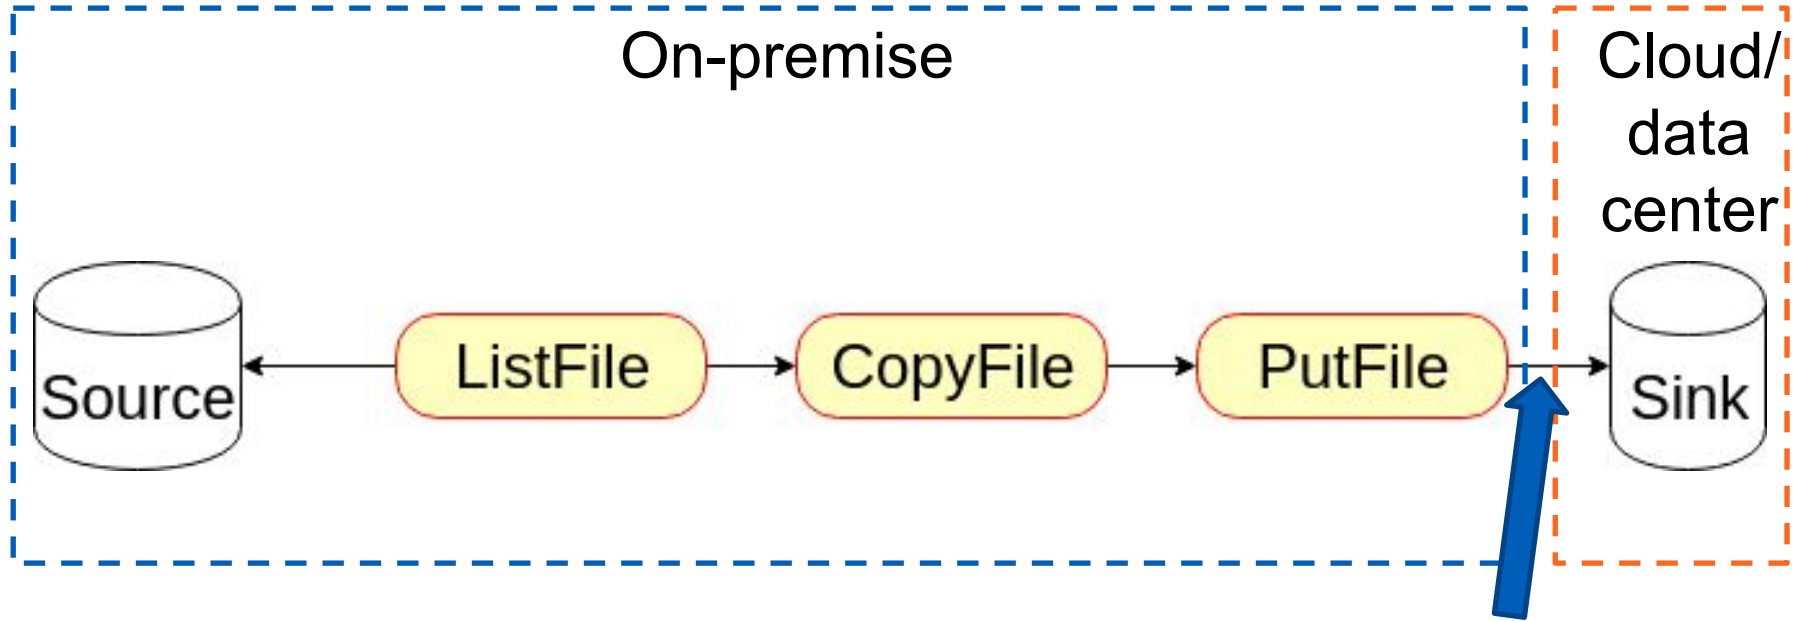


**Tenant/  
user**

**Ingestion pipeline developer  
( for whom?)**

**Data  
store/platform  
provider**

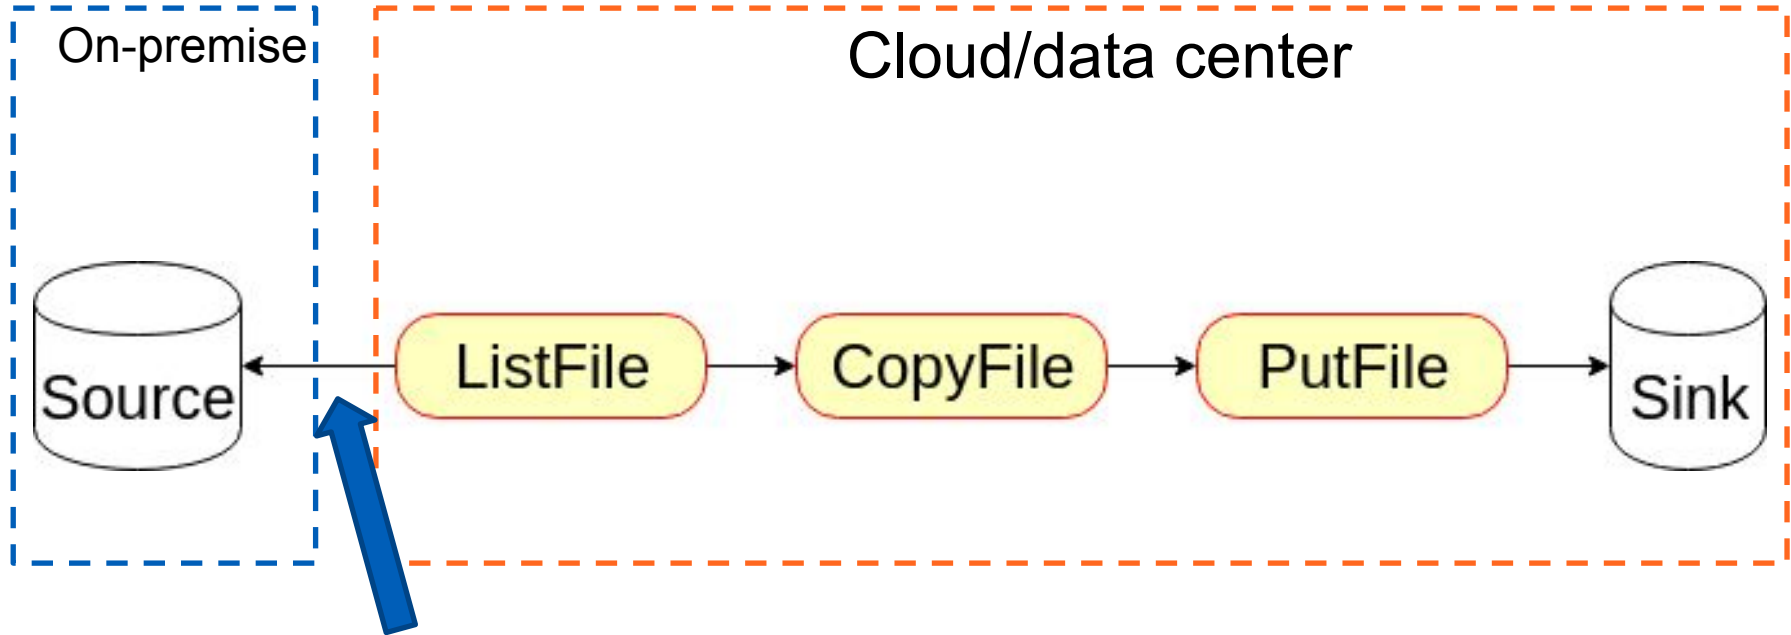
# Complex deployment and composition models



**APIs, protocols and deployment issues?**

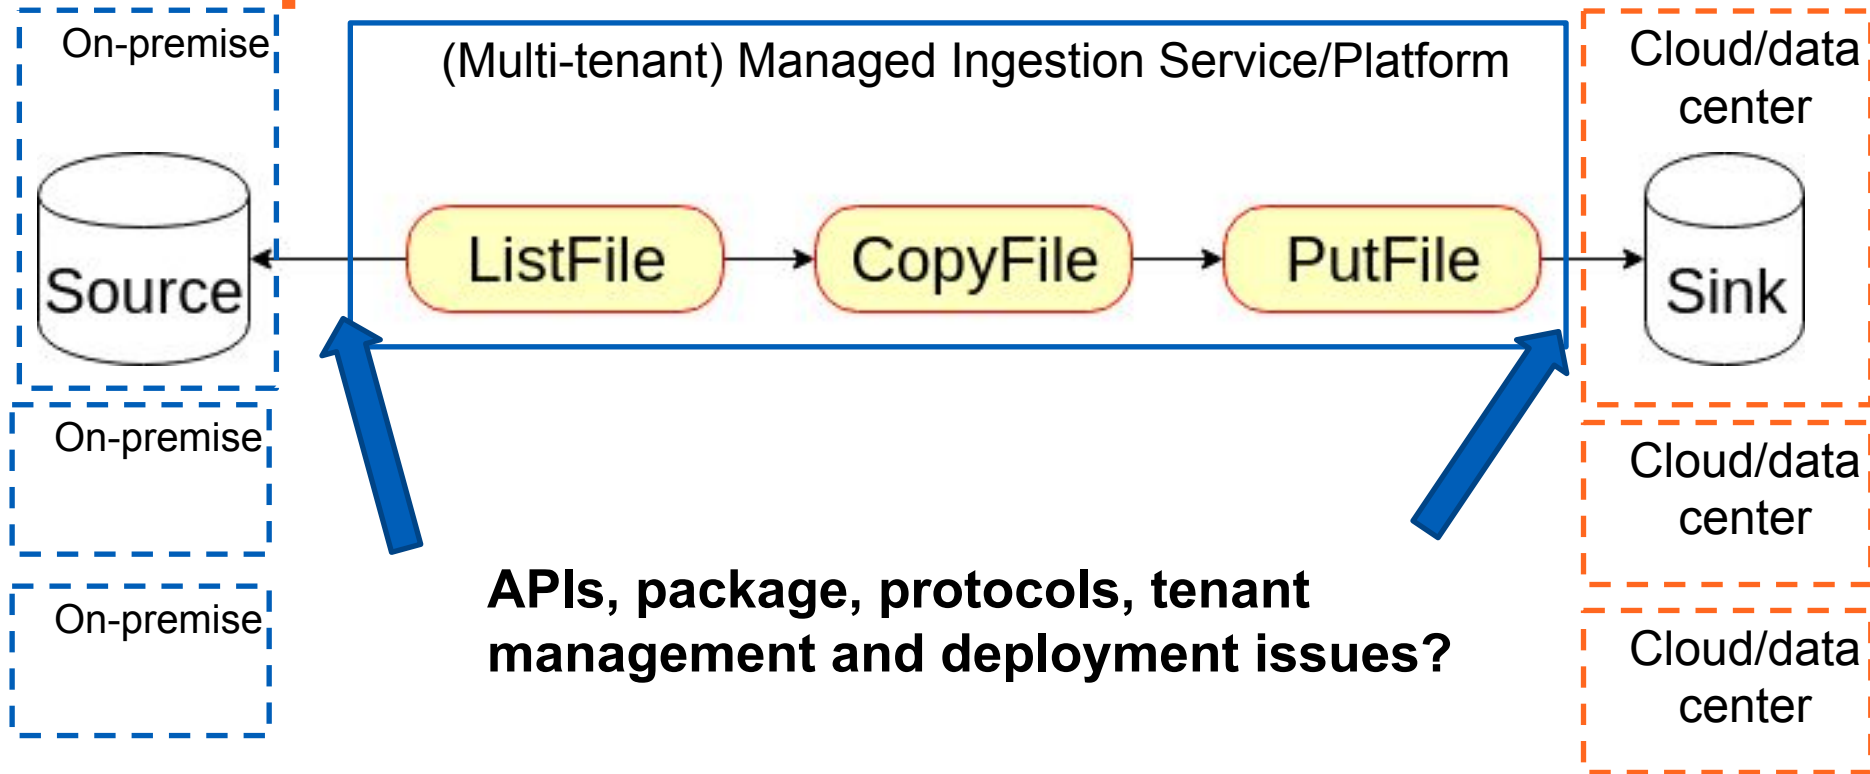


# Complex deployment and composition models



**APIs, protocols and deployment issues?**

# Complex deployment and composition models

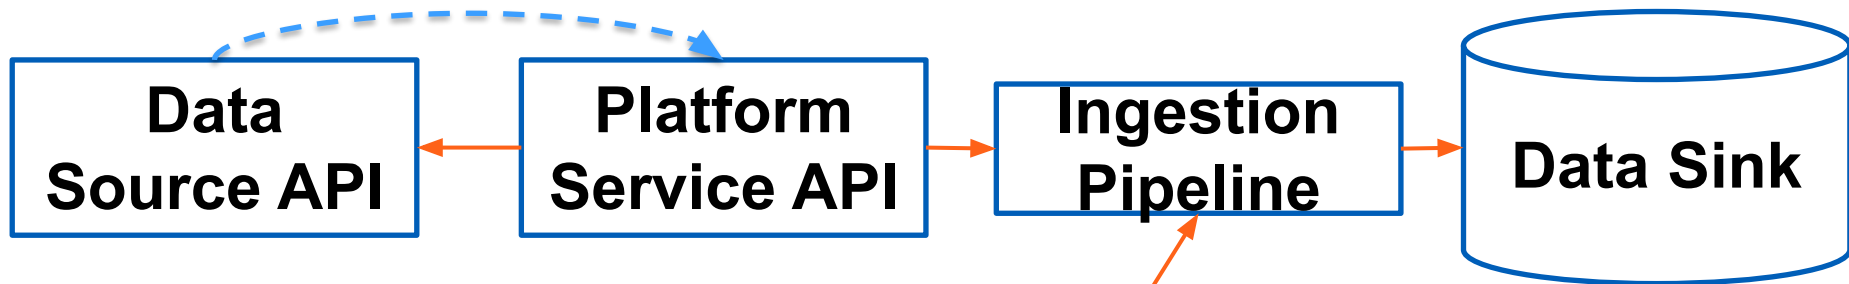


# Orchestrating batch ingestion pipelines


- **Data to be ingested is bounded**
  - files or messages are finite
- **Ingestion architectural styles**
  - (1) Direct APIs, (2) reactive pipelines, (3) tasks/workflows
- **Incremental ingestion**
  - dealing with the same data source but the data in the source has been changed over the time (related to change data capture)
- **Parallel and distributed execution**
  - use workflows and distributed processing engines

# Simple, direct APIs for ingestion

**Pull model:** register webhook/API

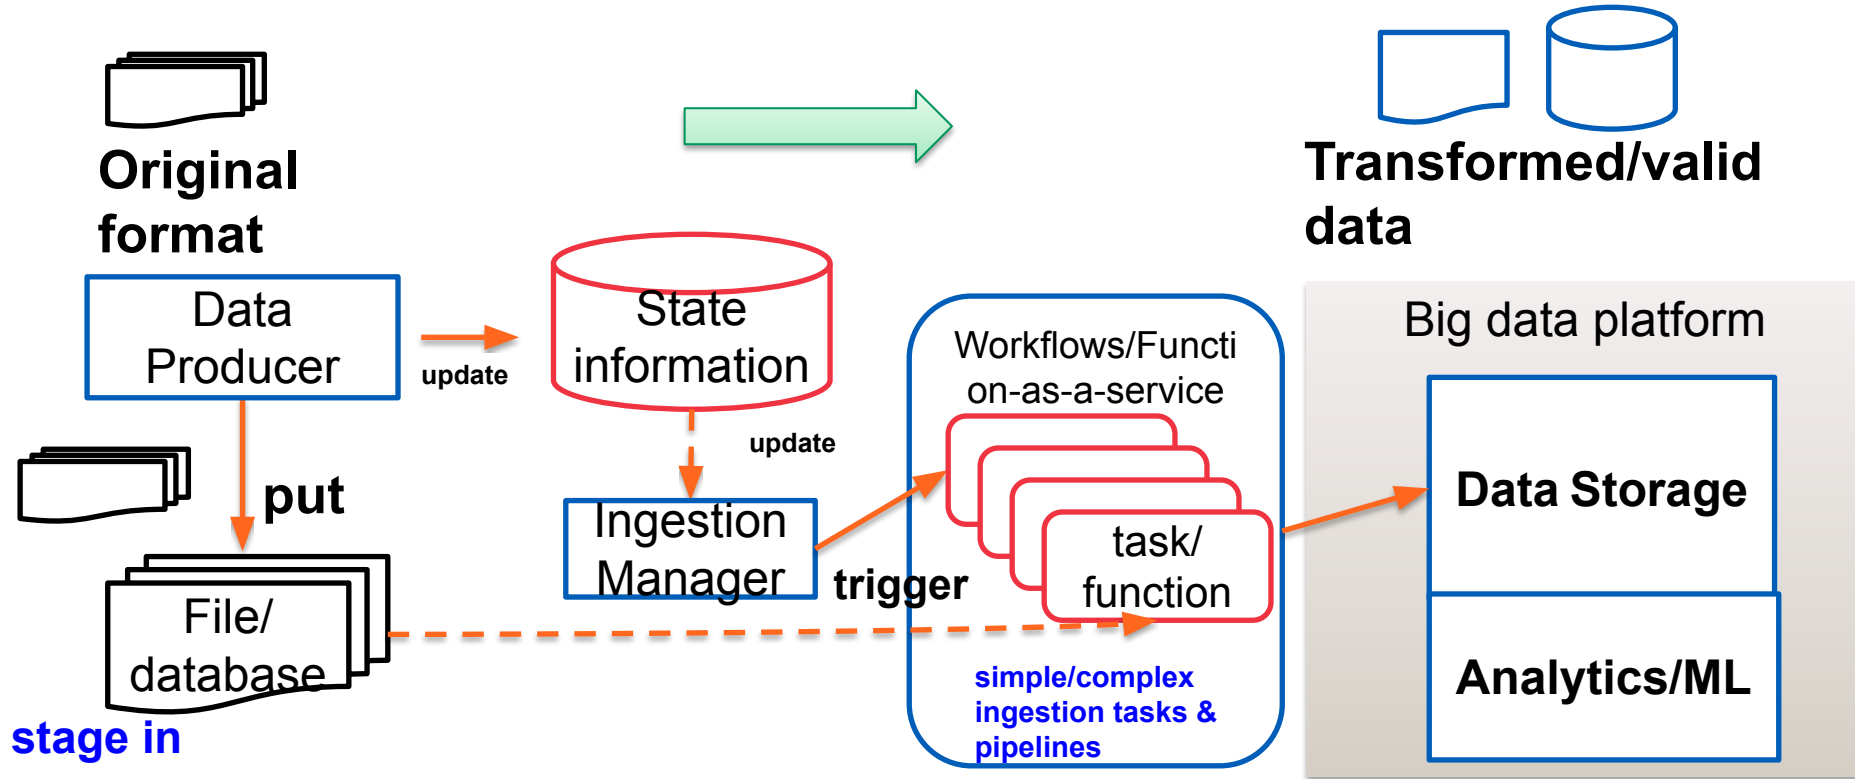


**Push model**



## Try to analyze pros and cons for your platform!

# Reactive pipelines with functions/workflows/containers



# Ingestion workflows orchestration

- **Different tasks**

- access and copy, extract, covert, quality check, and write data
- tasks are connected based on data or control flows


- **Workflows**

- a set of connected tasks is executed by an engine
- tasks can be scheduled and executed in different places
- flexible designs

- **Different tenants have different service level agreements**

- Performance, reliability, and cost.

# E.g., workflow based on scheduled time, with Apache Airflow

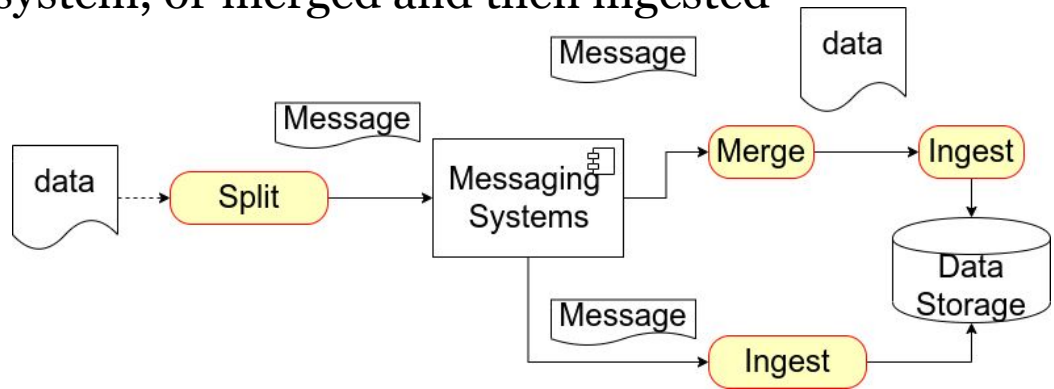


# Microbatching in data ingestion

- **Microbatching: we mean the strategy to deal with big dataset using batches of data (small chunks)**
  - not necessary the same as using batch systems to transfer small data in near realtime
- **Data is split into different chunks for ingestion**
  - using streaming or batch systems to transfer data chunks
  - chunks are ingested into the system, or merged and then ingested

**Example: with streaming system**

**Be careful with the data semantics/integrity!**

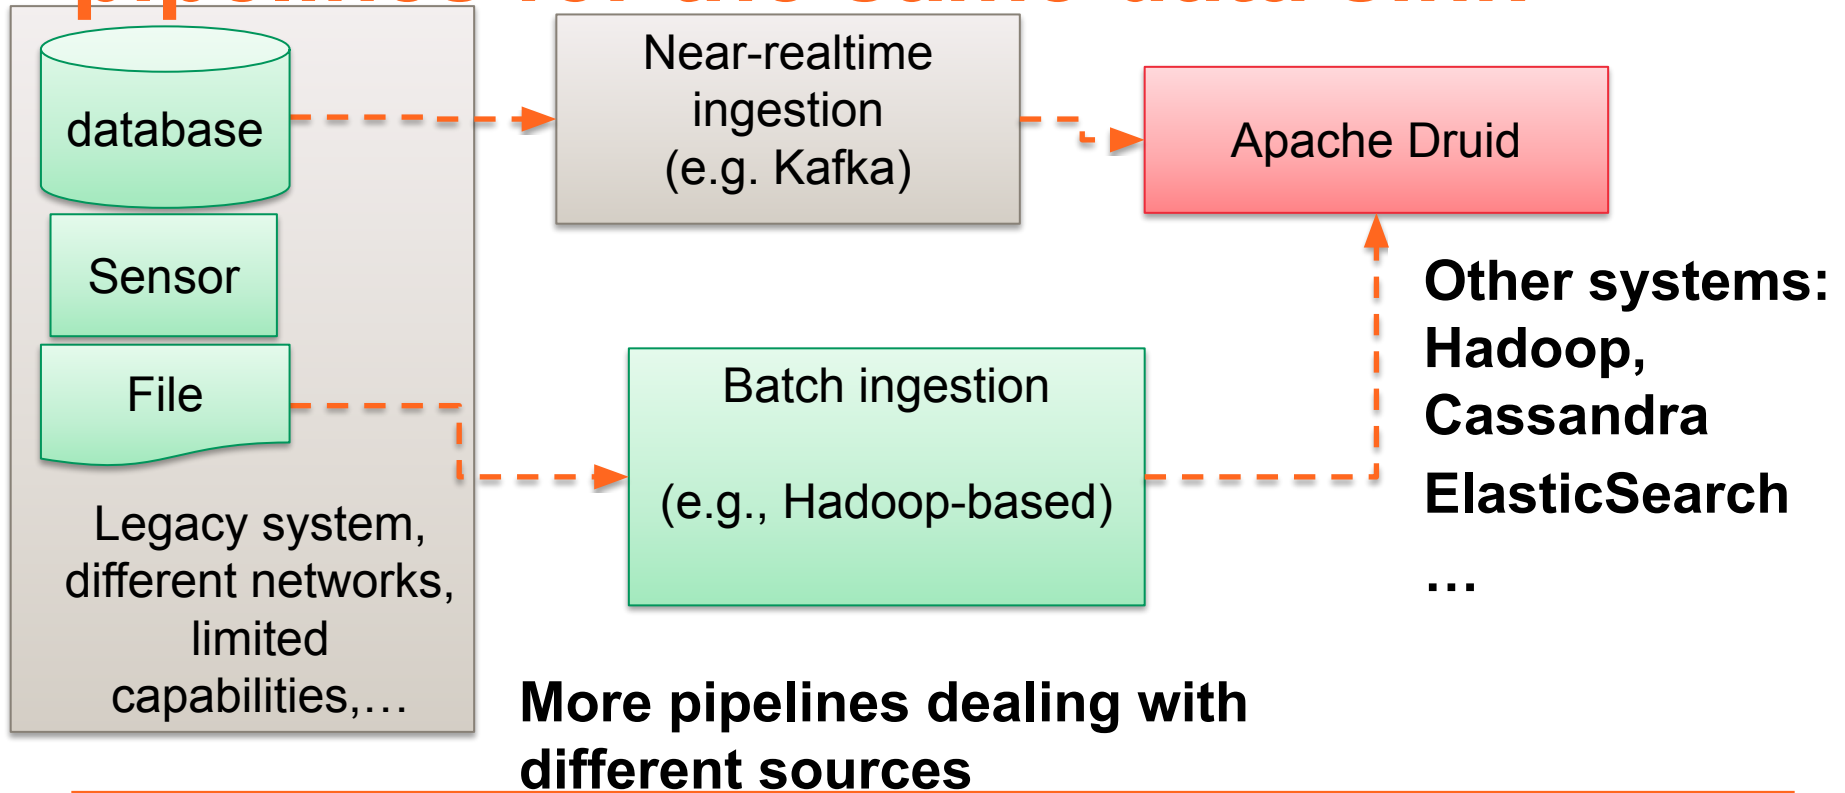




# Combining ingestion pipelines in big data platforms

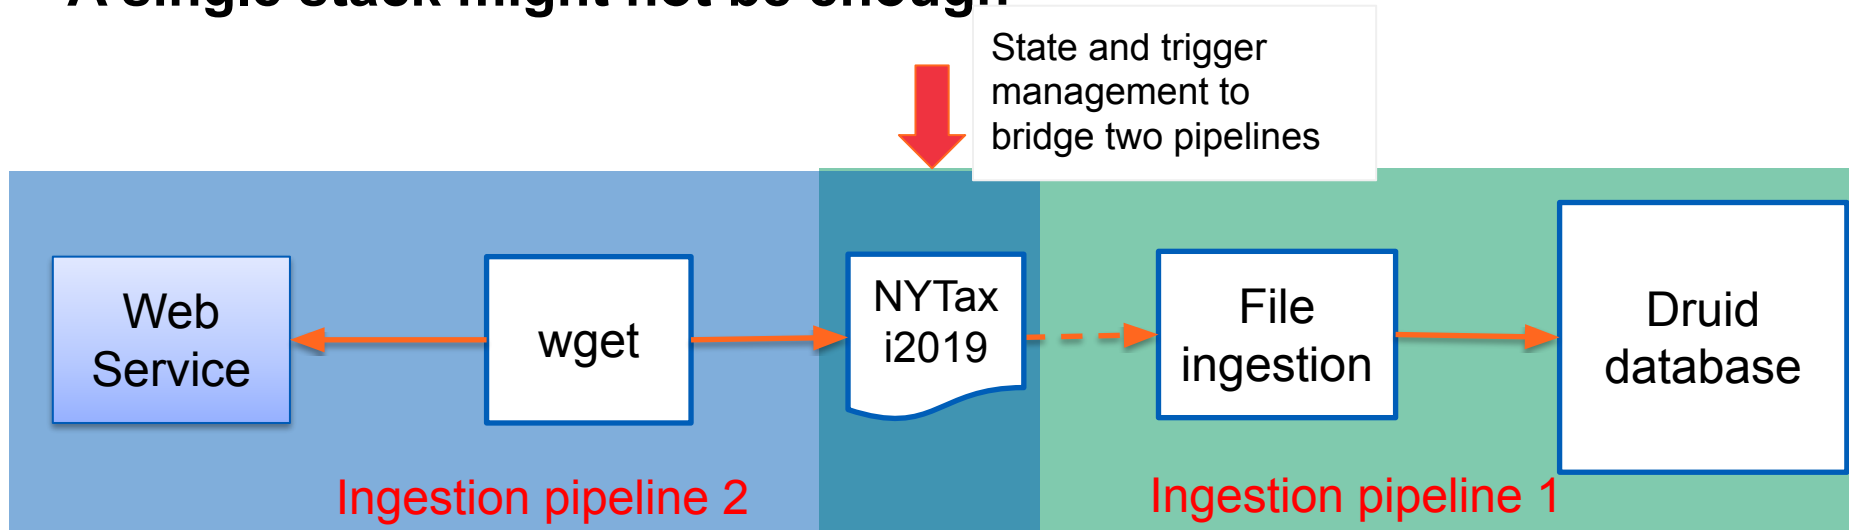
- **Multiple types of pipelines for multiple types of tenants/users**
  - **A tenant/user might need different integrated pipelines**
- ⇒ **Both batch and near-realtime ingestion are supported**
- **Complex architectural designs**
    - ingestion pipeline-to-pipeline needs “bridges”

# Supporting multiple types of pipelines for the same data sink



# Connecting different ingestion pipelines

A single stack might not be enough



**Real-world:**  
**both pipelines and their connections are complex**

# Quality control/data regulation assurance

## Data sources

Responsible data: profiling, sampling, measuring quality and inspecting data  
⇒ implications on data products

Log file

...

Transaction records

User-provided data

Access data

Process & profile data

Data Sinks

Data observability

data testing

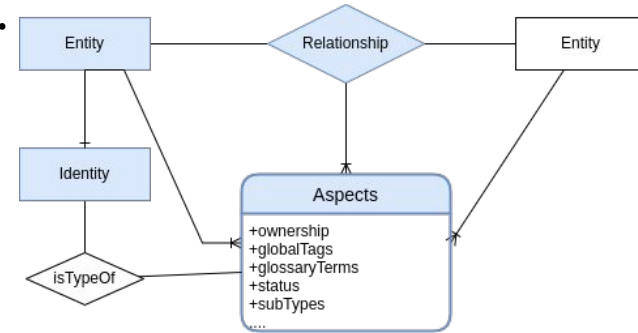
patterns/rules/AI

duplication detection

**Challenging issues: misinformation, GDPR, data quality, inappropriate content**

# Data lineage and observability (1)

- **FAIR principles** (<https://www.nature.com/articles/sdata201618>)
  - findable, accessible, interoperable, and reusable
- **Lineage/Provenance**
  - capture relevant information for understanding how data has been moved, transferred, processed, etc.
  - metadata models: W3C Provenance Model, DataHub, etc.
- **Key issues**
  - which metadata must be captured?
  - based on existing tools or your own?
- **Instrumentation/logging processes and automated data lineage → performance overhead!**



High level view of datahub  
see

<https://datahubproject.io/docs/metadata-modeling/metadata-model>

# Data lineage and observability (2)

- **Data observability: the health about data**
  - near-real time metrics, offline checks and possible dashboards
  - similar to service observability, relying on traces, logs, metrics, etc.
- **Focus on data**
  - data metrics (volumes, data quality, schemas, lineage)
  - issues due to data problems
  - data ingestion processes/workflows
- **Some solutions**
  - validation of data against design schemas (e.g., Schema Registry in Kafka)
  - checks of realtime and offline data quality attributes → integrate with data ingestion processes or offline data profiling
  - integrated data quality tests in pipelines (e.g., data testing)

Tools:



Microsoft Presidio

YData Profiling




Aalto University  
School of Science

CS-E4640 Big Data Platforms, Spring 2025, Hong-Linh Truong  
29/01/2025


# Quality control/data regulation assurance (1)

## Design: different evaluation modes


Pull,  
pass-by-reference  
model for  
evaluating data  
concerns



Pull, pass-by-value  
model for  
evaluating data  
concerns



Push model for evaluating data  
concerns of active data sources



**Source:** H. -L. Truong and S. Dustdar, "On Evaluating and Publishing Data Concerns for Data as a Service," *2010 IEEE Asia-Pacific Services Computing Conference*, doi: 10.1109/APSCC.2010.54.

# Quality control/data regulation assurance (2)

- **Before, after or during the ingestion/transformation**
- **In-process vs out-process**
  - in process: using libraries doing data quality very fast
  - out-process: a separate task in the workflow or external programs/services
- **Profiling, sampling, ML techniques for data quality**
- **Examples:**
  - Using a separate program like pydeequ Spark to check quality
    - <https://github.com/rdsea/bigdataplatforms/tree/master/tutorials/dataquality>
  - Anonymizing data
    - <https://microsoft.github.io/presidio/anonymizer/>




# Tooling for ingestion pipelines

# Tooling

- **Given different ingestion models, how do you deliver your ingestion tools/services?**
- **(Traditional) ways of REST API/specific client libraries**
  - upload using put/get operations
- **Workflows**
  - self-developed workflows vs automatically generated workflows
- **Pipelines are bundled into containers**
  - self-developed vs generic pipelines based on user configurations

# Design tools for ingestion processes: Logstash



```
input {
  file {
    path => "${MY_INPUT_DIR}/bts-data-alarm-2017.csv"
    start_position => "beginning"
  }
}
filter {
  csv {
    separator => ",",
    columns => ["station_id", "datapoint_id", "alarm_id",
  ]
}
}
output {
  stdout {
    #...
  }
}
```

**Pipeline is defined in a configuration file**

## Pluggable approaches

Figure source:

<https://www.elastic.co/guide/en/logstash/current/getting-started-with-logstash.html> (from the previous version of Logstash)

# Design tools for ingestion processes: Apache Druid

Allow the user to build the plan: select tasks, configuration, etc. and then generate ingestion pipelines

Connect and parse raw data

Transform data and configure schema

Tune parameters

Verify and submit

Start

Connect

Parse data

Parse time

Transform

Filter

Configure schema

Partition

Tune

Publish

Edit spec

VendorID,tpep\_pickup\_datetime,tpep\_dropoff\_datetime,passenger\_count,trip\_distance,RatecodeID,store\_and\_fwd\_flag,PULocationID,DOLocationID,payment\_type,  
2,11/04/2084 12:32:24 PM,11/04/2084 12:47:41 PM,1,1.34,1,N,238,236,2,10,0,0,5,0,0,0,3,10.8  
2,11/04/2084 12:25:53 PM,11/04/2084 12:29:00 PM,1,0.32,1,N,238,238,2,4,0,0,5,0,0,0,3,4.8  
2,11/04/2084 12:08:33 PM,11/04/2084 12:22:24 PM,1,1.85,1,N,236,238,2,10,0,0,5,0,0,0,3,10.8  
2,11/04/2084 11:41:35 AM,11/04/2084 11:59:41 AM,1,1.65,1,N,68,237,2,12,5,0,0,5,0,0,0,3,13.3  
2,11/04/2084 11:27:28 AM,11/04/2084 11:39:52 AM,1,1.07,1,N,170,68,2,9,0,0,5,0,0,0,3,9.8  
2,11/04/2084 11:19:06 AM,11/04/2084 11:26:44 AM,1,1.3,1,N,107,170,2,7,5,0,0,5,0,0,0,3,8.3  
2,11/04/2084 11:02:59 AM,11/04/2084 11:15:51 AM,1,1.85,1,N,113,137,2,10,0,0,5,0,0,0,3,10.8  
2,11/04/2084 10:46:05 AM,11/04/2084 10:50:09 AM,1,0.62,1,N,231,231,2,4,5,0,0,5,0,0,0,3,5.3  
2,07/11/2053 01:25:33 PM,07/11/2053 01:25:33 PM,1,0,1,N,264,264,2,0,0,0,0,0,0,0,0  
2,12/04/2042 08:51:43 AM,12/04/2042 08:54:47 AM,1,0.29,1,N,162,162,2,4,0,0,5,0,0,0,3,4.8  
2,06/25/2041 08:46:37 PM,06/25/2041 08:52:37 PM,1,1.34,1,N,239,151,2,7,0,5,0,5,0,0,0,3,8.3  
2,11/17/2037 09:24:28 PM,11/17/2037 09:46:03 PM,1,2.99,1,N,170,143,1,15,0,5,0,5,1,7,0,0,3,18  
2,02/02/2032 12:39:23 PM,02/02/2032 01:11:39 AM,4,23.21,1,N,132,228,2,62,0,5,0,5,0,0,0,3,63.3  
2,02/13/2031 05:36:35 PM,02/13/2031 05:45:36 PM,1,1.44,1,N,236,237,2,8,1,0,5,0,0,0,3,9.8  
2,02/13/2031 05:21:28 PM,02/13/2031 05:35:36 PM,1,1.69,1,N,141,236,2,8,1,0,5,0,0,0,3,9.8  
2,05/06/2029 08:43:14 PM,05/06/2029 09:03:14 PM,4,4.47,1,N,162,80,1,17,5,0,5,0,5,4,91,5,76,0,3,29.47  
2,05/05/2029 11:22:18 PM,05/06/2029 02:02:00 AM,1,11.51,1,N,148,244,1,34,5,0,5,0,5,0,0,3,35.8  
2,02/13/2026 11:53:54 AM,02/13/2026 11:58:02 AM,2,0.85,1,N,161,43,2,5,1,0,5,0,0,0,3,6.8  
2,02/13/2026 11:06:18 AM,02/13/2026 06:26:09 PM,2,3.14,1,N,163,246,2,20,1,0,5,0,0,0,3,21.8  
2,09/13/2021 12:19:52 PM,09/13/2021 12:22:07 PM,1,0,1,N,193,193,2,0,0,0,0,0,0,0,0  
2,12/10/2020 08:34:26 PM,12/10/2020 08:54:46 PM,1,4.62,1,N,50,231,2,17,5,0,5,0,5,0,0,3,18.8  
2,12/10/2020 08:23:43 PM,12/10/2020 08:32:35 PM,1,2.44,1,N,90,50,1,9,0,5,0,5,2,06,0,0,3,12.36  
2,08/01/2020 12:20:58 AM,08/01/2020 12:47:09 AM,1,16.71,1,N,143,138,1,45,5,0,5,0,5,10,41,5,76,0,3,62.47  
2,08/01/2020 12:07:04 AM,08/01/2020 12:20:28 AM,1,2.3,1,N,238,143,2,11,0,0,5,0,0,0,3,11.8  
2,03/05/2020 06:44:16 PM,03/06/2020 03:14:32 PM,1,2.39,1,N,125,161,2,11,0,0,5,0,0,0,3,11.8  
2,03/05/2020 06:33:57 PM,03/05/2020 06:40:39 PM,1,1.04,1,N,231,125,1,6,5,0,5,0,5,1,46,0,0,3,8.76

Druid ingests raw data and converts it into a custom, indexed format that is optimized for analytic queries.

To get started, please specify what data you want to ingest.

[Learn more](#)

Source type

local

Base directory

/opt/data/rawdata/bdp

File filter

\*.csv

This path must be available on the local filesystem of all Druid services.

Apply

**A** Aalto University  
School of Science

CS-E4640 Big Data Platforms, Spring 2025, Hong-Linh Truong  
29/01/2025

52

# Design tools for ingestion processes: Apache Nifi

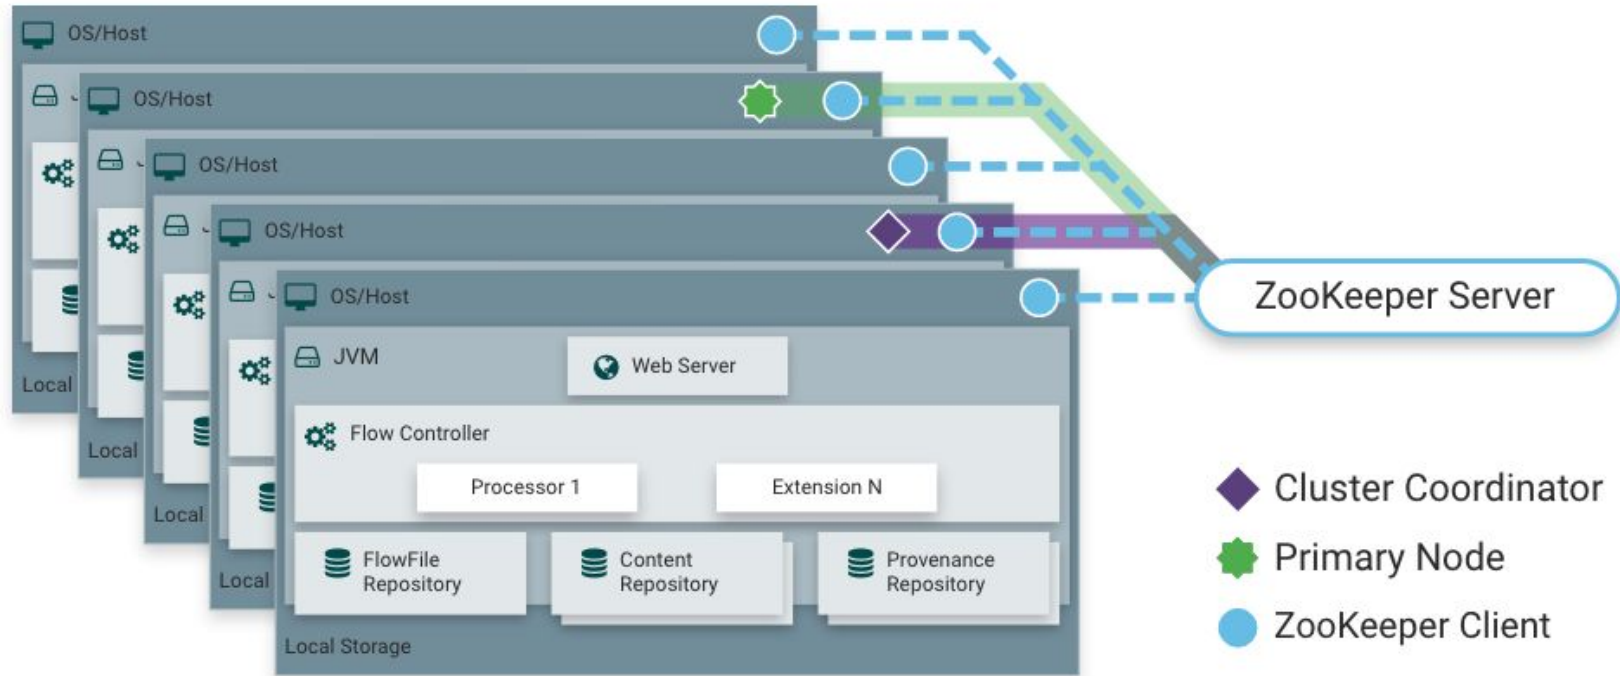


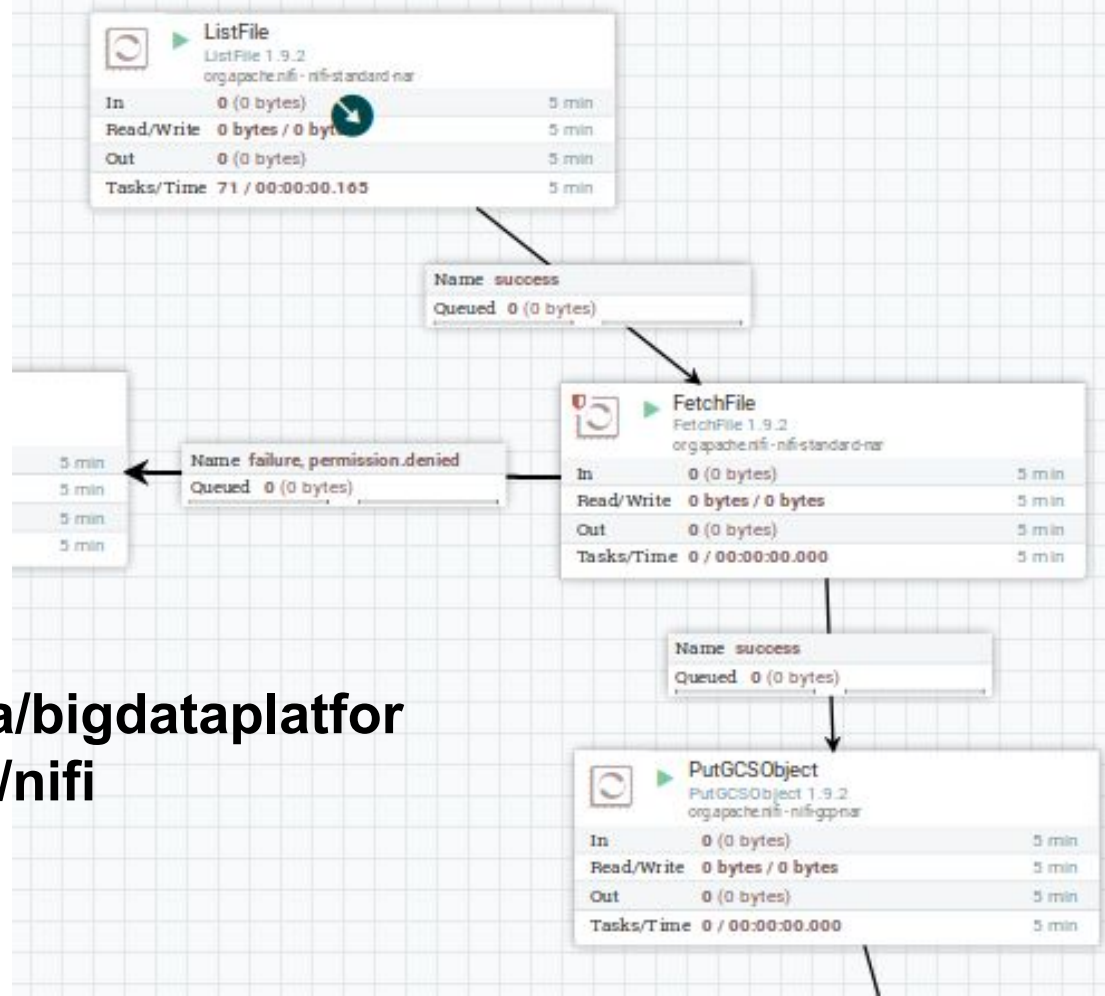
Figure source: <https://nifi.apache.org/docs.html>

# Design tools for ingestion processes: Apache Nifi - key concept

- Data is encapsulated into “FlowFile”
- **Processor** (Component) performs tasks
- **Processor** handle FlowFile and has different states
  - each state indicates the results of processing that can be used for establishing relationships to other components
- **Processors** are connected by **Connection**
- **Connection** can have many **relationships** based on states of upstream Processors

# Design tools for ingestion processes: Apache Nifi

See the tutorial:  
<https://github.com/rdsea/bigdataplatfor ms/tree/master/tutorials/nifi>



# Summary

- **Different designs of data ingestion for batch and streaming**
- **Ingestion is a complex pipeline**
  - many different sub tasks
  - complex requirements w.r.t performance, scale, failure handling
- **Different tools/stacks/services available**
  - share composable design principles, but different software models and deployments → explore them for your work
- **Do real-world designs**
  - hands-ons
  - complex designs but we do not need to “reinvent the wheel” → stay with core concepts and requirements to find the right tools!



# Thanks!

**Hong-Linh Truong**  
**Department of Computer Science**

**rdsea.github.io**