
Big Data Ingestion

Hong-Linh Truong
Department of Computer Science
linh.truong@aalto.fi



28.1.2026

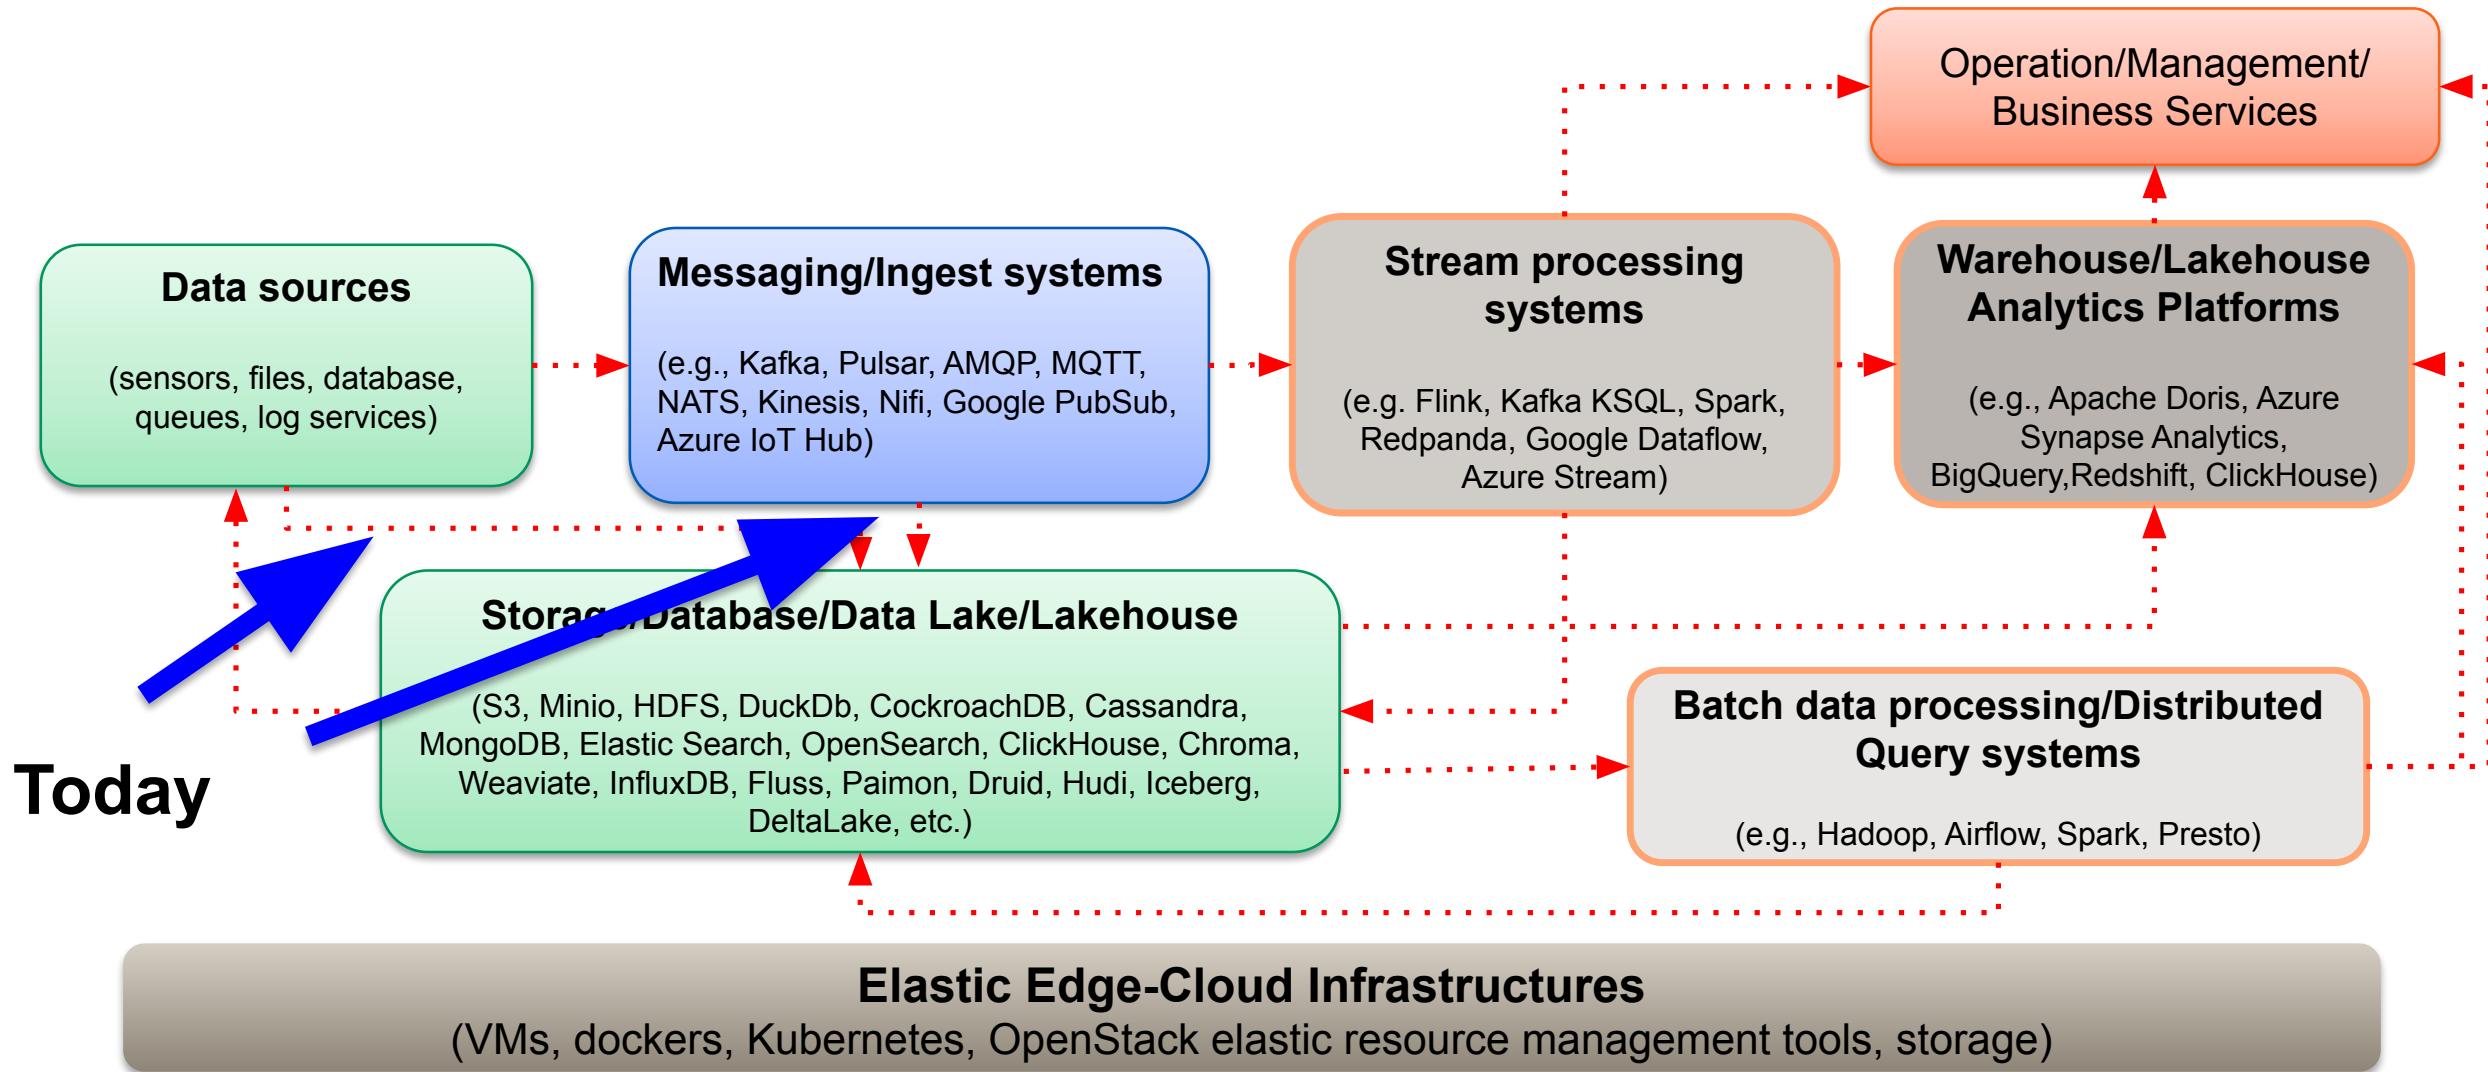


Content is available under
CC BY-SA 4.0 unless otherwise stated

Learning objectives

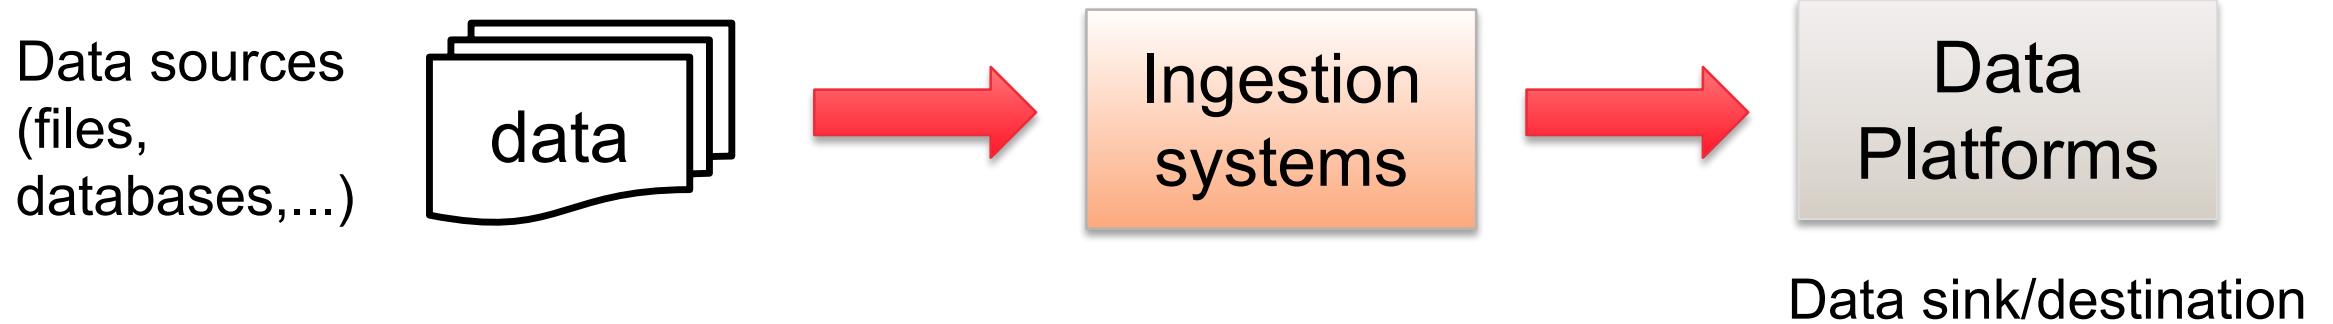
- Understand the overall design of data ingestion
- Study common tasks in data ingestion
- Understand and design efficient, robust data ingestion pipelines/processes
- Learn existing technologies/frameworks for your own design

Our big data at large-scale: the big picture in this course



Ingestion systems

Data ingestion: move data from different sources into data sinks/destinations within data platforms



The basic goals in terms of data operations:

- **insert** new data and **upsert** (insert and update) data into sinks
 - insert can be “**append-only**” and big data files can be **immutable**
- at very large scale!

Data sources and sinks

Data sources

File systems

REST Services

Messaging Systems
(MQTT, Kafka, etc.)

Databases

Data sinks

Storage/File Systems/Data Lake

Files

Big Database Services

Messaging Systems

Ingestion Pipelines

Rich, diverse types of
Connectors (libraries)
for source/sink
Connections (runtime)

Examples of data sinks

Hadoop File systems
Google Storage
Amazon Storage

Druid
Poison
Google BigQuery
Hive
MongoDB
ElasticSearch
Cassandra
InfluxDB
Hudi
Kafka, Pulsar

A!

Design and engineering aspects

- Tasks, pipelines and service models
- Non-functional requirements and service level agreement (SLA) between sources and platforms
- Deployment architectures

Performance and consistency tradeoffs

Reusability and extensibility

Diverse requirements from data sources

- Requirements based on data characteristics
 - multimodal data with structured, unstructured and semi-structured
 - speed, volume, accuracy, confidentiality, data regulation
- Interact with data sources:
 - Access APIs and protocols
 - REST API, ODBC, SFTP, specific client libs
 - MQTT, AMQP, CoAP, NATS, Kafka,...
 - Connection management:
 - performance, reliability and security
- How deep can a platform support complex requirements?
 - e.g., able to go into inside of data elements (understanding the syntax and semantics of data)?

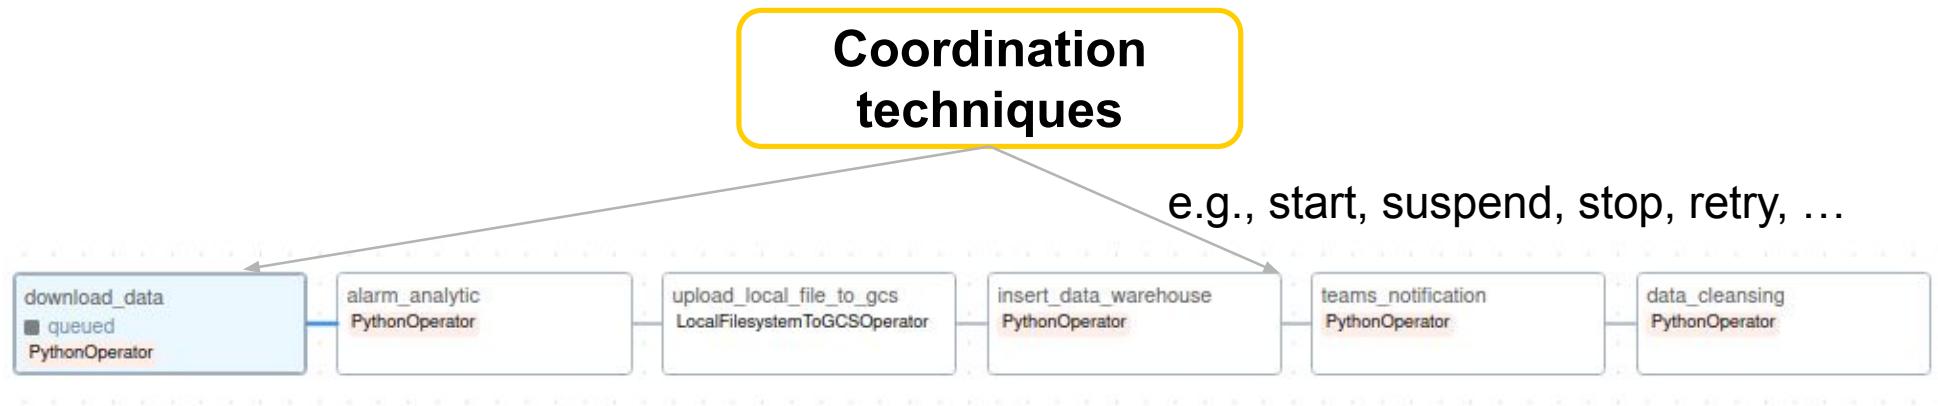
Data transformation

- Data transformation:
 - convert data from an existing form to another form for an (analytic) purpose
- Extract, Transform, Load (ETL)
 - Extract data from a source, Transform data and Load (save/store) data into a sink:
Extract → Transform → Load
 - ETL has many operations to deal with the semantics/syntax of data and the business of data
- ELT: Extract → Load → Transform
 - Data transformation done after, within the (target) platform
- **Modern ingestion**
 - data transformation together with ingestion tasks within a complex pipeline
 - both **Transform → Load** and **Load → Transform** designs

Performance, correctness and quality assurance

A!

Coordination of tasks in ingestion pipelines



- Big data ingestion involves
 - many tasks
 - multiple tenants/users
 - ad-hoc, on-demand vs scheduled task pipelines
 - data movement in a single vs cross data centers
- Complex coordination techniques
 - tasks are not in the same machine (executors)
 - → data exchange among tasks (data dependencies)
 - different ways to implement the coordination

When should the
ingestion be executed?
based on what?

Modern ingestion pipelines: common tasks and requirements

Tasks

- Common tasks
 - data extraction
 - change data capture
 - data wrangling/transformation
 - data storing
 - lineage and observability for quality assurance/governance (quality check)
 - backfilling
- Consumer/user defined tasks vs platform tasks
- Other supports within tasks
 - compression, encryption, end-to-end security

Differences: batch vs near real time (streaming) ingestions

Extensible, composable tasks as plugins

- Basic tasks for big data ingestion can be (re)used in different cases
- Support end-user goals
 - enables the user to do many tasks through configurations and extensions
- Enable pluggable approaches is important



Data access and extraction

- Data Access
 - obtaining/copy data from sources
 - including change data capture (CDC)
 - often built based on common protocols and APIs
 - connector library: strongly related to data storage/database/datalake sink/source)
 - runtime connection management: maintain list of connections created from connectors
 - reusability is important!
- Encryption, masking/anonymization
 - might need to be done when accessing and extracting data
 - also during transfers of data
 - data security requirements, personally identifiable information

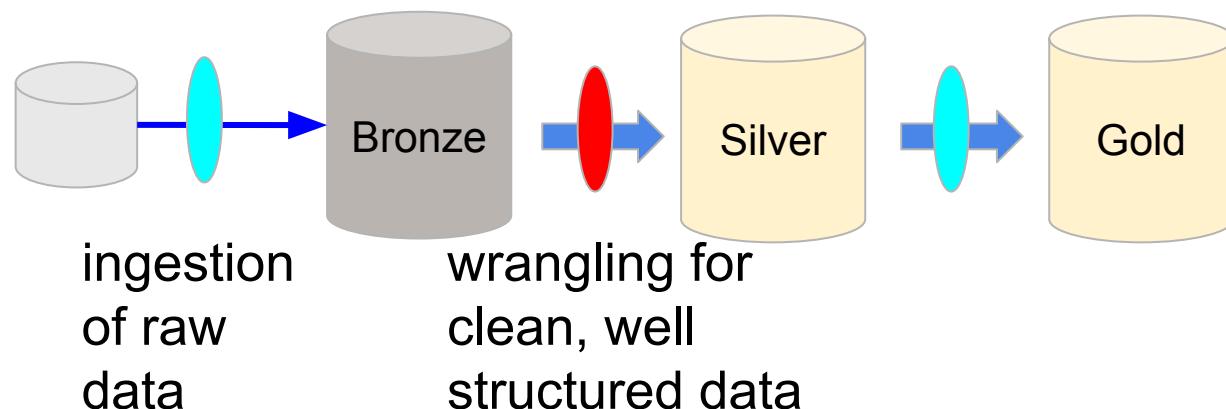
Change data capture (CDC)

It is important to capture new data, changes at immediately (as soon as possible) for continuous analysis to support decision making

- The principles:
 - capture and ingest only **new data** by listening data changes
 - “**new**”: application-specific, e.g., based on time, value, and version.
 - leverage many features of databases (update, query, insert operations), data stream offsets and status notification (e.g., the availability of new files)
 - customize detection mechanism
- Implementation in different tools
 - e.g., Redhat Debezium, Hudi DeltaStreamer, Kafka Connect

Data wrangling

- Convert/transform data from one form to another
 - cleansing, filtering, merging, enriching, inferring, and reshaping data
- Require access to the data content!
- Key design choices
 - when and where: during the ingestion vs after the ingestion
 - by whom: which features will be provided by a platform provider?



A!

Examples

**Write your own
code with
Pandas/Dask and
Dataframe?**



**Automatically
generate code for
wrangling, e.g.
using GenAI?**

```
Alarms={}
with open(sys.argv[1], 'rb') as csvfile:
    reader = csv.DictReader(csvfile)
    for row in reader:
        try:
            #print row['Started']
            alarm_time = datetime.strptime(row['Started'], '%d.%m.%Y %H:%M:%S')
            #diff =start_time - alarm_time
            #print "different time is ",diff
            if alarm_time >=start_time:
                #print(row['RNW Object Name'], row['Severity'])
                typeOfAlarm = 0
                cleanSeverity = re.sub('\W+', '',row['Severity'])
                if (cleanSeverity in mobifone.AlarmSeverity.keys()):
                    typeOfAlarm = mobifone.AlarmSeverity[cleanSeverity]
                #print ("Type of Alarm: ",typeOfAlarm)

                if row['RNW Object Name'] in Alarms:
                    #print "Again"
                    severies =Alarms[row['RNW Object Name']];
                    severies[typeOfAlarm]=serveries[typeOfAlarm]+1
                else:
                    severies =[row['RNW Object Name'],0,0,0,0,0,0]
                    severies[typeOfAlarm]=serveries[typeOfAlarm]+1
                    Alarms[row['RNW Object Name']] = severies;
            except:
                print "Entry has some problem"
                print row
            #timestamp =long(row['TIME'])
            #times.append(datetime.datetime.fromtimestamp(timestamp/1000))
            #times.append(long(row['TIME']))
            #signals.append(float(row['GSM_SIGNAL_STRENGTH']))
        datafram =pd.DataFrame(Alarms,index=mobifone.AlarmSeverityIndex).transpose()
        alarmdata =datafram.as_matrix();
        #TODO print Alarms to fine
        #only for debugging
        print datafram
        datafram.to_csv(outputFile, index=False)
```

A!

Complex data transformation and processing

- More complex data processing
 - extract only important data
 - feature engineering
 - enrich data on the fly with external sources

Example: extract vectors from images

```
class TowheeExtractor(BaseExtractor):  
    def __init__(self):  
        self.towhee_feature_extractor = (  
            pipe.input('file_path')  
                .map('file_path', 'img', ops.image_decode.cv2_rgb())  
                .map('img', 'embedding', ops.image_embedding.timm(model_name='resnet50'))  
                .map('embedding', 'embedding', ops.towhee.np_normalize())  
                .output('embedding')  
        )  
  
    def get_model_name(self):  
        return "resnet50"  
  
    def feature_extractor(self, image_file):  
        ##currently only a single figure so we have to get the first element  
        embedding=self.towhee_feature_extractor(image_file).get()[0]  
        return embedding
```

Complex code & libraries hide low level distributed/parallel tasks

- Complex distributed and parallel tasks for ingestion
 - Complex coordination
 - Underlying, internal task models:
 - MapReduce model
 - embarrassingly parallel model
 - full directed acyclic graph (DAG) task model

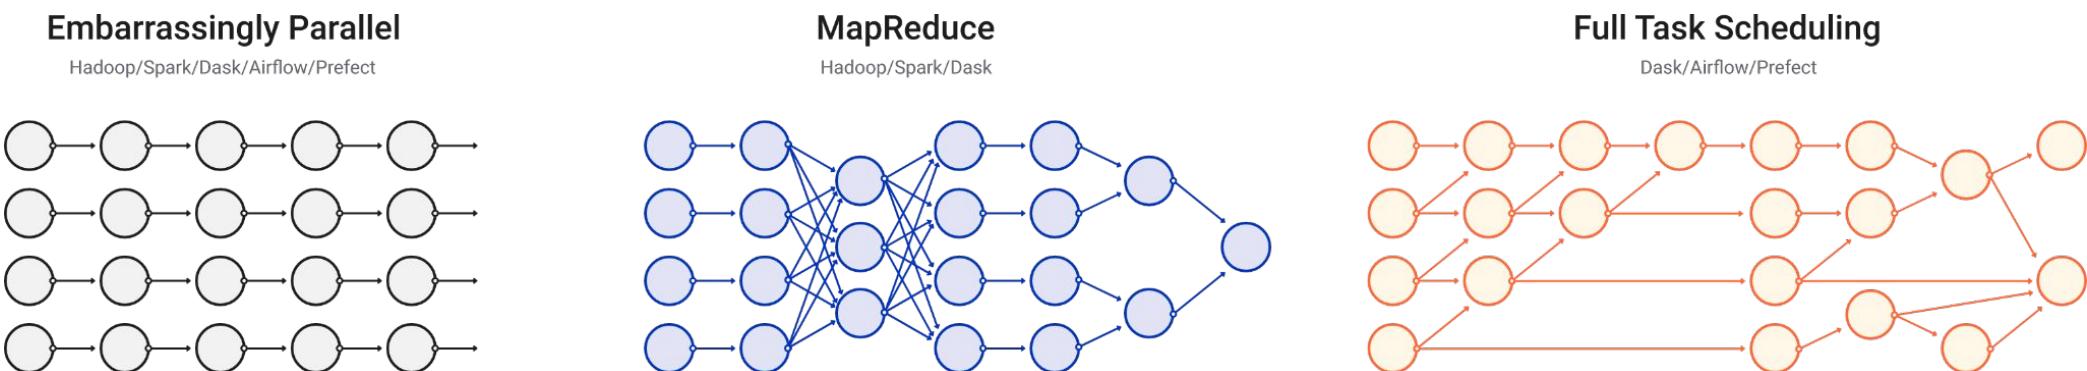
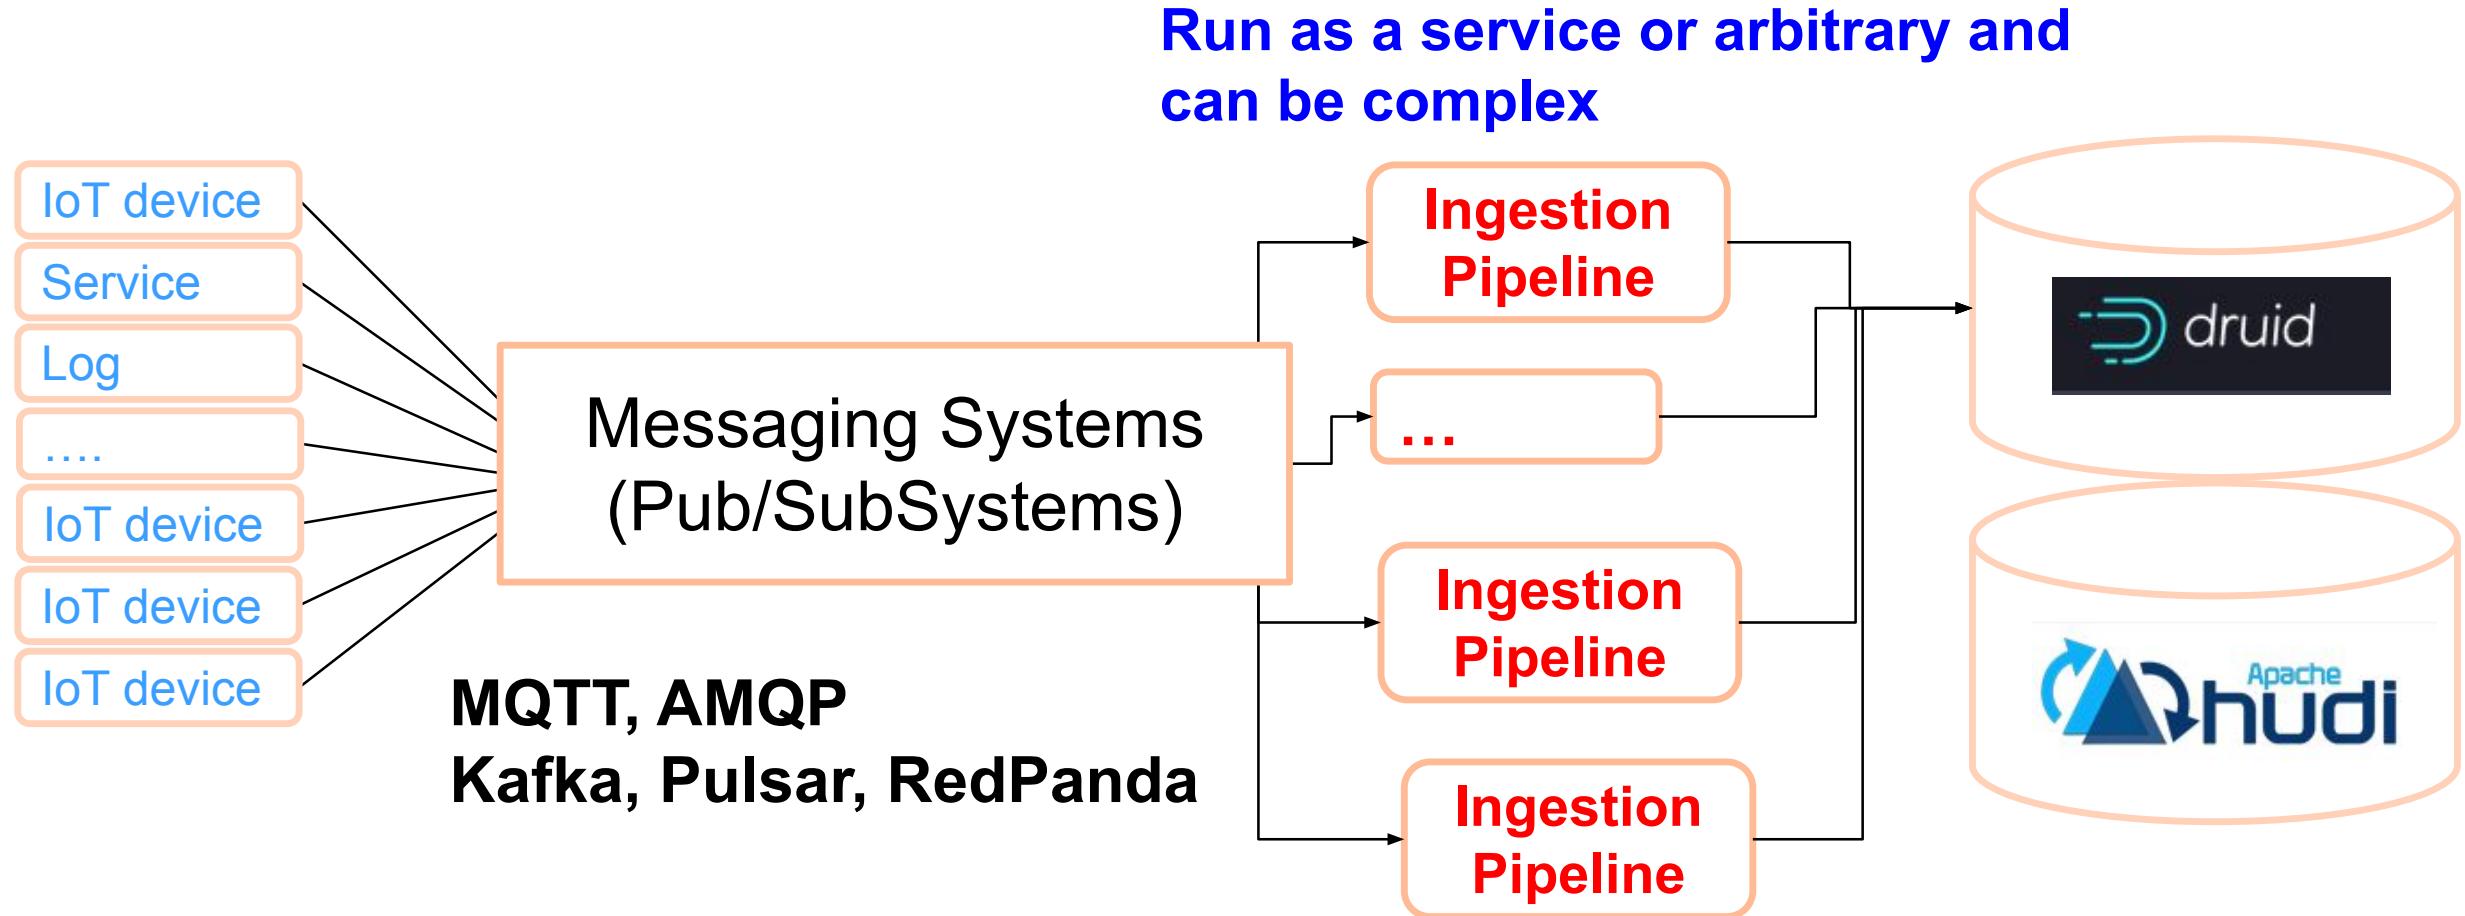


Figure source: <https://docs.dask.org/en/stable/graphs.html>

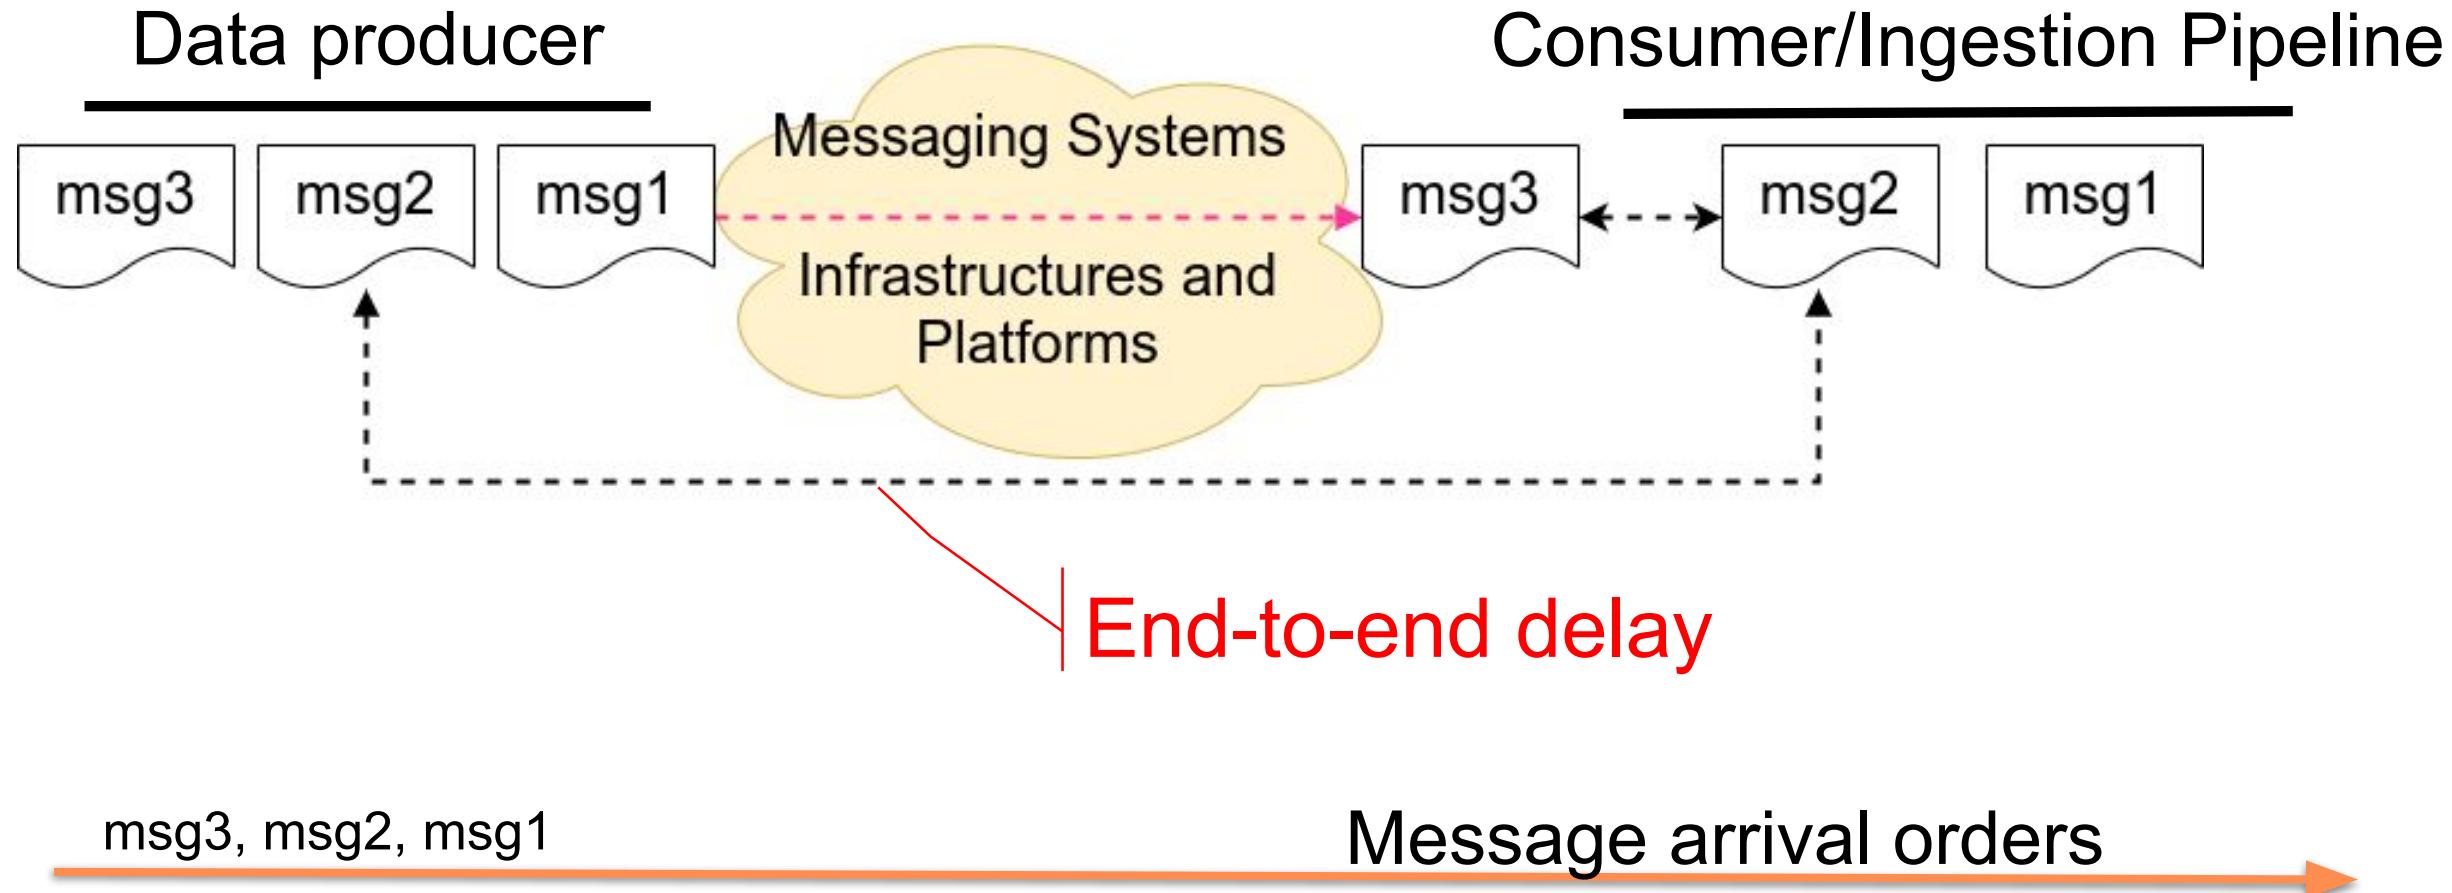
A!

Near real time streaming ingestion



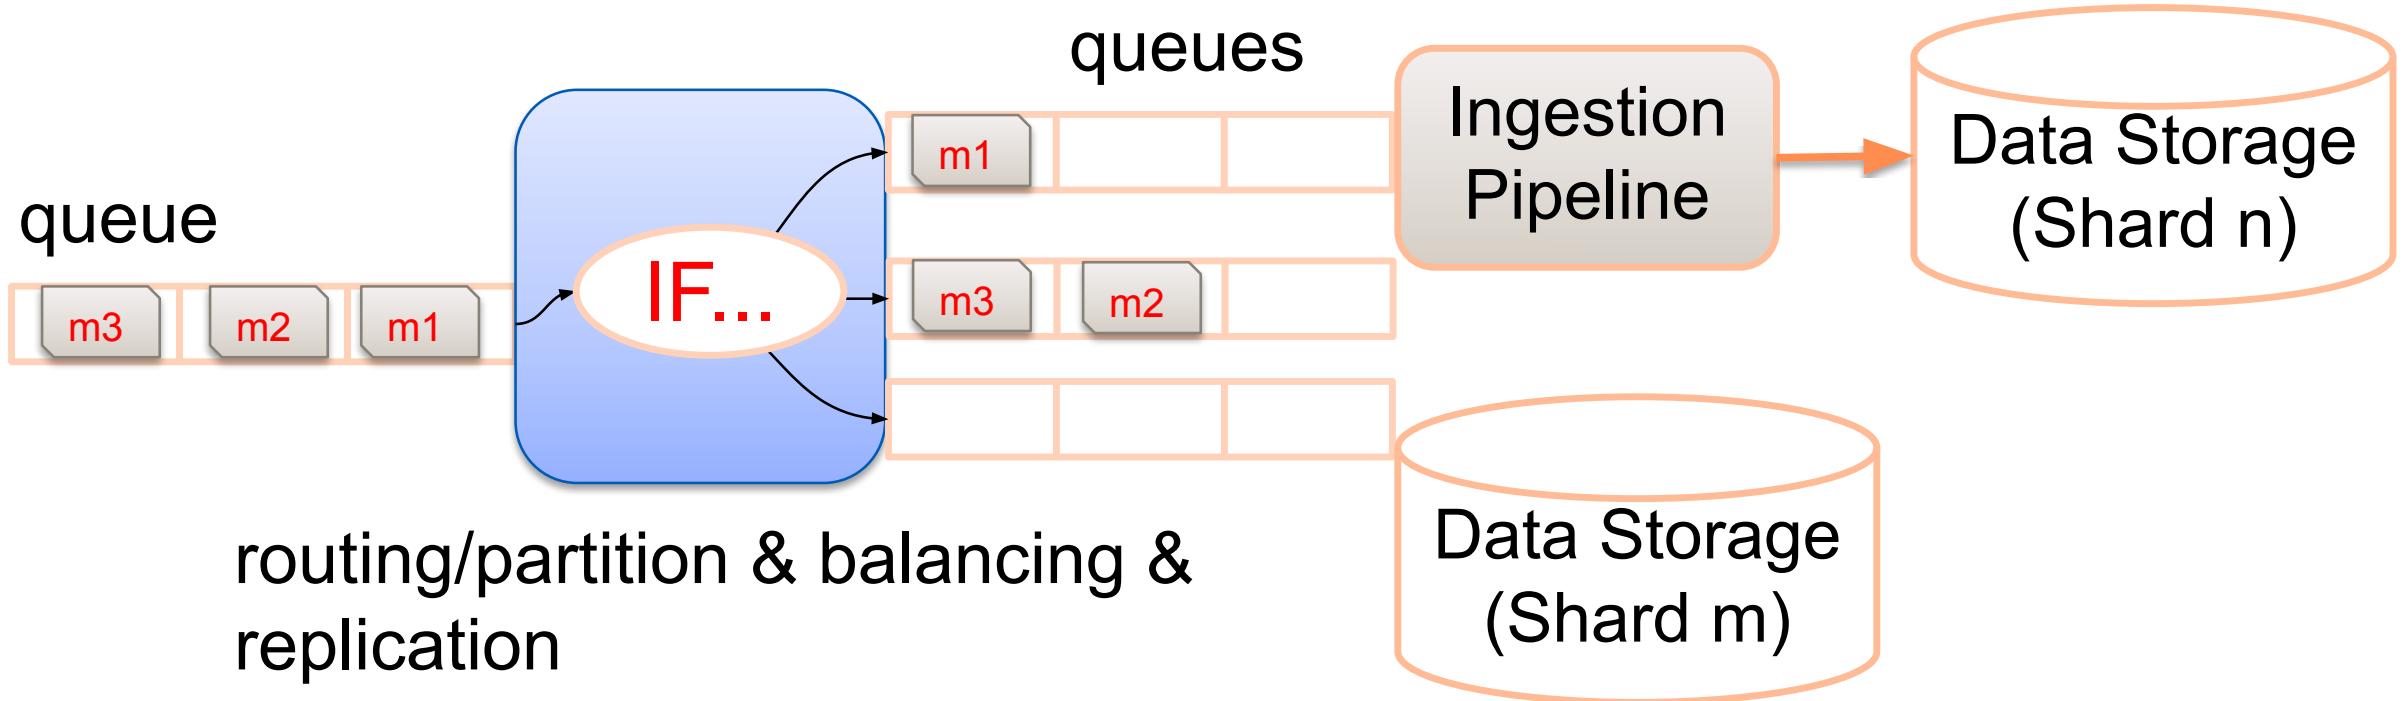
A!

Key issues in streaming data ingestion



A!

Split (pub/sub) and partition with ingestion



Some key issues for ingestion of streaming data

- Late data, data out of order
- Data loss or data duplication
 - e.g, at most once, at least once or exactly once
- Back pressure and data retention
 - for individual components or the whole pipelines
- Scalability and elasticity
 - changes in data streams can be unpredictable
- Data quality
 - how to do it with fast processing and minimum overhead

Dealing with diverse data structures

- The data sender/producer and the receiver/consumer are **diverse**
 - implemented with different languages and software technologies
 - need to guarantee the message syntax and semantics
 - performance overhead due to data format conversion
- Solutions: don't assume! agreed in advance
 - agreed in advance ⇒ within the implementation or with a standard
 - know and use tools to deal with **syntax differences**
- Understanding the syntax allows some automatic transformations/quality checks
- But semantics are domain/application-specific

Dealing with diverse data structures: example of interoperability in data transfer: Arvo

Syntax specification
<https://avro.apache.org/>

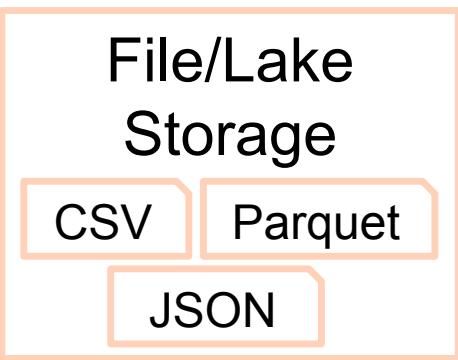


Relevant issues: data compression and security; schema validation and evolution

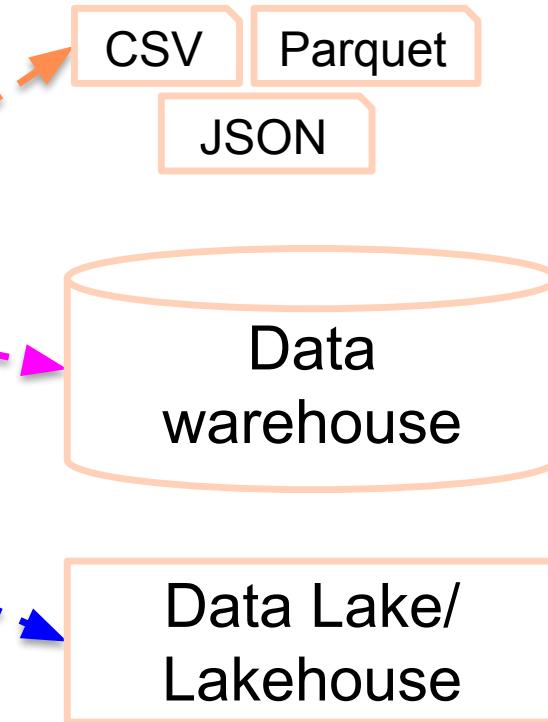
A!

Dealing with diverse data structures: interoperability in data processing in ETL

Data sources and formats



Data sinks and formats



Example: Apache Arrow

in-memory columnar format

rich ecosystems

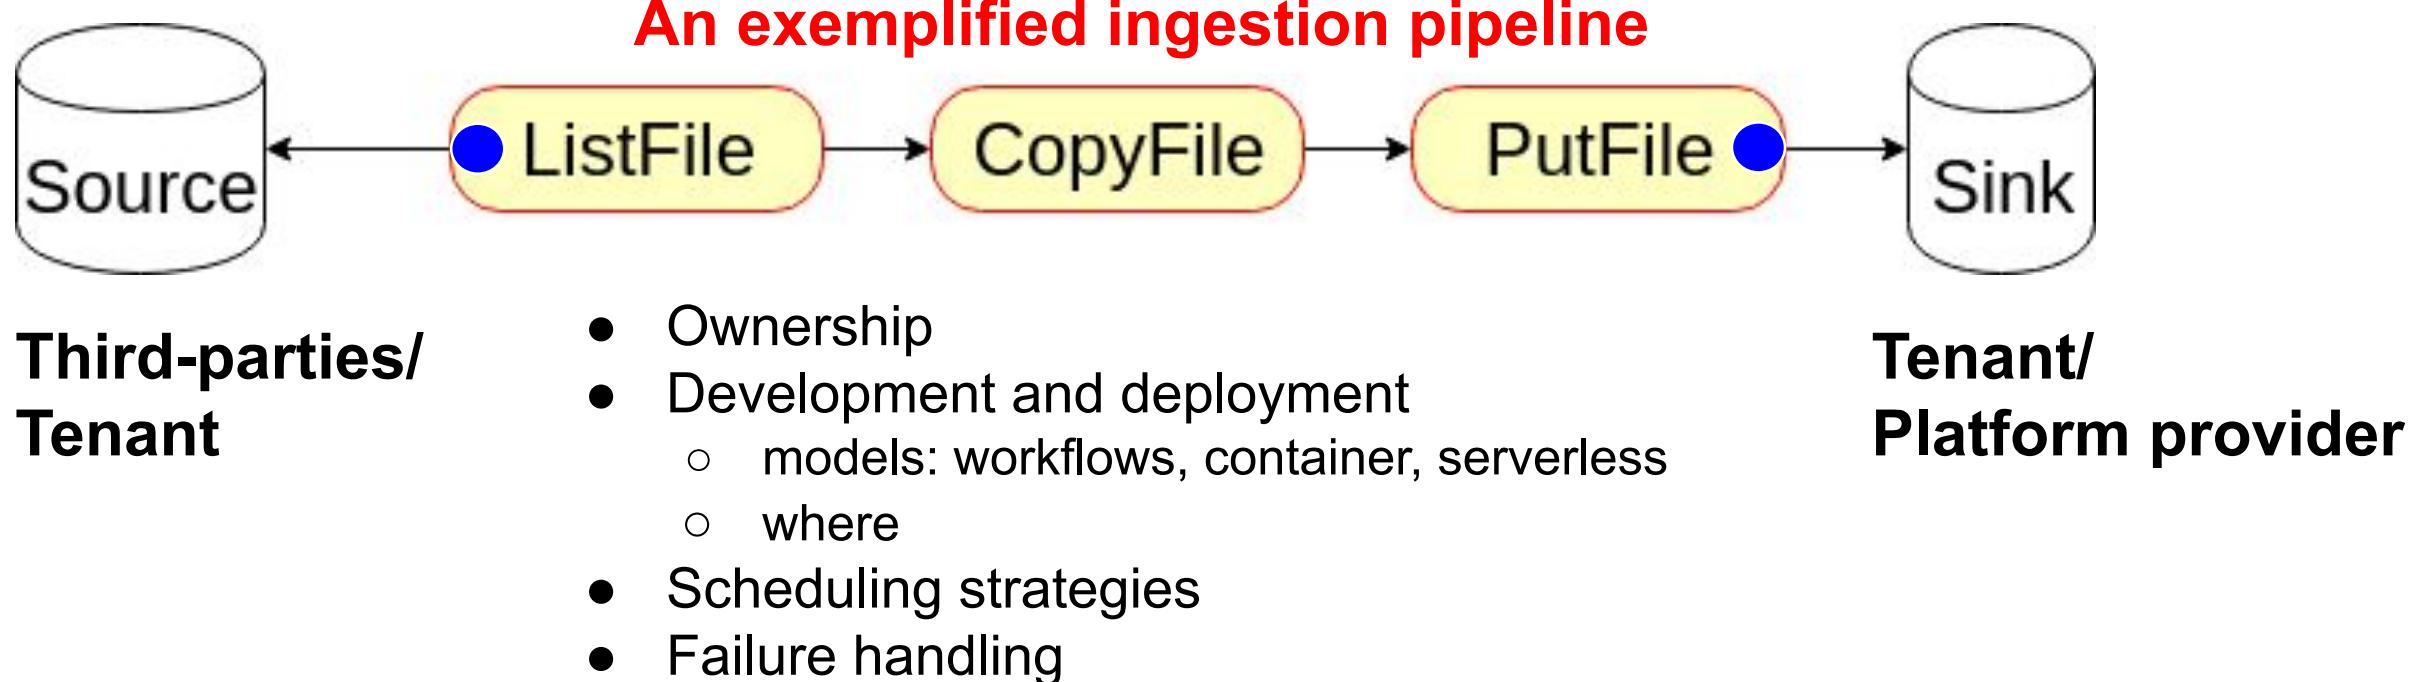
<https://arrow.apache.org/>

A!

Ingestion pipelines/processes: composition, deployment, orchestration, and quality assurance

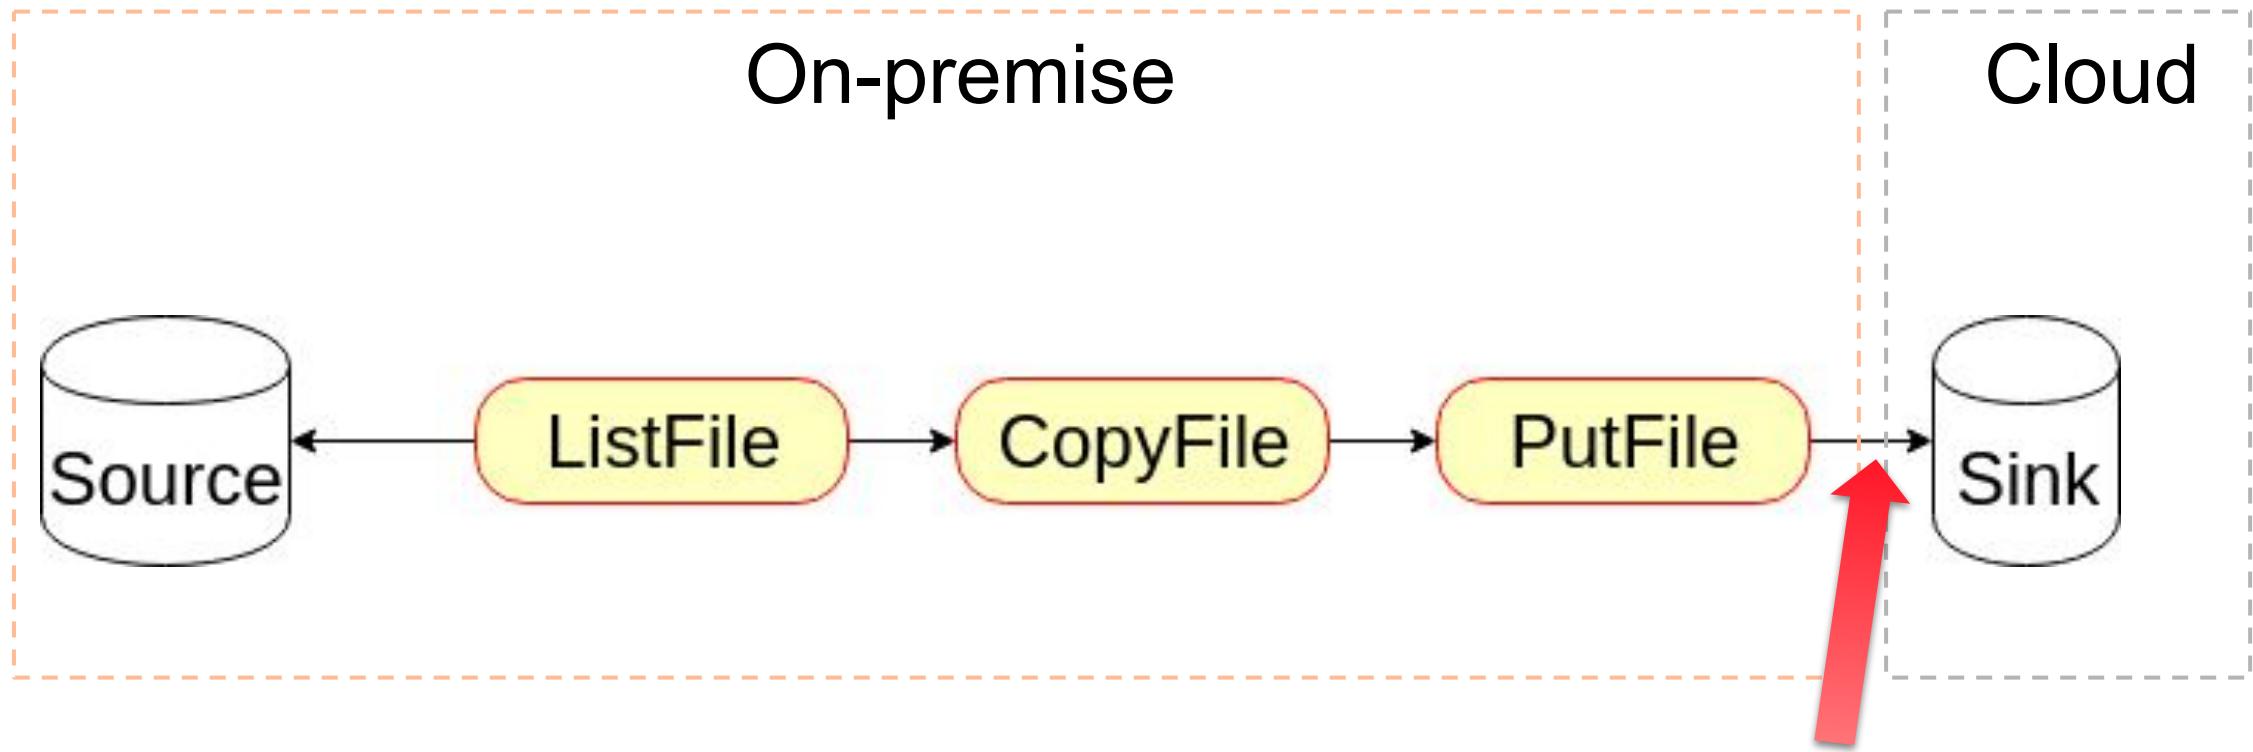
Complex deployment and composition models (1)

Understanding strong dependencies between protocols/APIs, security, performance, connection management, and service-level agreement (SLA)



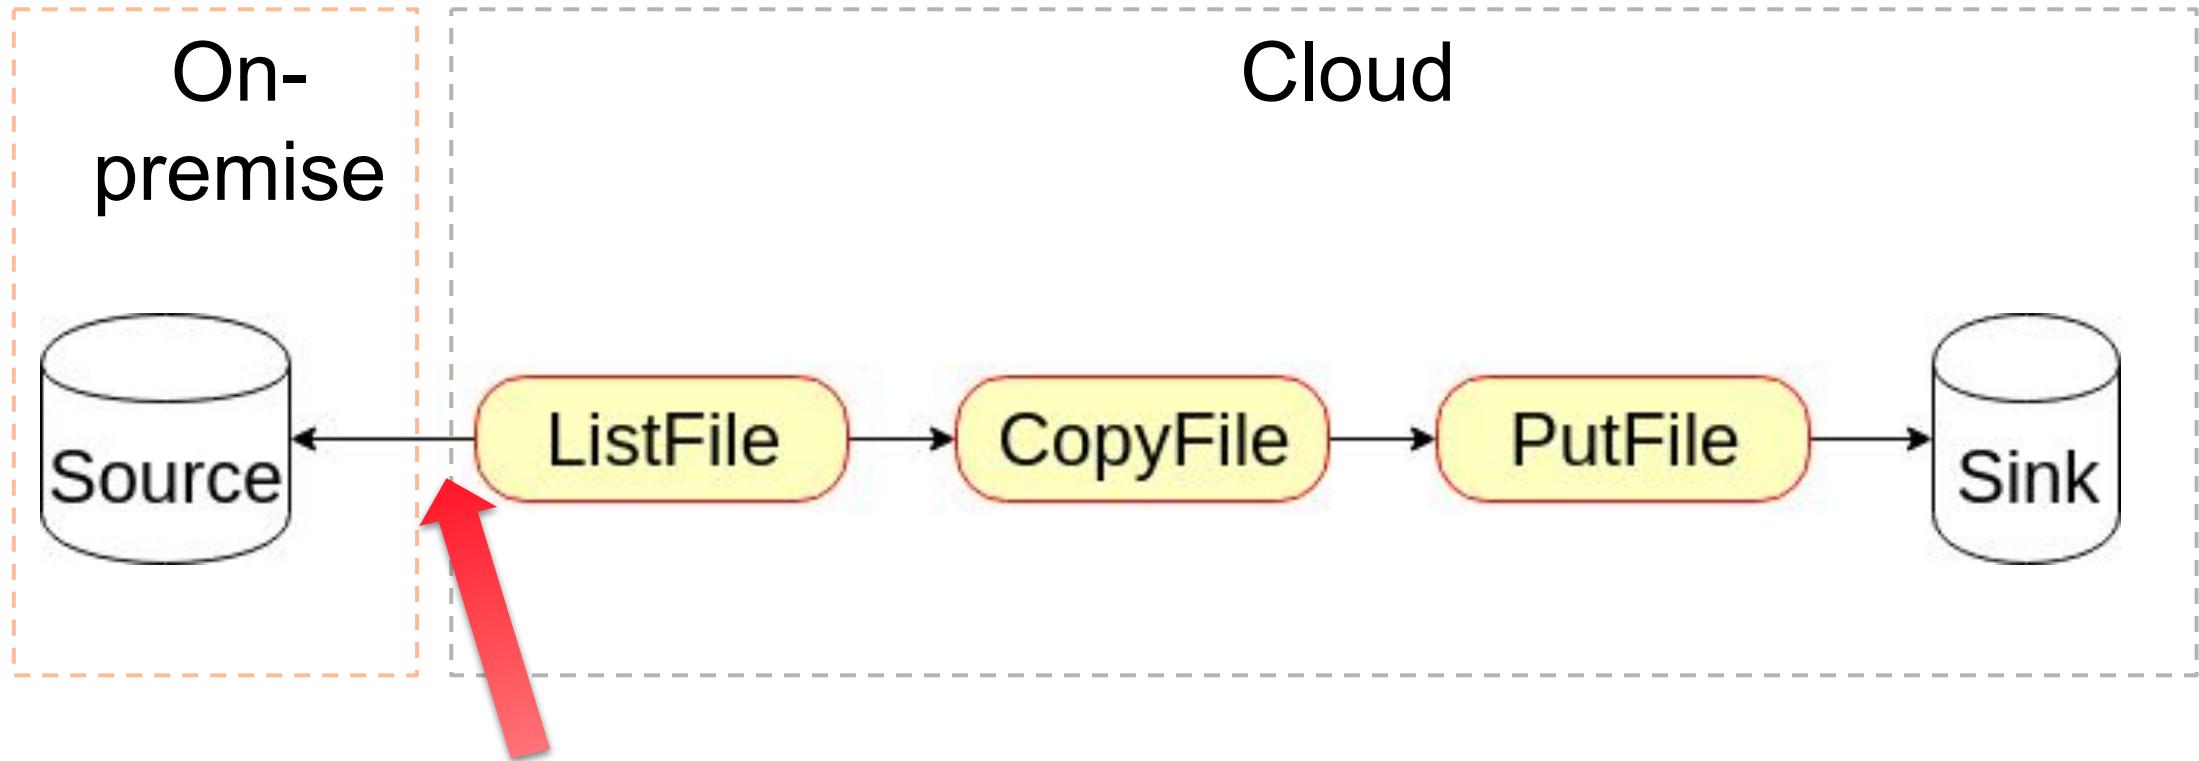
A!

Complex deployment and composition models (2)



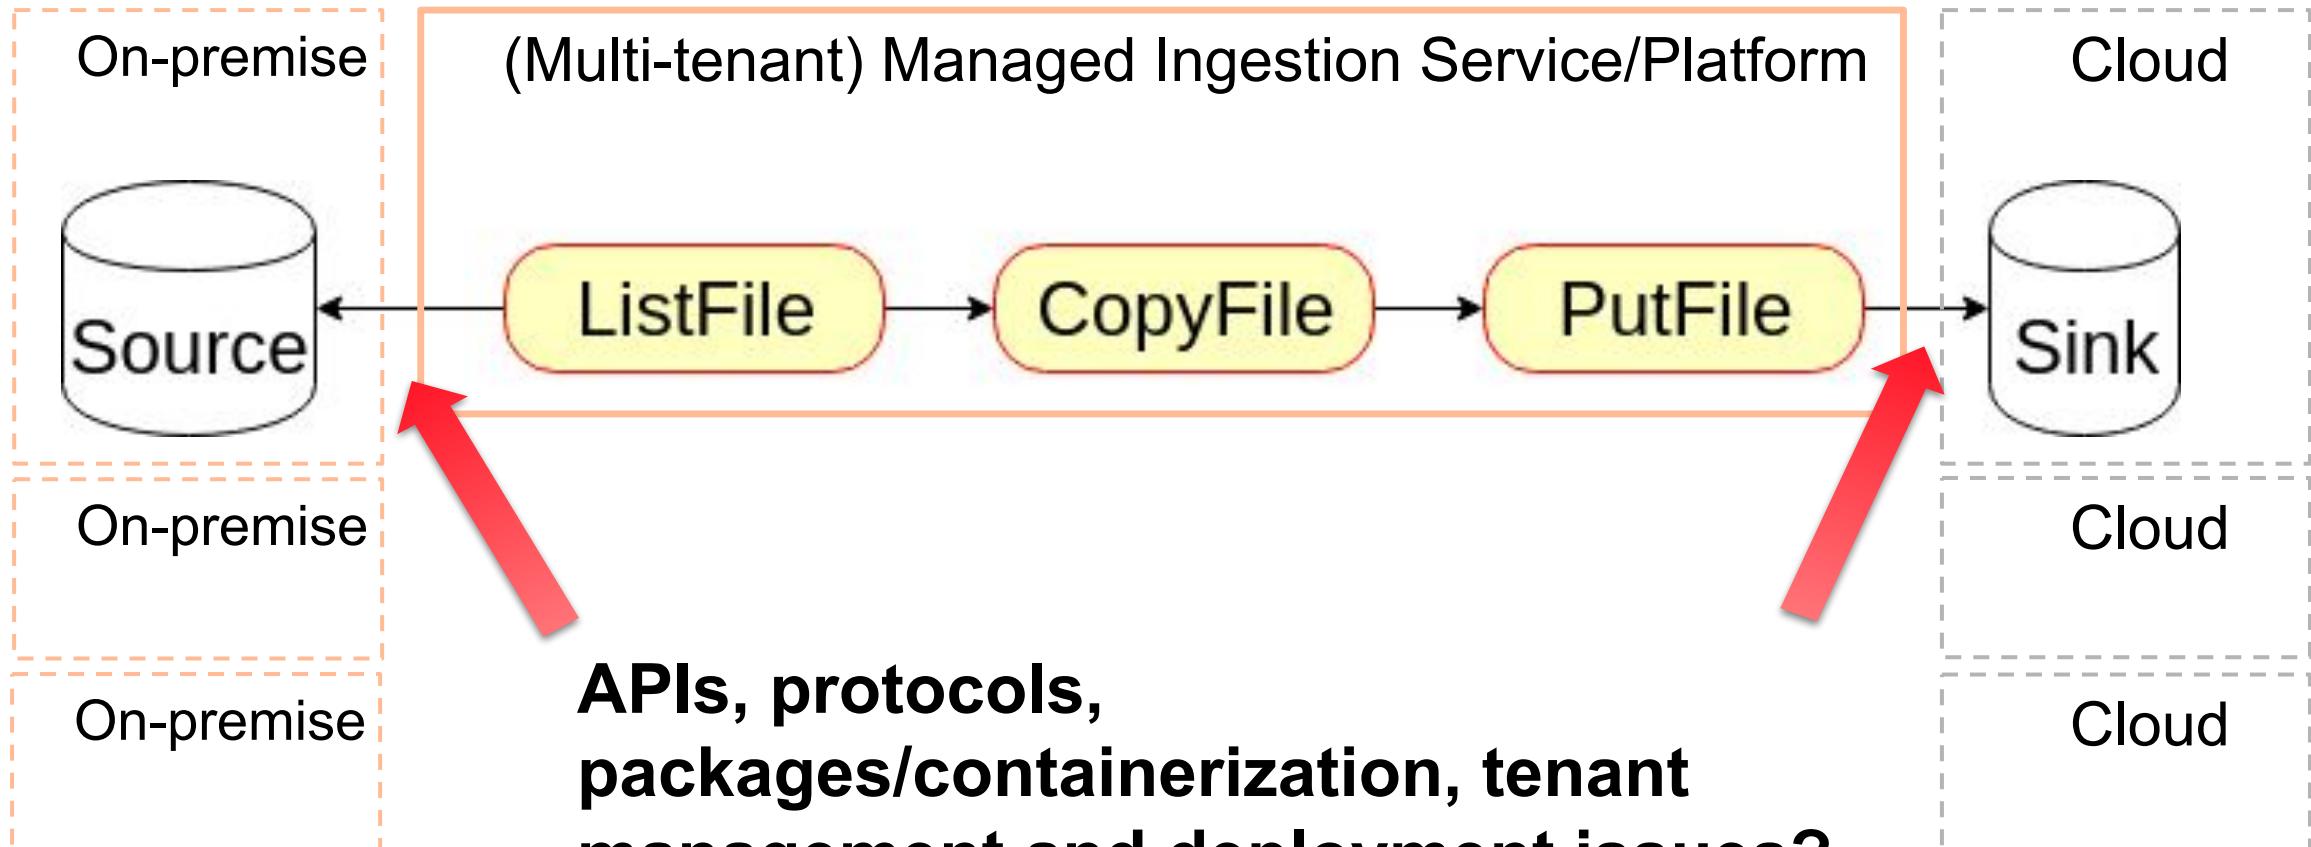
APIs, protocols and deployment issues?

Complex deployment and composition models (3)



APIs, protocols and deployment issues?

Complex deployment and composition models (4)



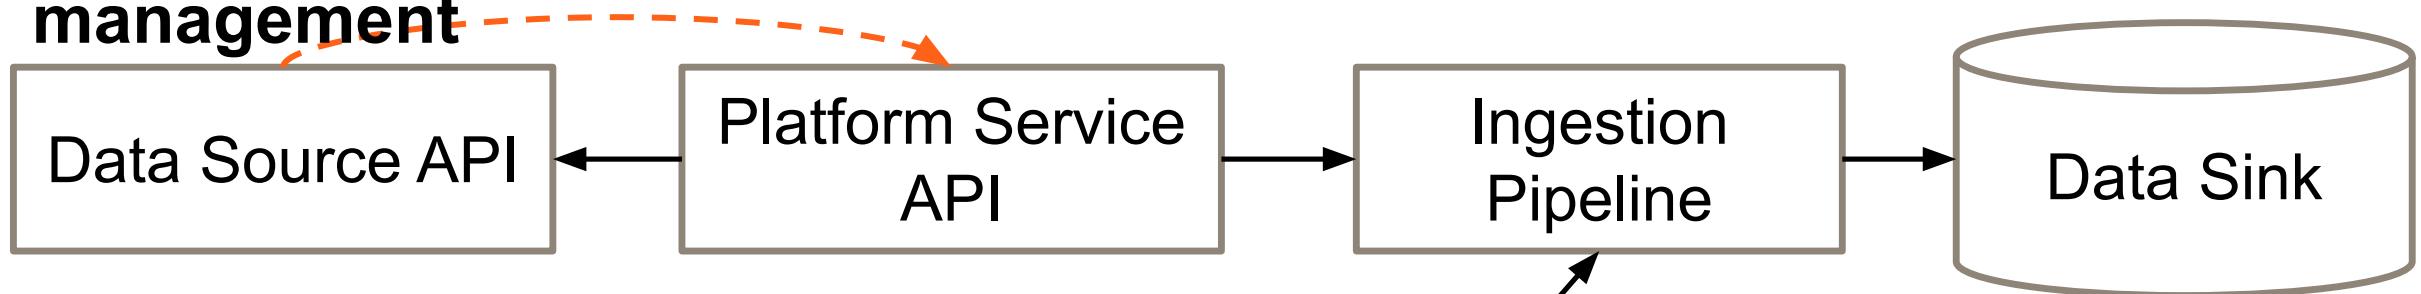
A!

Orchestrating batch ingestion pipelines

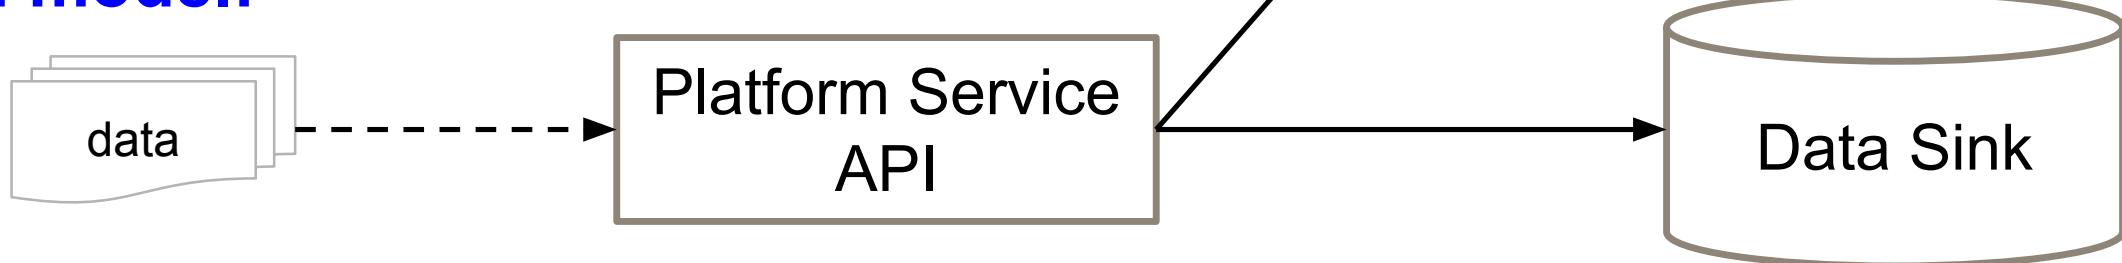
- Data to be ingested is bounded
 - files or messages are finite
- Ingestion architectural styles
 - (1) Direct APIs, (2) reactive pipelines, (3) tasks/workflows
- Incremental ingestion
 - dealing with the same data source but the data in the source has been changed over the time (related to change data capture)
- Parallel and distributed execution
 - use workflows and distributed processing engines

Simple, direct APIs for ingestion

Pull model: register webhook/API ⇒ Connection management



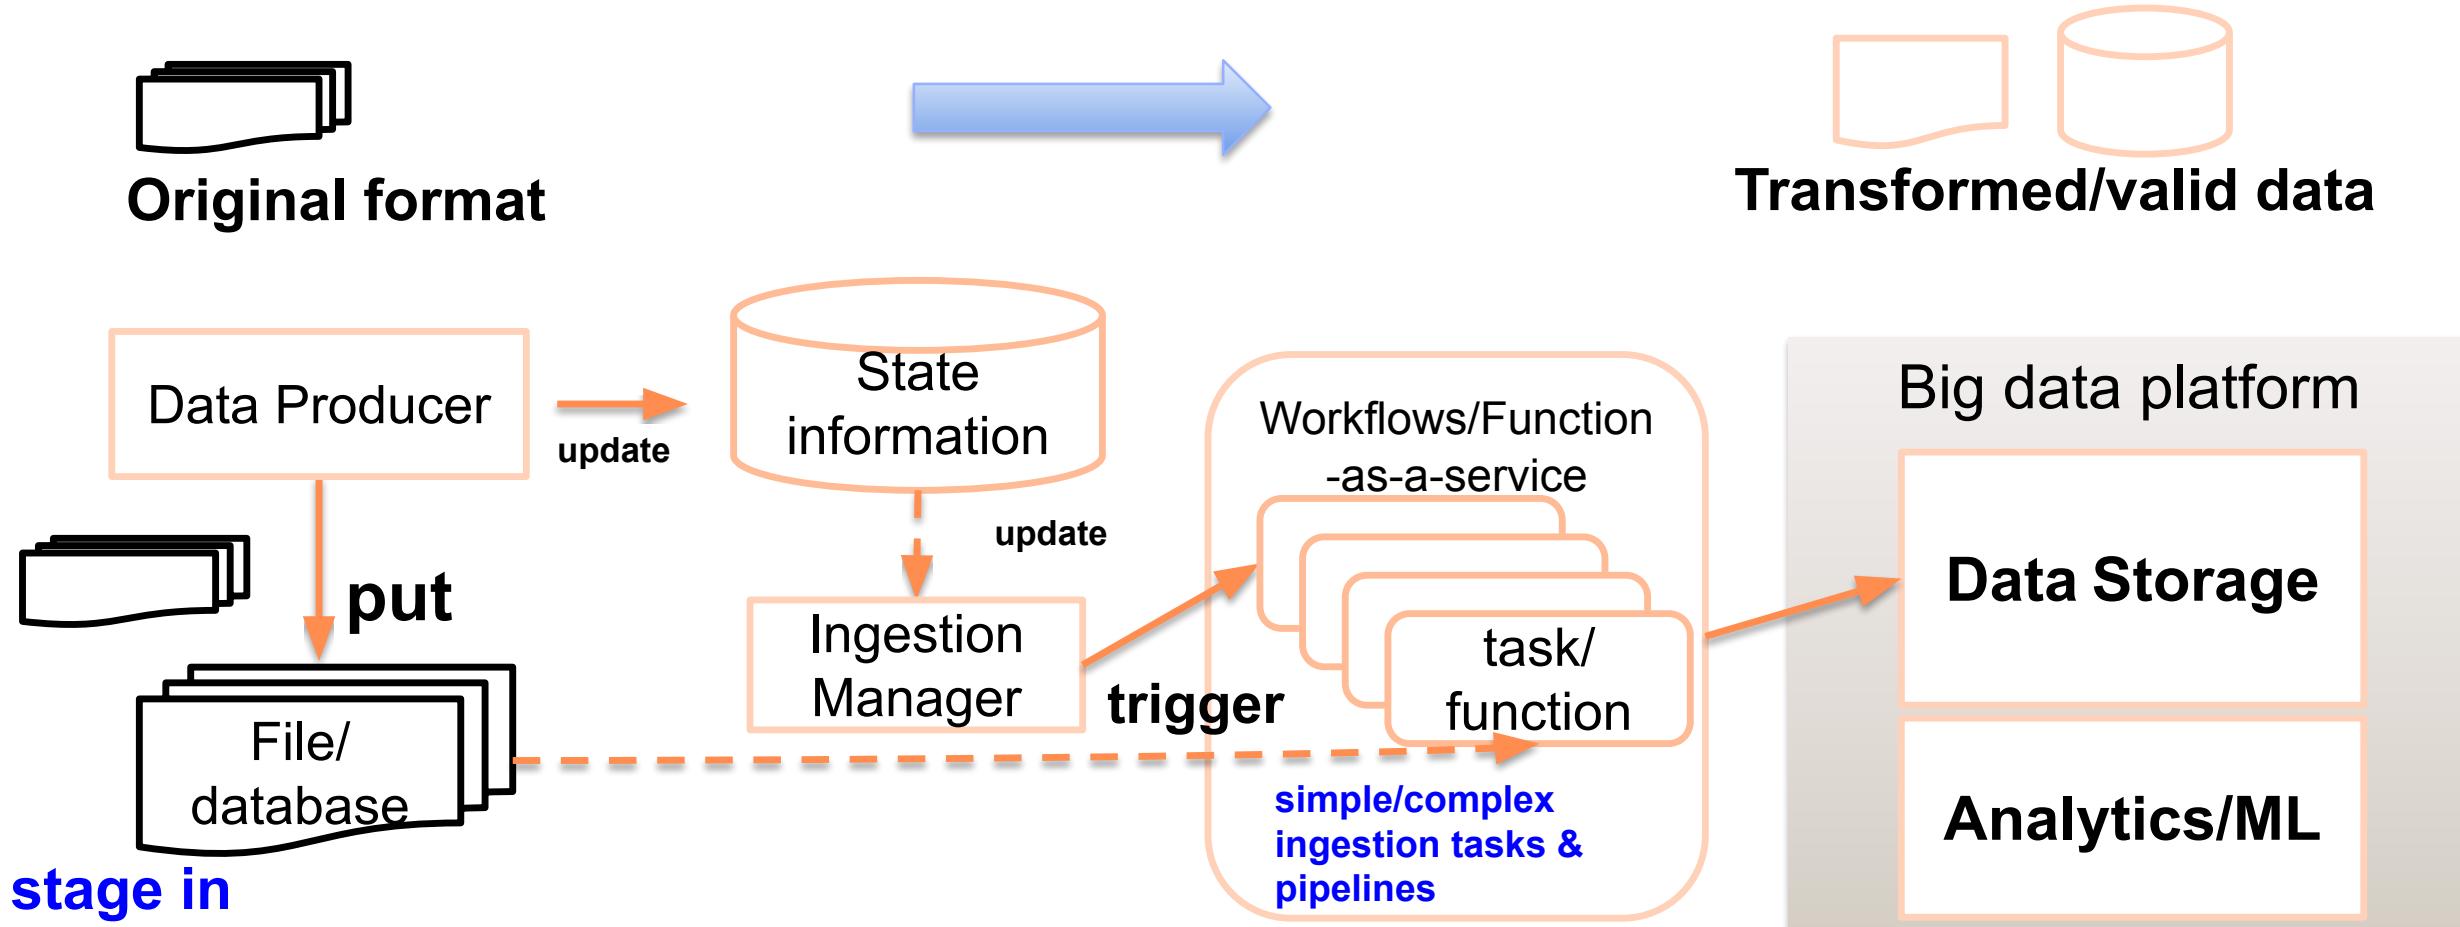
Push model:



Analyze pros and cons in the design

A!

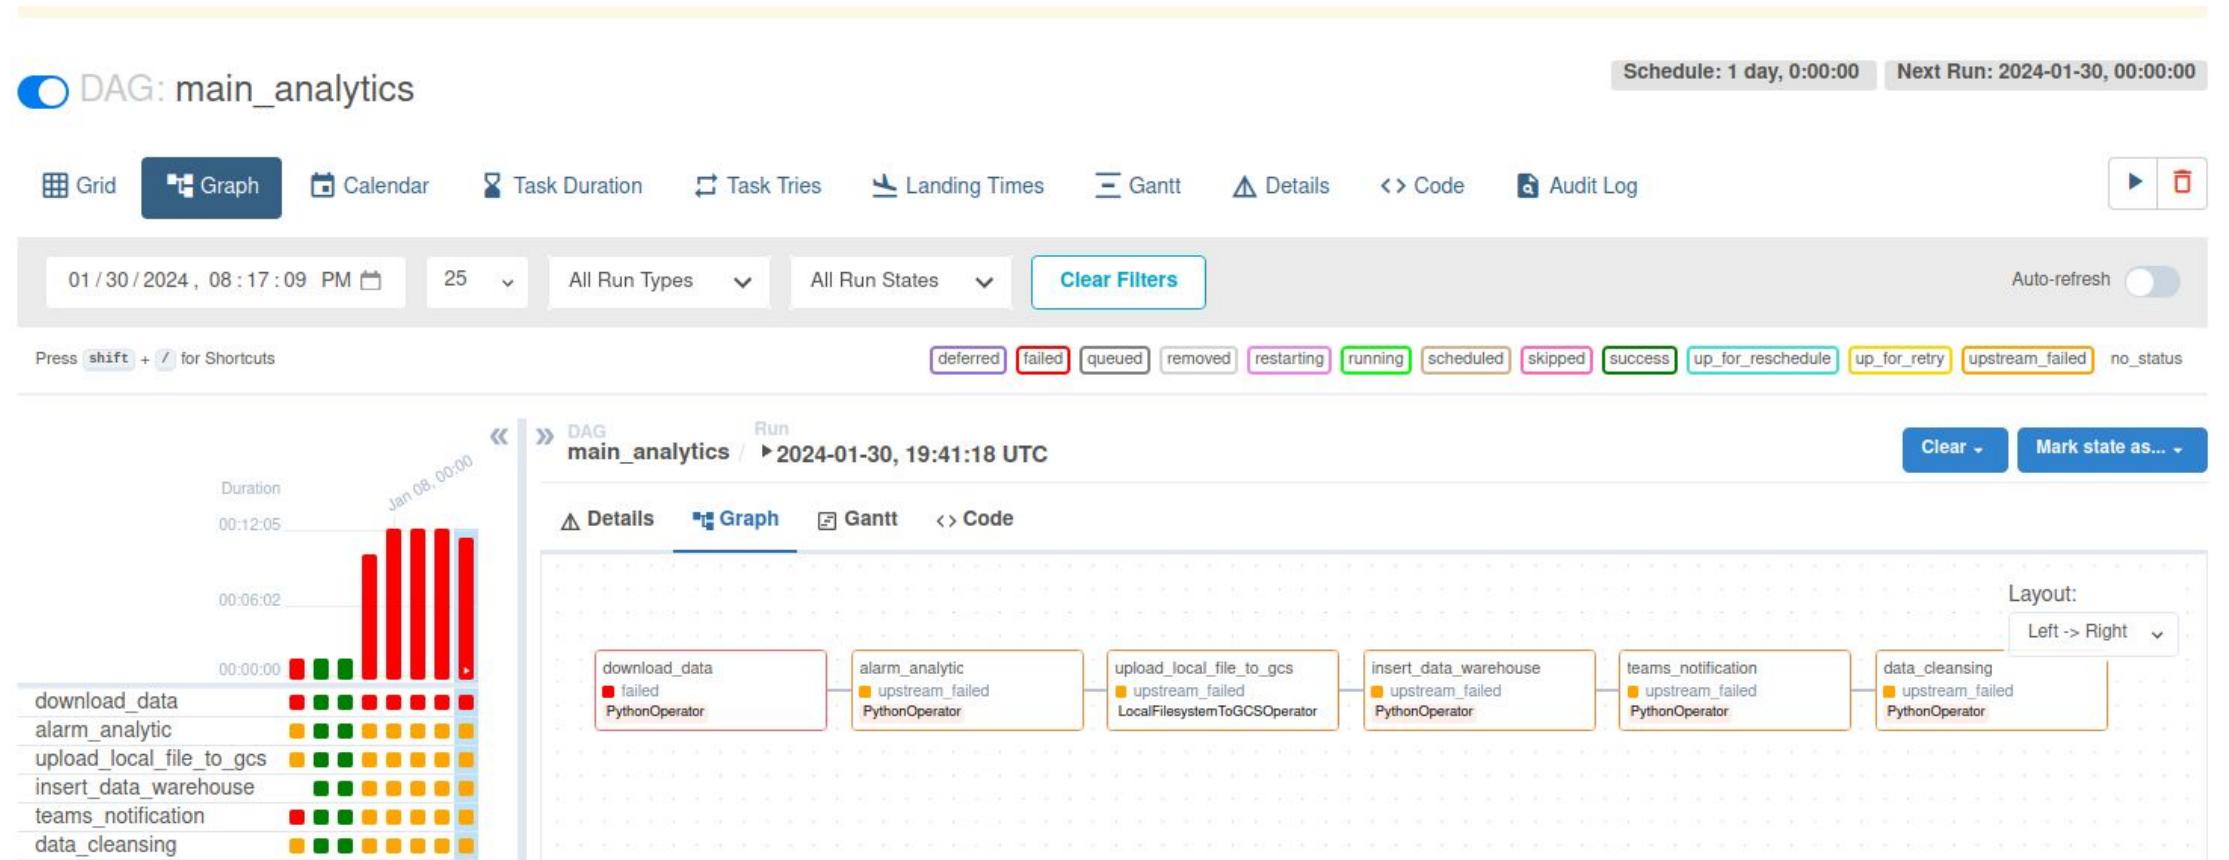
Reactive pipelines with functions/workflows/containers



Ingestion **workflows** orchestration

- Different tasks
 - access and copy, extract, covert, quality check, and write data
 - tasks are connected based on data or control flows
- Workflows
 - a set of connected tasks is executed by an engine
 - tasks can be scheduled and executed in different places
 - flexible designs
- Different tenants have different service level agreements
 - performance, reliability, and cost.

E.g., workflow based on scheduled time, with Apache Airflow

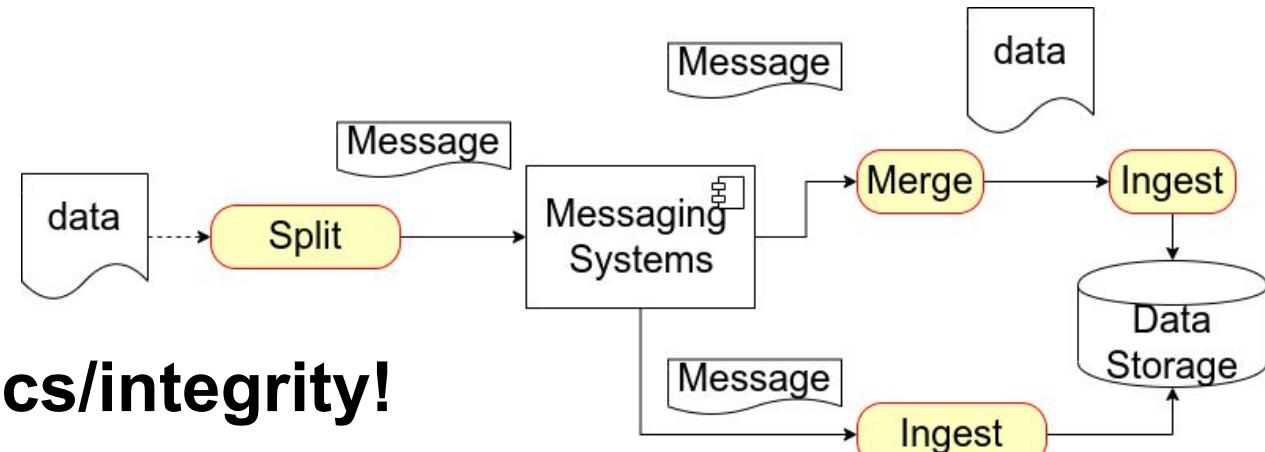


A!

Microbatching in data ingestion

- Microbatching: we mean the strategy to deal with big dataset using batches of data (small chunks)
 - not necessary the same as using batch systems to transfer small data in near real time
- Data is split into different chunks for ingestion
 - using streaming or batch systems to transfer data chunks
 - chunks are ingested into the system, or merged and then ingested

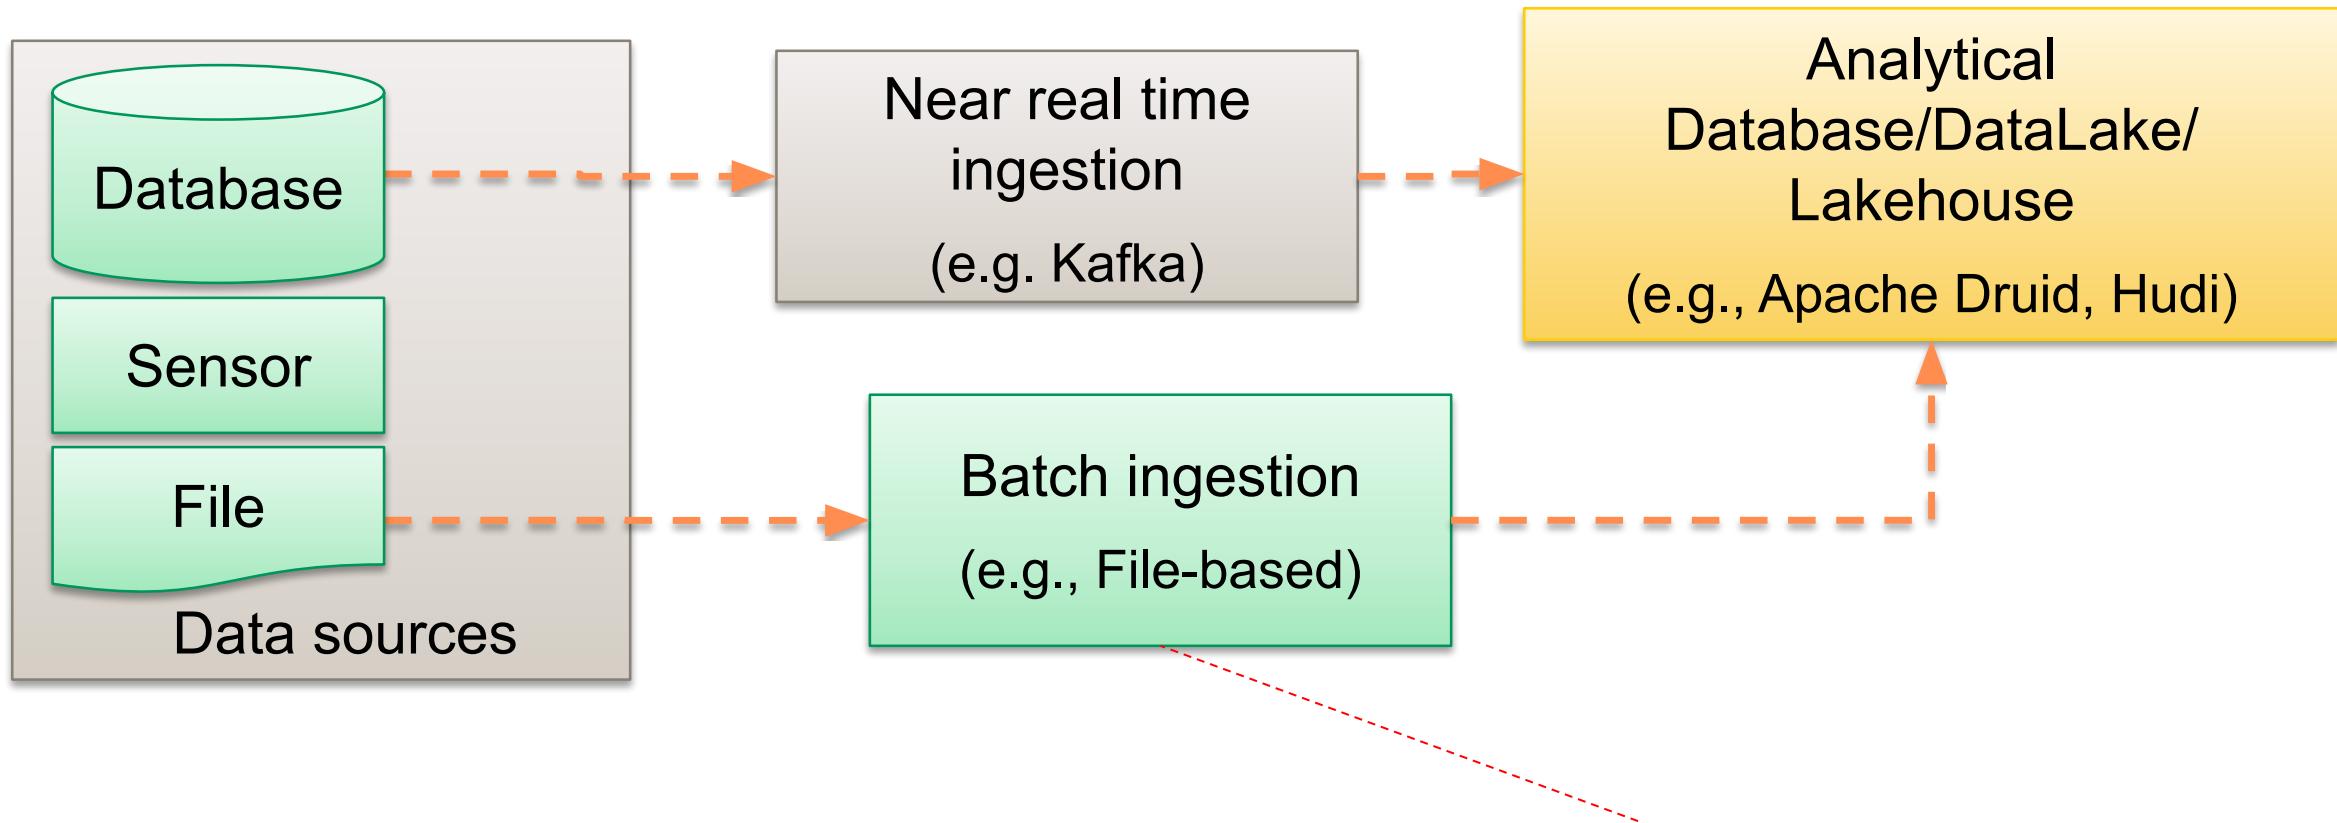
**Example: with
streaming system**



Challenge: the data semantics/integrity!

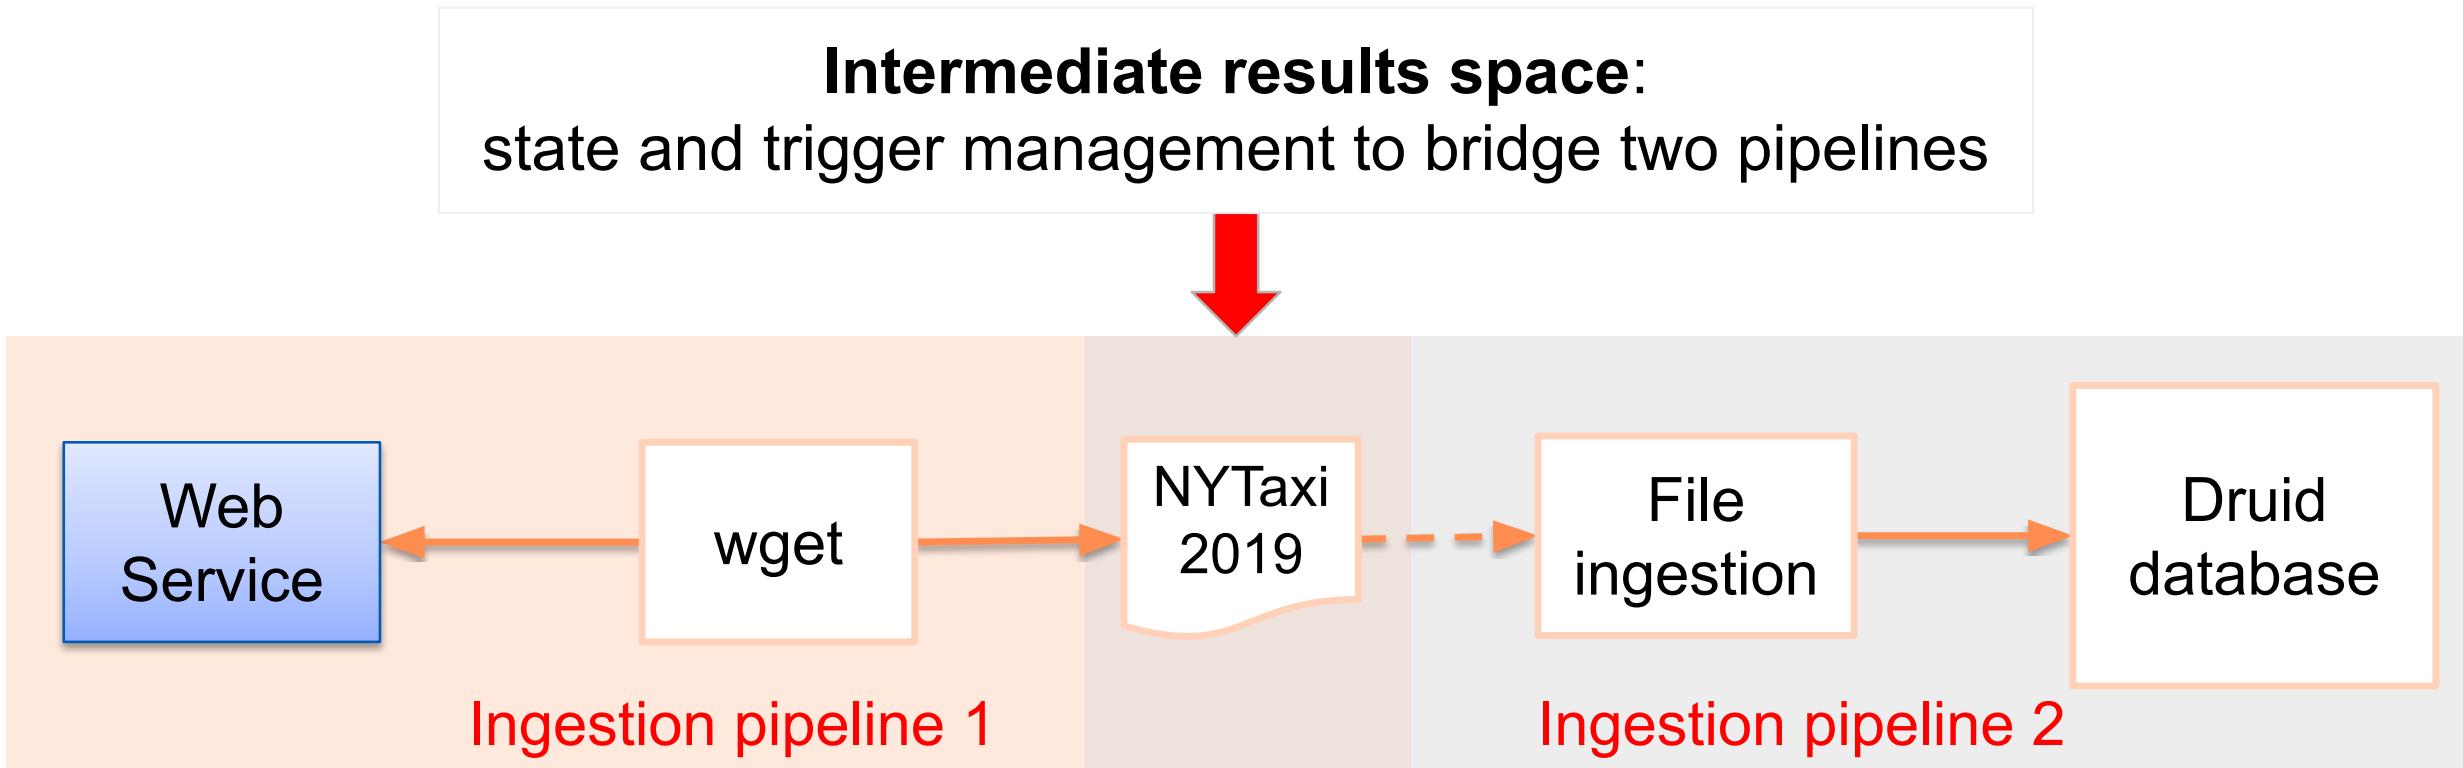
A!

Supporting multiple types of pipelines for **the same data sink**



Programming models, orchestration and scheduling are very diverse

Connecting different ingestion pipelines: bridge designs

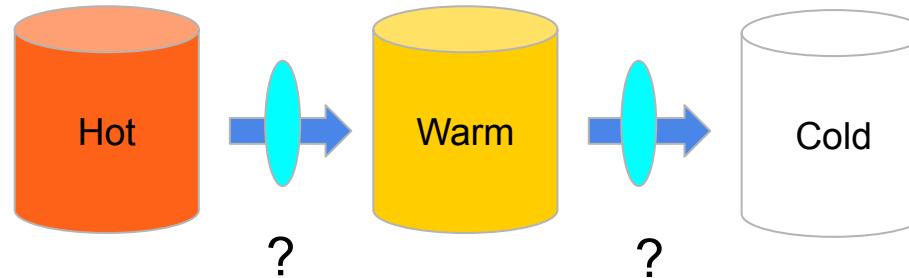


Real-world: a single stack might not be enough, both pipelines and their connections are complex

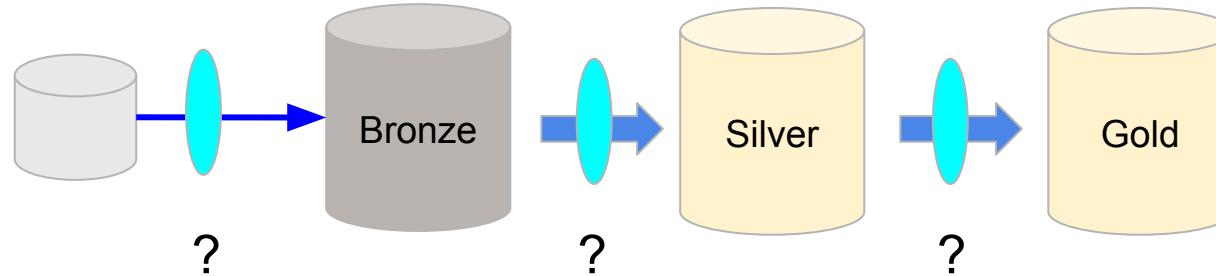
A!

Discussion: streaming or batch ingestion

Access frequency: Hot → Warm → Cold

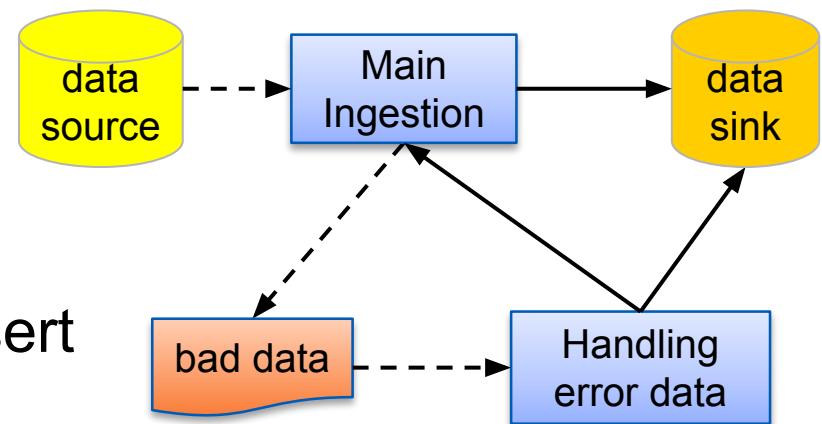


Medallion architecture



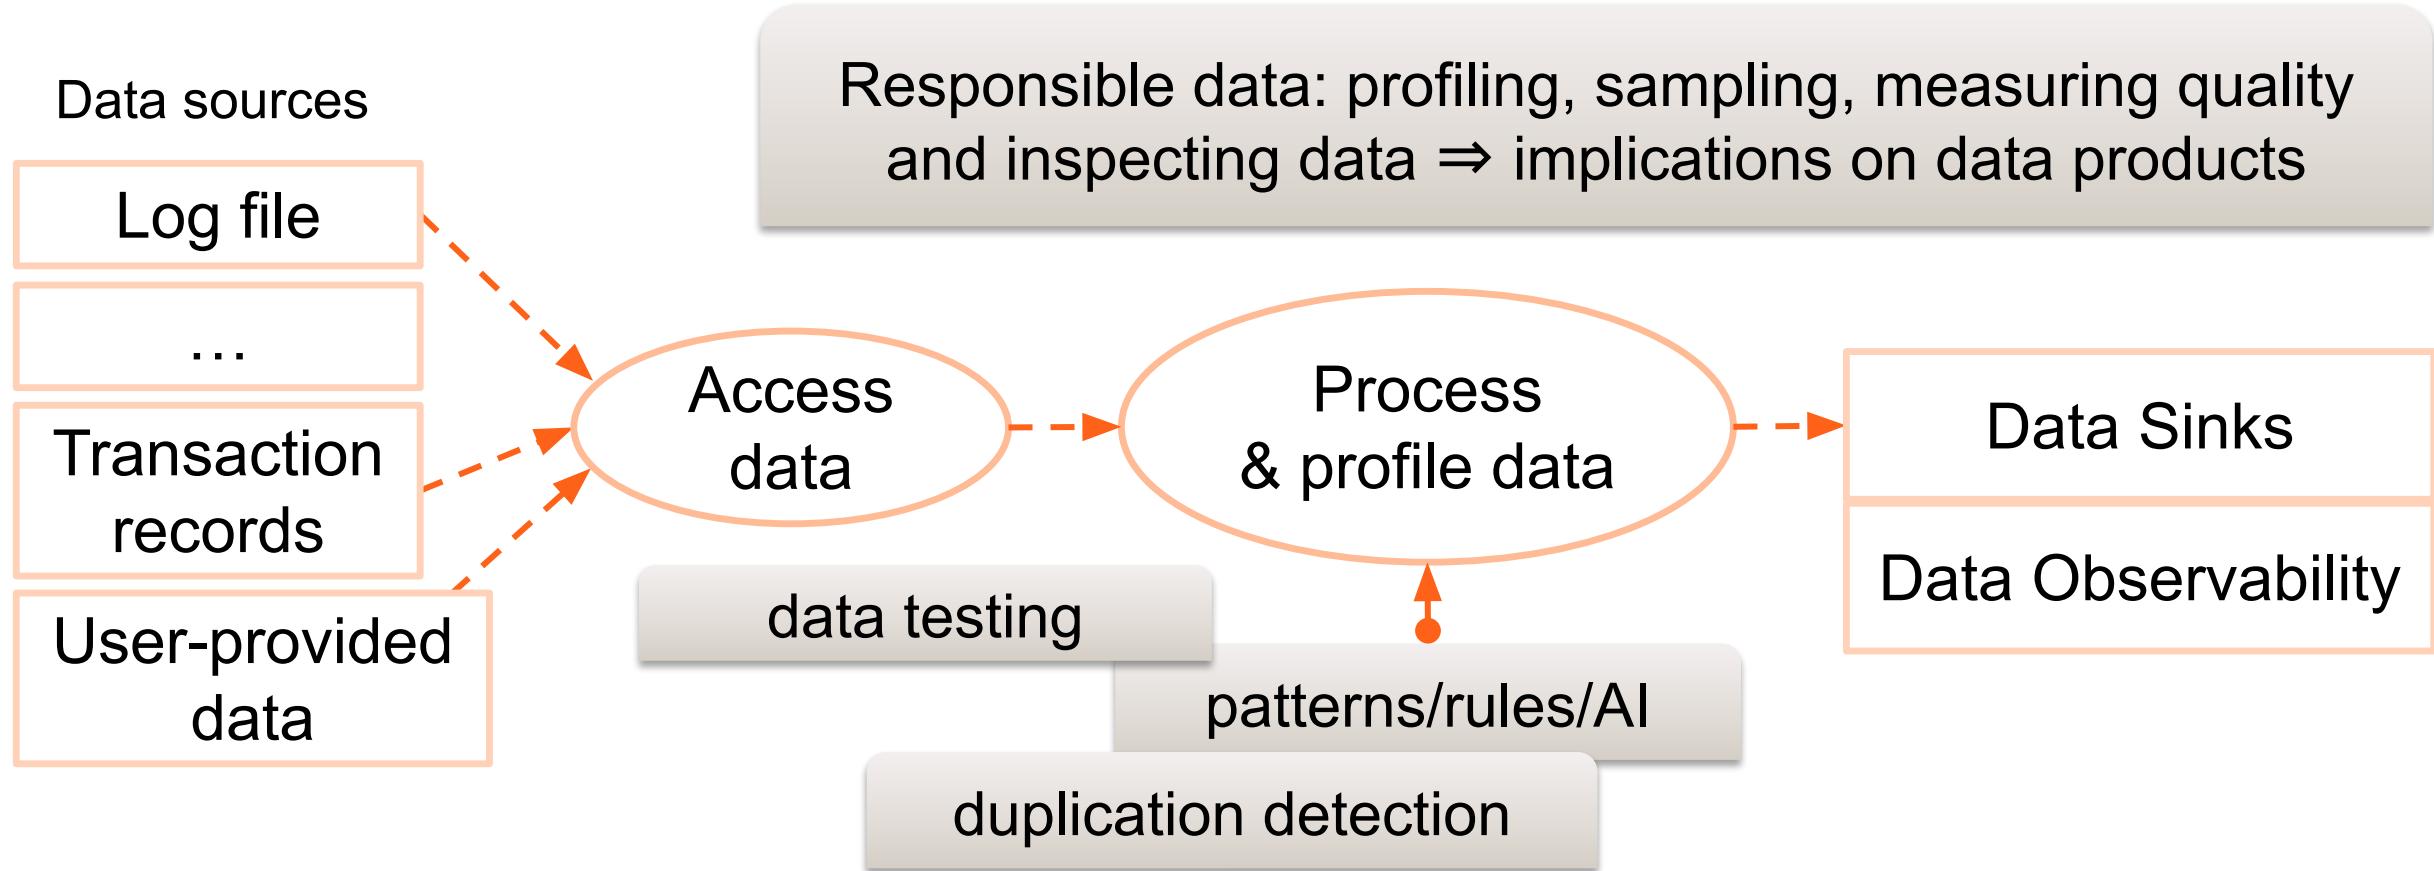
Failure handling

- Data error records
 - abort completely vs ingest qualified records and return errors
 - strategies for duplicated data
 - discard, use last value, etc.
 - acceptance error threshold
- Idempotent design
 - not causing a problem/introducing new data, even if we repeat the same ingestion
- Differences for insert (including append-only) and upsert
- Task failure handling
 - Rollback and retry → costly
 - Avoid big data chunk
 - small data ingestion → performance overhead



A!

Quality control/data regulation assurance



Challenging issues: misinformation, GDPR, data quality, inappropriate content

A!

Data lineage and observability (1)

- FAIR principles (<https://www.nature.com/articles/sdata201618>)
 - findable, accessible, interoperable, and reusable
- Lineage/Provenance
 - capture relevant information for understanding how data has been moved, transferred, processed, etc.
 - metadata models: W3C Provenance Model, DataHub, etc.
- Key issues
 - which metadata must be captured?
 - based on existing tools or your own?
- Instrumentation/logging processes and automated data lineage → performance overhead!

Check: <https://datahubproject.io/docs/metadata-modeling/metadata-model>

Data lineage and observability (2)

- Data observability: the health about data
 - near-real time metrics, offline checks and possible dashboards
 - similar to service observability, relying on traces, logs, metrics, etc.
- Focus on data
 - data metrics (volumes, data quality, schemas, lineage)
 - issues due to data problems
 - data ingestion processes/workflows
- Some solutions
 - validation of data against design schemas (e.g., Schema Registry in Kafka)
 - checks of realtime and offline data quality attributes → integrate with data ingestion processes or offline data profiling
 - integrated data quality tests in pipelines (e.g., data testing)

Tools:



Microsoft Presidio

YData Profiling

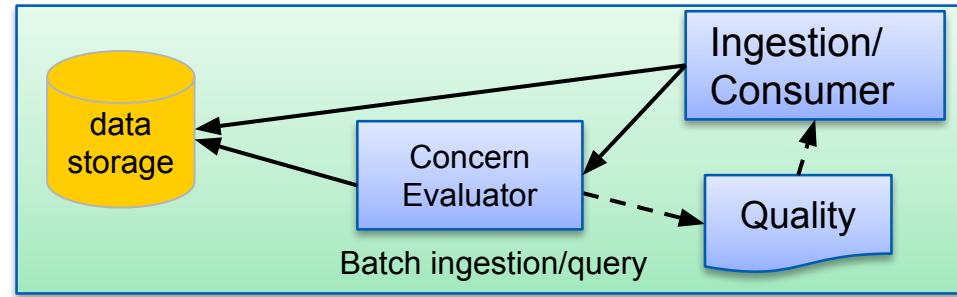


A!

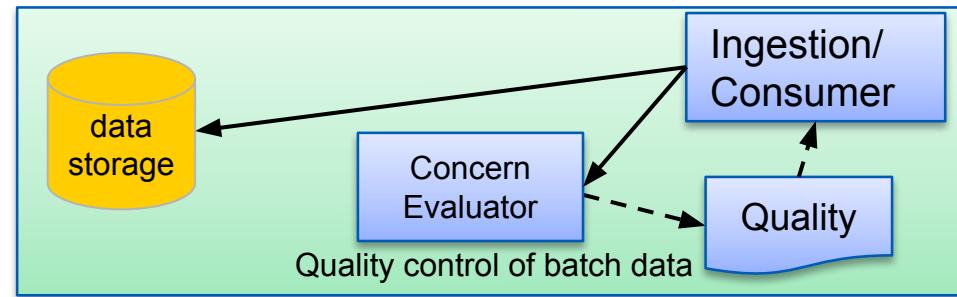
Quality control and regulation assurance (1)

Design: different evaluation mechanisms

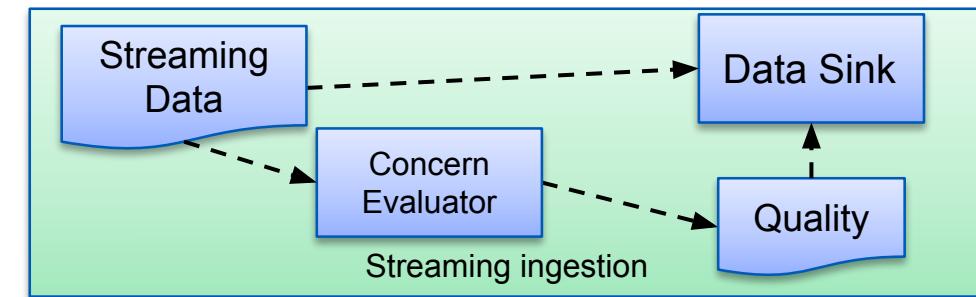
**Pull,
pass-by-reference
model for
evaluating data
concerns**



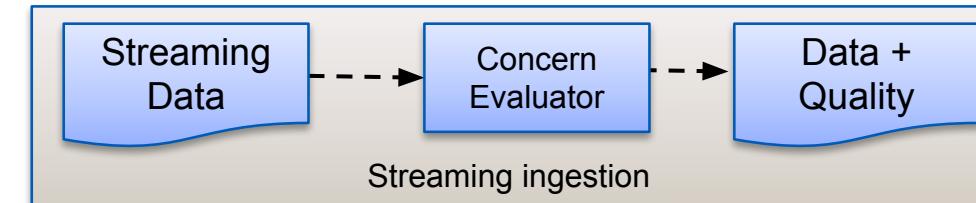
**Pull, pass-by-value
model for
evaluating data
concerns**



**Push model for evaluating data concerns
of active data sources: sampling + delay**



Potential performance bottleneck



Reading Source: H. -L. Truong and S. Dustdar, "On Evaluating and Publishing Data Concerns for Data as a Service," 2010 IEEE Asia-Pacific Services Computing Conference, doi: 10.1109/APSCC.2010.54.

A!

Quality control and regulation assurance (2)

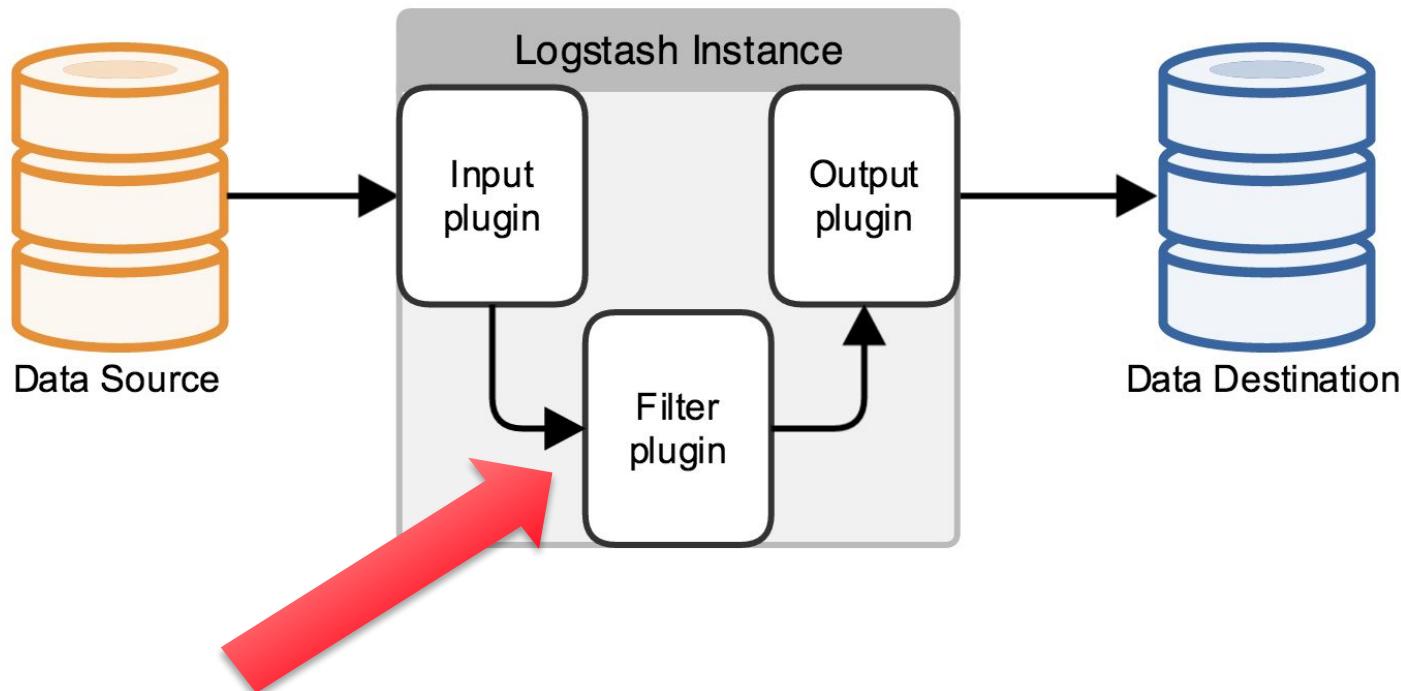
- Before, after or during the ingestion/transformation
- In-process vs out-process
 - in process: using libraries doing data quality → must be very fast
 - out-process: a separate task in the workflow or external programs/services
- Profiling, sampling, AI/ML techniques for data quality
- Examples:
 - Using a separate program like PyDeequ Spark to check quality
 - <https://github.com/rdsea/bigdataplatforms/tree/master/tutorials/dataquality>
 - Anonymizing data
 - <https://microsoft.github.io/presidio/anonymizer/>

Example tooling for ingestion pipelines

Study existing tools

- Different ways to deliver ingestion pipelines
- (Traditional) ways of REST API/specific client libraries
 - upload using put/get operations
- Workflows
 - self-developed workflows vs automatically generated workflows
- Pipelines are bundled into containers
 - self-developed vs generic pipelines based on user configurations

Design tools for ingestion processes 1: Logstash



Pluggable approaches

Figure source:

<https://www.elastic.co/guide/en/logstash/current/getting-started-with-logstash.html> (from the previous version of Logstash)

```
input {  
  file {  
    path => "${MY_INPUT_DIR}/bts-data-alarm-2017.csv"  
    start_position => "beginning"  
  }  
}  
filter {  
  csv {  
    separator => ","  
    columns => ["station_id","datapoint_id","alarm_id"]  
  }  
}  
output {  
  stdout {  
    #...  
  }  
}
```

Pipeline is defined in a configuration file

Design tools for ingestion processes 2: Apache Druid

Allow the user to build the plan: select tasks, define configuration, etc. and then generate ingestion pipelines

The screenshot shows the Apache Druid Ingestion UI interface. At the top, there are four tabs: "Connect and parse raw data", "Transform data and configure schema", "Tune parameters", and "Verify and submit". Below these tabs, there are several sub-options: "Start", "Connect", "Parse data", "Parse time", "Transform", "Filter", "Configure schema", "Partition", "Tune", "Publish", and "Edit spec".

The main area displays a preview of raw data from a CSV file. The data includes columns such as VendorID, tpep_pickup_datetime, tpep_dropoff_datetime, passenger_count, trip_distance, RatecodeID, store_and_fwd_flag, PULocationID, DOLocationID, payment_type, and various timestamp and numerical values.

To the right of the data preview, there is a descriptive text block: "Druid ingests raw data and converts it into a custom, indexed format that is optimized for analytic queries. To get started, please specify what data you want to ingest." Below this text is a "Learn more" link.

Configuration settings on the right side include:

- Source type:** local
- Base directory:** /opt/data/rawdata/bdp
- File filter:** *.csv

A warning message in an orange box states: "⚠ This path must be available on the local filesystem of all Druid services." Below the configuration section is a blue "Apply" button.

A!

Design tools for ingestion processes 3: Apache Nifi

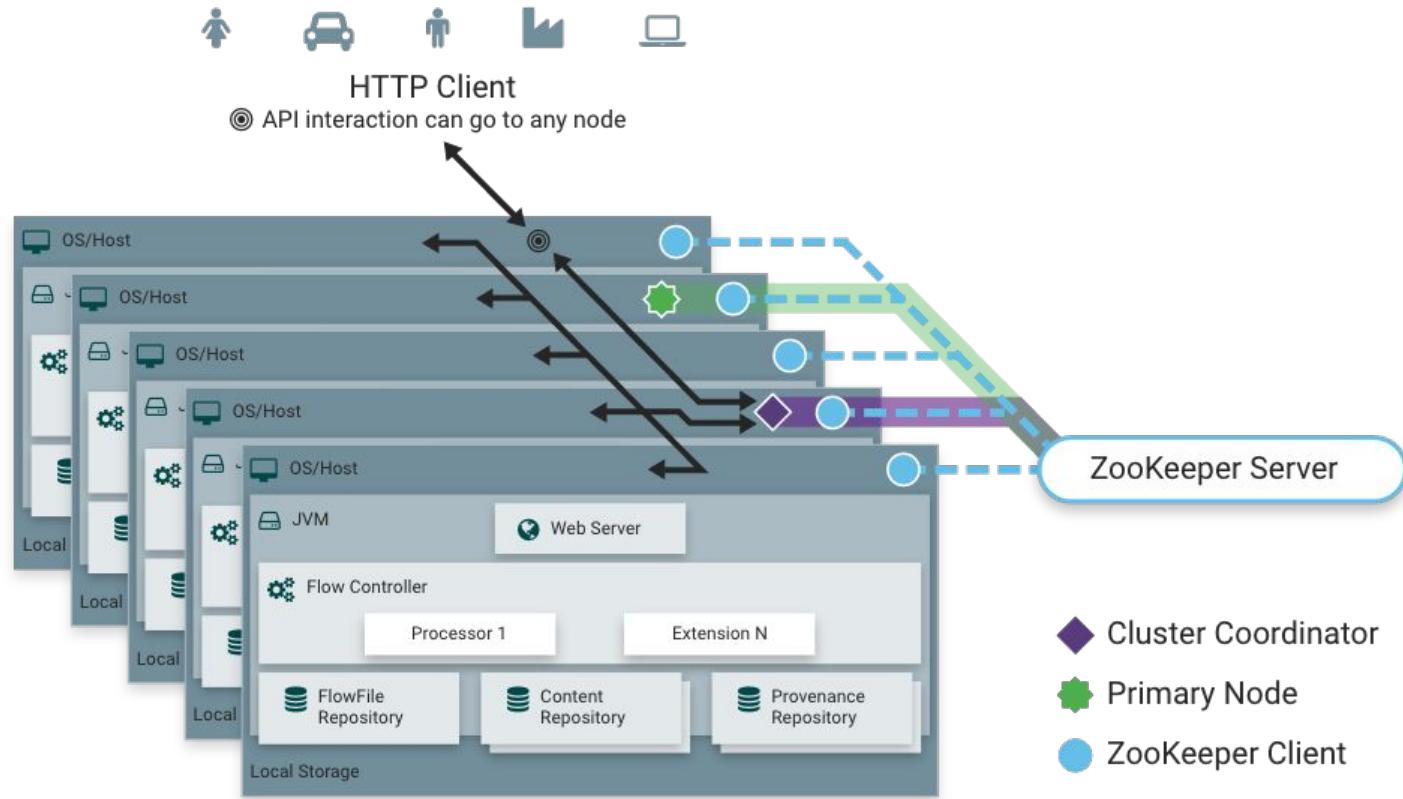


Figure source: <https://nifi.apache.org/nifi-docs/administration-guide.html#clustering>

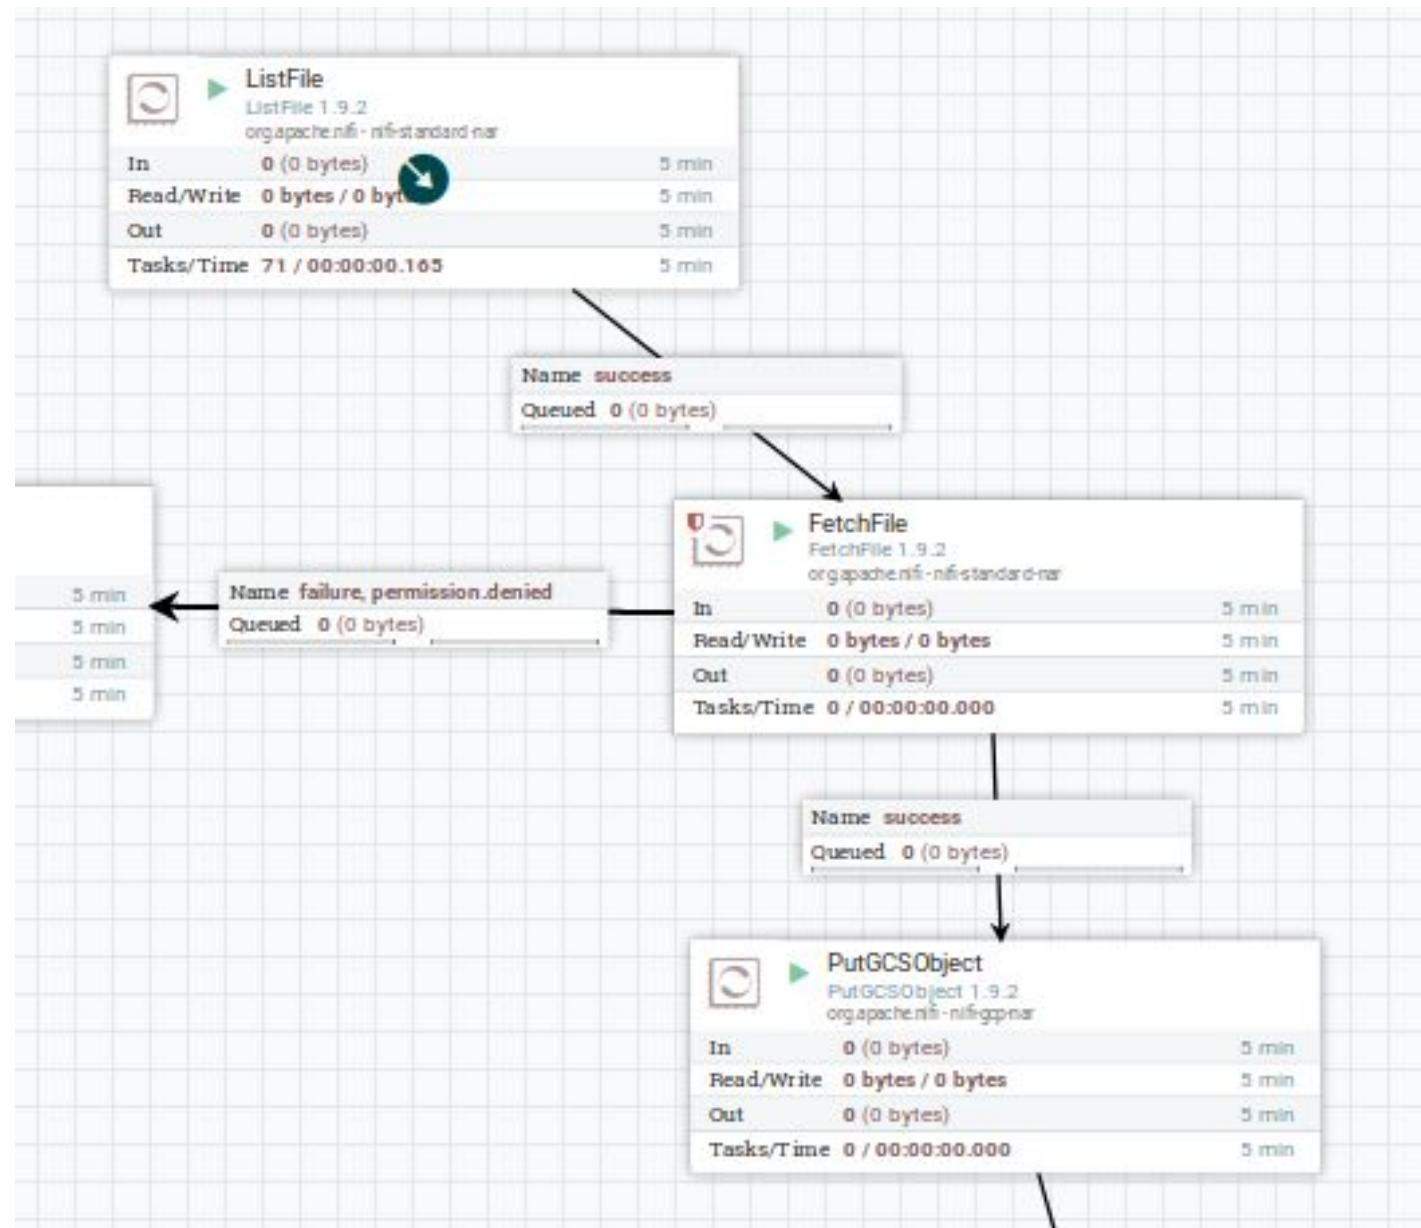
Design tools for ingestion processes: Apache Nifi - key concept

- Data is encapsulated into “**FlowFile**”
- **Processor** (Component) performs tasks
- Processor handle FlowFile and has different states
 - each state indicates the results of processing that can be used for establishing relationships to other components
- **Processors** are connected by **Connection**
- Connection can have many **relationships** based on states of upstream Processors

Design tools for ingestion processes: Apache Nifi

See the tutorial:

<https://github.com/rdsea/bigdataplatforms/tree/master/tutorials/nifi>



A!

Summary

- Different designs of data ingestion for batch and streaming
- Ingestion is a complex pipeline
 - many different sub tasks
 - complex requirements w.r.t performance, scale, failure handling
- Different tools/stacks/services available
 - share composable design principles, but different software models and deployments → explore them for your work
- Do real-world designs
 - hands-ons with widely used tools
 - complex designs but we do not need to “reinvent the wheel” → stay with core concepts and requirements to find the right tools!

Thanks!

Hong-Linh Truong
Department of Computer Science

rdsea.github.io



—
Kiitos
aalto.fi