



Aalto University  
School of Science

# Programming Models for Big Data Processing

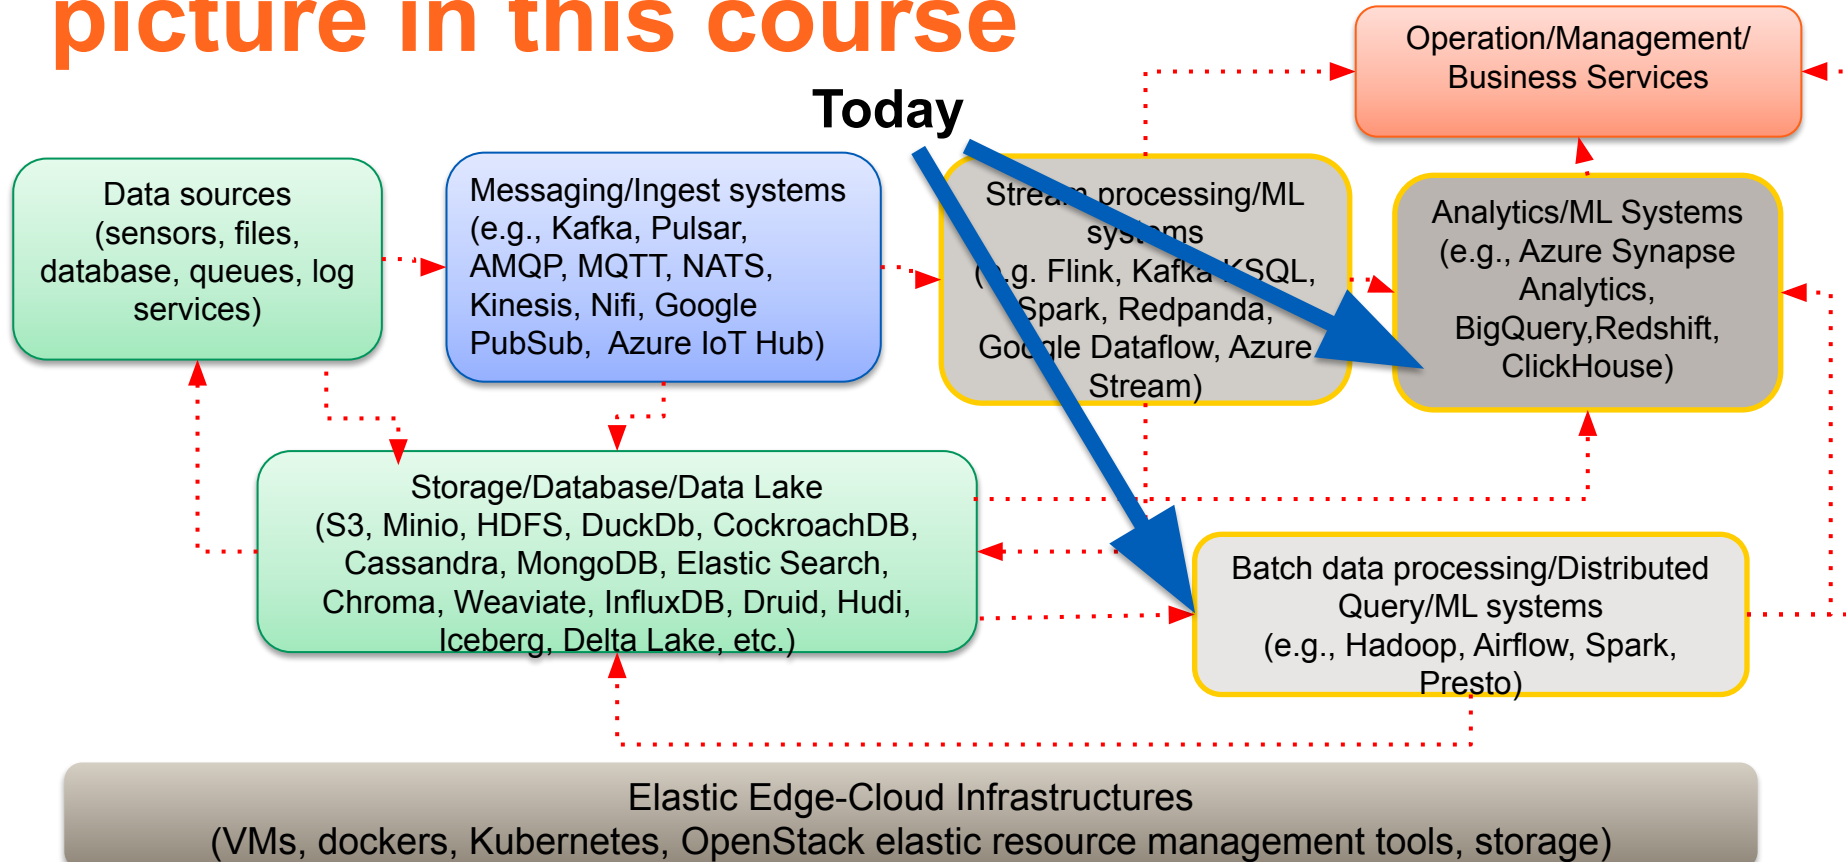
*Hong-Linh Truong*  
*Department of Computer Science*  
*[linh.truong@aalto.fi](mailto:linh.truong@aalto.fi)*, *<https://rdsea.github.io>*

CS-E4640 Big Data Platforms, Spring 2025, Hong-Linh Truong  
12/02/2025

# Learning objectives

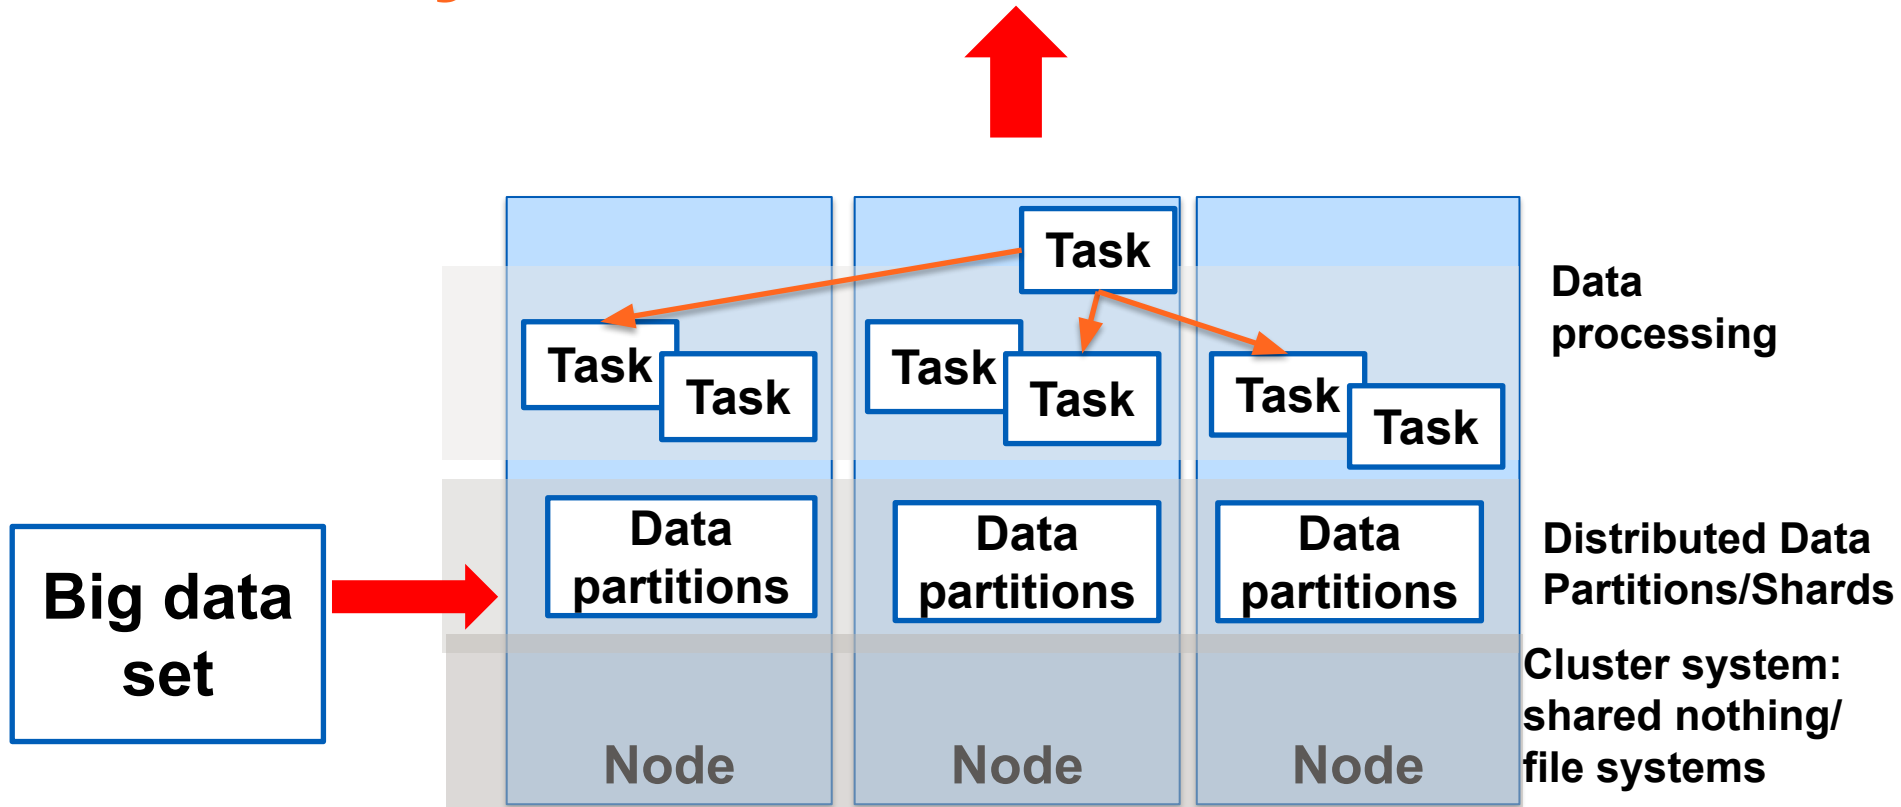
- **Be familiar with key processing models and common techniques using multiple nodes/clusters for data processing**
- **Understand programming models and supports in Dask and Spark for data processing**
- **Able to perform practical programming features for data ingestion, transformation and analysis**

# Big data at large-scale: the big picture in this course



# Understanding common aspects

# Analytics with cluster systems



# Parallel processing of data in distributed file systems/storage

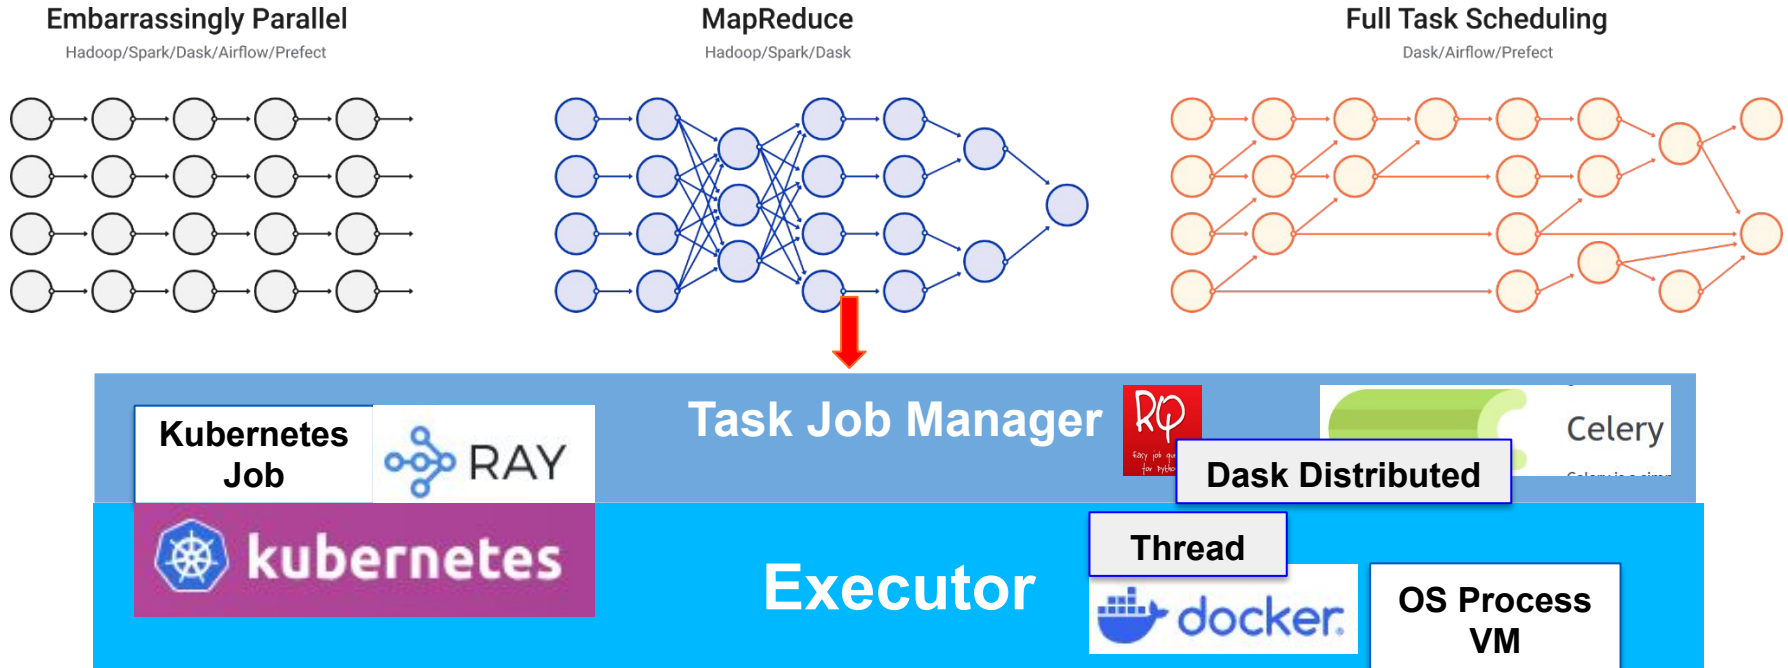
- **Distributed data with different data formats**
  - multiple types of data transformation/analytics with high concurrent/parallel data writes/reads
- **Explore parallel processing at different levels**
  - **data organization**: different data access/analytics frequencies, e.g., data organized into hot, warm and cold data
  - **individual data collection**: items in a collection, e.g., a set of data files/tables, can be processed in parallel
  - **parts of individual data file/table** can be processed in parallel
- **Leverage multiprocessing features from modern compute resources speedup data processing**

# Parallel processing of data for Data Lake/Lakehouse

- **Modern lake data: cloud or on-premise**
  - *multiple types of data* from different sources (databases, files, sensors, etc)
  - *different forms in storage*: raw data, enriched/processed/cleansing, application-/business curated data, sandbox data (for testing, collaboration)
  - *common, standard, cost optimal storage*: object storage (S3, Azure), (distributed) file storage (Hadoop FS), ...
- **Data Lake/Lakehouse core**
  - data tables, metadata and catalogs
  - open standards: Parquet, ORC, Iceberg tables, Delta Lake formats
  - many processing and governance tasks

# Parallel and distributed data processing models

Figure source: <https://docs.dask.org/en/stable/graphs.html>





# DataFrame/Table view of data

## Example taxi records: named columns

| passenger_count | trip_distance | RatecodeID | store_and_fwd_flag | PULocationID | DOLocationID | payment_type | fare_amount | extra | mta_tax | tip_amount | tolls_amount | improvement_surcharge | total_amount |
|-----------------|---------------|------------|--------------------|--------------|--------------|--------------|-------------|-------|---------|------------|--------------|-----------------------|--------------|
| 1               | 1.34          | 1          | N                  | 238          | 236          | 2            | 10.0        | 0.0   | 0.5     | 0.0        | 0.0          | 0.3                   | 10.8         |
| 1               | 1.34          | 1          | N                  | 238          | 236          | 2            | 10.0        | 0.0   | 0.5     | 0.0        | 0.0          | 0.3                   | 10.8         |
| 1               | 0.32          | 1          | N                  | 238          | 238          | 2            | 4.0         | 0.0   | 0.5     | 0.0        | 0.0          | 0.3                   | 4.8          |
| 1               | 0.32          | 1          | N                  | 238          | 238          | 2            | 4.0         | 0.0   | 0.5     | 0.0        | 0.0          | 0.3                   | 4.8          |
| 1               | 1.85          | 1          | N                  | 236          | 238          | 2            | 10.0        | 0.0   | 0.5     | 0.0        | 0.0          | 0.3                   | 10.8         |
| 1               | 1.85          | 1          | N                  | 236          | 238          | 2            | 10.0        | 0.0   | 0.5     | 0.0        | 0.0          | 0.3                   | 10.8         |
| 1               | 1.65          | 1          | N                  | 68           | 237          | 2            | 12.5        | 0.0   | 0.5     | 0.0        | 0.0          | 0.3                   | 13.3         |
| 1               | 1.65          | 1          | N                  | 68           | 237          | 2            | 12.5        | 0.0   | 0.5     | 0.0        | 0.0          | 0.3                   | 13.3         |
| 1               | 1.07          | 1          | N                  | 170          | 68           | 2            | 9.0         | 0.0   | 0.5     | 0.0        | 0.0          | 0.3                   | 9.8          |
| 1               | 1.07          | 1          | N                  | 170          | 68           | 2            | 9.0         | 0.0   | 0.5     | 0.0        | 0.0          | 0.3                   | 9.8          |
| 1               | 1.3           | 1          | N                  | 107          | 170          | 2            | 7.5         | 0.0   | 0.5     | 0.0        | 0.0          | 0.3                   | 8.3          |
| 1               | 1.3           | 1          | N                  | 107          | 170          | 2            | 7.5         | 0.0   | 0.5     | 0.0        | 0.0          | 0.3                   | 8.3          |
| 1               | 1.85          | 1          | N                  | 113          | 137          | 2            | 10.0        | 0.0   | 0.5     | 0.0        | 0.0          | 0.3                   | 10.8         |
| 1               | 1.85          | 1          | N                  | 113          | 137          | 2            | 10.0        | 0.0   | 0.5     | 0.0        | 0.0          | 0.3                   | 10.8         |
| 1               | 0.62          | 1          | N                  | 231          | 231          | 2            | 4.5         | 0.0   | 0.5     | 0.0        | 0.0          | 0.3                   | 5.3          |
| 1               | 0.62          | 1          | N                  | 231          | 231          | 2            | 4.5         | 0.0   | 0.5     | 0.0        | 0.0          | 0.3                   | 5.3          |
| 1               | 0.0           | 1          | N                  | 264          | 264          | 2            | 0.0         | 0.0   | 0.0     | 0.0        | 0.0          | 0.0                   | 0.0          |
| 1               | 0.29          | 1          | N                  | 162          | 162          | 2            | 4.0         | 0.0   | 0.5     | 0.0        | 0.0          | 0.3                   | 4.8          |
| 1               | 0.29          | 1          | N                  | 162          | 162          | 2            | 4.0         | 0.0   | 0.5     | 0.0        | 0.0          | 0.3                   | 4.8          |
| 1               | 1.34          | 1          | N                  | 239          | 151          | 2            | 7.0         | 0.5   | 0.5     | 0.0        | 0.0          | 0.3                   | 8.3          |

- Very common we analyze **big data files** based on this view
- Streaming data can be also represented as **unbounded tables**



<https://pandas.pydata.org/docs/>



<https://github.com/pola-rs/polars>

```

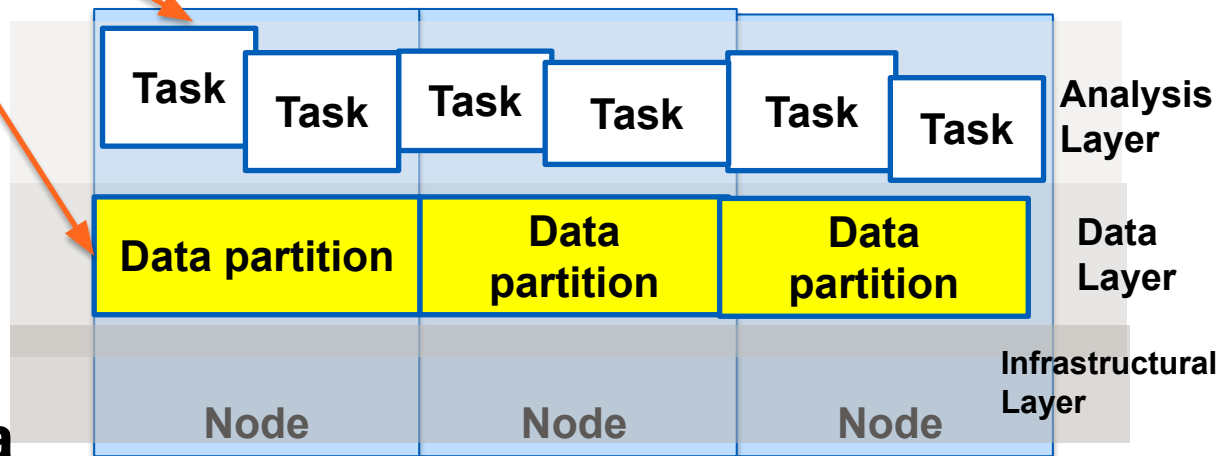
inputFile =args.input_file
## hadoop inputFile="hdfs://"
df =spark.read.csv(inputFile,header=True,inferSchema=True)
#df.show()
print("Number of trips", df.count())
#number of passenger count per vendor and total amount of money
passenger_exprs = {"passenger_count":"sum","total_amount":"sum"}
df2 = df.groupBy('VendorID').agg(passenger_exprs)
# Where do you want to write the output
df2.repartition(1).write.csv(args.output_dir,header=True)

```

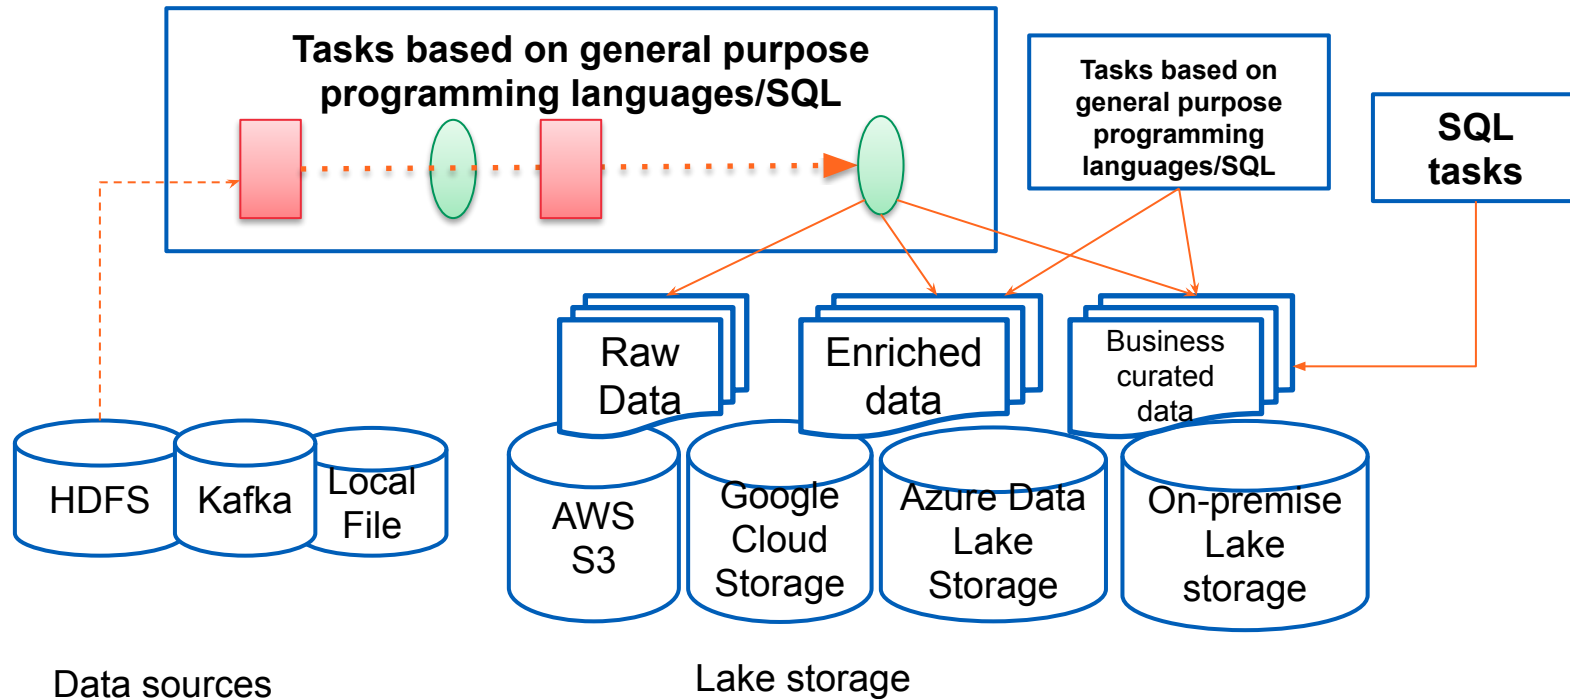
Python/Java/Rust/...

Result

**General purpose programming languages:** what we want when developing analysis programs for big data



# General purpose programming languages +SQL for Data Lake/Lakehouse



# Example

**Spark program with Spark Delta for processing data and store the processed data into a cloud data lake storage**

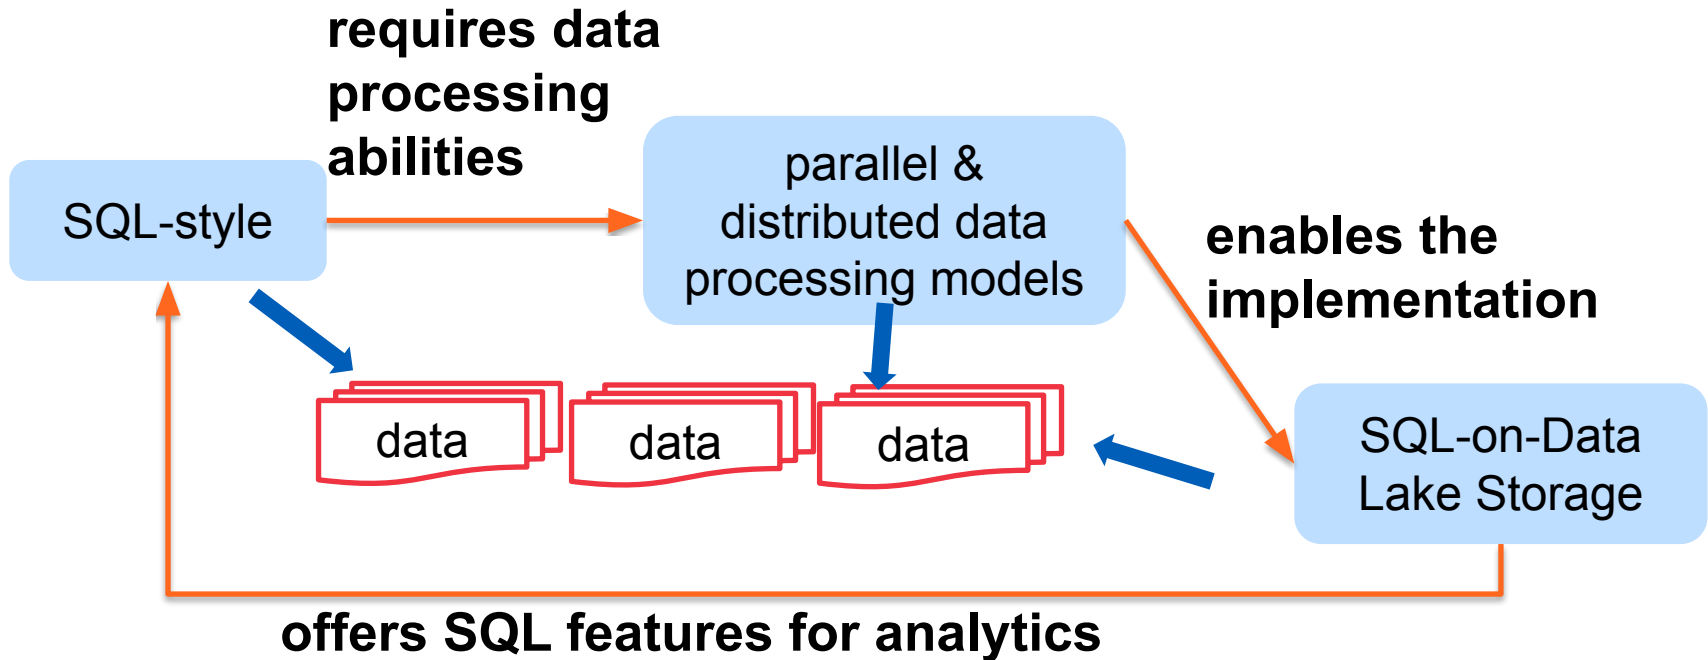
```
## hadoop inputFile="hdfs://"
spark_df =spark.read.csv(inputFile,header=True,inferSchema=True)
print(spark_df.head(10))
#do many things, before producing data for datalake
spark_df.write.format("delta").mode("append").save(lake_table_path)
```

E.g., Data lake storage  
based on Google  
Cloud Storage  
(<https://delta.io/>)

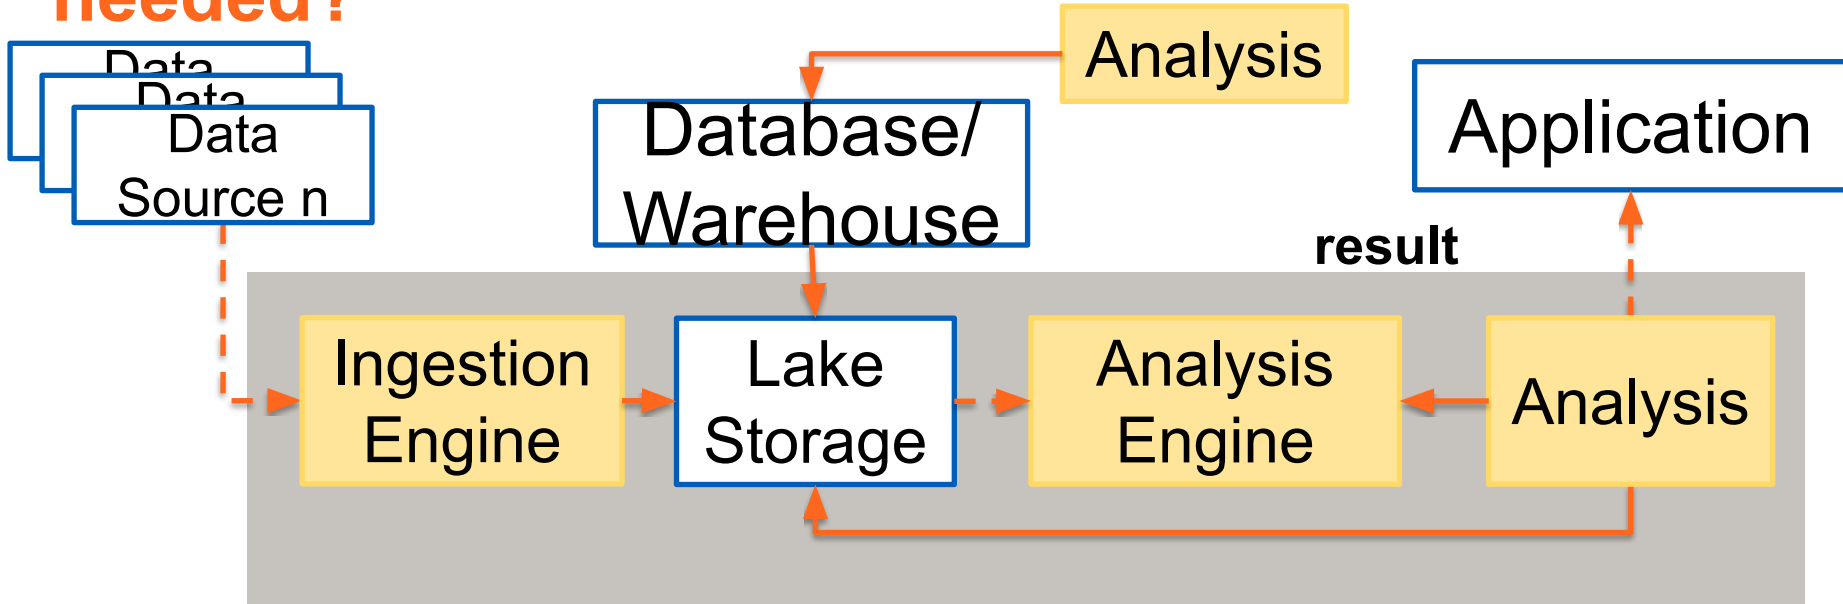
**Data lake storage with  
Iceberg tables,  
Pyarrow an Pylceberg**  
(<https://py.iceberg.apache.org/>)

```
37 catalog = SqlCatalog(
38     catalog_name,
39     **catalog_config["catalog"][catalog_name],
40 )
41
42 if data_type == ".parquet":
43     df = pq.read_table(input_data)
44 else:
45     df = csv.read_csv(input_data)
46 catalog.create_namespace_if_not_exists(namespace)
47 logger.info(f'Existing namespaces: {catalog.list_namespaces()}')
48 full_tablename=f'{namespace}.{table_name}'
49 if not catalog.list_namespaces((namespace)):
50     catalog.create_namespace(namespace)
51 table = catalog.create_table_if_not_exists(
52     full_tablename,
53     schema=df.schema,
54 )
55 table.append(df)
```

# Enabling SQL-style with parallel/distributed data processing



# Where is distributed/parallel data processing needed?



## Data Ingestion:

- Spark Streaming
- Kafka Connect
- Apache Nifi
- etc.

## Storage

- HDFS, AWS S3, Google Storage, Azure Data Lake Storage, Iceberg tables, etc., as storage

## Computing/Data Processing Frameworks

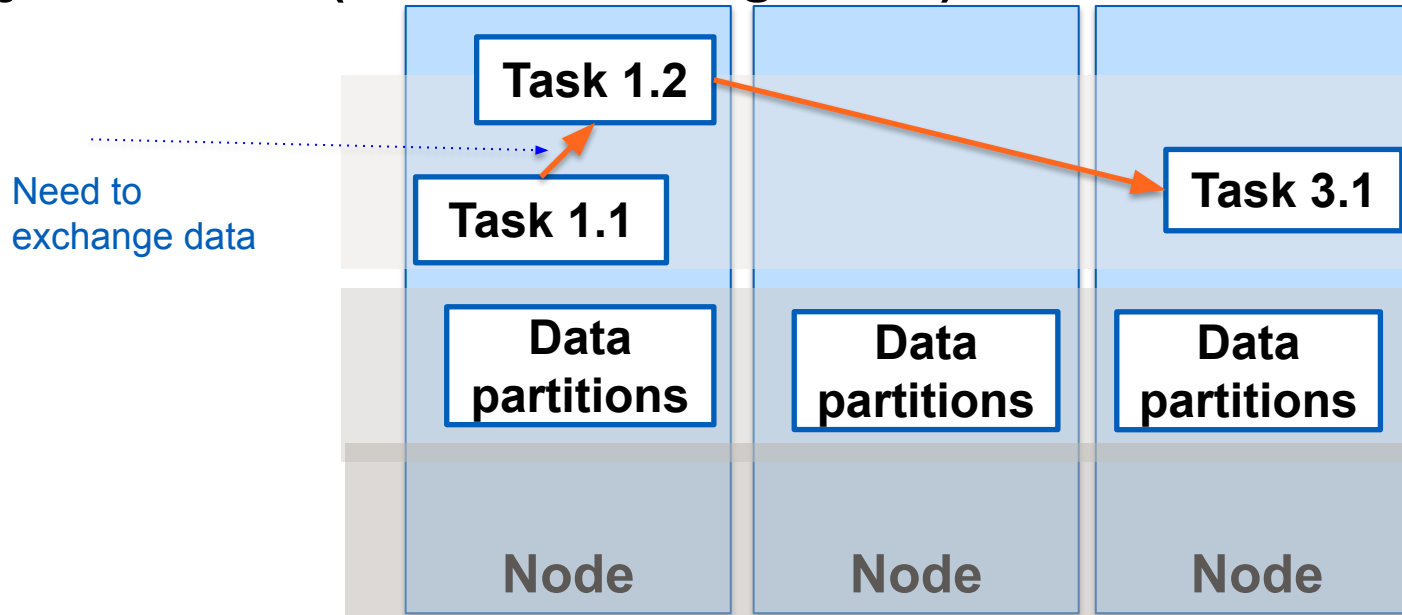
- Apache Spark
- Hadoop MapReduce
- Dask, Ray, etc.

# Common aspects

- **Data input/output connectors**
  - for reading data from sources and loading data into data sinks
- **Data collections as abstract (big/distributed) data structures**
  - for modeling/representing data in suitable views for processing
- **Data operations**
  - operations applied to data in data collections
- **Execution models**
  - tasks, jobs, workflows
  - scheduling; lazy vs eager execution; delayed and future execution
- **Task fault-tolerance and data exchange among tasks in distributed processes/machines**

# Exchange data possibilities

Try to analyze and identify this problem with a framework of your choice (and fault management)





# Programing frameworks in our focus

- **Programming models**

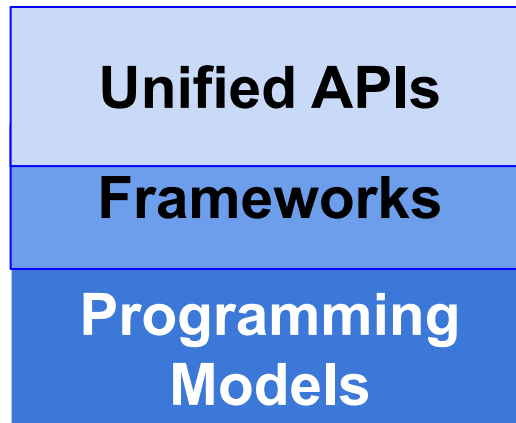
- OS-based multi-threads/processes
- embarrassingly parallel programming
- MapReduce/Spark
- workflows
- (distributed) SQL processing with MPP (Massive Parallel Processing)

- **Programming frameworks**

- Apache Hadoop/Spark, Dask, Polars, Apache Airflow

- **Not in our focus:**

- HPC MPI (Message Passing Interface)



e.g., IBIS  
<https://ibis-project.org/>

# Embarrassingly parallel processing workload and task graph dependencies with Dask

<https://www.dask.org/>

# Key features

- **Data input/output connectors**

- file types: CSV, Parquet, HDF, ORC, Json
- source: Cloud storage (S3, Google), HDFS, Snowflake, BigQuery, Delta Lake

- **Data Collections**

- Array (like numpy array), Bag/Multiset (suitable for unstructured data, like text), DataFrame

- **Operations:**

- joins, concatenation, aggregation (first, sum, ...)
- grouping/resampling, SQL-alike support
- functions/computation suitable for arrays

- **Execution modes**

- lazy by default; and support specific delayed and future tasks

# Key features

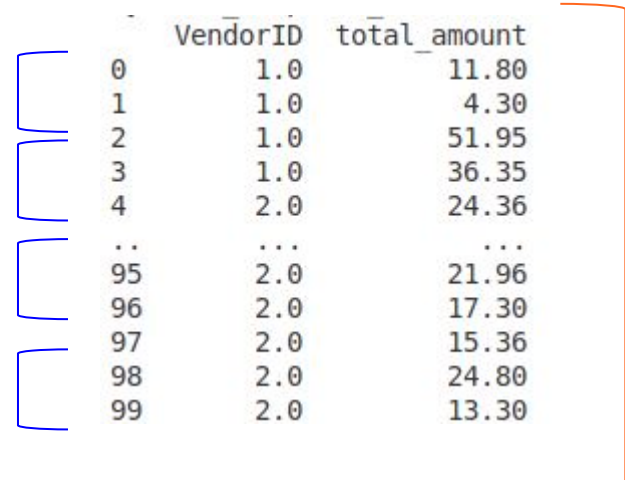
- **Data can be splitted and processed in parallel tasks**
  - many operations on dataframes/tables can be parallelized, with little/without dependency among tasks
  - using directed acyclic graph (DAG) to represent tasks
  - little communication among them, little data shuffle between tasks
- **Single and multiple nodes as computing resources**
  - multiprocessing in single node vs distributed nodes
  - scheduling graphs using OS threads and processes to execute tasks
  - data exchange among tasks using shared memory, direct communication or disk
  - using different resource management systems: Kubernetes, SLURM, PBS, etc.

# Parallelizing dataframe → embarrassingly

- A big dataset can be presented as a Dask dataframe
  - a Dask dataframe can be partitioned into different partitions
- Perform operations on data partitions with lazy principles
  - explicitly call compute() method → computation

Dataframe  
in partitions

Dask  
Dataframe

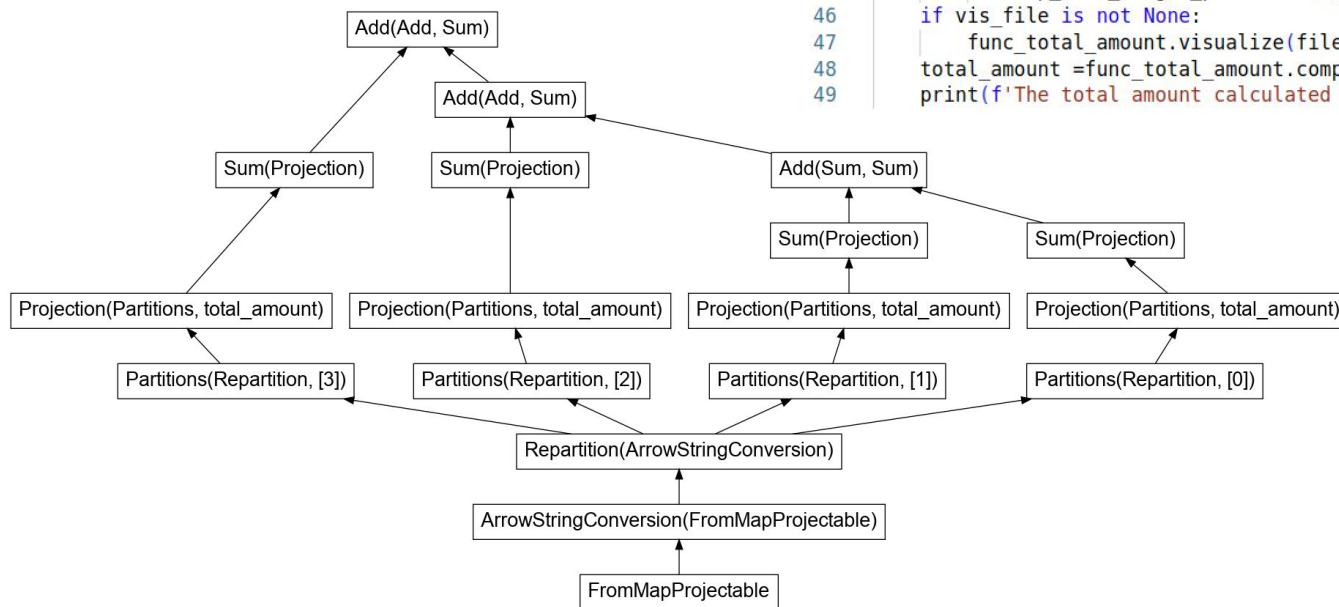


|    | VendorID | total_amount |
|----|----------|--------------|
| 0  | 1.0      | 11.80        |
| 1  | 1.0      | 4.30         |
| 2  | 1.0      | 51.95        |
| 3  | 1.0      | 36.35        |
| 4  | 2.0      | 24.36        |
| .. | ...      | ...          |
| 95 | 2.0      | 21.96        |
| 96 | 2.0      | 17.30        |
| 97 | 2.0      | 15.36        |
| 98 | 2.0      | 24.80        |
| 99 | 2.0      | 13.30        |

```
Data records: 1369765
Data has 4 partitions
Partition 0 has 342441
Partition 1 has 342441
Partition 2 has 342441
Partition 3 has 342442
```

# Example

```
34 from dask.distributed import Client
35 # make sure that dask scheduler and worker running
36 client = Client(f'{dask_scheduler_host}:{dask_scheduler_port}')
37 taxi_df = dd.read_csv(input_file, dtype = dtype,
38                       assume_missing=True,
39                       low_memory=False)
40 print(f'Total records: {len(taxi_df)}')
41 p_taxi_df = taxi_df.repartition(npartitions=num_partitions)
42 func_total_amount = p_taxi_df.get_partition(0)["total_amount"].sum()
43 for i in range(0,num_partitions):
44     func_total_amount = func_total_amount \
45         + p_taxi_df.get_partition(i)["total_amount"].sum()
46 if vis_file is not None:
47     func_total_amount.visualize(filename=vis_file)
48 total_amount = func_total_amount.compute()
49 print(f'The total amount calculated from this file is {total_amount}')
```

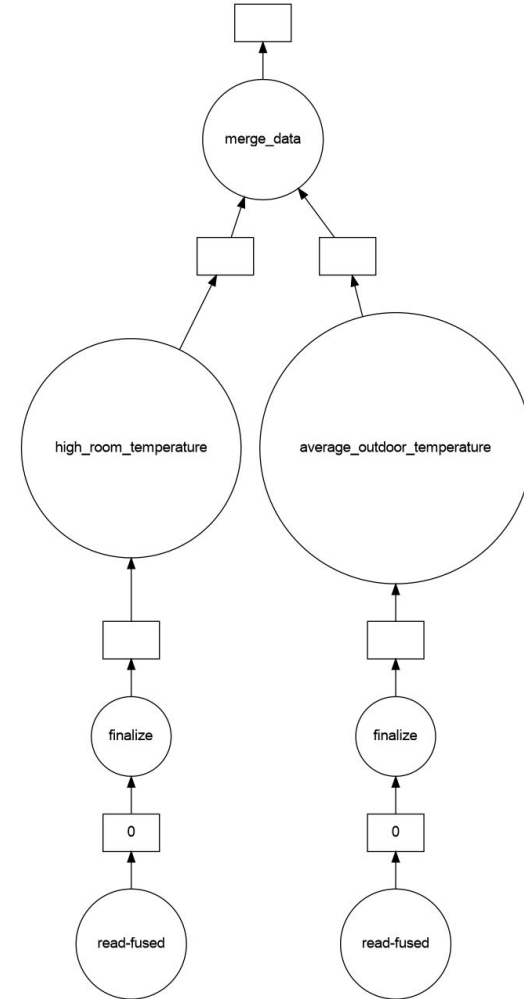


# Task dependency based on DAG

- **Flexible to define task graphs**
  - as a directed acyclic graph
- **Explicitly lazy, deferred execution**
  - using `dask.delayed()/@dask.delayed` to declare delayed tasks
- **Concurrent, asynchronous eager execution**
  - using future tasks
- **Suitable for problems cannot be solved with Dask Dataframe**

# Task dependency based on DAG

```
60 if delayed_mode:
61     # delayed tasks
62     task11 = dask.delayed(high_room_temperature)(bts_alarm_df)
63     task12 = dask.delayed(average_outdoor_temperature)(bts_parameter_df)
64     final_task = dask.delayed(merge_data)(task11, task12)
65     if vis_file is not None:
66         final_task.visualize(filename=vis_file)
67     final_result = final_task.compute()
68     print(f'First 100 elements\n: {final_result.head(100)}')
```





# Apache Spark

<https://spark.apache.org/>

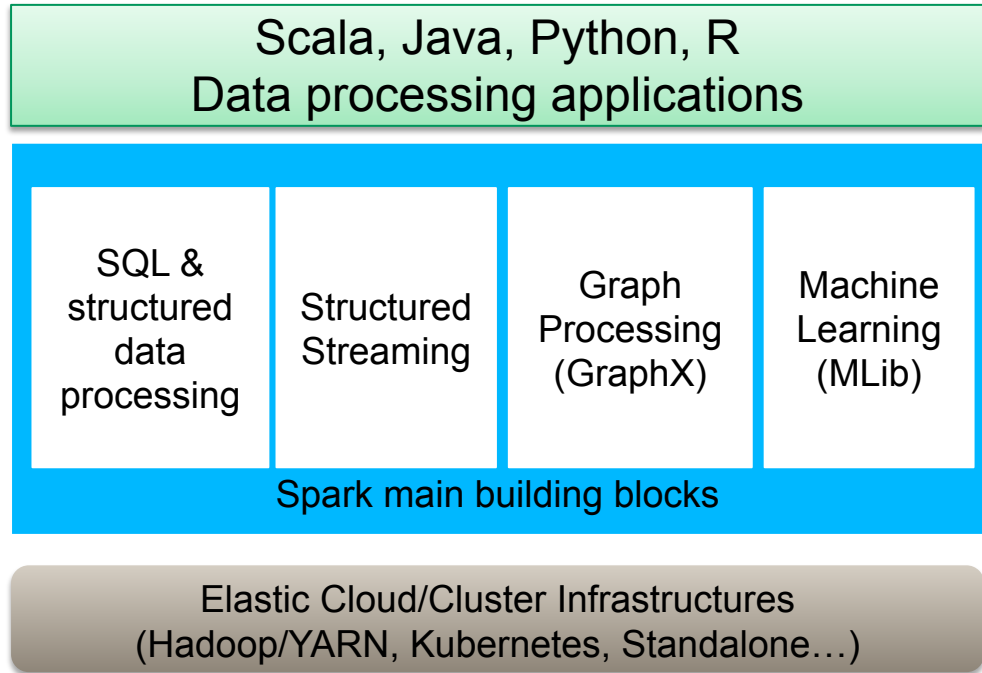
# Apache Spark

- **Cluster-based high-level computing framework**
- **“unified engine” for different types of big data processing**
  - SQL/structured data processing
  - Machine learning
  - Graph processing
  - Streaming processing
- **It is a powerful computing framework and system  $\Rightarrow$  an important service for a big data platform**
  - public cloud: Google DataProc, Azure HDInsight, Amazon EMR
  - data lake systems: e.g., Hudi and Delta Lake

# Apache Spark

Can be run a top

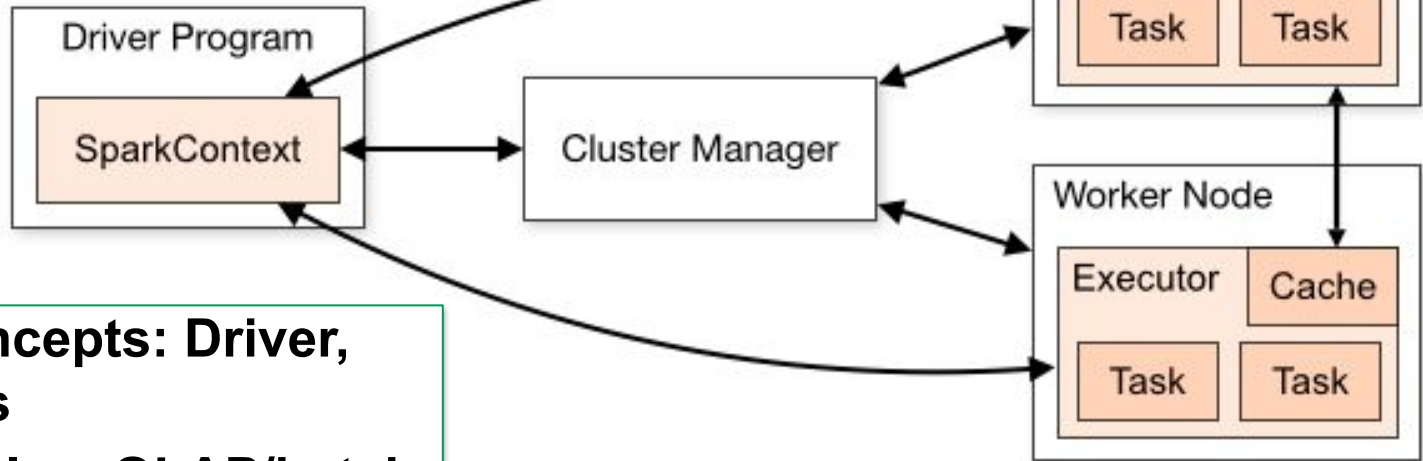
- Hadoop (using HDFS and YARN)
- Kubernetes
- Standalone machines



**Source:** <http://spark.apache.org/>

# Execution model in a cluster system

**Driver** manages operations and tasks in nodes



Computing resources  
in a cluster node

**Common concepts: Driver, Nodes, Tasks**

**Workload styles: OLAP/batch jobs with a lot of data**

Figure source:

<http://spark.apache.org/docs/latest/cluster-overview.html>

# Spark application management: high-level view

- **Submission/Request**

- submit the Spark application for running
- resource is provided for running the Driver

- **Launch**

- the Driver requests resources for executors (through SparkContext)
- establish executors across worker nodes

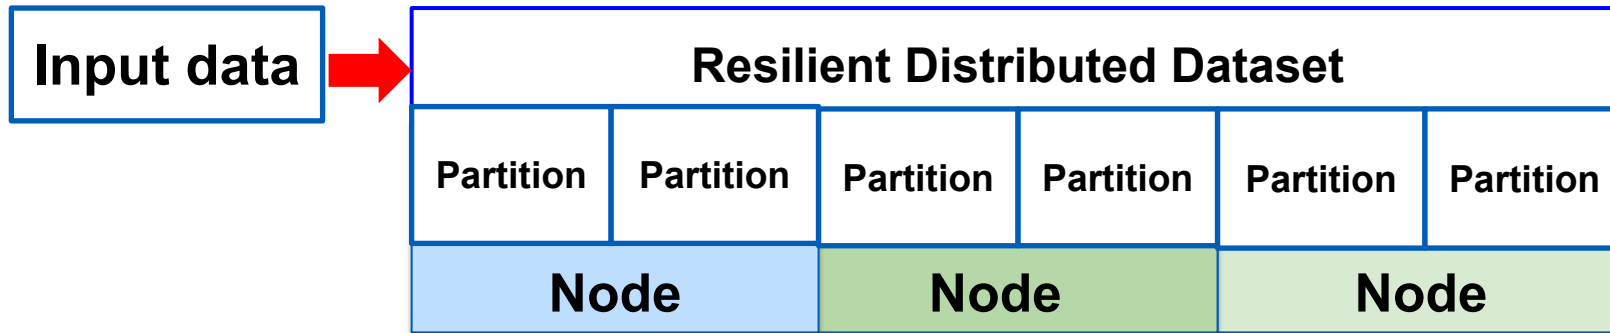
- **Execution**

- the Driver starts to execute code and move data

- **Finish/Completion:**

- finish, release executors

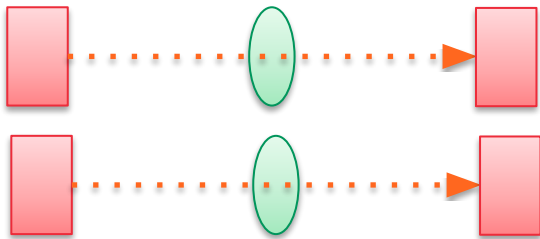
# Key features



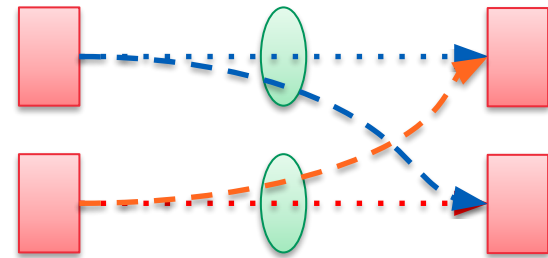
- Input data is **distributed** in different nodes for processing
  - support partitions for data processing: a node keeps one or n partitions, a partition resides only in a node  $\Rightarrow$  for computing
- Key operations: **transformations** and **actions** on data
- Leverage parallel computing concepts to run **multiple tasks**
  - data operation  $\rightarrow$  task executed by executor
  - parallel tasks, task pipeline, DAG of processing stages
- Persistent data in memory/disk for operations

# Transformation operations

- **Transformation**
  - instructions about how to transform a data in a form to another form  $\Rightarrow$  it will not change the original data (immutability)
- **Only tell what to do: to build a DAG (direct acyclic graph)  $\rightarrow$  a lineage of what to do**
- **lazy approach  $\Rightarrow$  real transformation will be done at **action operators****



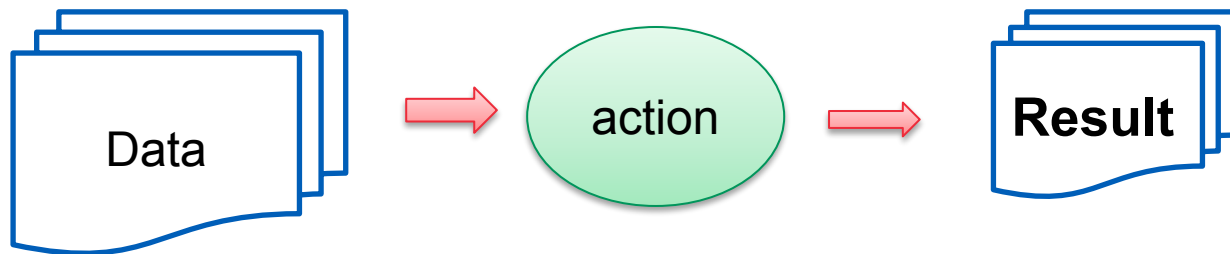
**Narrow transformation,  
no data shuffle**



**Wide transformation, cross data  
partitions, requires a shuffle**

# Action operations

- **Compute the results for a set of transformations**
  - examples: count or average
- **Actions: view, collect, write, calculation**



**Lazy approach: an action triggers execution of transformation operations  $\Rightarrow$  enable various types of optimization**



# Spark program: programming elements

- **SparkSession**

- acts as a program driver to manage the execution of tasks
- SparkContext: manages connection to a cluster and internal services

- **Data APIs**

- low-level Resilient Distributed Dataset (RDD) & shared variables
- high-level DataFrames/DataSets
- load and hold distributed data
- transformation and action functions

- **ML, graph and streaming functions and pipelines**

---

# Spark program logic: typical steps

- **Load data and distribute data**
  - data is **immutable** after created
  - data partition in Spark: a partition is allocated in a node
- **Perform **transformations and actions** operations**
  - *transformations*: build plans for transforming data models
  - *actions*: perform computation on data

# Resilient distributed dataset (RDD)

- **Low-level data structure**

- collection of data elements partitioned across nodes in the cluster
- with data sharing, parallel operations, fault-tolerant features

- **Create RDD**

- created by loading data from files (text, sequence file) including local file systems, HDFS, Cassandra, HBase, Amazon S3, etc.

- **Persist RDD**

- in memory or to files

# Example with RDD

VendorID,tpep\_pickup\_datetime,tpep\_dropoff\_datetime,passenger\_count,trip\_distance,RatecodeID,store\_and\_fwd\_flag,PULocationID,DOLocationID,payment\_type,fare\_amount,extra,mta\_tax,tip\_amount,tolls\_amount,improvement\_surcharge,total\_amount

2,11/04/2084 12:32:24 PM,11/04/2084 12:47:41 PM,1,1.34,1,N,238,236,2,10,0,0.5,0,0,0.3,10.8

2,11/04/2084 12:32:24 PM,11/04/2084 12:47:41 PM,1,1.34,1,N,238,236,2,10,0,0.5,0,0,0.3,10.8

2,11/04/2084 12:25:53 PM,11/04/2084 12:29:00 PM,1,0.32,1,N,238,238,2,4,0,0.5,0,0,0.3,4.8

as a text file

```
conf = SparkConf().setAppName("cse4640-rddshow").setMaster(args.master)
sc = SparkContext(conf=conf)
##modify the input data
rdd=sc.textFile(args.input_file)
## if there is a header we can filter it otherwise comment two lines
csvheader = rdd.first()
rdd = rdd.filter(lambda csventry: csventry != csvheader)
## using map to parse csv text entry
rdd=rdd.map(lambda csventry: csventry.split(","))
rdd.repartition(1)
rdd.saveAsTextFile(args.output_dir)
```

# Shared variables

- **A function is executed a remote and various tasks running in parallel**
  - how do tasks share variables? common patterns in parallel computing: *broadcast and global variable/counter*
- **Variables used in parallel operations**
  - variables are copied among parallel tasks
  - shared among tasks or between tasks and the driver
- **Types of variables**
  - broadcast variables: cache a value in all nodes
  - accumulators: a global counter shared across processes

# Examples

```
conf = SparkConf().setAppName("CS4040-Broadcast").setMaster("ygs:master")
sc = SparkContext(conf=conf)
bVar = sc.broadcast([5,10])
print("The value of the broadcast",bVar.value,sep=" ")
counter = sc.accumulator(0)
sc.parallelize([1, 2, 3, 4]).foreach(lambda x: counter.add(bVar.value[0]))
print("The value of the counter is ",counter.value,sep=" ")
```

## Use cases:

- **Broadcast variables: lookup tables**
- **Accumulators: monitoring/checkpoint counters**

# Spark SQL and DataFrames

- **High-level APIs**
  - design with common programming patterns in data analysis, multi-language support
- **SparkSQL: enable dealing with structured data**
  - SQL query execution, Hive, JDBC/ODBC
- **DataFrame**
  - distributed data organized into named columns, similar to a table in relational database
  - Pandas and Spark DataFrames have similar design concepts

# DataFrame

```
inputFile =args.input_file
df =spark.read.csv(inputFile,header=True,inferSchema=True)
print("Number of partition",df.rdd.getNumPartitions())
df.show()
```

| PROVINCECODE | DEVICEID | IFINDEX | FRAME | SLOT | PORT | ONUINDEX | ONUID | TIME                | SPEEDIN | SPEEDOUT |
|--------------|----------|---------|-------|------|------|----------|-------|---------------------|---------|----------|
| YN 1         | 3023     | 528     | 1     | 2    | 7    | 39 10    | 07039 | 01/08/2019 00:04:07 | 148163  | 49018    |
| YN 1         | 3023     | 528     | 1     | 2    | 7    | 38 10    | 07038 | 01/08/2019 00:04:07 | 1658    | 1362     |
| YN 1         | 3023     | 528     | 1     | 2    | 7    | 9 10     | 07009 | 01/08/2019 00:04:07 | 6693    | 5185     |
| YN 1         | 3023     | 528     | 1     | 2    | 7    | 8 10     | 07008 | 01/08/2019 00:04:07 | 640     | 544      |
| YN 1         | 3023     | 528     | 1     | 2    | 7    | 11 10    | 07011 | 01/08/2019 00:04:07 | 118     | 114      |
| YN 1         | 3023     | 528     | 1     | 2    | 7    | 10 10    | 07010 | 01/08/2019 00:04:07 | 28514   | 12495    |
| YN 1         | 3023     | 528     | 1     | 2    | 7    | 13 10    | 07013 | 01/08/2019 00:04:07 | 868699  | 23400    |
| YN 1         | 3023     | 528     | 1     | 2    | 7    | 15 10    | 07015 | 01/08/2019 00:04:07 | 1822    | 1120     |
| YN 1         | 3023     | 528     | 1     | 2    | 7    | 17 10    | 07017 | 01/08/2019 00:04:07 | 998069  | 117345   |
| YN 1         | 3023     | 528     | 1     | 2    | 7    | 16 10    | 07016 | 01/08/2019 00:04:07 | 22402   | 1804     |
| YN 1         | 3023     | 528     | 1     | 2    | 7    | 19 10    | 07019 | 01/08/2019 00:04:07 | 640     | 791      |
| YN 1         | 3023     | 760     | 1     | 1    | 10   | 49 10    | 10049 | 01/08/2019 00:04:07 | 662     | 494      |
| YN 1         | 3023     | 760     | 1     | 1    | 10   | 48 10    | 10048 | 01/08/2019 00:04:07 | 2158    | 759      |
| YN 1         | 3023     | 528     | 1     | 2    | 7    | 21 10    | 07021 | 01/08/2019 00:04:07 | 0       | 0        |
| YN 1         | 3023     | 760     | 1     | 1    | 10   | 51 10    | 10051 | 01/08/2019 00:04:07 | 2600890 | 54153    |
| YN 1         | 3023     | 528     | 1     | 2    | 7    | 20 10    | 07020 | 01/08/2019 00:04:07 | 330     | 184      |



# Create DataFrame

**DataFrames can be created from a Hive table, from Spark data sources, or another DataFrame**

## Load and save

- From Hive, JSON, CSV
- HDFS, cloud object storage (AWS S3, Google Cloud Storage, Azure Blob Storage), Delta Lake, local files, etc.



**Figure source:**

<https://databricks.com/blog/2015/02/17/introducing-dataframe-s-in-spark-for-large-scale-data-science.html>

# DataFrame Transformations & Actions

- **Several transformations can be done**
  - *Think transformation for relational database or matrix*
- **Select**
  - *df.select*
- **Filter**
  - *df.filter*
- **Groupby**
  - *df.groupBy*
- **Handle missing data**
  - *Drop duplicate rows, drop rows with NA/null data*
  - *Fill NA/null data*

## Actions

- return values calculated from DataFrame

## Examples

- reduce, max, min, sum, variance and stdev

**⇒ Distributed and parallel processing but it is done by the framework**

# Example of a Spark

```
#!/usr/bin/env python2
#encoding: UTF-8
# CS-E4640
import csv
import sys
from datetime import datetime
from pyspark.sql import SparkSession
import numpy as np
from pyspark.sql import functions as F
import argparse
parser = argparse.ArgumentParser()
parser.add_argument('--input_file', help='input data file')
parser.add_argument('--output_dir', help='output dir')
args = parser.parse_args()

##define a context
spark = SparkSession.builder.appName("cse4640-onu").getOrCreate()
#NOTE: using hdfs:///..... for HDFS file or file:///
inputFile = args.input_file
df = spark.read.csv(inputFile, header=True, inferSchema=True)
#df.show()
print("Number of records", df.count())
exprs = {"SPEEDIN": "avg"}
df2 = df.groupBy('ONUID').agg(exprs)
df2.repartition(1).write.csv(args.output_file, header=True)
```

Session/Driver



Read data



Apply operations



# Spark application runtime view

- **Tasks:**

- a unit of work executed in an executor: e.g., performing transformations of a data partition

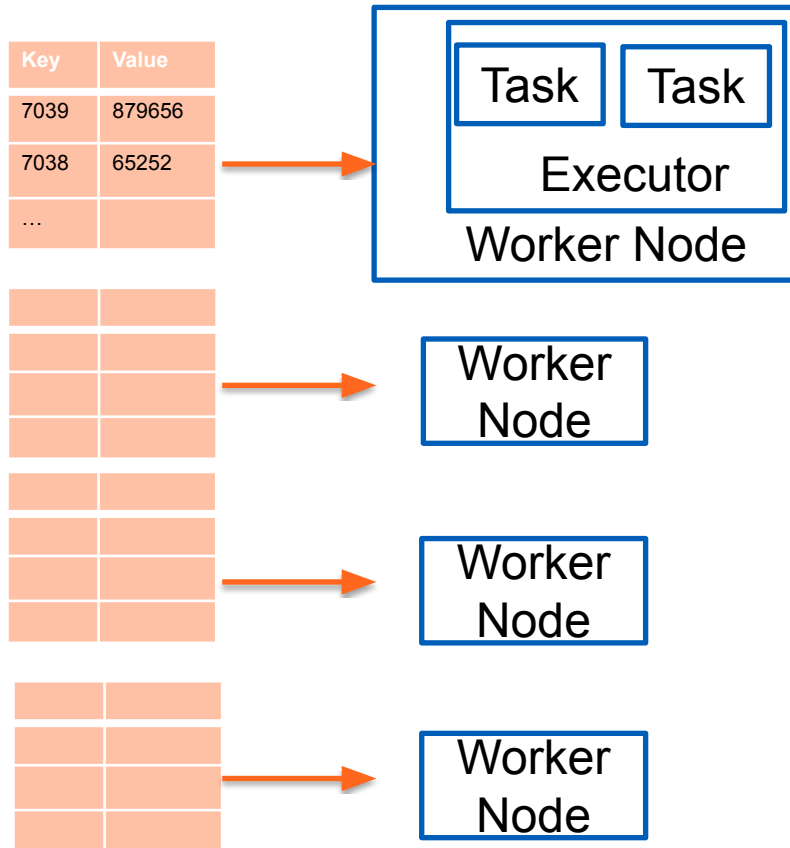
- **Stage: Shuffle Map Stage & Result Stage**

- a set of tasks executed in many nodes for performing the same operation
- move to a new stage: through a shuffle to produce output partitions or an action to produce results

- **Job**

- runtime view of an action operation (actual computation produces a result), includes many stages of tasks

# Data Distribution



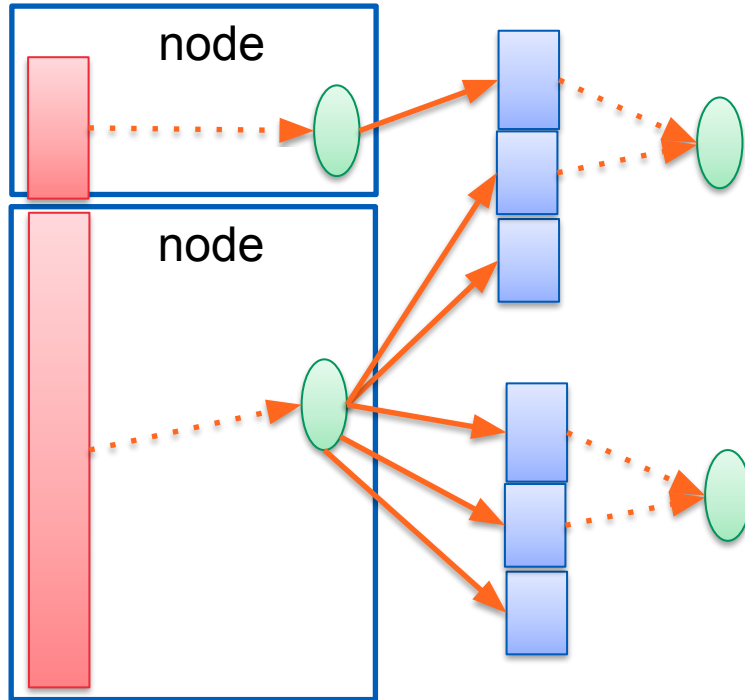
**One task works on a partition at a time**

**⇒ Parallelism and performance are strongly dependent on number of partitions, tasks, CPU cores**

# Data Distribution: load balance

Imbalance  
processing

more data shuffle



- It is important to have well-balanced data distribution across nodes
- Detection:
  - look at runtime execution time to see problems or check your data
- Examples of solution:
  - repartition
  - change group keys

# Pipelining, Shuffle and DAG

- **Operations work in a pipeline without moving data across nodes**
  - e.g., map->filter, select->filter
- **Shuffle persistent**
  - shuffle needs move data across nodes
  - source tasks save shuffle files into local disks for data shuffle, then the target tasks will read data from source nodes
    - *Save time, recovery, fault tolerance*

# Other important support of Spark

- **MLlib - Machine learning**

- Distributed and parallel machine learning algorithms with big data and clusters

- **Streaming: data processing in near real-time**

- *Related to our topic: stream data processing*

- **Graph Processing: Spark GraphX**

- Parallel computation for graphs

- **Many third-party frameworks, e.g.,**

- SparkOCR (<https://www.johnsnowlabs.com/spark-ocr/>), SparkNLP (<https://nlp.johnsnowlabs.com/>)
- PyDeequ (<https://pydeequ.readthedocs.io/en/latest/README.html#>) Data quality
  - check our example:  
<https://github.com/rdsea/bigdataplatforms/tree/master/tutorials/dataquality>



# Massive parallel processing for distributed query engines

# Massive parallel processing employed by distributed query engine

- **Key concepts**
  - using SQL as a way to query different types of data sources like data lake, warehouse, and databases
  - the query engine is decoupled from data sources/storage
  - using massive parallel processing (MPP) to support parallel tasks accessing different data sources at a large-scale with many compute nodes
- **Complex fault tolerance and optimization:**
  - failure management, query and data movement costs, ...
- **Mostly for analytics: interactive analytics, seconds – minutes**

# Example distributed SQL engine: Presto/Trino

Small exercise  
(see course git)

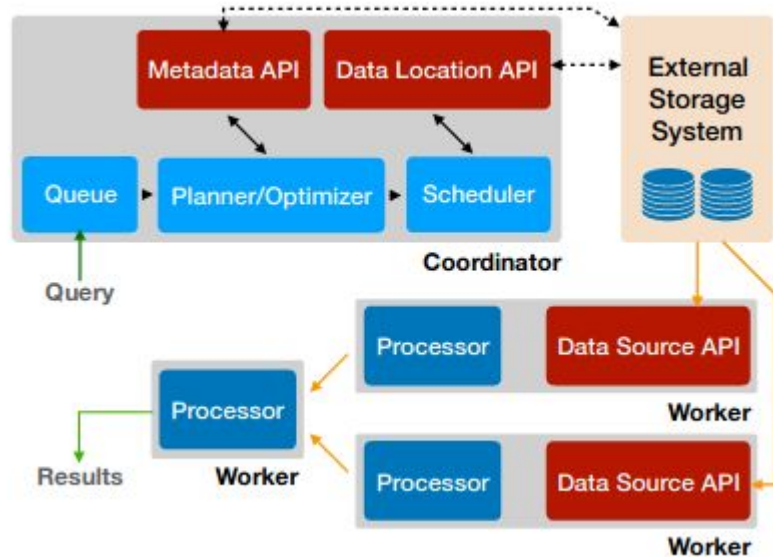
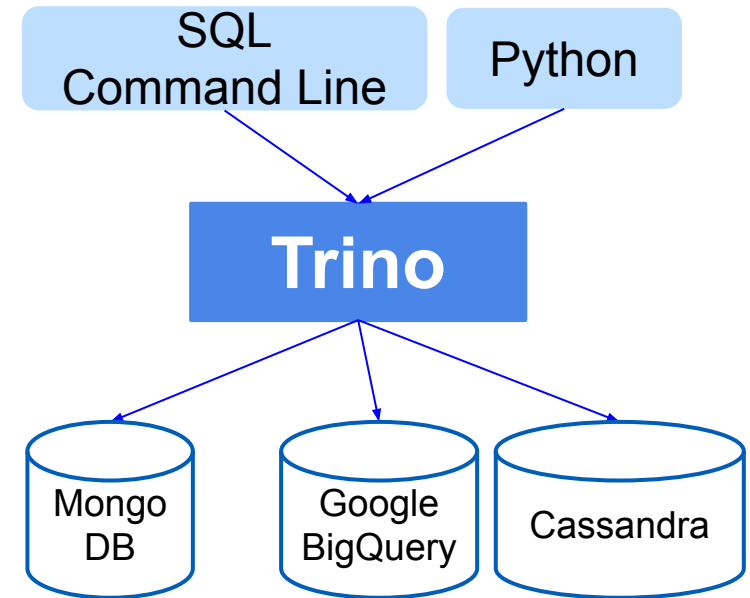


Figure source: Presto: SQL on Everything  
<https://ieeexplore.ieee.org/stamp/stamp.jsp?tp=&arnumber=8731547&tag=1>



Trino (<https://trino.io/>): from a fork of Presto

# Summary

- **Different programming models for data processing**
  - models and tools selected based on data workload and ecosystems, including underlying compute resource management
  - both developers and platform operator/provider must carefully decide the programming models for data processing
- **Effects of modernization and composability in data platforms:**
  - Spark is powerful but many emerging ones, e.g., Polars and DuckDB, which may be suitable due to learning curves, management, data load
- **Thoughts**
  - think about the success of Apache Spark: rich ecosystems!
  - think if you combine data, different distributed programming supports for your big data platform

# Thanks!

**Hong-Linh Truong**  
**Department of Computer Science**

**rdsea.github.io**

# Monitoring Spark: executors and tasks

