
Programming Models for Big Data Processing

Hong-Linh Truong
Department of Computer Science
linh.truong@aalto.fi



11.2.2026




Content is available under
CC BY-SA 4.0 unless otherwise stated

Learning objectives


- Be familiar with key processing models and common techniques using multiple nodes/clusters for data processing
- Understand programming models and supports in Dask and Spark for data processing
- Able to perform practical programming features for data ingestion, transformation and analysis

Our big data at large-scale: the big picture in this course



Understanding common aspects

Analytics with cluster systems




A!

Parallel and distributed processing of data from distributed file systems/storage

- Process distributed data with different data formats
 - multiple types of data transformation/analytics with high concurrent/parallel data writes/reads
- Explore parallel and concurrent processing at different levels
 - **data organization**: different data access/analytics frequencies, e.g., data organized into hot, warm and cold data
 - **individual data collection**: items in a collection, e.g., a set of data files/tables, can be processed in parallel
 - **parts of individual data file/table** can be processed in parallel
- Leverage multiprocessing features from modern compute resources to speed up data processing
 - multi/many-cores and accelerators

Parallel and distributed data processing models

Figure source: <https://docs.dask.org/en/stable/graphs.html>



Choosing suitable programming models is based on use cases and ecosystem!

Using general purpose programming languages: DataFrame/Table view of data

Example taxi records: named columns

passenger_count	trip_distance	RatecodeID	store_and_fwd_flag	PULocationID	DOLocationID	payment_type	fare_amount	extra	mta_tax	tip_amount	tolls_amount	improvement_surcharge	total_amount
1	1.34	1	N	238	236	2	10.0	0.0	0.5	0.0	0.0	0.3	10.8
1	1.34	1	N	238	236	2	10.0	0.0	0.5	0.0	0.0	0.3	10.8
1	0.32	1	N	238	238	2	4.0	0.0	0.5	0.0	0.0	0.3	4.8
1	0.32	1	N	238	238	2	4.0	0.0	0.5	0.0	0.0	0.3	4.8
1	1.85	1	N	236	238	2	10.0	0.0	0.5	0.0	0.0	0.3	10.8
1	1.85	1	N	236	238	2	10.0	0.0	0.5	0.0	0.0	0.3	10.8
1	1.65	1	N	68	237	2	12.5	0.0	0.5	0.0	0.0	0.3	13.3
1	1.65	1	N	68	237	2	12.5	0.0	0.5	0.0	0.0	0.3	13.3
1	1.07	1	N	170	68	2	9.0	0.0	0.5	0.0	0.0	0.3	9.8
1	1.07	1	N	170	68	2	9.0	0.0	0.5	0.0	0.0	0.3	9.8
1	1.3	1	N	107	170	2	7.5	0.0	0.5	0.0	0.0	0.3	8.3
1	1.3	1	N	107	170	2	7.5	0.0	0.5	0.0	0.0	0.3	8.3
1	1.85	1	N	113	137	2	10.0	0.0	0.5	0.0	0.0	0.3	10.8
1	1.85	1	N	113	137	2	10.0	0.0	0.5	0.0	0.0	0.3	10.8
1	0.62	1	N	231	231	2	4.5	0.0	0.5	0.0	0.0	0.3	5.3
1	0.62	1	N	231	231	2	4.5	0.0	0.5	0.0	0.0	0.3	5.3
1	0.0	1	N	264	264	2	0.0	0.0	0.0	0.0	0.0	0.0	0.0
1	0.29	1	N	162	162	2	4.0	0.0	0.5	0.0	0.0	0.3	4.8
1	0.29	1	N	162	162	2	4.0	0.0	0.5	0.0	0.0	0.3	4.8
1	1.34	1	N	239	151	2	7.0	0.5	0.5	0.0	0.0	0.3	8.3

- Very common we analyze big data files based on this view
- Streaming data can be also represented as unbounded tables



<https://pandas.pydata.org/docs/>



<https://github.com/pola-rs/polars>



<https://github.com/modin-project/modin>


A!

```
inputFile = args.input_file
## hadoop inputFile="hdfs://"
df = spark.read.csv(inputFile,header=True,inferSchema=True)
#df.show()
print("Number of trips", df.count())
#number of passenger count per vendor and total amount of money
passenger_expressions = {"passenger_count": "sum", "total_amount": "sum"}
df2 = df.groupBy('VendorID').agg(passenger_expressions)
# Where do you want to write the output
df2.repartition(1).write.csv(args.output_dir,header=True)
```


Python/Java/Rust/...

Using general purpose programming languages:

what we want when
developing analysis
programs for big data



General purpose programming languages +SQL for Data Lake/Lakehouse



Example with writing data to data lakes

Spark program with Spark Delta for processing data and store the processed data into a cloud data lake storage


```
## hadoop inputFile="hdfs://"
spark_df = spark.read.csv(inputFile,header=True,inferSchema=True)
print(spark_df.head(10))
#do many things, before producing data for datalake
spark_df.write.format("delta").mode("append").save(lake_table_path)
```

E.g., Data lake storage based on Google Cloud Storage (<https://delta.io/>)


Data lake storage with Iceberg tables, Pyarrow and Pylceberg
(<https://py.iceberg.apache.org/>)

```
37     catalog = SqlCatalog(
38         catalog_name,
39         **catalog_config["catalog"][catalog_name],
40     )
41
42     if data_type == ".parquet":
43         df = pq.read_table(input_data)
44     else:
45         df = csv.read_csv(input_data)
46     catalog.create_namespace_if_not_exists(namespace)
47     logger.info(f'Existing namespaces: {catalog.list_namespaces()}')
48     full_tablename=f'{namespace}.{table_name}'
49     if not catalog.list_namespaces((namespace)):
50         catalog.create_namespace(namespace)
51     table = catalog.create_table_if_not_exists(
52         full_tablename,
53         schema=df.schema,
54     )
55     table.append(df)
```

Enabling **SQL-style** with parallel/distributed data processing



Where is distributed/parallel data processing needed?



Data Ingestion:

- Spark Streaming
- Kafka Connect
- Apache Nifi
- etc.

Storage:

- HDFS, AWS S3, Google Storage, Azure Data Lake Storage, Iceberg tables, etc., as storage

Programming Frameworks for data processing:

- Apache Spark
- Hadoop MapReduce
- Dask, Ray, etc.


A!

Common principles and techniques

- Data input/output connectors
 - for reading data from sources and writing data into data sinks
- Data collections as abstract (big/distributed) data structures
 - for modeling/representing data in suitable views for processing
- Data partitioning
 - in the view of processing, a similar principle in data storage
- Data operations
 - operations applied to data in data collections
- Execution models
 - tasks and workflows
 - job scheduling; lazy vs eager execution; future execution
- Task fault-tolerance and data exchange among tasks in distributed processes/machines

A!

Data Shuffle



Which operation/analysis could lead to the data shuffle?

Shuffling data from one node to another node for another task is expensive!


- via memory and network?
- via local disk and network?
- via distributed/shared file systems and network?

Which one we shall avoid?

Exchange data possibilities

Try to analyze and identify this problem with a framework of your choice
(and fault management)

Need to
exchange
data



Programming frameworks in our focus

- Programming models
 - OS-based multi-threads/processes
 - embarrassingly parallel programming
 - MapReduce/Spark
 - workflows
 - (distributed) SQL processing with MPP (Massive Parallel Processing)
- Programming frameworks
 - Apache Hadoop/Spark, Dask, Polars, Apache Airflow
- Not in our focus:
 - HPC MPI (Message Passing Interface), NCCL (NVIDIA Collective Communications Library)

e.g., IBIS
<https://ibis-project.org/>

Unified APIs

Programming
Frameworks

Programming
Models

Data processing workload and task graph dependencies with Dask

<https://www.dask.org/>

Key features

- Data input/output connectors
 - file types: CSV, Parquet, HDF5, ORC, Json
 - source: Cloud storage (S3, Google), HDFS, Snowflake, BigQuery, Delta Lake
- Data Collections
 - Array (like numpy array), Bag/Multiset (suitable for unstructured data, like text), DataFrame
- Operations:
 - joins, concatenation, aggregation (`first`, `sum`, ...)
 - grouping/resampling, SQL-alike support
 - functions/computation suitable for arrays
- Execution modes
 - lazy by default; and support specific delayed and future tasks

Key features

- Data can be splitted and processed in parallel tasks
 - many operations on dataframes/tables can be parallelized, with little/without dependency among tasks
 - using directed acyclic graph (DAG) to represent tasks
 - little communication among them, little data shuffle between tasks
- Single and multiple compute nodes for processing
 - multiprocessing in single node vs distributed nodes
 - scheduling graphs using OS threads and processes to execute tasks
 - data exchange among tasks using shared memory, direct communication or network file systems
 - using different resource management systems: Kubernetes, SLURM, PBS, etc.
 - spill data into disks when running out of memory

Parallelizing dataframe → embarrassingly

- A big dataset can be presented as a Dask dataframe
 - a Dask dataframe can be partitioned into different partitions
- Perform operations on data partitions with lazy principles
 - explicitly call `compute()` method
→ computation


Dataframe
in partitions

	VendorID	total_amount
0	1.0	11.80
1	1.0	4.30
2	1.0	51.95
3	1.0	36.35
4	2.0	24.36
..
95	2.0	21.96
96	2.0	17.30
97	2.0	15.36
98	2.0	24.80
99	2.0	13.30

Dask
Dataframe

Data records: 1369765
Data has 4 partitions
Partition 0 has 342441
Partition 1 has 342441
Partition 2 has 342441
Partition 3 has 342442

Example




```
34 from dask.distributed import Client
35 # make sure that dask scheduler and worker running
36 client = Client(f'{dask_scheduler_host}:{dask_scheduler_port}')
37 taxi_df = dd.read_csv(input_file, dtype = dtype,
38                         assume_missing=True,
39                         low_memory=False)
40 print(f'Total records: {len(taxi_df)}')
41 p_taxi_df = taxi_df.repartition(npartitions=num_partitions)
42 func_total_amount = p_taxi_df.get_partition(0)["total_amount"].sum()
43 for i in range(1,num_partitions):
44     func_total_amount = func_total_amount \
45         + p_taxi_df.get_partition(i)["total_amount"].sum()
46 if vis_file is not None:
47     func_total_amount.visualize(filename=vis_file)
48 total_amount =func_total_amount.compute()
49 print(f'The total amount calculated from this file is {total_amount}'')
```

Task dependency based on DAG

- Flexible to define task graphs
 - as a directed acyclic graph
- Explicitly lazy, deferred execution
 - using `dask.delayed()` / `@dask.delayed` to declare delayed tasks
- Concurrent, asynchronous eager execution
 - using future tasks
- Suitable for problems cannot be solved with Dask Dataframe

Task dependency based on DAG

```
60 if delayed_mode:  
61     # delayed tasks  
62     task11 = dask.delayed(high_room_temperature)(bts_alarm_df)  
63     task12 = dask.delayed(average_outdoor_temperature)(bts_parameter_df)  
64     final_task = dask.delayed(merge_data)(task11,task12)  
65     if vis_file is not None:  
66         final_task.visualize(filename=vis_file)  
67     final_result = final_task.compute()  
68     print(f'First 100 elements\n: {final_result.head(100)}')  
69
```



A!

Apache Spark

<https://spark.apache.org/>

Apache Spark

- Cluster-based high-level computing framework
- “unified engine” for different types of big data processing
 - SQL/structured data processing
 - Machine learning
 - Graph processing
 - Streaming processing
- It is a powerful computing framework and system ⇒ an important service for a big data platform
 - public cloud: Google DataProc, Azure HDInsight, Amazon EMR
 - data lake systems: e.g., Hudi and Delta Lake

Apache Spark

Can be run a top

- Clusters of shared-nothing compute nodes with Hadoop (using HDFS and YARN)
- Kubernetes
- A set of standalone machines in a master-worker architecture





Figure source: <http://spark.apache.org/>

Execution model in a cluster system

Driver manages operations and tasks in nodes



Common concepts: **Driver, Node, Task and Executor**

Computing resources in a cluster node





Figure source:
<http://spark.apache.org/docs/latest/cluster-overview.html>

Spark application management: high-level view

- Submission/Request
 - submit the Spark application for running
 - resource is provided for running the Driver
- Launch
 - the Driver requests resources for executors (through SparkContext)
 - establish executors across worker nodes
- Execution
 - the Driver starts to execute code and move data
- Finish/Completion:
 - finish, release executors

Key features




- Input data is **distributed** in different nodes for processing
 - support partitions for data processing: a node keeps one or n partitions, a partition resides only in a node ⇒ for computing
- Key operations: **transformations** and **actions** on data
- Leverage parallel computing concepts to run **multiple tasks**
 - data operation -> task executed by executor
 - parallel tasks, task pipeline, DAG of processing stages
- Persistent data in memory/disk for operations


A!

Transformation operations

- Transformation: instructions about how to transform a data in a form to another form ⇒ it will not change the original data (immutability)
- Only tell what to do: to build a DAG (directed acyclic graph) → a lineage of what to do
- Lazy approach ⇒ real transformations will be done when **action operations are triggered**



**Narrow transformation,
no data shuffle**



**Wide transformation, cross data partitions,
requires a shuffle**

Action operations

- Compute the results for a set of transformations
 - examples: count or average
- Actions: view, collect, write, calculation



Lazy approach: an action triggers execution of transformation operations ⇒ enable various types of optimization

Spark program: programming elements

- **SparkSession**
 - acts as a program driver to manage the execution of tasks
 - `SparkContext`: manages connection to a cluster and internal services
- **Data APIs**
 - low-level Resilient Distributed Dataset (RDD) & shared variables
 - high-level `DataFrames`/`DataSets`
 - load and hold distributed data
 - transformation and action functions
- **ML, graph and streaming functions and pipelines**

Spark program logic: typical steps

- Load data and distribute data
 - data is **immutable** after created
 - data partition in Spark: a partition is allocated in a node
- Perform **transformations and actions** operations
 - *transformations*: build plans for transforming data models
 - *actions*: perform computation on data

Resilient distributed dataset (RDD)

- Low-level data structure
 - collection of data elements partitioned across nodes in the cluster
 - with data sharing, parallel operations, fault-tolerant features
- Create RDD
 - created by loading data from files (text, sequence file) including local file systems, HDFS, Cassandra, HBase, Amazon S3, etc.
- Persist RDD
 - in memory or to files

Example with RDD

```
VendorID,tpep_pickup_datetime,tpep_dropoff_datetime,passenger_count,trip_distance,RatecodeID,store_and_fwd_flag,PULocationID,DOLocationID,payment_type,fare_amount,extra,mta_tax,tip_amount,tolls_amount,improvement_surcharge,total_amount  
2,11/04/2014 12:32:24 PM,11/04/2014 12:47:41 PM,1,1.34,1,N,238,236,2,10,0,0.5,0,0,0.3,10.8  
2,11/04/2014 12:32:24 PM,11/04/2014 12:47:41 PM,1,1.34,1,N,238,236,2,10,0,0.5,0,0,0.3,10.8  
2,11/04/2014 12:25:53 PM,11/04/2014 12:29:00 PM,1,0.32,1,N,238,238,2,4,0,0.5,0,0,0.3,4.8
```

as a text file

```
conf = SparkConf().setAppName("cse4640-rddshow").setMaster(args.master)
sc = SparkContext(conf=conf)
##modify the input data
rdd=sc.textFile(args.input_file)
## if there is a header we can filter it otherwise comment two lines
csvheader = rdd.first()
rdd = rdd.filter(lambda csventory: csventory != csvheader)
## using map to parse csv text entry
rdd=rdd.map(lambda csventory: csventory.split(","))
rdd.repartition(1)
rdd.saveAsTextFile(args.output_dir)
```

A!

Shared variables

- A function is executed a remote and various tasks running in parallel
 - how do tasks share variables? common patterns in parallel computing:
broadcast and global variable/counter
- Variables used in parallel operations
 - variables are copied among parallel tasks
 - shared among tasks or between tasks and the driver
- Types of variables
 - broadcast variables: cache a value in all nodes
 - accumulators: a global counter shared across processes

Examples

```
conf = sparkContext.getConf().setMaster("local[4]").setAppName("myapp")
sc = SparkContext(conf=conf)
bVar = sc.broadcast([5,10])
print("The value of the broadcast",bVar.value,sep=" ")
counter = sc.accumulator(0)
sc.parallelize([1, 2, 3, 4]).foreach(lambda x: counter.add(bVar.value[0]))
print("The value of the counter is ",counter.value,sep=" ")
```

Use cases:

- Broadcast variables: lookup tables
- Accumulators: monitoring/checkpoint counters

Spark SQL and DataFrames

- High-level APIs
 - design with common programming patterns in data analysis, multi-language support
- SparkSQL: enable dealing with structured data
 - SQL query execution, Hive, JDBC/ODBC
- DataFrame
 - distributed data organized into named columns, similar to a table in relational database
 - Pandas and Spark DataFrames have similar design concepts

DataFrame

```
inputFile =args.input_file
df =spark.read.csv(inputFile,header=True,inferSchema=True)
print("Number of partition",df.rdd.getNumPartitions())
df.show()
```

PROVINCECODE	DEVICEID	IFINDEX	FRAME	SLOT	PORT	ONUINDEX	ONUID	TIME	SPEEDIN	SPEEDOUT
N 10	23 26		1	2	7	39 100560530		'08/2019 00:04:07	148163	49018
N 10	23 26		1	2	7	38 100560530		'08/2019 00:04:07	1658	1362
N 10	23 26		1	2	7	9 100560530		'08/2019 00:04:07	6693	5185
N 10	23 26		1	2	7	8 100560530		'08/2019 00:04:07	640	544
N 10	23 26		1	2	7	11 100560530		'08/2019 00:04:07	118	114
N 10	23 26		1	2	7	10 100560530		'08/2019 00:04:07	28514	12495
N 10	23 26		1	2	7	13 100560530		'08/2019 00:04:07	868699	23400
N 10	23 26		1	2	7	15 100560530		'08/2019 00:04:07	1822	1120
N 10	23 26		1	2	7	17 100560530		'08/2019 00:04:07	998069	117345
N 10	23 26		1	2	7	16 100560530		'08/2019 00:04:07	22402	1804
N 10	23 26		1	2	7	19 100560530		'08/2019 00:04:07	640	791
N 10	23 26		1	1	10	49 100560530		'08/2019 00:04:07	662	494
N 10	23 26		1	1	10	48 100560530		'08/2019 00:04:07	2158	759
N 10	23 26		1	2	7	21 100560530		'08/2019 00:04:07	0	0
N 10	23 26		1	1	10	51 100560530		'08/2019 00:04:07	2600890	54153
N 10	23 26		1	2	7	20 100560530		'08/2019 00:04:07	330	184

A!

Create DataFrame

DataFrames can be created from a Hive table, from Spark data sources, or another DataFrame

Load and save

- From Hive, JSON, CSV
- HDFS, cloud object storage (AWS S3, Google Cloud Storage, Azure Blob Storage), Delta Lake, local files, etc.



Formats and Sources supported by DataFrames

Figure source:

<https://databricks.com/blog/2015/02/17/introducing-dataframes-in-spark-for-large-scale-data-science.html>

DataFrame Transformations & Actions

Transformations

- Several operations, think transformation for a relational table or a matrix

Example

- Select: `df.select`
- Filter: `df.filter`
- Groupby: `df.groupBy`
- Handle missing data
 - Drop duplicate rows, drop rows with NA/null data
 - Fill NA/null data

Actions

- Return values calculated from DataFrame

Examples

- `reduce`, `max`, `min`, `sum`, `variance` and `stdev`

⇒ Distributed and parallel processing but it is done by the framework

Example of a Spark program

```
#!/usr/bin/env python2
#encoding: UTF-8
# CS-E4640
import csv
import sys
from datetime import datetime
from pyspark.sql import SparkSession
import numpy as np
from pyspark.sql import functions as F
import argparse
parser = argparse.ArgumentParser()
parser.add_argument('--input_file', help='input data file')
parser.add_argument('--output_dir', help='output dir')
args = parser.parse_args()

##define a context
spark = SparkSession.builder.appName("cse4640-onu").getOrCreate()
#NOTE: using hdfs:///..... for HDFS file or file:///
inputFile = args.input_file
df = spark.read.csv(inputFile,header=True,inferSchema=True)
#df.show()
print("Number of records", df.count())
exprs = {"SPEEDIN": "avg"}
df2 = df.groupBy('ONUID').agg(exprs)
df2.repartition(1).write.csv(args.output_file,header=True)
```


Session/Driver

Read data


Apply operations

Spark application runtime view

- Tasks:
 - a unit of work executed in an executor: e.g., performing transformations of a data partition
- Stage: Shuffle Map Stage & Result Stage
 - a set of tasks executed in many nodes for performing the same operation which does not lead to a data shuffle
 - move to a new stage: through a shuffle to produce output partitions or an action to produce results
- Job
 - runtime view of an action operation (actual computation produces a result), includes many stages of tasks



Data Distribution




One task works on a partition at a time

⇒ Parallelism and performance are strongly dependent on number of partitions, tasks, CPU cores

Data Distribution: load balance

Imbalance processing

more data shuffle



- It is important to have well-balanced data distribution across nodes
- Detection:
 - look at runtime execution time to see problems or check your data
- Examples of solution:
 - repartition
 - change group keys

Pipelining, Shuffle and DAG

- Operations work in a pipeline without moving data across nodes
 - e.g., map → filter, select → filter
- Shuffle persistent
 - shuffle needs move data across nodes
 - source tasks save shuffle files into local disks for data shuffle, then the target tasks will read data from source nodes
 - Save time, recovery, fault tolerance

A!

Massive parallel processing for distributed query engines

Massive parallel processing employed by distributed query engine

- Key concepts
 - using SQL as a way to query different types of data sources like data lake, warehouse, and databases
 - managing catalogs and schemas about data sources
 - the query engine is decoupled from data sources/storage
 - using massive parallel processing (MPP) to support parallel tasks accessing different data sources at a large-scale with many compute nodes
- Complex fault tolerance and optimization:
 - failure management, query and data movement costs, ...
- Mostly for analytics: interactive analytics, seconds – minutes

Example distributed SQL engine: Presto/Trino





Figure source: Presto: SQL on Everything
<https://ieeexplore.ieee.org/stamp/stamp.jsp?tp=&arnumber=8731547&tag=1>

Small exercise
(see course git)



Trino (<https://trino.io/>): from a fork of Presto

Summary

- Different programming models for data processing
 - models and tools selected based on data workload and ecosystems, including underlying compute resource management
 - both developers and platform operator/provider must carefully decide the programming models for data processing
- Effects of modernization and composability in data platforms
 - Spark is powerful but many emerging ones, e.g., Polars and DuckDB, which may be suitable due to learning curves, management, and data load
- Rich ecosystems
 - combine data, different distributed programming supports for big data platforms

Thanks!

Hong-Linh Truong
Department of Computer Science

rdsea.github.io



—
Kiitos
aalto.fi