



Aalto University  
School of Science

# Stream Processing and Big Data Platforms

*Hong-Linh Truong*

*Department of Computer Science*

*[linh.truong@aalto.fi](mailto:linh.truong@aalto.fi)*, *<https://rdsea.github.io>*

# Learning objectives

- **Understand fundamental concepts and techniques in stream processing in big data**
- **Able to design streaming analytics in big data platforms and applications**
- **Able to select and use common stream processing frameworks**

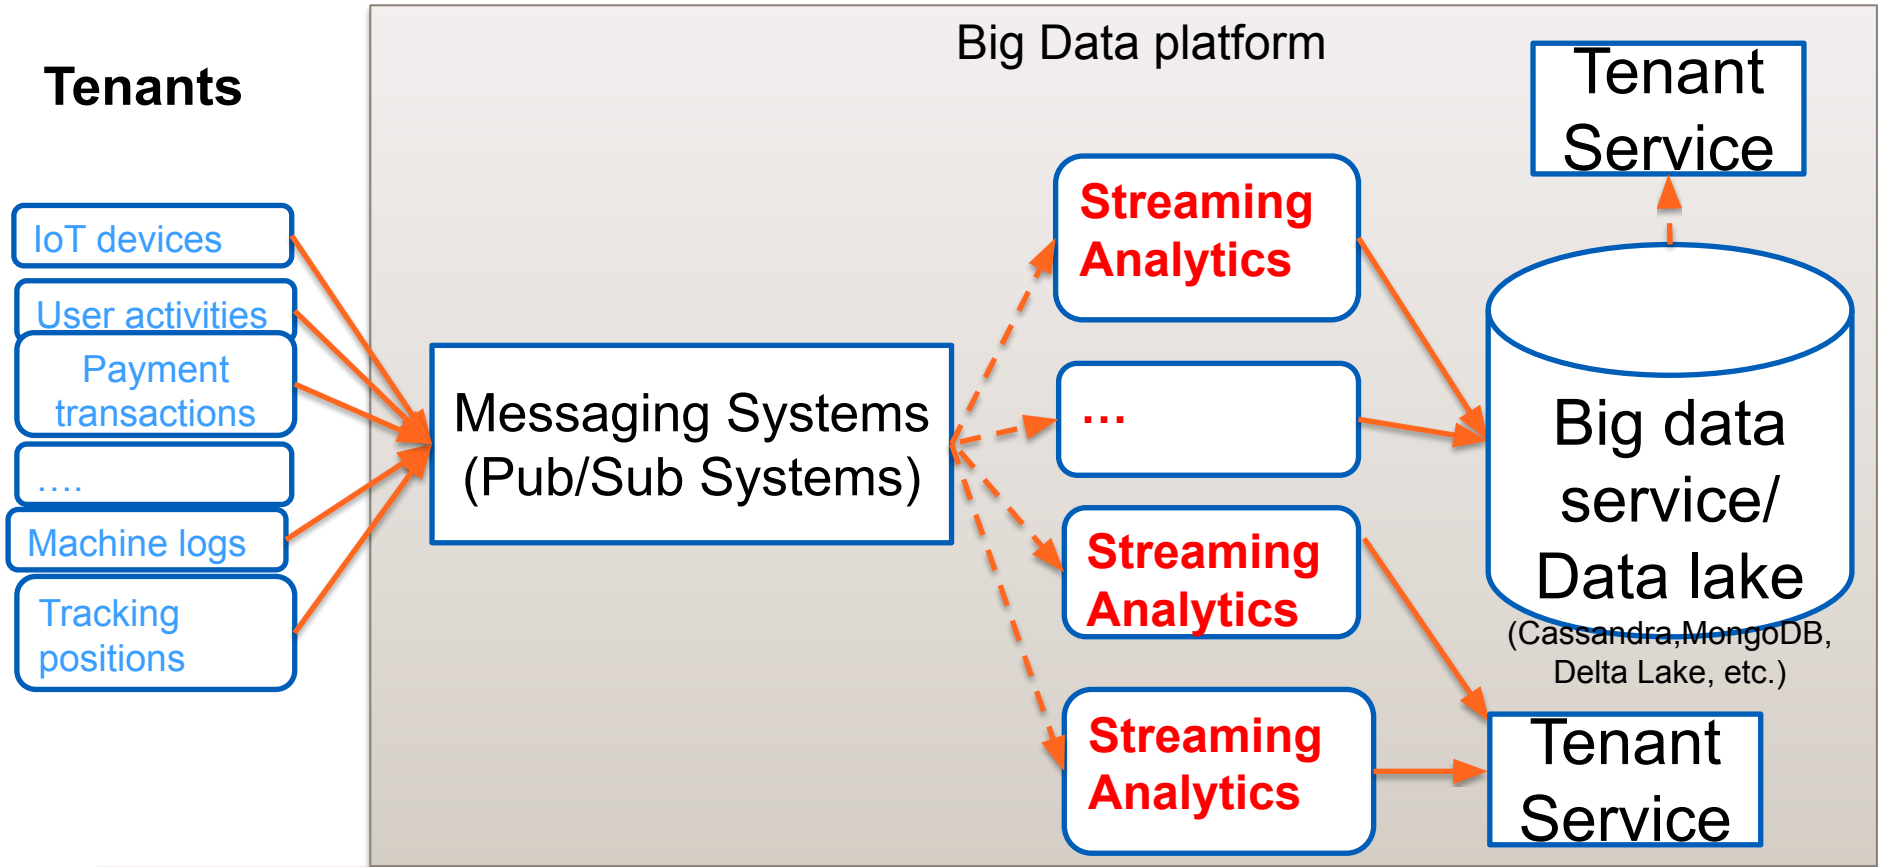


# Stream analytics for data in motion

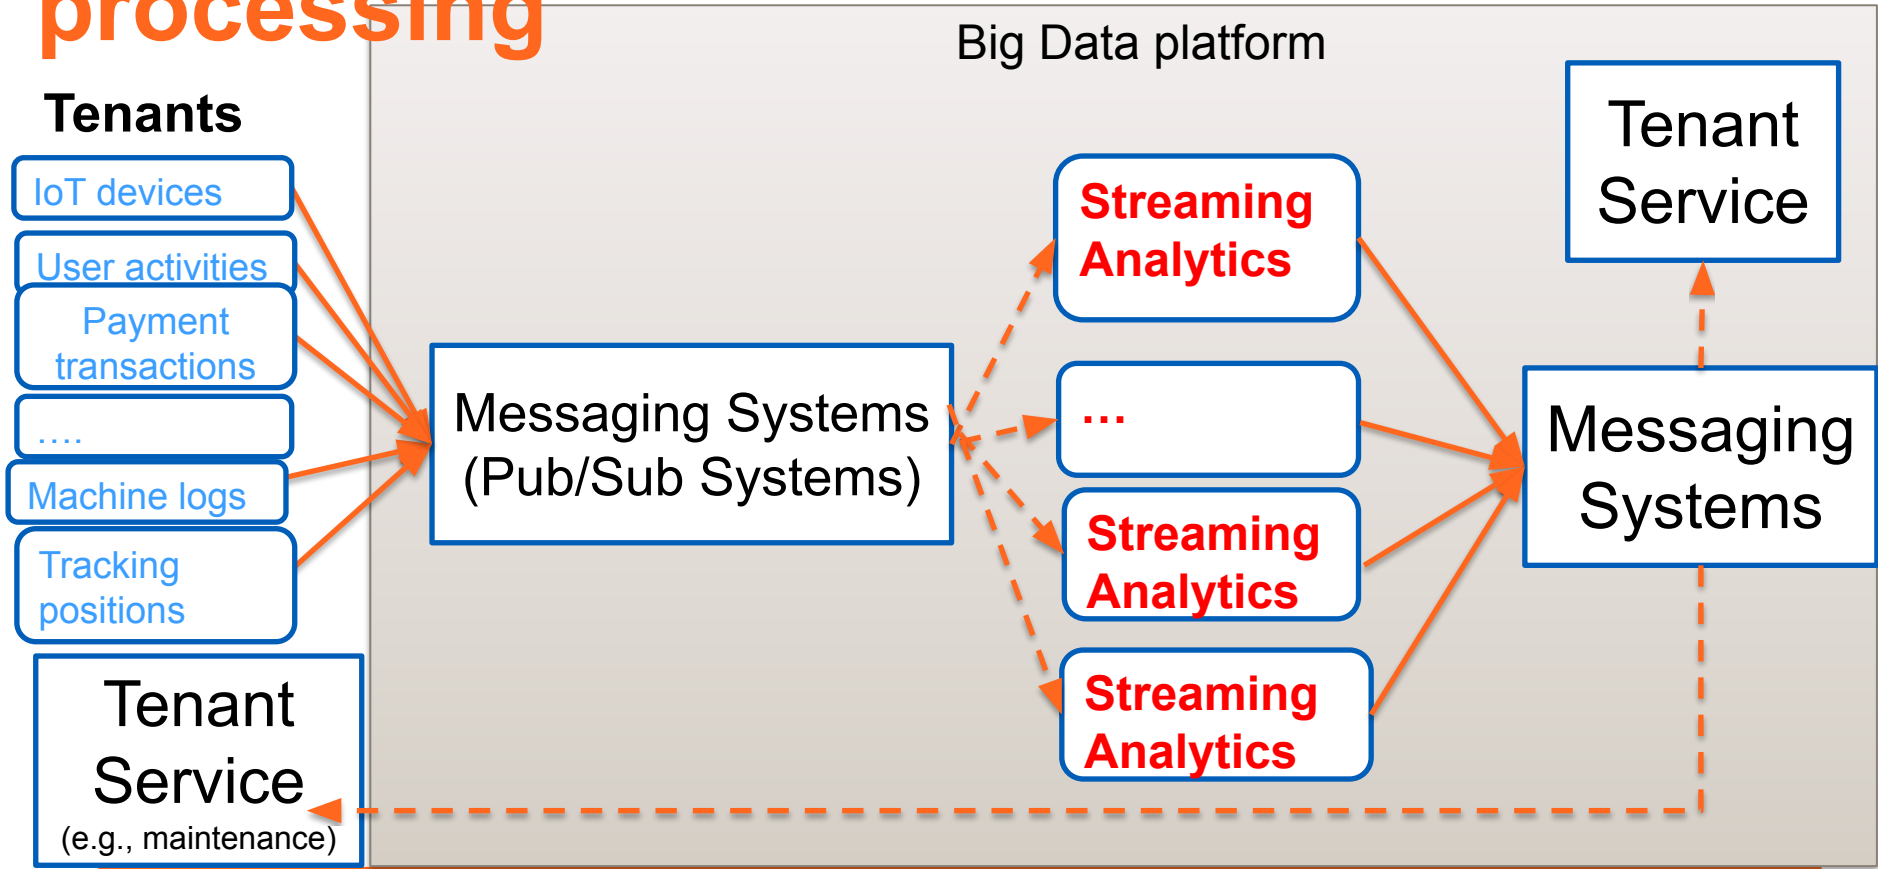
# Stream processing in big data

- **Big data coming from streams at near real-time**
  - the data element/unit may be “small” but voluminous and delivered in a near real-time manner
  - high and volatile throughput, but low processing time expected
  - not just *“take a record and store it into a database”*
- **Require large-scale computing infrastructures and many other platform services**
  - *task parallelism*: multiple tasks for processing data
  - *data parallelism*: data is partitioned into concurrent/parallel data streams  
⇒ distributed, parallel processing tasks
  - *stateful analytics*: processing needs state information across multiple data records and time

# Near real-time streaming data processing



# Near real-time streaming data processing



# Example in the cloud – Azure stream analytics for stream processing and big data platforms

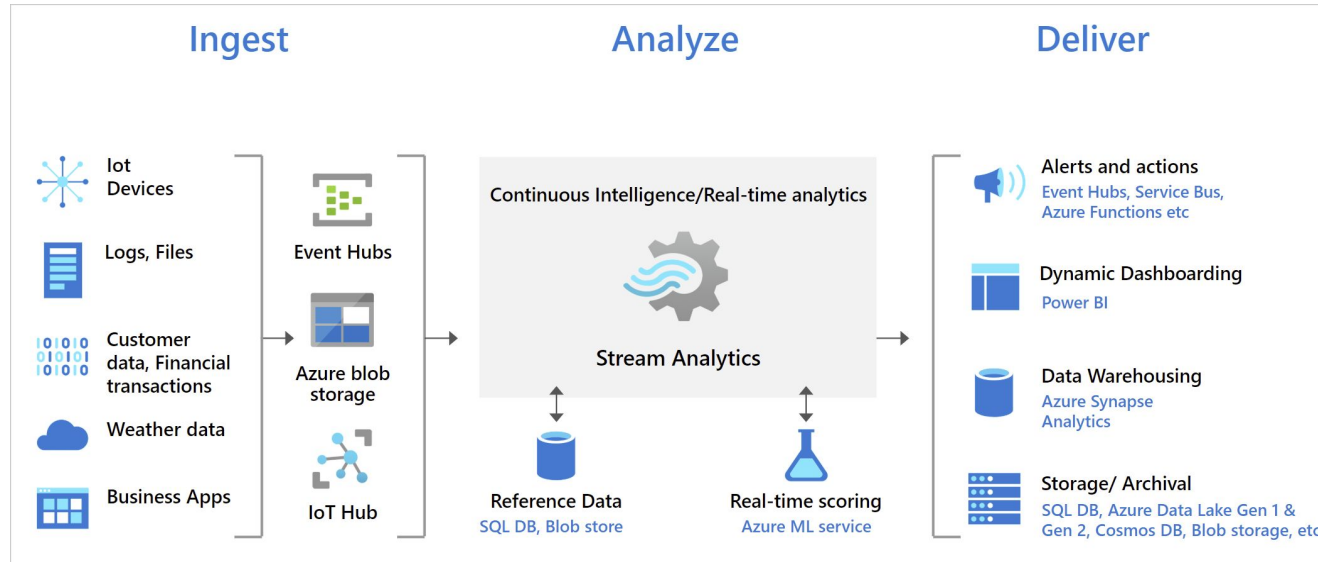


Figure source: <https://docs.microsoft.com/en-us/azure/stream-analytics/stream-analytics-introduction>

Known public cloud services: Amazon Kinesis, Google Dataflow, Alibaba Cloud DataHub



# From complex event processing (CEP) in the age of enterprise computing



Esper CEP



Our practices focus on modern technologies like: Apache Flink, Apache Spark and Arroyo, which are used intensively in business systems and big cloud platforms

# Stream processing and big data platforms

- **Stream processing services as a component of data platforms**
  - a big data technology for pre-processing, ingestion and high-level analytics, including near-real time machine learning
- **Stream processing services as data platforms**
  - a big data platform offers mainly stream processing services for streaming analytics
  - analytics on the fly as the first class feature
  - e.g., IoT analytics, e-commerce user activities, fraud detection, real time AI/ML

# Stream Processing – key concepts

# Common building blocks

- **The way to connect data to streams and obtain data records (messages) from the streams**
  - focusing very much on *connector concepts* and well-defined message structures (JSON, Avro, customized binary format, etc.)
  - connectors implement complex data handling mechanisms (low level session management, message retainment, delivery quality of service)
- **The way to specify/program the “analytics” logic**
  - *analytics functions, statements* and how they are glued together to process flows of messages
  - high-level, easy to use
- **The distributed engine to process analytics tasks**
  - handle complex task processing atop multiple compute nodes
- **The way to push the result to external components (sink databases, new streams, files)**

# Data stream programming

**Data stream:** a sequence/flow of data units


Data units are defined by applications: a data unit can be data described by a primitive data type or by a complex data type, a serializable object, etc.

**Streaming data:** produced by (near)realtime data sources as well as (big) static data sources  $\Rightarrow$  *unbounded* and *bounded*

- Examples of data streams
  - Continuous media (e.g., video for video analytics)
  - Discrete media (e.g., stock market events, twitter events, system monitoring events, comments, notifications, log records)

# Messages of events/data records

- Messages encapsulating real-world events, data records and other types of data
- Data to be sent/processed can be in a simple or complex structure



We focus on **unbounded discrete** messages of data

# Message representations and streams

- **Data Sources**

- via message brokers, databases, websocket, different IO adapters/connectors, etc.

- **Data Sinks**



- messaging systems, databases, file storage/systems (S3, HDFS), etc.

- **Data representations**



- text/CSV, JSON, Arvo format, etc.
- serialization and deserialization (short name: SerDe) are required
- data format validation
- data schema registry for registered schemas

# Publisher view: how messages are published

## Messaging system



**topic=queue;  
no partition**




**topic = n partitions = n  
queues**

Topic and topic partitions




# Handling messages for consumption (processing)

Messages in systems




simple straight forward



fan-out/broadcast


**different mechanisms for "routing"**

complex mapping/routing



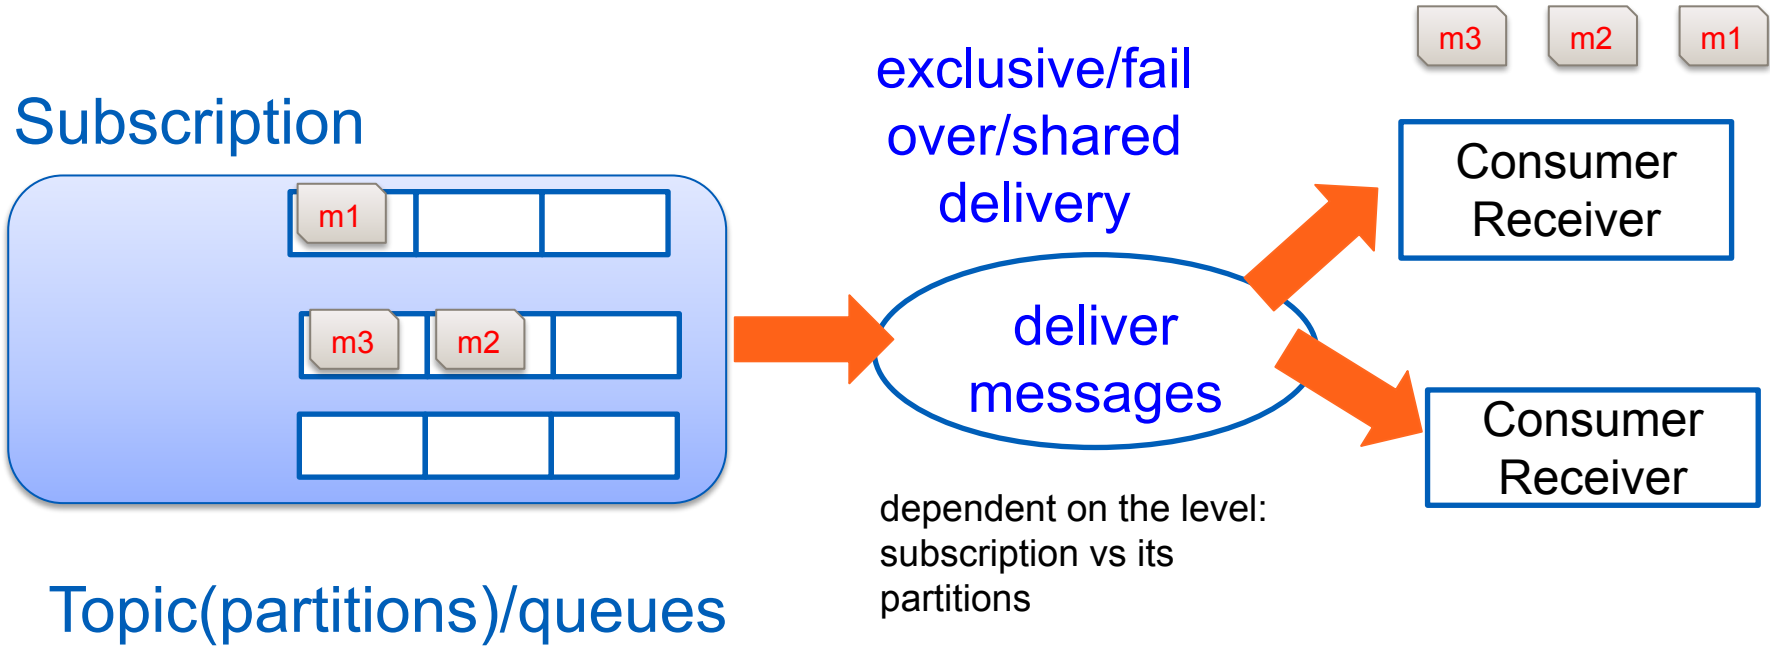
allowing parallel processing of the same messages

**messages for consumption**



parallel processing of different messages

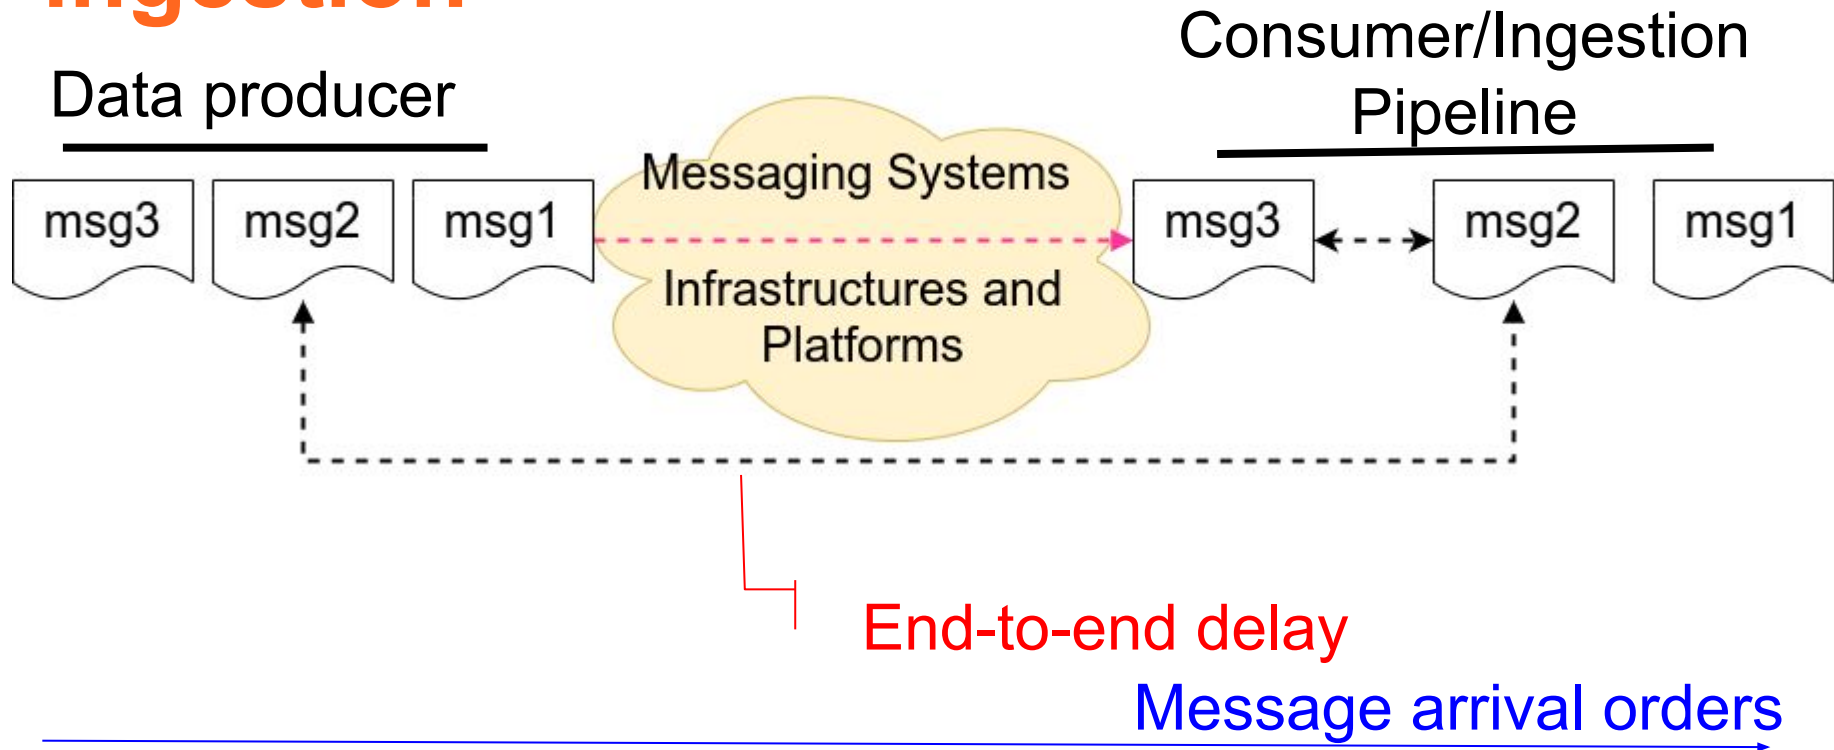
# Consumer view in accessing messages: subscription and delivery



# Some key issues

- **Data order & delivery**
  - late data, out of order data
- **Times associated with messages and processing**
- **Data parallelism**
  - key-based data processing
- **Task parallelism**
  - stateful vs stateless processing

# Key issues in streaming data ingestion

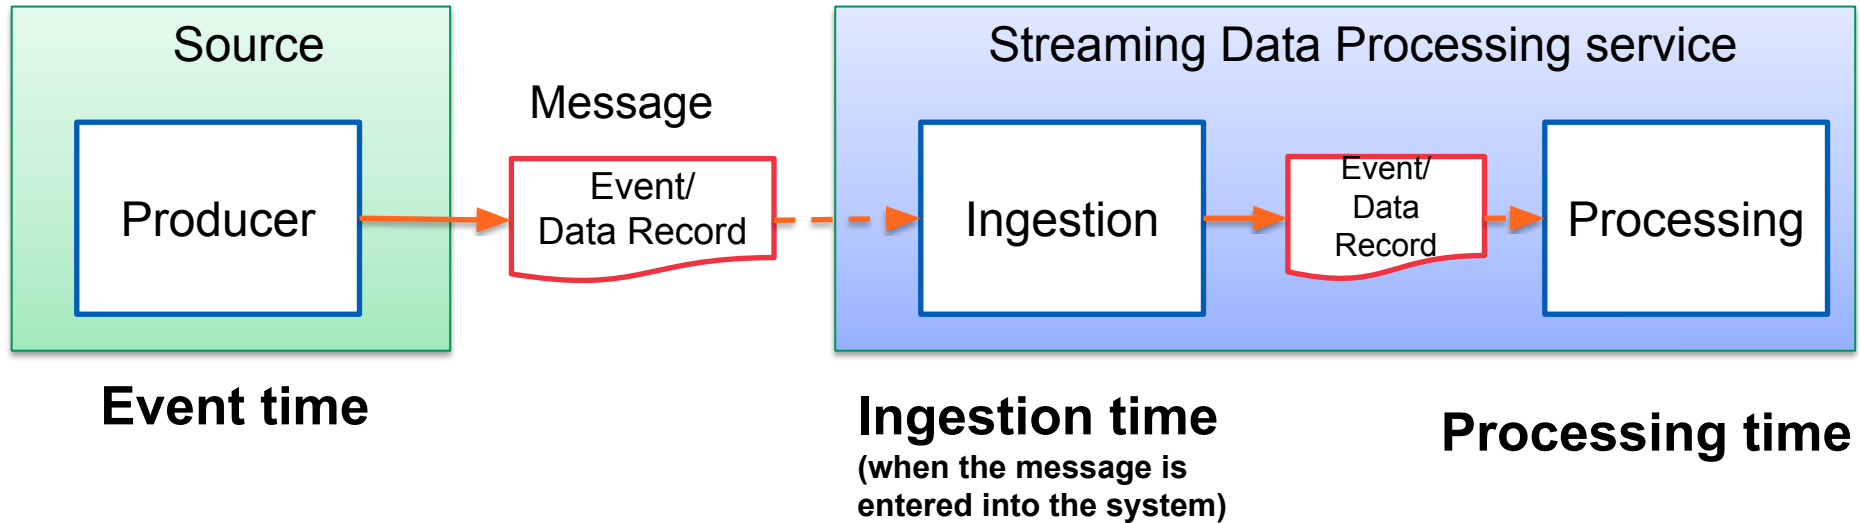


**Without a timestamp associated to a message, do we know the delay or out of order?**

**What is the consequence of delay/out of order for processing?**


# Key issues in streaming data: the notion of times

## Times associated with data and processing



**Which time is important for analytics (from business viewpoint)?**

# Data parallelism: partition stream data based on some keys for analytics



Can be processed in parallel using multiple nodes


With **keyed data**: enable parallel processing based on the keys

```
12 //  
13 public class BTSAlarmEvent {  
14     public String station_id; ← keys?  
15     public String datapoint_id; ←  
16     public String alarm_id; ←  
17     public Date event_time;  
18     public float value;  
19     public float valueThreshold;  
20     BTSAlarmEvent() {  
21
```

# Windows of data

Window is used to group data for processing:

Which constraints are used to determine a window?



a stream of events

sliding/tumble window size: period of time or number of events/records

Arrival order



# Windowing

- **Windows size:**
  - time or number of records
- **Tumbling window:**
  - identified by size, no gap between windows
- **Sliding window:**
  - identified by size and a sliding interval
- **Session Window:**
  - identified by “gap” between windows (e.g., the gap of events is used to mark “sessions”)

# Functions applied to Windows of data

If we

- **specify a set of conditions**  $\Rightarrow$  windows will be created according to the conditions to store message in corresponding windows

then we can

- **Apply functions to messages in** the window that match these conditions

**Task parallelism: we can have a lot of such functions executed in parallel in multiple compute nodes**

# Functions

- **Can be simple or complex!**
  - built-in and user-defined functions
- **Core for analytics and ML**
- **Examples**
  - individual threshold/alarm based alerting, atypical events monitoring
  - data rollup
  - anomaly detection based on statistical functions, like quantile/T-digest, ...
  - real time AI/machine learning

# Example

**Monitoring working hours of (taxi/truck) drivers (assume events about pickup/drop captured at near real-time):**

- Windows: **12 hours**
- Partitioning data/Keyed streams: **licenseID**
- Function: determine **working and break times and check with the law/regulation**

**Source:**

<https://www.infoworld.com/article/3293426/how-to-build-stateful-streaming-applications-with-apache-flink.html>


What if events/records come late into the windows?

Do we need to deal with late, out of order events/records?

*correctness and completeness* issues

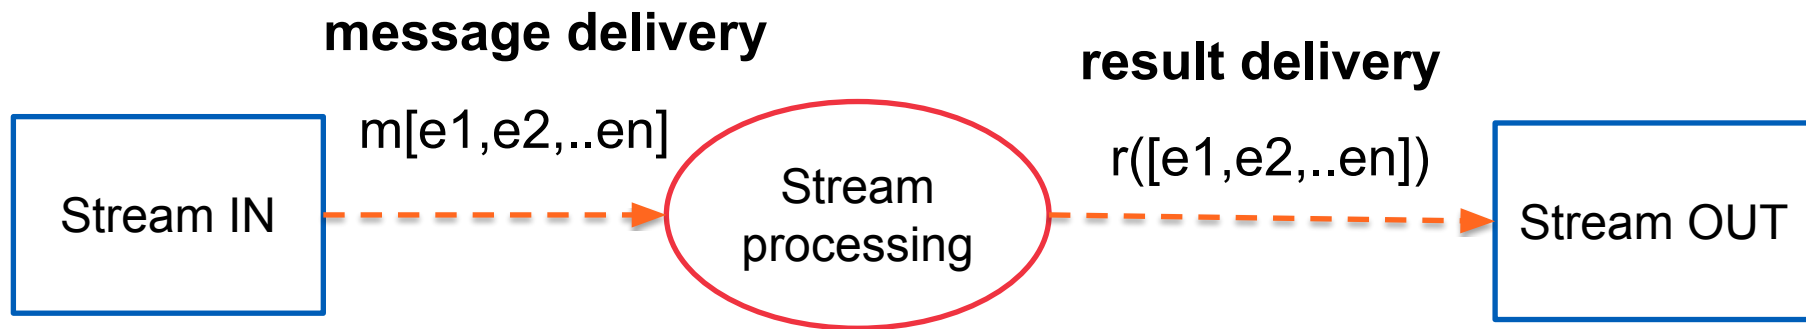
# Support lateness

- Identify timestamp of events/data records
- Identify watermark in streams
  - a watermark is a timestamp
  - a watermark indicates that no events which are older than the watermark should be processed
  - enable the delay of processing functions to wait for late events
- Using watermark to ignore late data  $\Rightarrow$  computing under “**incompleteness** assumption”



# Delivery guarantees

Exactly once? at least once? or at-most-once  
End-to-end?



**What if the stream processing fails and restarts**

# Message and processing guarantees

- **Message guarantees are the job of the broker/messaging system**
- **Processing guarantees are the job of the stream processing frameworks**
- **They are highly connected if messaging systems and processing frameworks are tightly coupled (e.g., Kafka case)**



# End-to-end exactly once

- **Exactly once for processing is not enough**
- **Messaging systems must support**
  - redeliver messages/data, recoverable data
- **Sink and output must support exactly one**
  - idempotent results, roll back
- **Coordination among various components**

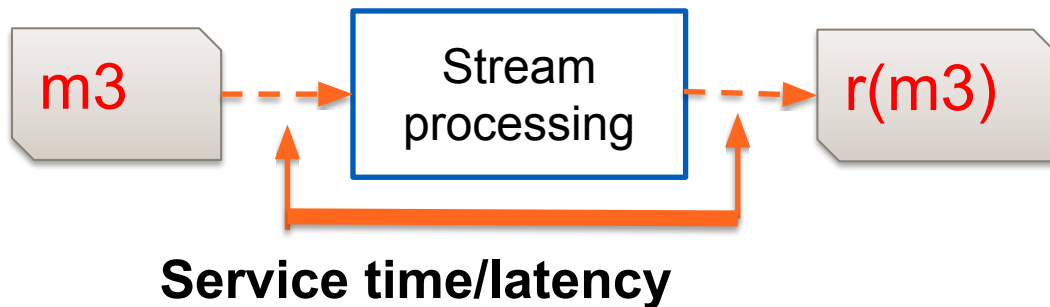

Further reading:

<https://flink.apache.org/features/2018/03/01/end-to-end-exactly-once-apache-flink.html>

<https://www.confluent.io/blog/simplified-robust-exactly-one-semantics-in-kafka-2-5/>

<https://docs.microsoft.com/en-us/azure/hdinsight/spark/apache-spark-streaming-exactly-once>

# Performance metrics



# Latency and throughput

- **Service latency**

- subseconds! e.g., milliseconds
- max, min or percentile?  $\Rightarrow$  up to application requirements

- **Throughput**

- how many messages can be processed per second?

- **Goal: low latency and high throughput!**

# Structure of streaming data processing programs (1)

- **We have multiple streams of data, different functions for processing data, multiple computing nodes**
- **Data exchange between tasks**
  - links in task graphs reflect data flows
- **Stream processing**
  - centralized or distributed (in terms of computing resources)
  - simple functions vs complex ones

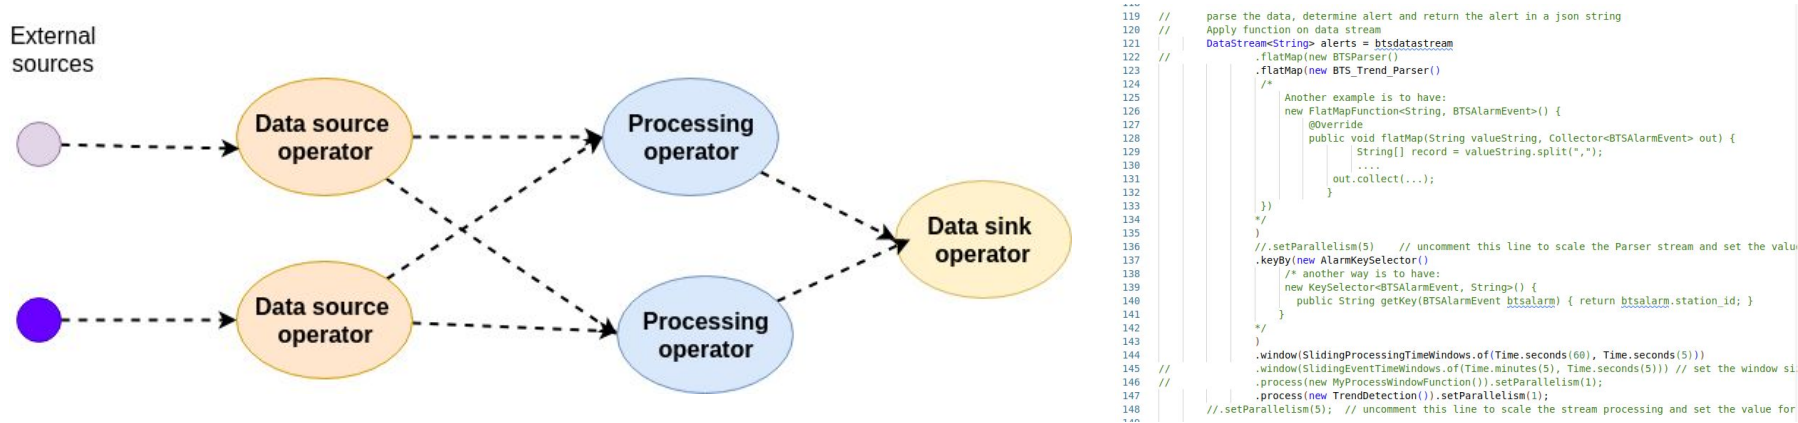
# Structure of streaming data processing programs (2) - examine a simple example

```
126 WAIT AND PROCESS DATA
127 '''
128 while True:
129     '''
130     Receive the data from source
131     '''
132     msg = consumer.receive()
133     '''
134     when should we do this?
135     consumer.acknowledge(msg)
136     '''
137     try:
138         '''
139         MAIN TRANSFORMATION, HERE IS WITH A FUNCTION
140         '''
141         ## assume that the selected data schema is json
142         result = dt_process_json_style(msg, op_processor)
143         ##store the result to the right data sink
144         dt_store_to_sink(result)
145     except Exception as ex:
146         logging.warn(f'{ex}')
147         logging.info("Continue to wait")
148 '''
```

**How to handle  
possible errors**

Note: Example with an external Pulsar consumer for data transformation

# Structure of streaming data processing programs (3)

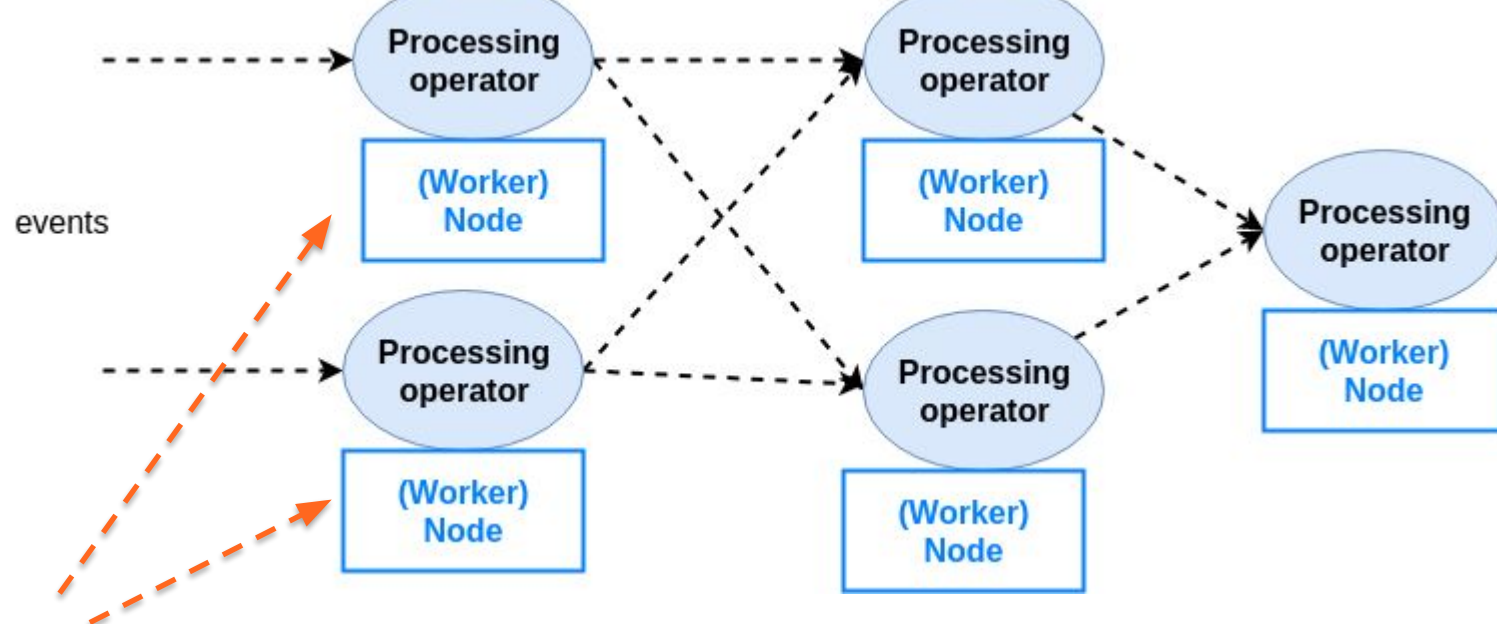


- **Dataflows:**

- *Data source operators: represent sources of streams*
- *Processing operators: represent processing functions*

# Distributed processing topology in a cluster

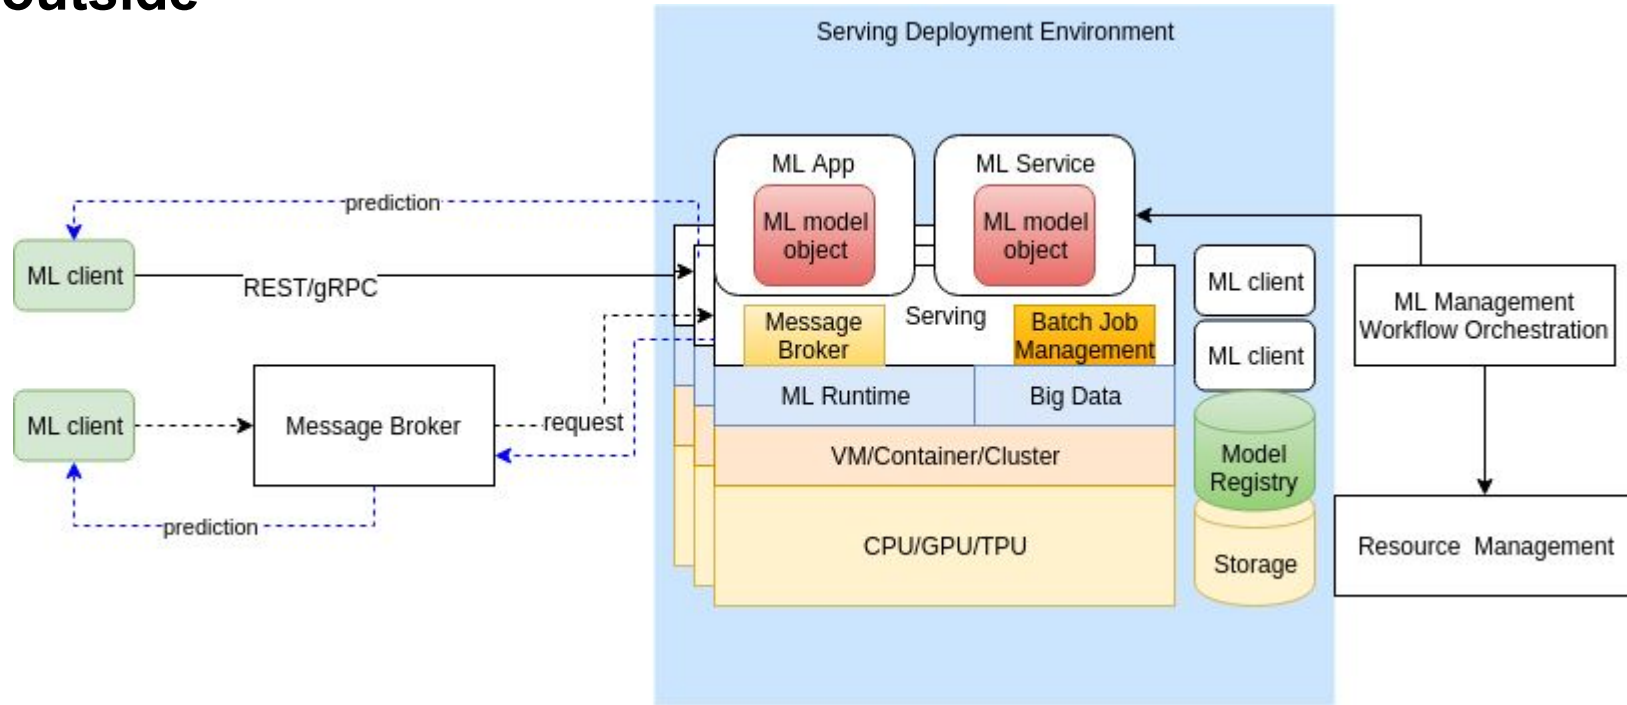
A graph of tasks (running operators); all tasks are running



Nodes of a cluster (VMs, containers, Kubernetes)

# Integration with ML

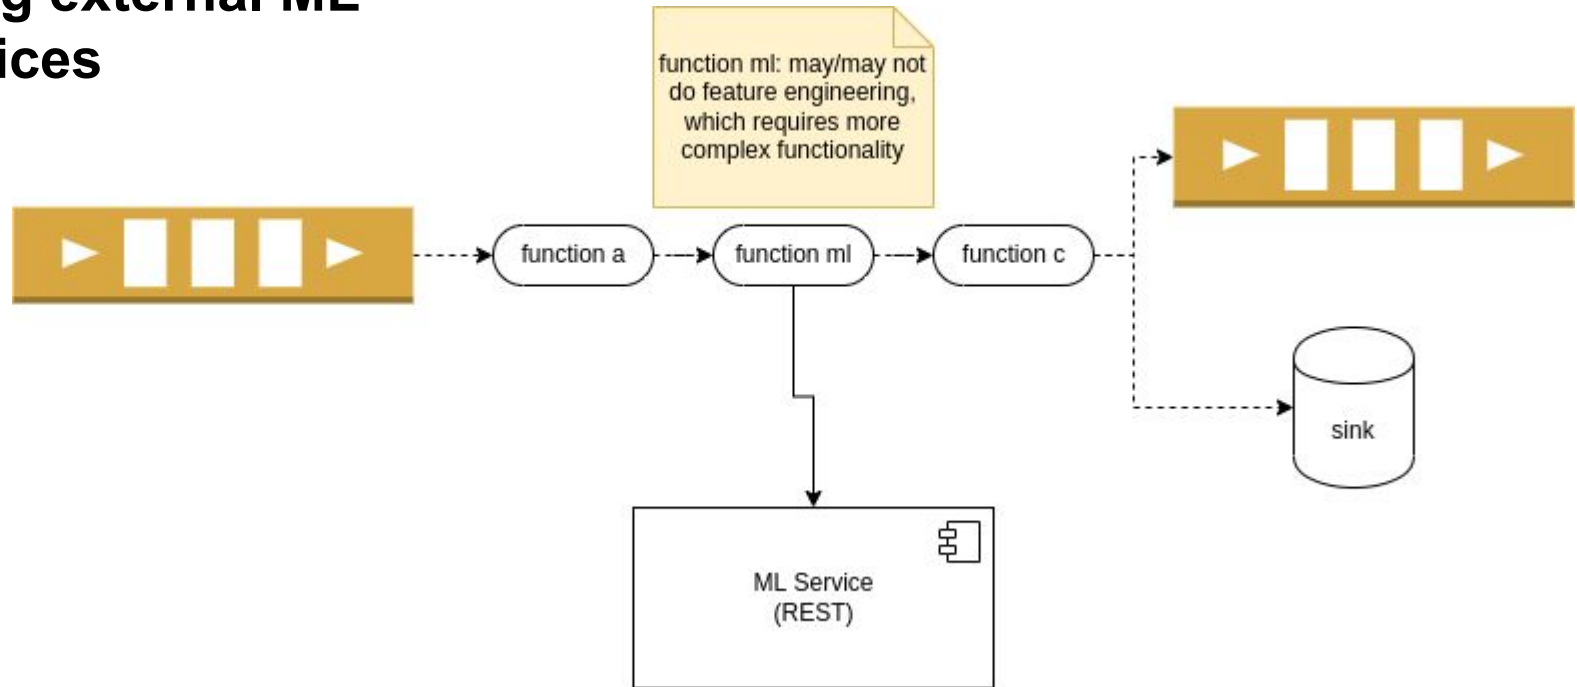
## From outside





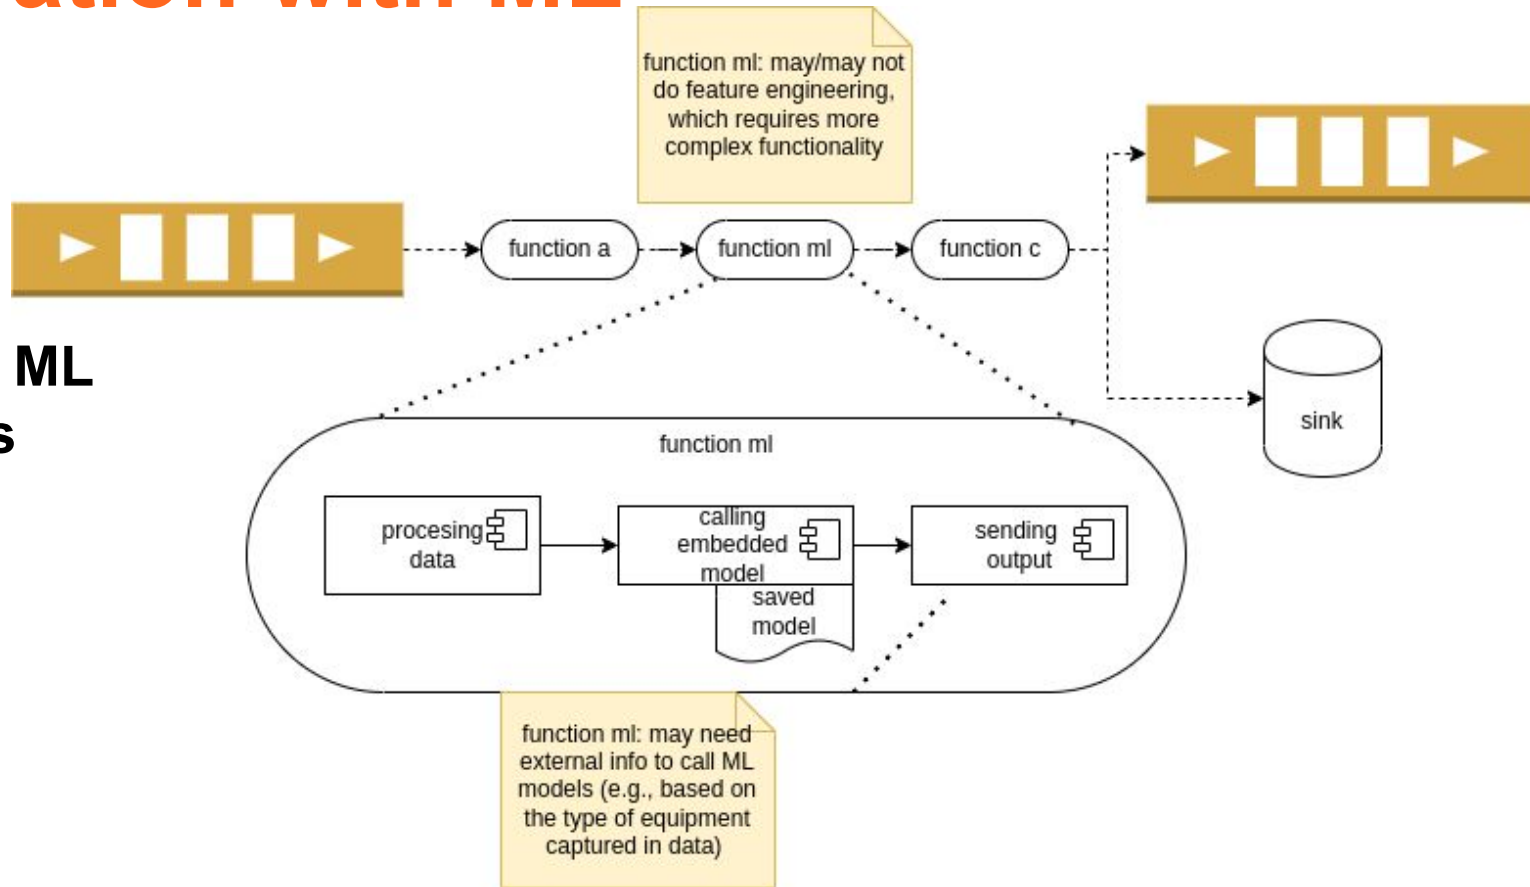
# Integration with ML

## Using external ML services




# Integration with ML

## Embedded ML capabilities



# Distributed, composable processing topologies in cross distributed sites



# Common concepts in existing frameworks - programming level

- **How to write streaming program?**
- **With programming languages**
  - low level APIs
  - DSL
  - Java, Scala, Python (Spark, Flink, Kafka)
- **High-level data models**
  - KSQL
- **Flow/pipeline description**
  - Node-RED/GUI-based flow editors

# Common concepts in existing frameworks - key common concepts

- **Abstraction of streams**
- **Connector library for data sources/sinks**
  - very important for application domains
- **Runtime elasticity**
  - add/remove (new) operators
  - add/remove underlying computing nodes
- **Fault tolerance**

# Where do you find most of concepts that we have discussed

- **Apache Storm**
  - <https://storm.apache.org/>
- **Apache Spark (Structured Streaming)**
  - <https://spark.apache.org/>
- **Apache Kafka Streams and KSQL**
  - strongly bounded to Kafka messaging
- **Apache Flink (Stream Analytics)**
  - native, clustered, better data sources/sinks
- **Apache Beam (<https://beam.apache.org/>)**
  - unifying programming models for batch and stream processing

# Practical learning paths

- **Path 1: if you don't have a preference and need challenges**
  - Apache Flink Stream API (e.g., with RabbitMQ/Kafka connectors)
- **Path 2: many of you have worked with Kafka**
  - Kafka Streams DSL (everything can be done with Kafka)
- **Path 3: for those of you who are working with Spark (and Python is the main programming language)**
  - Apache Spark Structured Streaming
- **Path 4: for those who deal with MQTT brokers**
  - Apache Storm (but also Kafka, ...): Spout and Bolt API or Stream API

# Summary

- **Focus:**

- Practical programming with one of the stacks:
  - *Apache Flink Stream API (with different connectors)*
  - *Kafka Streams*
- Check the common concepts in other tools/systems

- **Action:**

- Work on use cases where you can use stream analytics (as a user/developer)  $\Rightarrow$  there are many interesting analytics
- Provision services for stream processing (as a platform)



# Thanks!

**Hong-Linh Truong**  
**Department of Computer Science**

**rdsea.github.io**