



Aalto University
School of Science

Apache Flink for Stream Processing in Big Data Platforms

Hong-Linh Truong

Department of Computer Science

linh.truong@aalto.fi, *<https://rdsea.github.io>*

Stateful computations

- **Many computations on streaming data need to “memorize” data and use such data for computations**
 - calculate the number of records over a period of time
 - detect anomalies in data

⇒ **stateful computations**

- **Requires to manage and store data for real-time processing**
- **Apache Flink supports stateful computations with streaming data at very large scale, high performance, low latency**

Some impression

Four Billion Records per Second! What is Behind Alibaba Double 11 - Flink Stream-Batch Unification Practice during Double 11 for the Very First Time

Alibaba Clouder

December 2, 2020

7,274


0

Amazon Managed Service for Apache Flink

Build and run fully managed Apache Flink applications

Get started

Request more information



Apache Flink




Figure source: <https://flink.apache.org/>

Flink runtime view

- Parallelism
- Checkpointing
- Monitoring

**Remember:
stream applications
running 24/7**

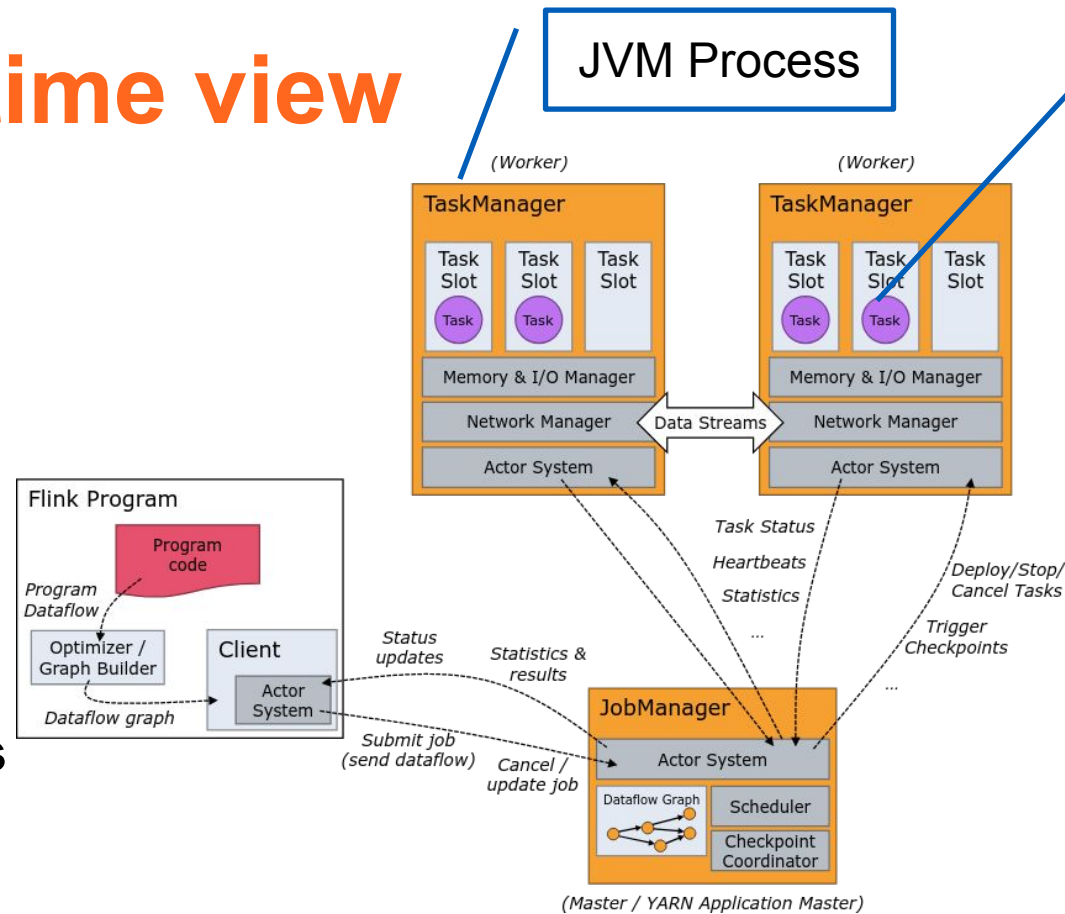



Figure source: <https://nightlies.apache.org/flink/flink-docs-release-1.20/docs/concepts/flink-architecture/>

Main elements in Flink applications



- Rich set of sources and sinks via many connectors

Connectors

- **Major systems in big data ecosystems**
- **We have used many of them in our study**
 - Apache Kafka
 - Apache Cassandra
 - Elasticsearch (sink)
 - Hadoop FileSystem
 - RabbitMQ
 - Apache NiFi
 - Google PubSub

Main programming primitives

- **Setting environments**
- **Handling inputs and outputs via data streams**
- **Key functions for processing data**
- **Stream processing flows**



Bounded and unbounded streams

Stream processing flows

Split streaming data into different windows using a key for analytics purposes

Keyed data/Keyed window: if we can separate data via keys

```
stream
  .keyBy(...)           <- keyed versus non-keyed windows
  .window(...)          <- required: "assigner"
  [.trigger(...)]       <- optional: "trigger" (else default trigger)
  [.evictor(...)]       <- optional: "evictor" (else no evictor)
  [.allowedLateness(...)] <- optional: "lateness" (else zero)
  [.sideOutputLateData(...)] <- optional: "output tag" (else no side output for late data)
  .reduce/aggregate/apply() <- required: "function"
  [.getSideOutput(...)] <- optional: "output tag"
```

Source: <https://nightlies.apache.org/flink/flink-docs-master/docs/dev/datastream/operators/windows/>

Stream processing flows

Handling streaming data without a key for analytics purposes

```
stream
  .windowAll(...)          <- required: "assigner"
  [.trigger(...)]          <- optional: "trigger" (else default trigger)
  [.evictor(...)]          <- optional: "evictor" (else no evictor)
  [.allowedLateness(...)]  <- optional: "lateness" (else zero)
  [.sideOutputLateData(...)] <- optional: "output tag" (else no side output for late data)
  .reduce/aggregate/apply() <- required: "function"
  [.getSideOutput(...)]    <- optional: "output tag"
```

Source: <https://nightlies.apache.org/flink/flink-docs-master/docs/dev/datastream/operators/windows/>

Dataflow vs Execution Graph

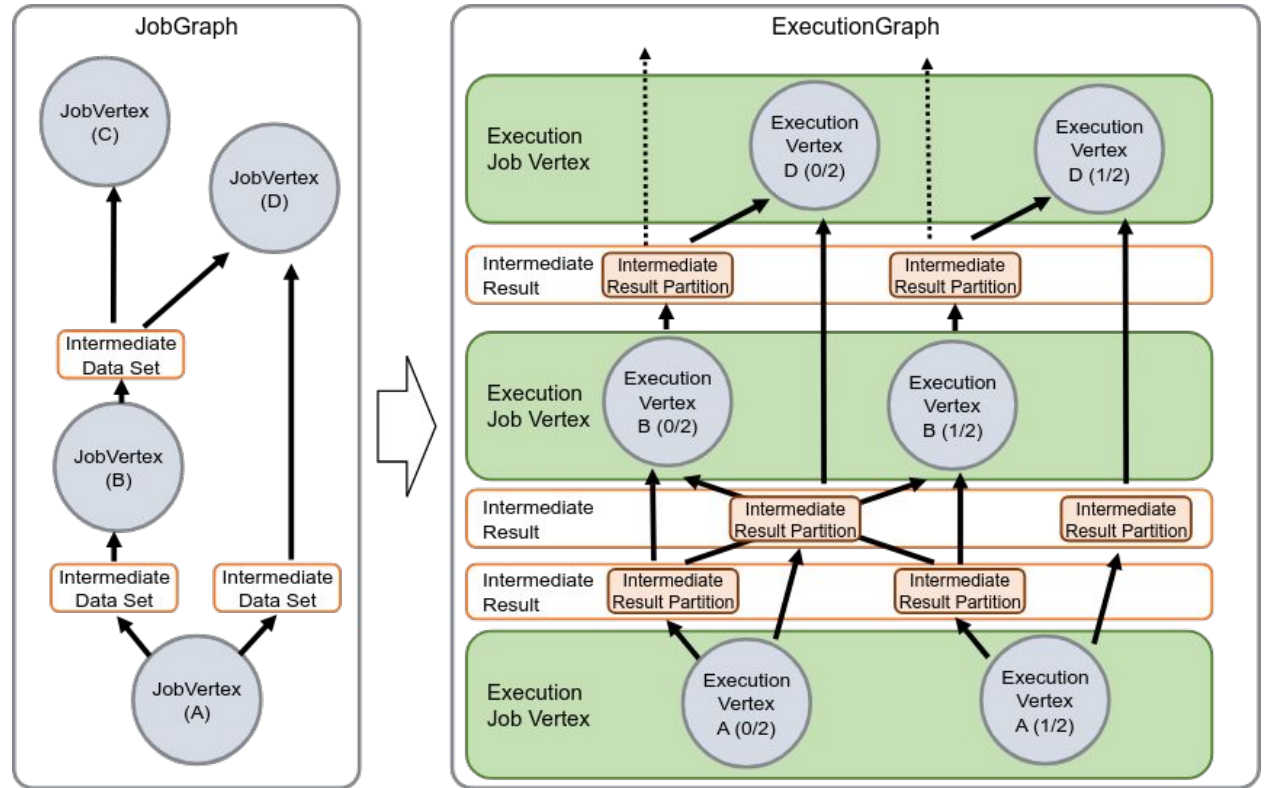



Figure source:

https://nightlies.apache.org/flink/flink-docs-master/docs/internals/job_scheduling/

Windows and Times

Windows



Times





Figure source: <https://nightlies.apache.org/flink/flink-docs-master/docs/concepts/time/>


Batch/Tumbling Windows



Use cases:
period computation
(e.g. stock,
temperature, IoT
data)

Figure source: <https://nightlies.apache.org/flink/flink-docs-master/docs/dev/datastream/operators/windows/>

Sliding windows



Use cases:
Moving average

Figure source: <https://nightlies.apache.org/flink/flink-docs-master/docs/dev/datastream/operators/windows/>

Session Windows

Use cases:
Web/user activities
clicks




Figure source: <https://nightlies.apache.org/flink/flink-docs-master/docs/dev/datastream/operators/windows/>

Window Functions

- **Reduce Function**

- Reduce through the combination of two inputs

- **Aggregate Function**

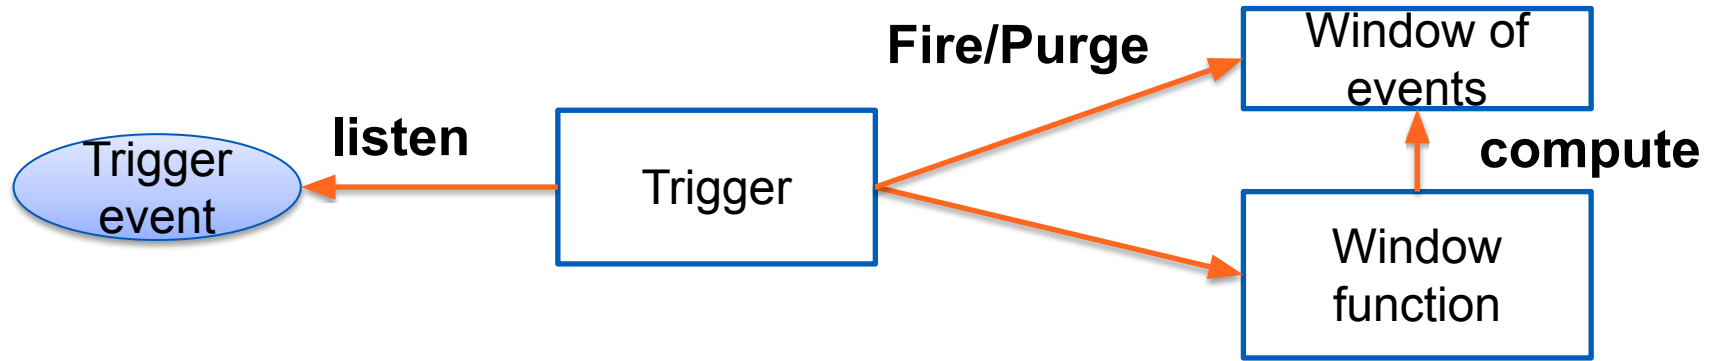
- Add an input into an accumulator

- **ProcessWindow Function**

- Get all elements of the windows and many other information so that you can do many tasks

Triggers & Evictor

- **Trigger:** determine if a window is ready for window functions



Evictor: actions **after** the trigger fires and **before** **and/or after** the windows function is called

Fault tolerance

- Principles: checkpointing, restarts operators from the latest successful checkpoints
- Need support from data stream sources/sinks w.r.t. (end-to-end) exactly once message receiving and result delivery

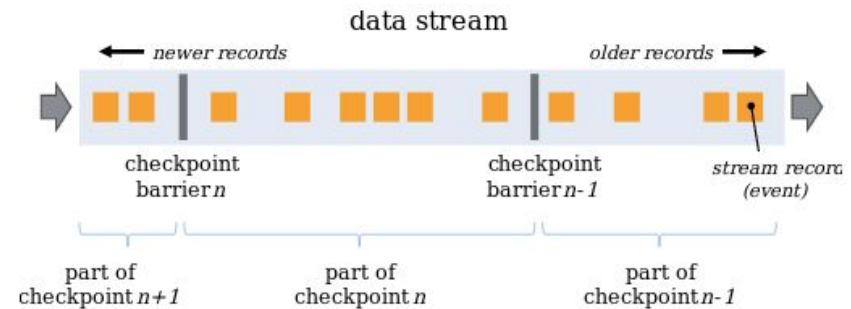


Figure source:

<https://nightlies.apache.org/flink/flink-docs-release-1.18/docs/concepts/stateful-stream-processing/>

Example with Base Transceiver Station


Data in our git

```
station_id,datapoint_id,alarm_id,event_time,value,valueThreshold,isActive,storedtime
1161115016,121,308,2017-02-18 18:28:05 UTC,240,240,false,
1161114050,143,312,2017-02-18 18:56:20 UTC,28.5,28,true,
1161115040,141,312,2017-02-18 18:22:03 UTC,56.5,56,true,
1161114008,121,308,2017-02-18 18:34:09 UTC,240,240,false,
1161115040,141,312,2017-02-18 18:20:49 UTC,56,56,false,
```

See the code in our git:

<https://github.com/rdsea/bigdataplatforms/tree/master/tutorials/streamingwithflink/>

Simple example



Monitoring

Apache Flink Dashboard

Overview

Jobs

Running Jobs

Completed Jobs

Task Managers

Job Manager

Submit New Job

Simple CS-E4640 BTS Flink Application RUNNING 2

ID: 81efb959d02448b7c44328ad75a824af | Start Time: 2019-11-04 14:00:14 | Duration: 55s

[Overview](#) | [Exceptions](#) | [TimeLine](#) | [Checkpoints](#) | [Configuration](#)

Source: Custom Source -> Flat Map
Parallelism: 1

HASH

Window(SlidingProcessingTimeWindows(60000, 5000), ProcessingTimeTrigger, MyProcessWindowFunction) -> (Sink: Unnamed, Sink: Print to Std. Out)
Parallelism: 1

Detail

SubTasks

TaskManagers

Watermarks

Accumulators

BackPressure

Metrics

Window(SlidingProcessingTimeWindows(60000, 5000), ProcessingTimeTrigger, MyProcessWindowFunction) -> (Sink: Unnamed, Sink: Print to Std. Out)

Status: RUNNING

Task: 1

Parallelism: 1

Records Sent: 0

Start Time: 2019-11-04 14:00:14

Bytes Received: 1.64 KB

End Time: -

Records Received: 29

Duration: 55s

Bytes Sent: 0 B

Name	Status	Bytes Received	Records Received	Bytes Sent	Records Sent	Parallelism	Start Time	Duration	End Time	Tasks
Window(SlidingProcessingTimeWindows(60000, 5000), ProcessingTimeTrigger, MyProcessWindowFunction) -> (Sink: Unnamed, Sink: Print to Std. Out)	RUNNING	1.64 KB	29	0 B	0	1	2019-11-04 14:00:14	55s	-	1
Source: Custom Source -> Flat Map	RUNNING	0 B	0	1.61 KB	29	1	2019-11-04 14:00:14	55s	-	1

Cancel Job

Summary

- **Focus:**

- Practical programming with one of the stacks:
 - *Apache Flink Stream API (with different connectors)*
 - *Kafka Streams*
- Check the common concepts in other tools/systems

- **Action:**

- Work on use cases where you can use stream analytics (as a user/developer) \Rightarrow there are many interesting analytics
- Provision services for stream processing (as a platform)

Thanks!

Hong-Linh Truong
Department of Computer Science

rdsea.github.io